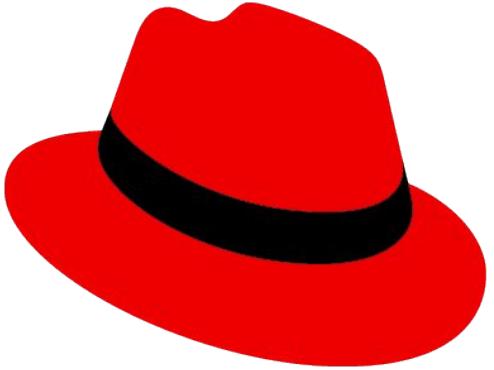


# Kubernetes Patterns

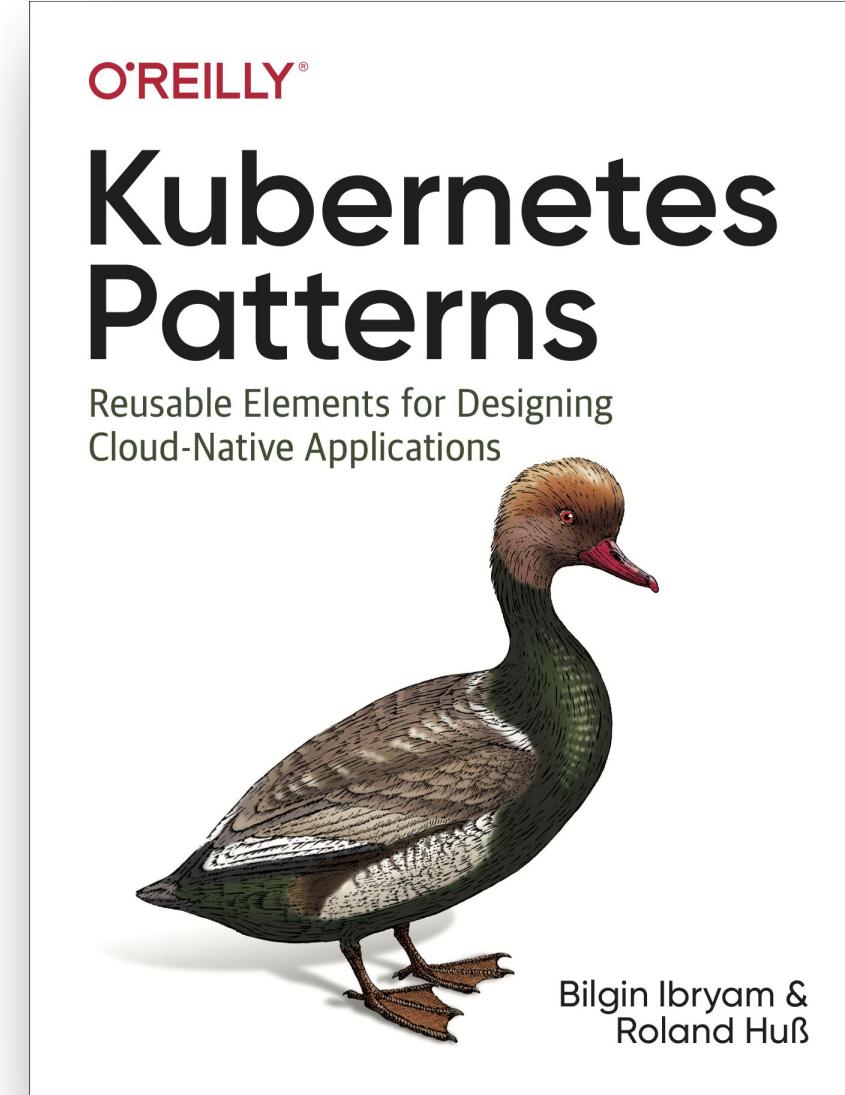
Reusable Elements for Designing  
Cloud-Native Applications

Dr. Roland Huß  
Principal Software Engineer  
@ro14nd  
<https://k8spatterns.io>



# Red Hat

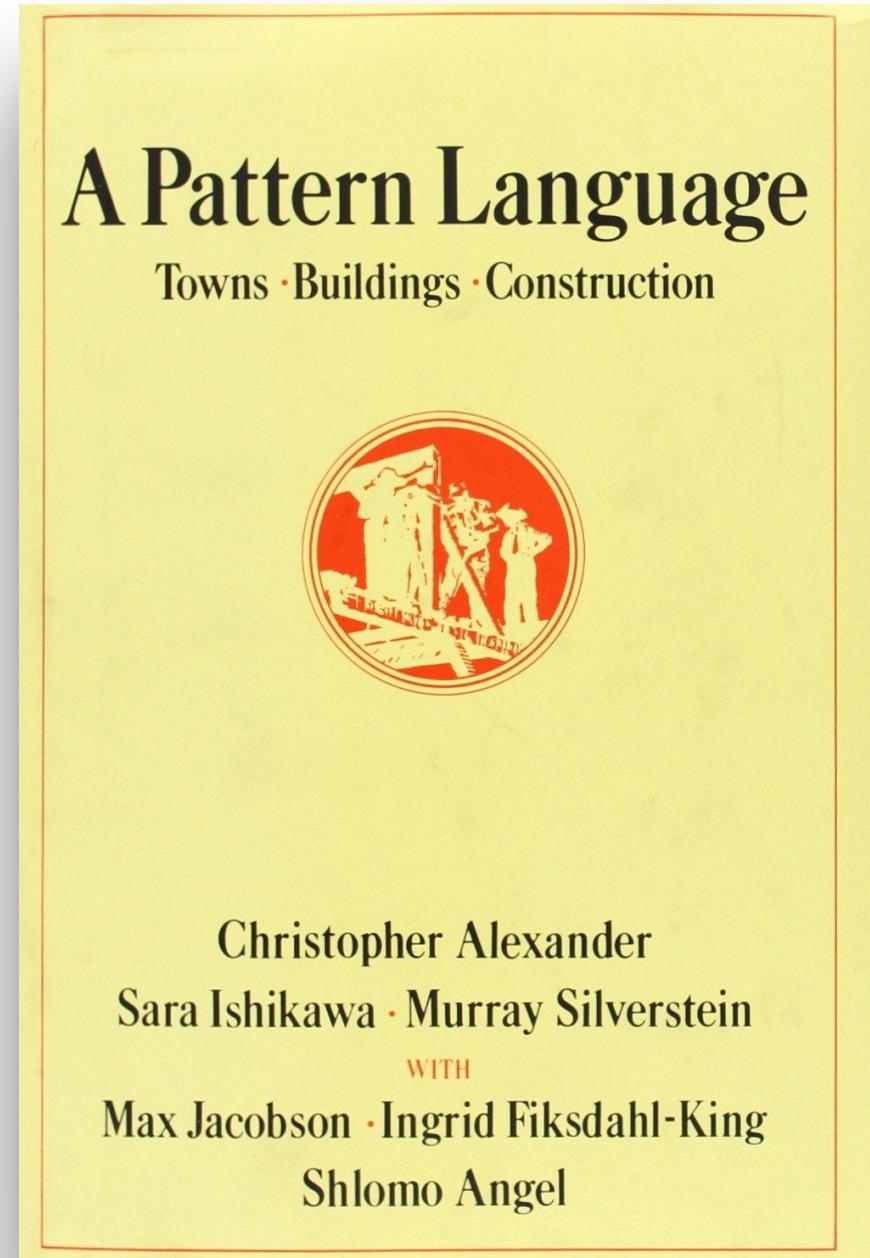
## New Logo !



<https://k8spatterns.io>

The background of the image is a intricate geometric pattern. It features a repeating motif of overlapping diamond shapes, creating a sense of depth and movement. The colors used are various shades of brown, tan, and black, which are arranged in a way that suggests a three-dimensional perspective. The overall effect is one of a sophisticated and artistic design.

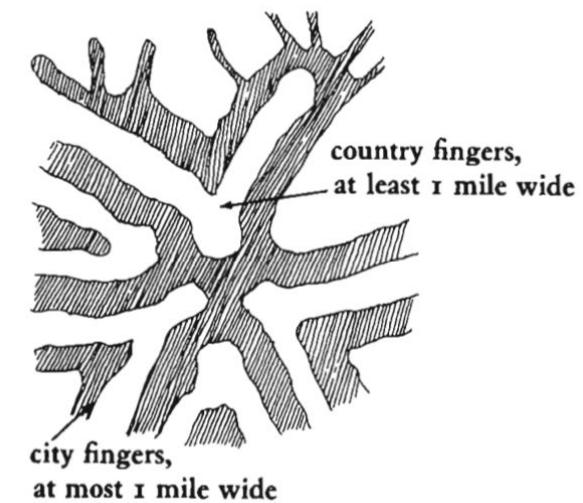
# Patterns

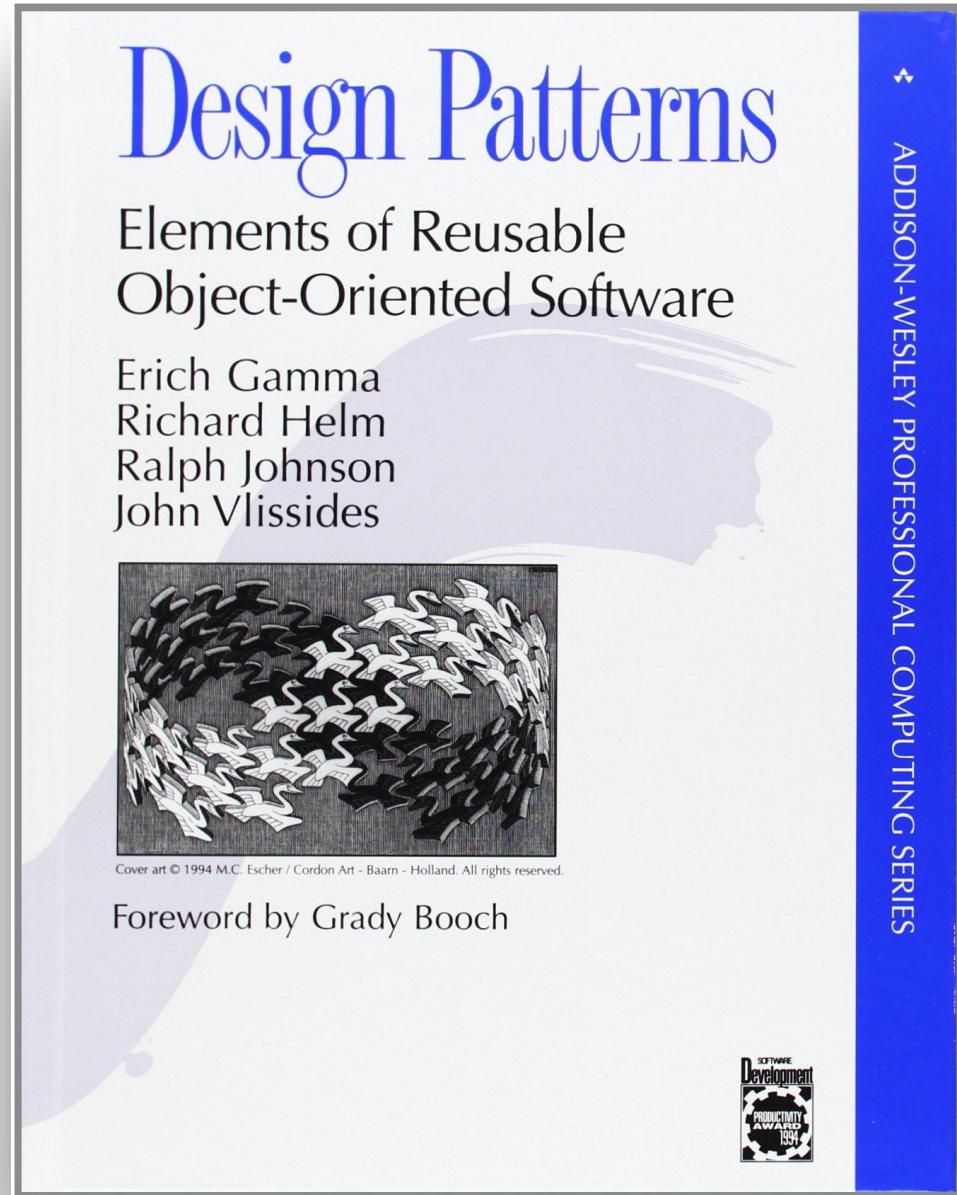


### 3 CITY COUNTRY FINGERS\*\*



*When the countryside is far away  
the city becomes a prison.*





## Patterns

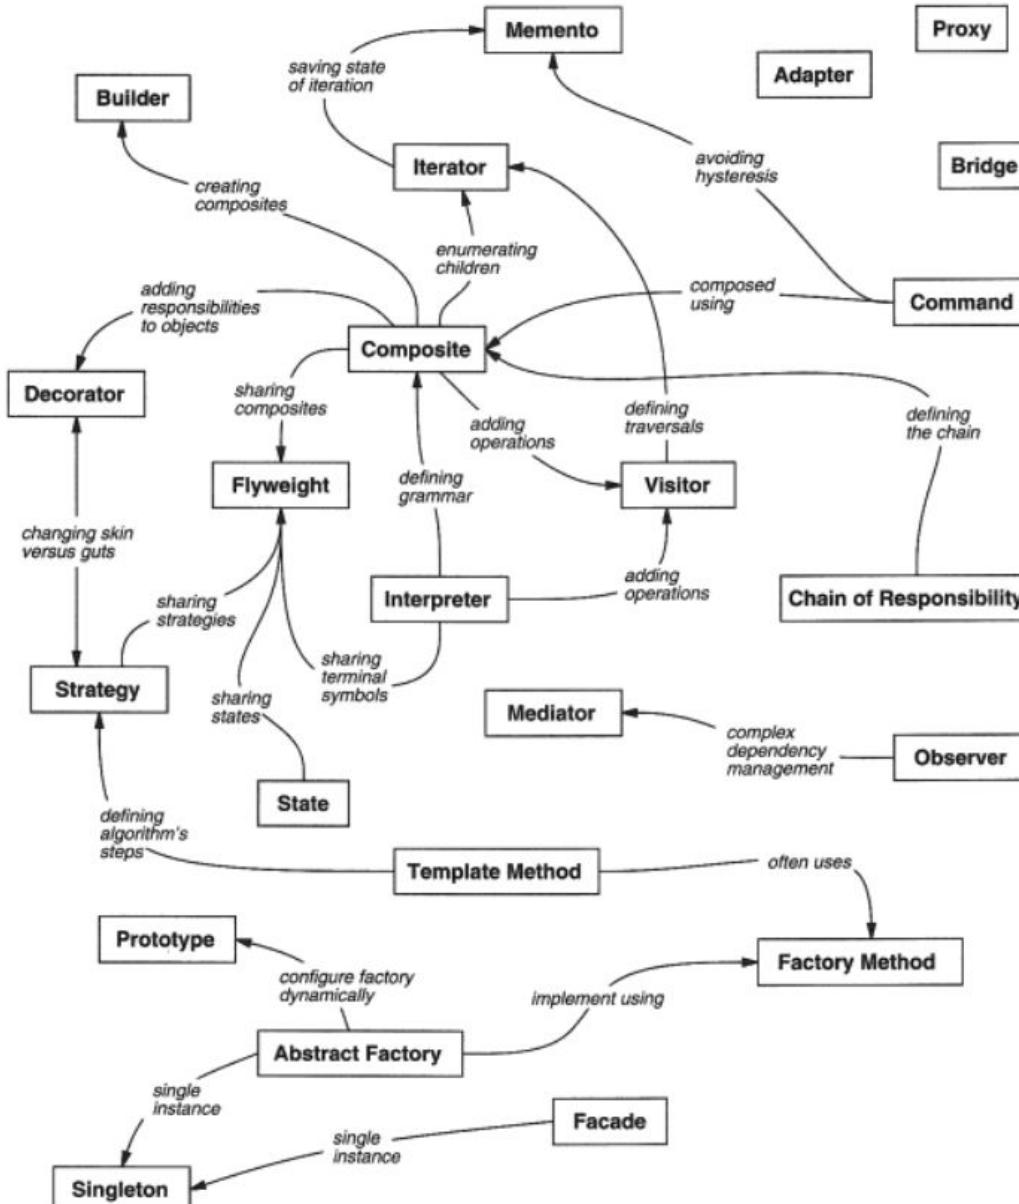
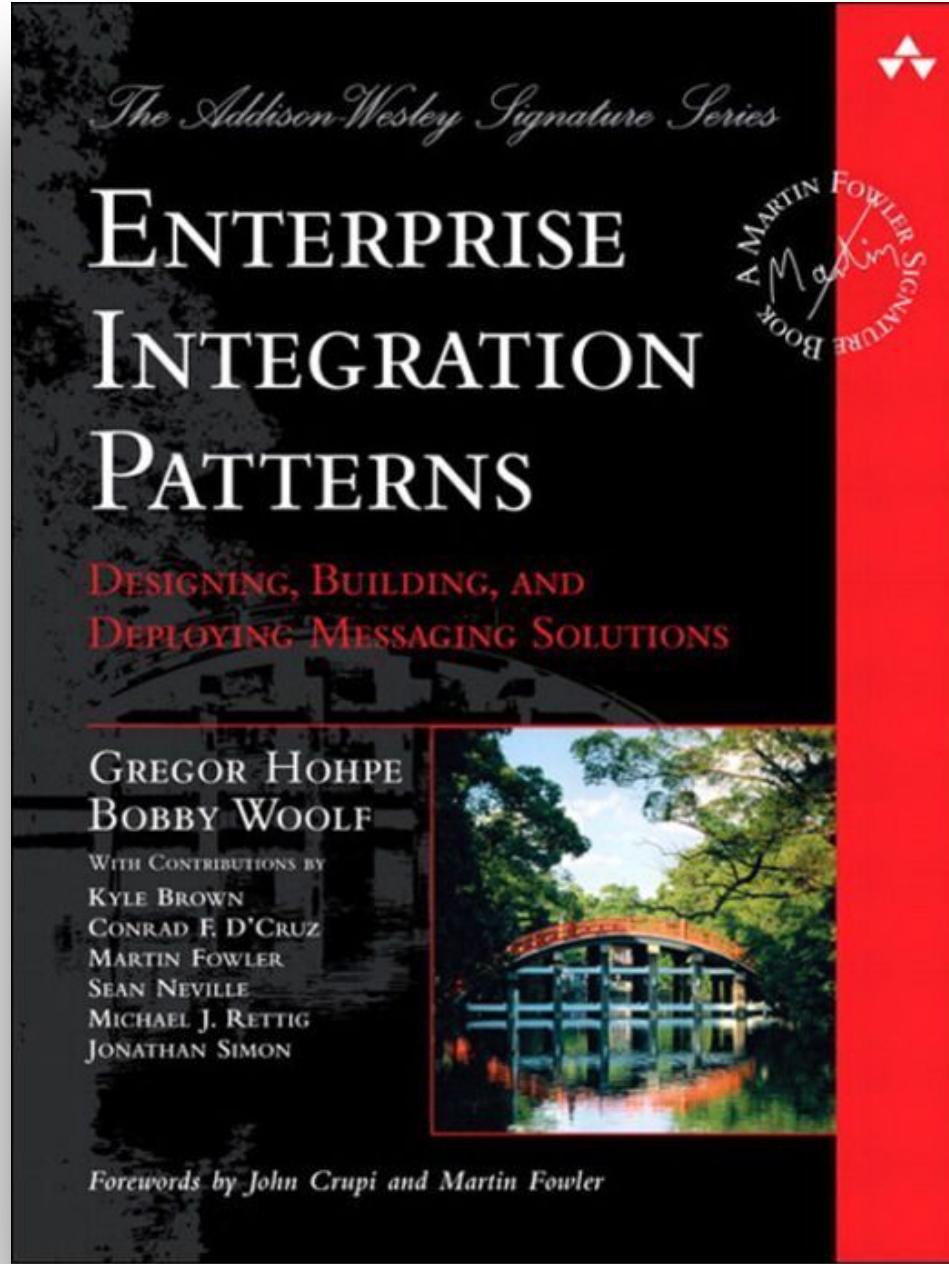


Figure 1.1: Design pattern relationships



# Kubernetes

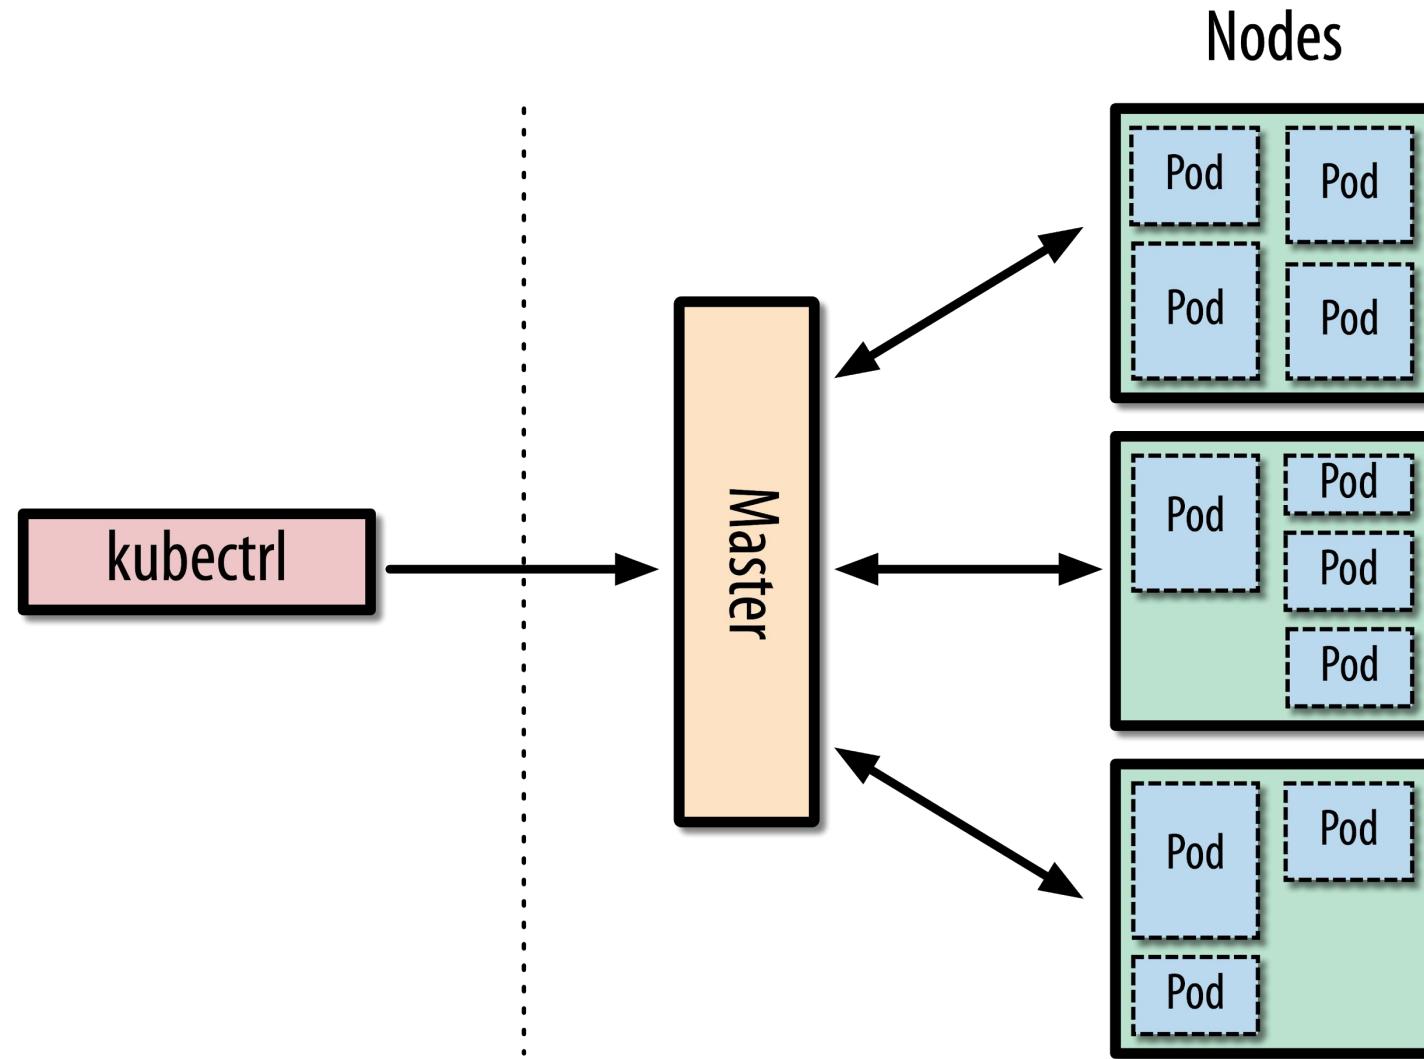


# Kubernetes

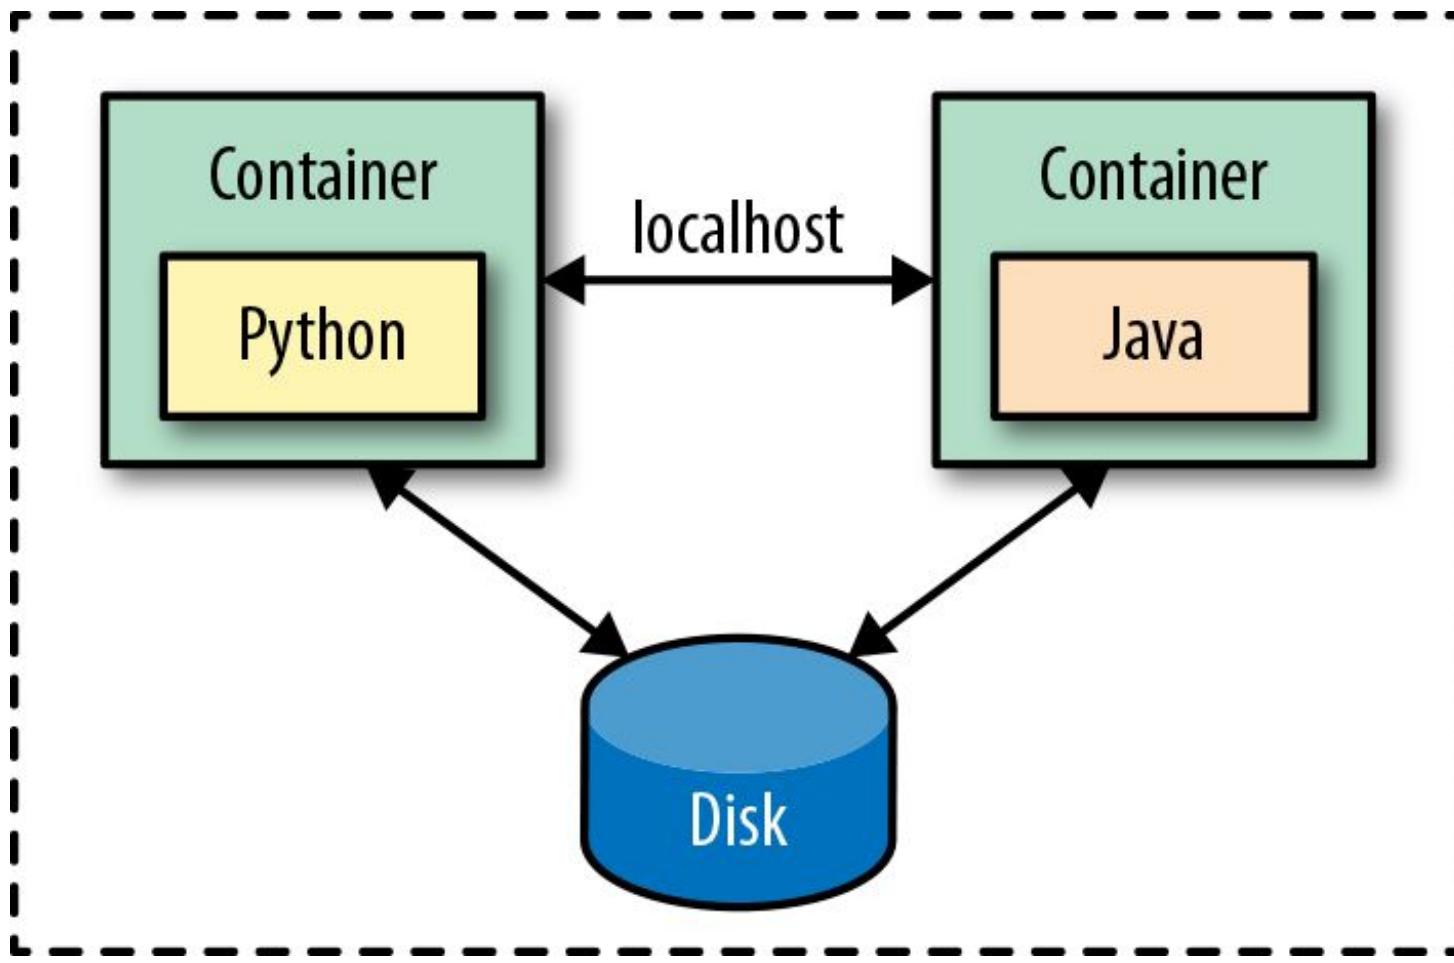


- Open Source container orchestration system started by Google in 2014
  - ※ Scheduling
  - ※ Self-healing
  - ※ Horizontal and vertical scaling
  - ※ Service discovery
  - ※ Rollout and Rollbacks
- Declarative resource-centric REST API

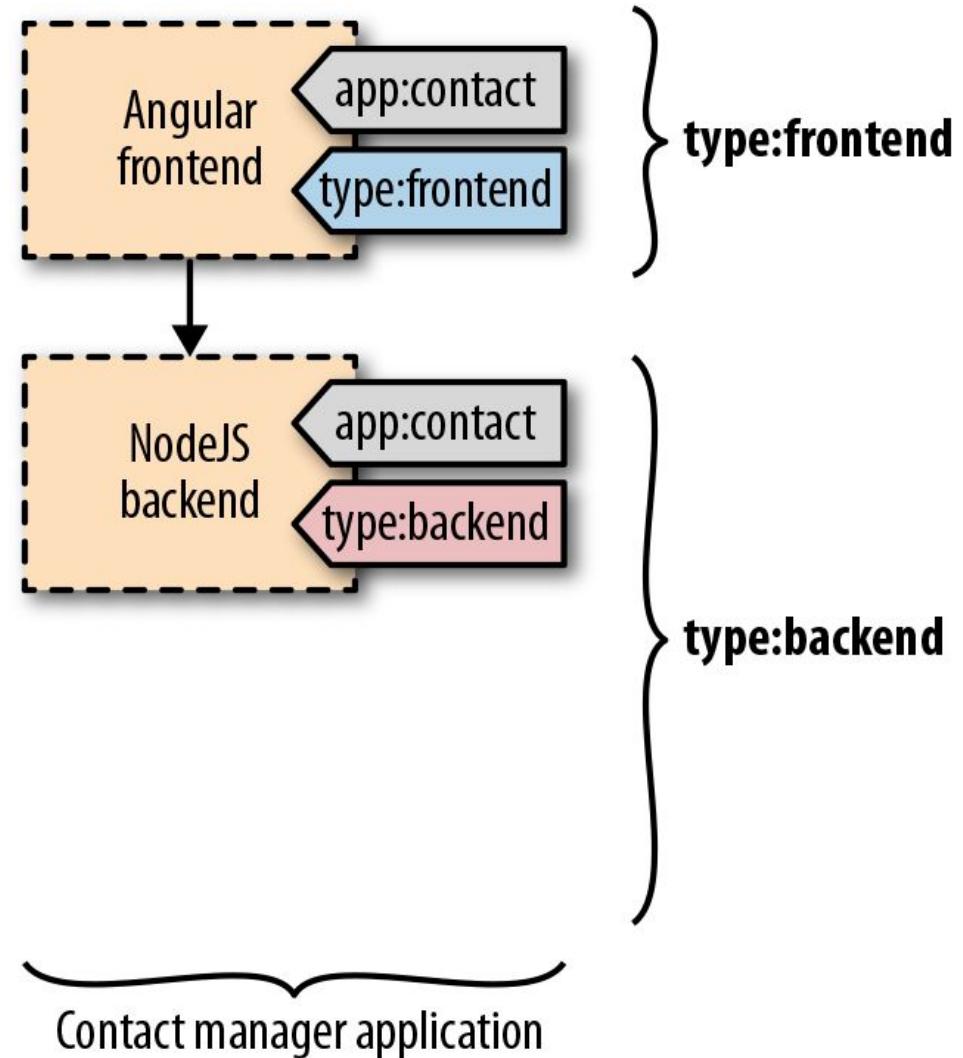
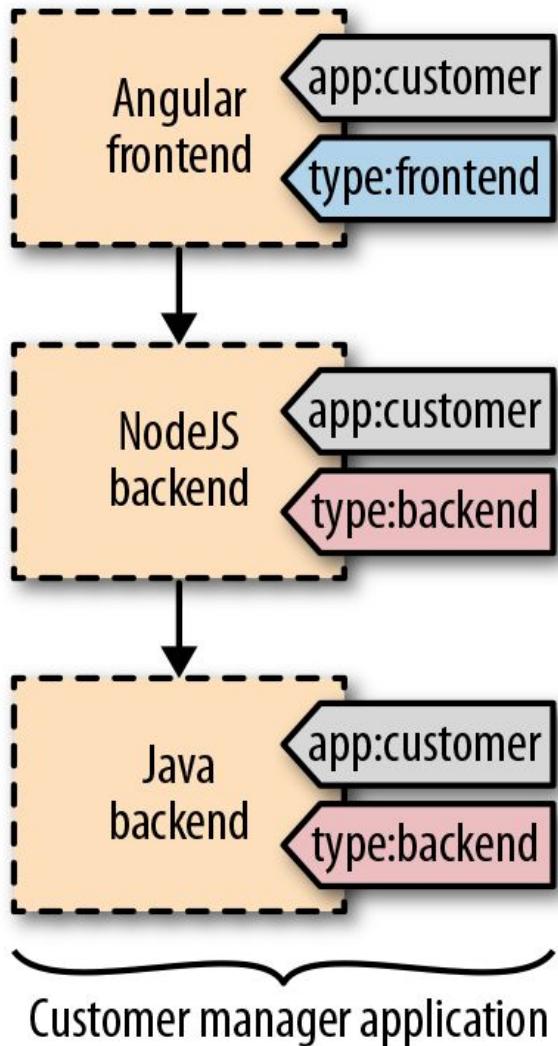
# Architecture

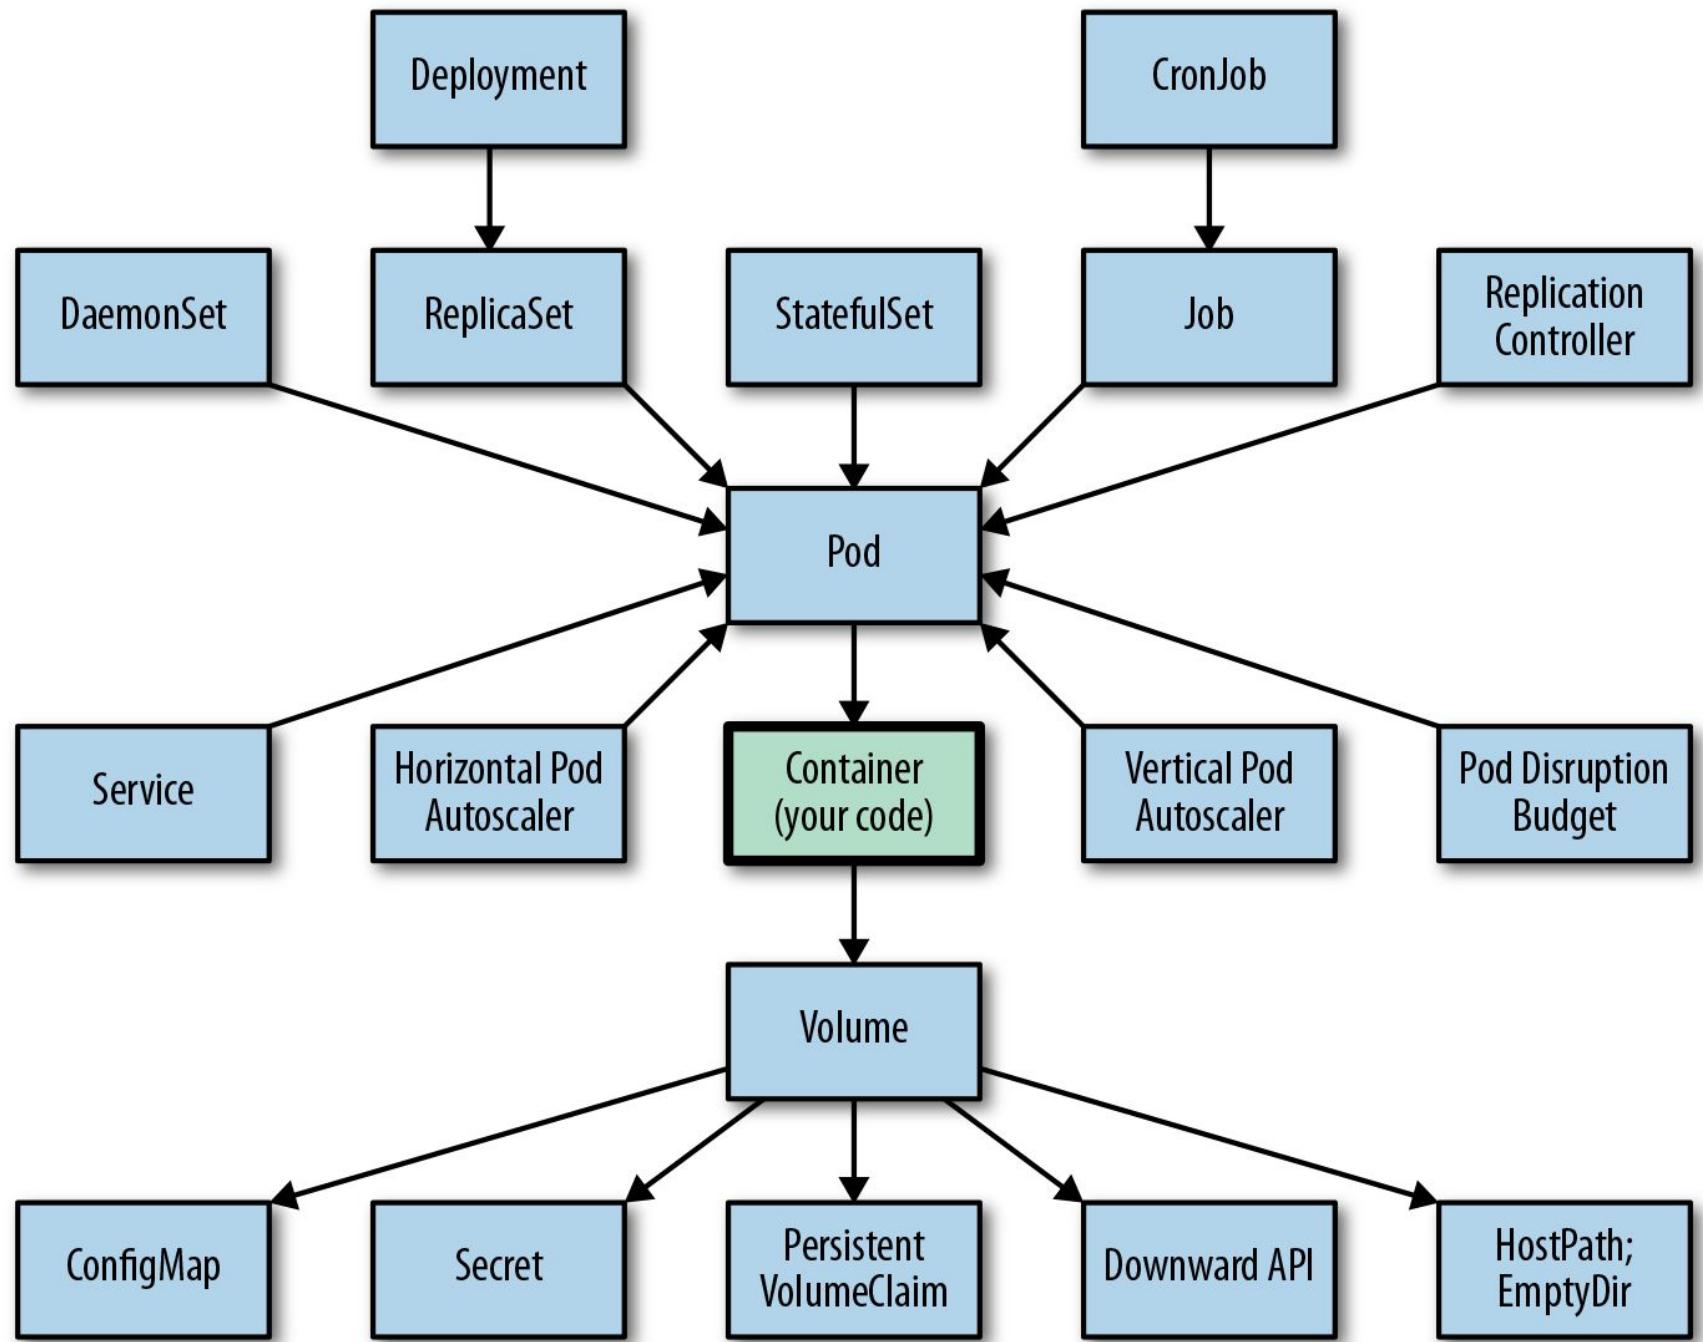


# Pod



# Labels







# Predictable Demands

---

# How to declare application requirements.

# Application Requirements

- Declared requirements help in
  - Scheduling decisions
  - Capacity planning
  - Matching infrastructure services
- Runtime dependencies
  - Persistent Volumes
  - Host ports
  - Dependencies on ConfigMaps and Secrets

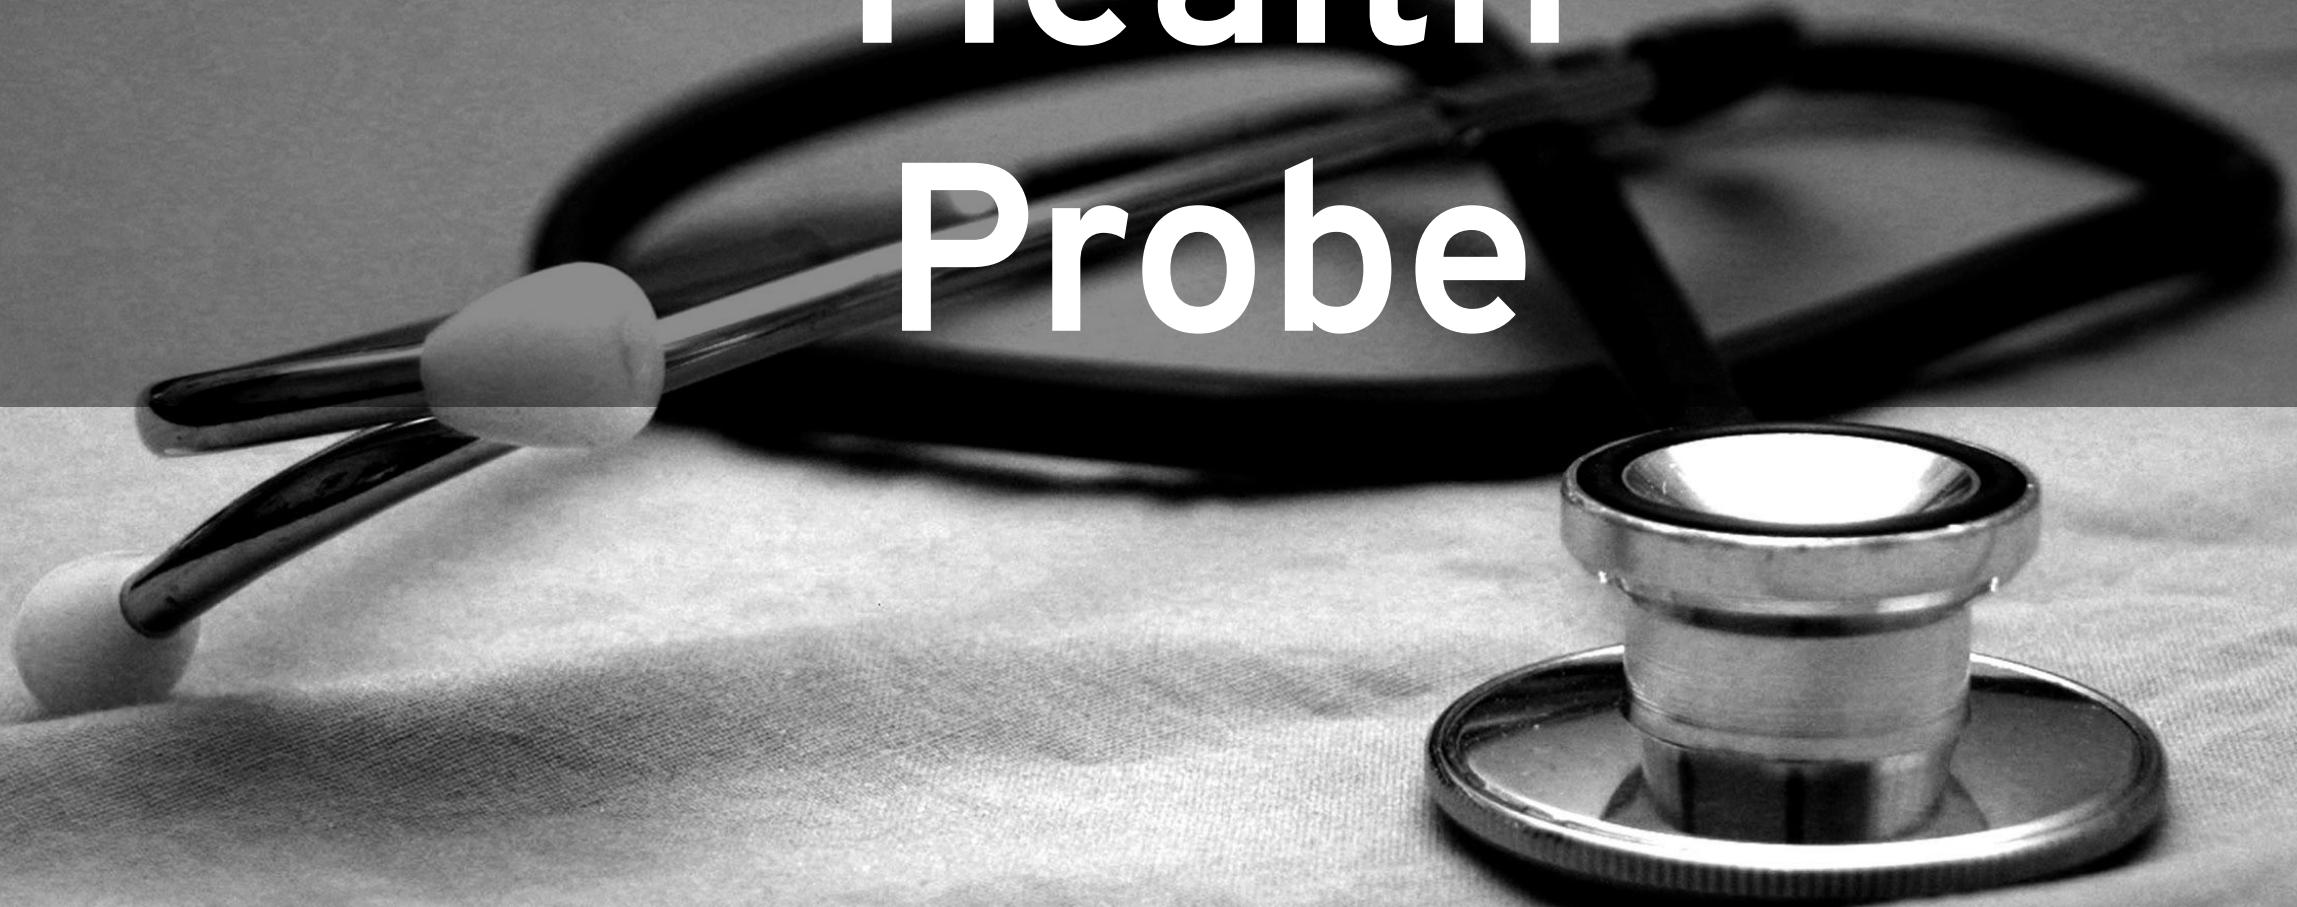
# Resource Profile

```
apiVersion: v1
kind: Pod
metadata:
  name: http-server
spec:
  containers:
  - image: nginx
    name: nginx
  resources:
    requests:
      cpu: 200m
      memory: 100Mi
    limits:
      cpu: 300m
      memory: 200Mi
```

# Quality-of-Service Classes

- **Best Effort**
  - No requests or limits
- **Burstable**
  - requests < limits
- **Guaranteed**
  - requests == limits

# Health Probe



---

# How to communicate an application's health state to Kubernetes.

# Monitoring Health

- Process health checks
  - Checks running process
  - Restarts container if container process died
- Application provided Health Probes
  - **Liveness Probe:** Check application health
  - **Readiness Probe:** Check readiness to process requests

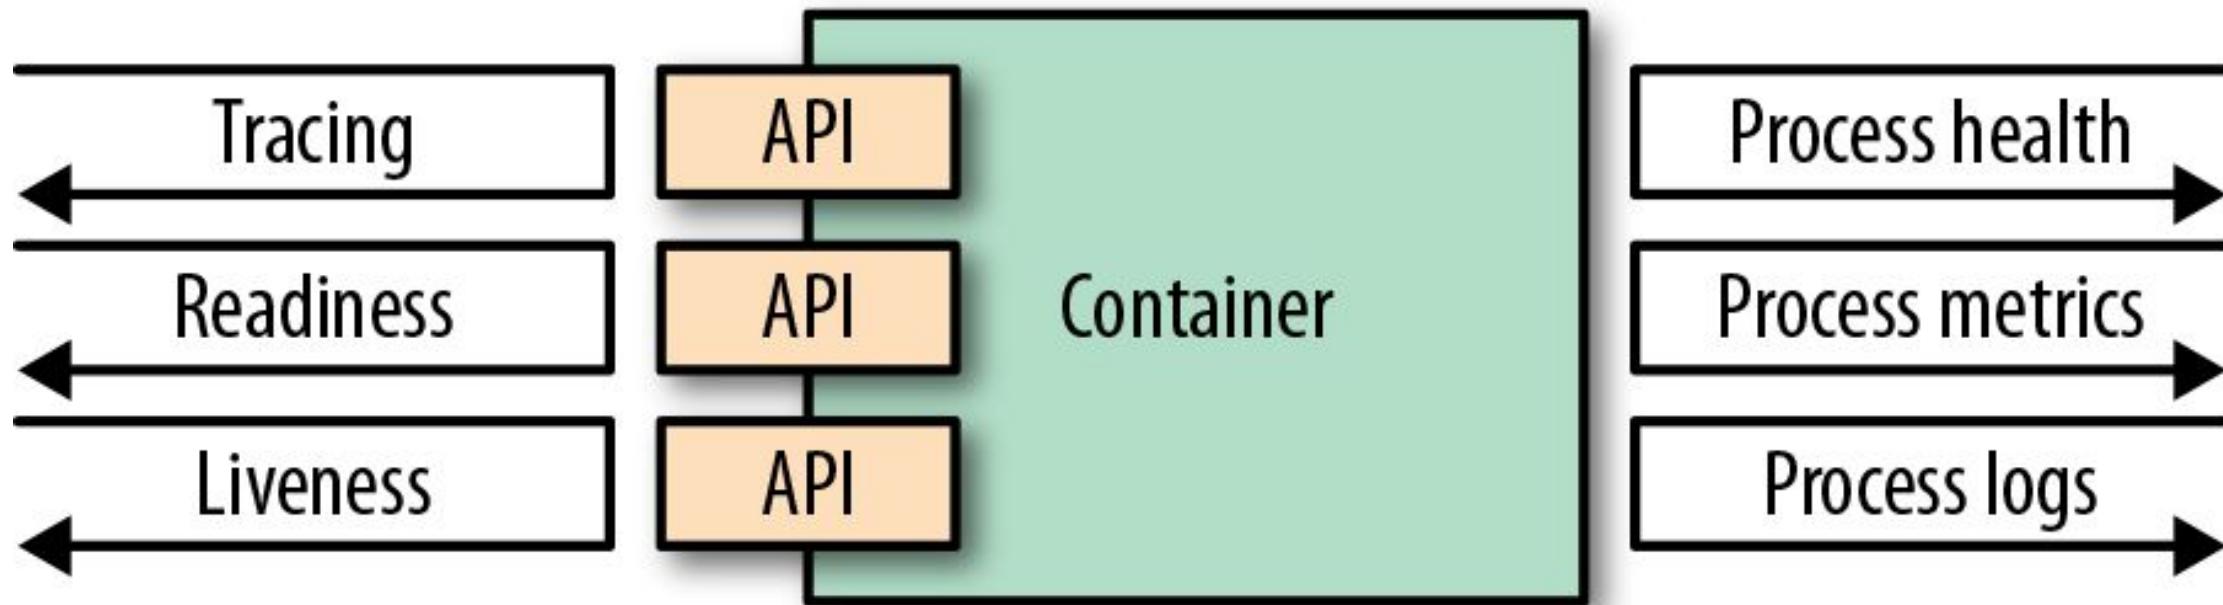
# Liveness & Readiness

- Liveness Probe
  - Restarting containers if liveness probes fail
- Readiness Probe
  - Removing from service endpoint if readiness probe fails
- Probe methods
  - HTTP endpoint
  - TCP socket endpoint
  - Unix command's return value

## Health Probe

```
apiVersion: v1
kind: Pod
metadata:
  name: pod-with-readiness-check
spec:
  containers:
  - image: k8spatterns/random-generator:1.0
    name: random-generator
    livenessProbe:
      httpGet:
        path: /actuator/health
        port: 8080
      initialDelaySeconds: 30
    readinessProbe:
      exec:
        command: [ "stat", "/var/run/random-generator-ready" ]
```

# Container Observability Options



# Declarative Deployment

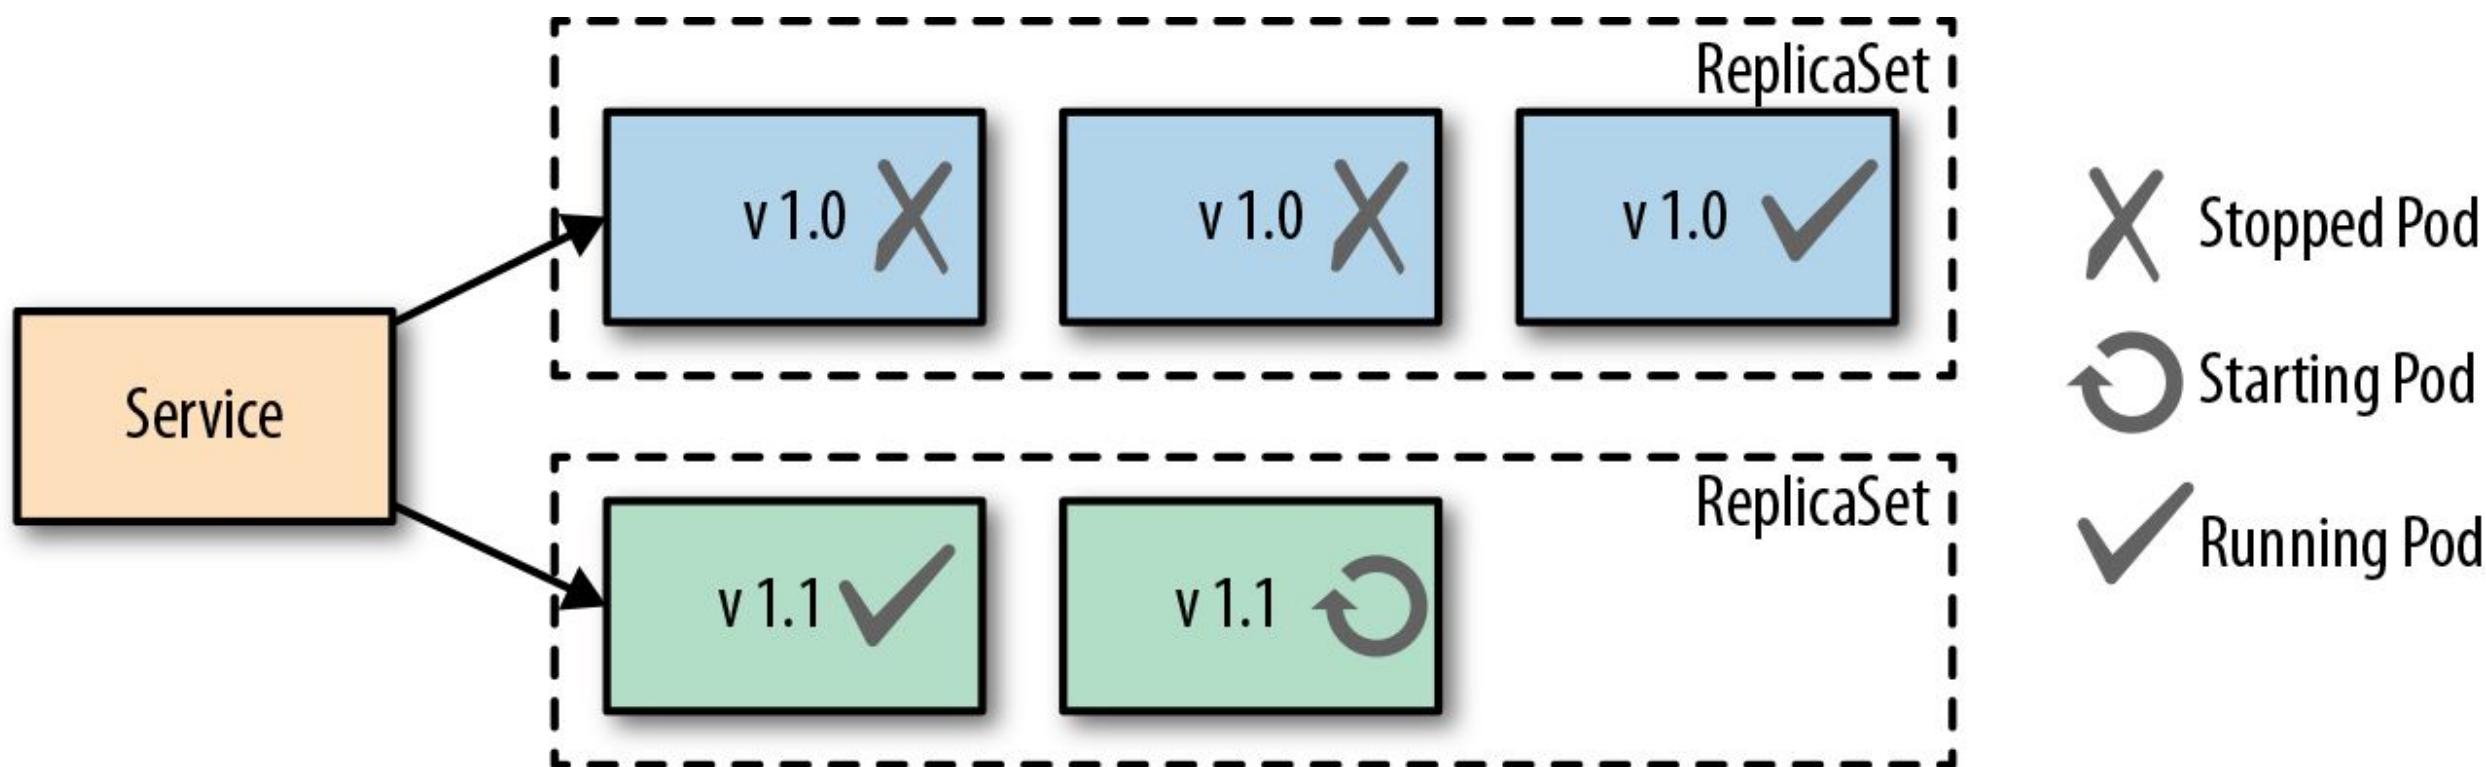
---

# How to perform application installations and upgrades by configuration.

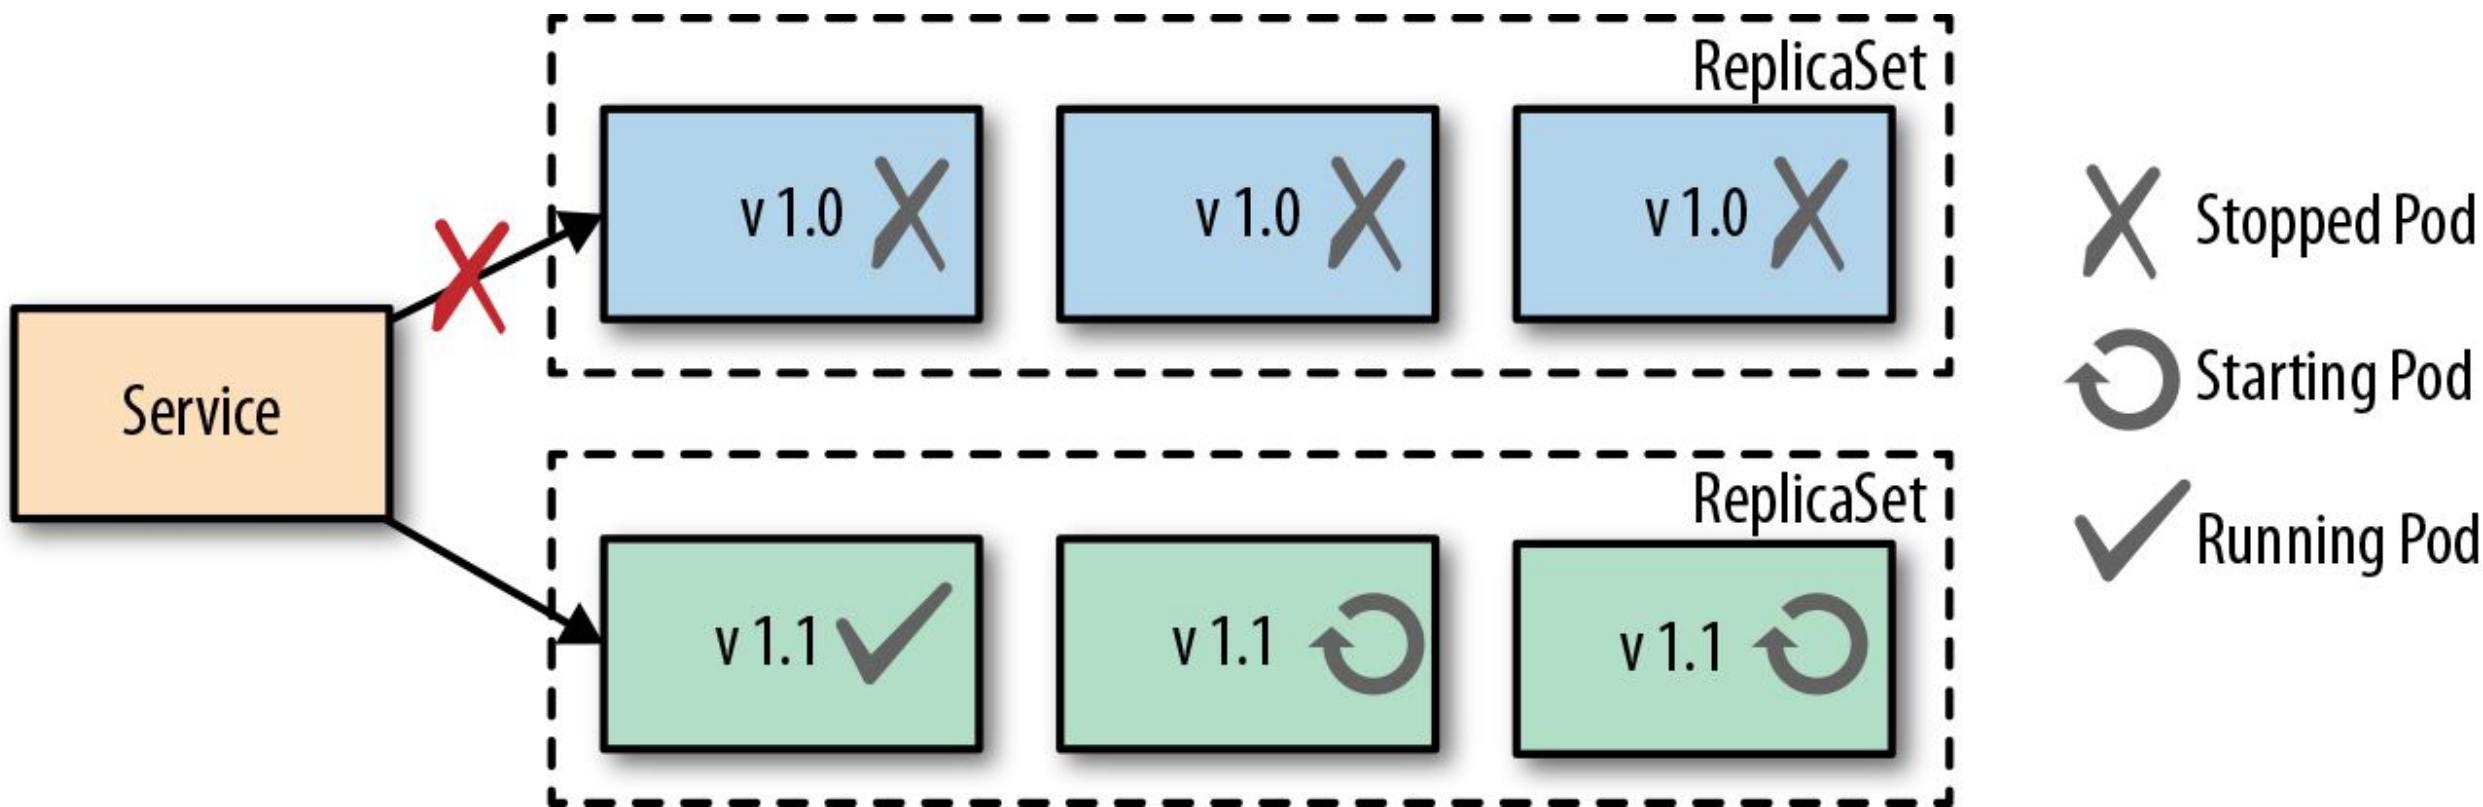
# Deployment

- **Declarative** vs. **Imperative** deployment
- Deployment Kubernetes Resource:
  - Holds template for Pod
  - Creates ReplicaSet on the fly
  - Allows rollback
  - Update strategies are declarable

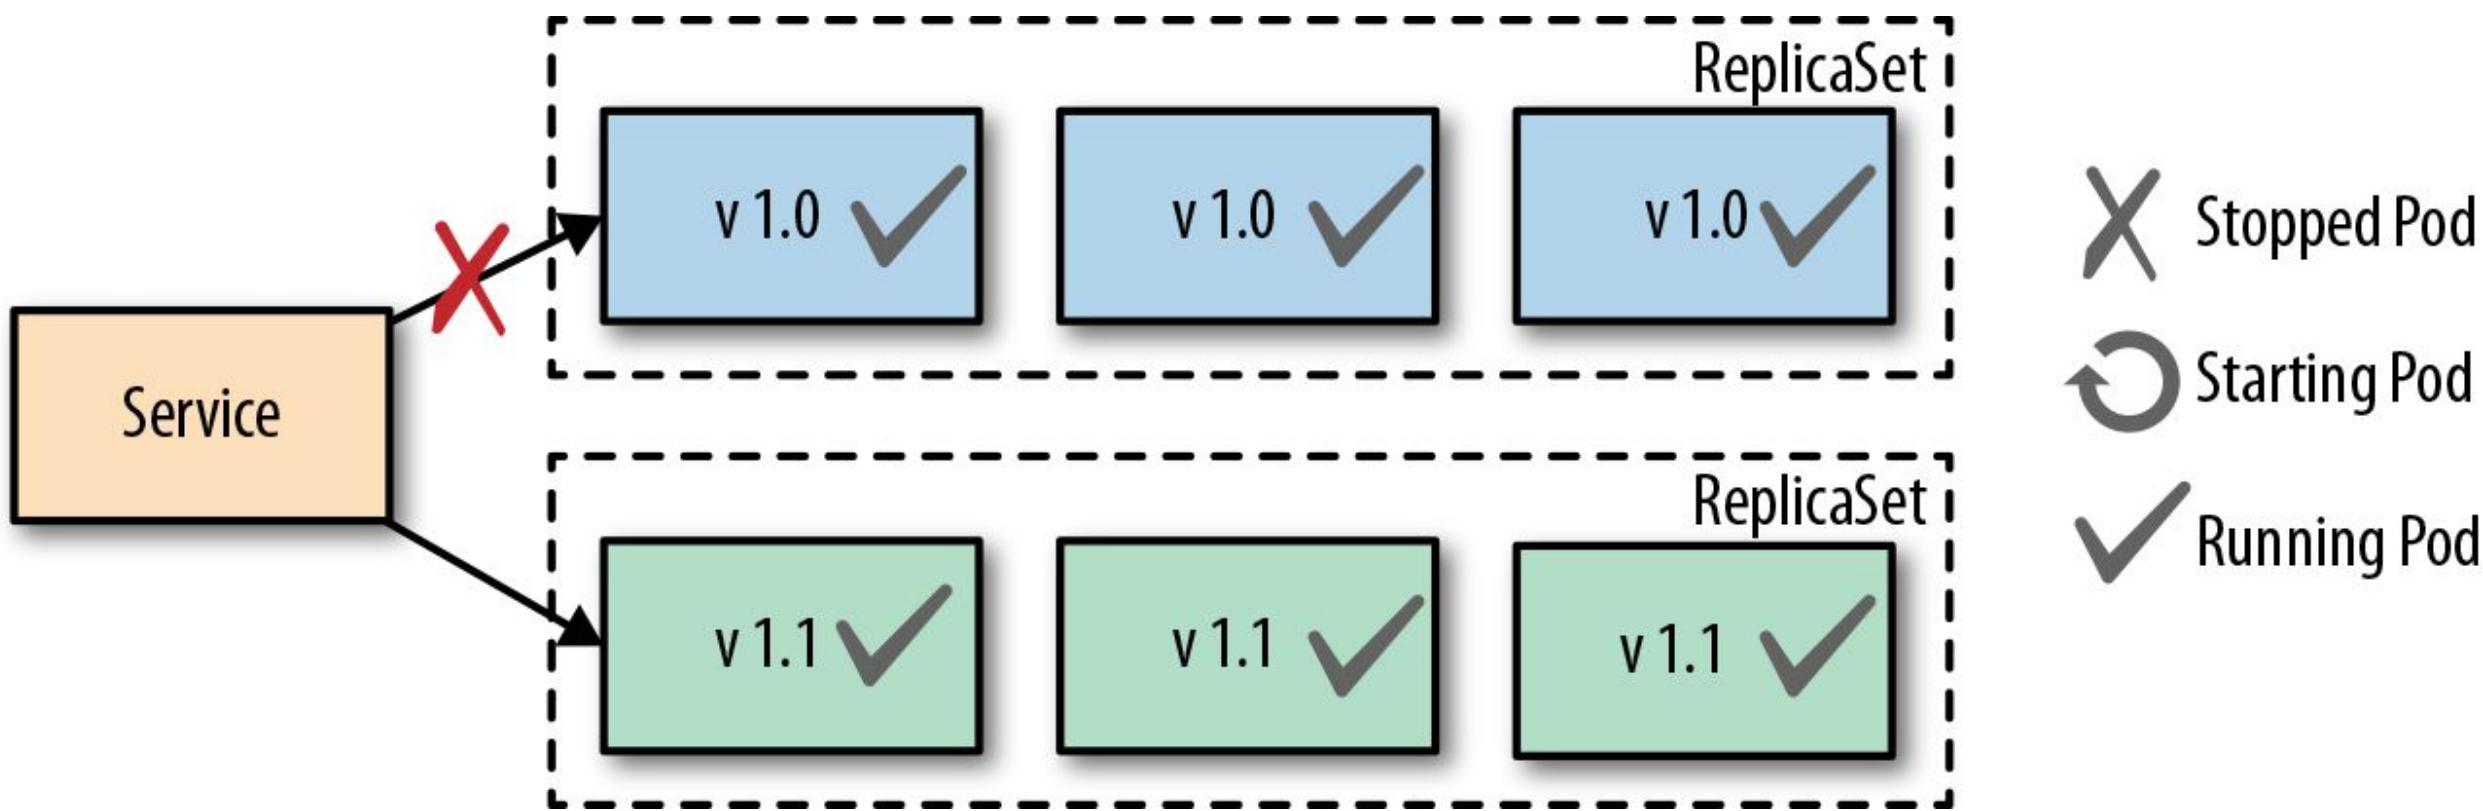
# Rolling Deployment



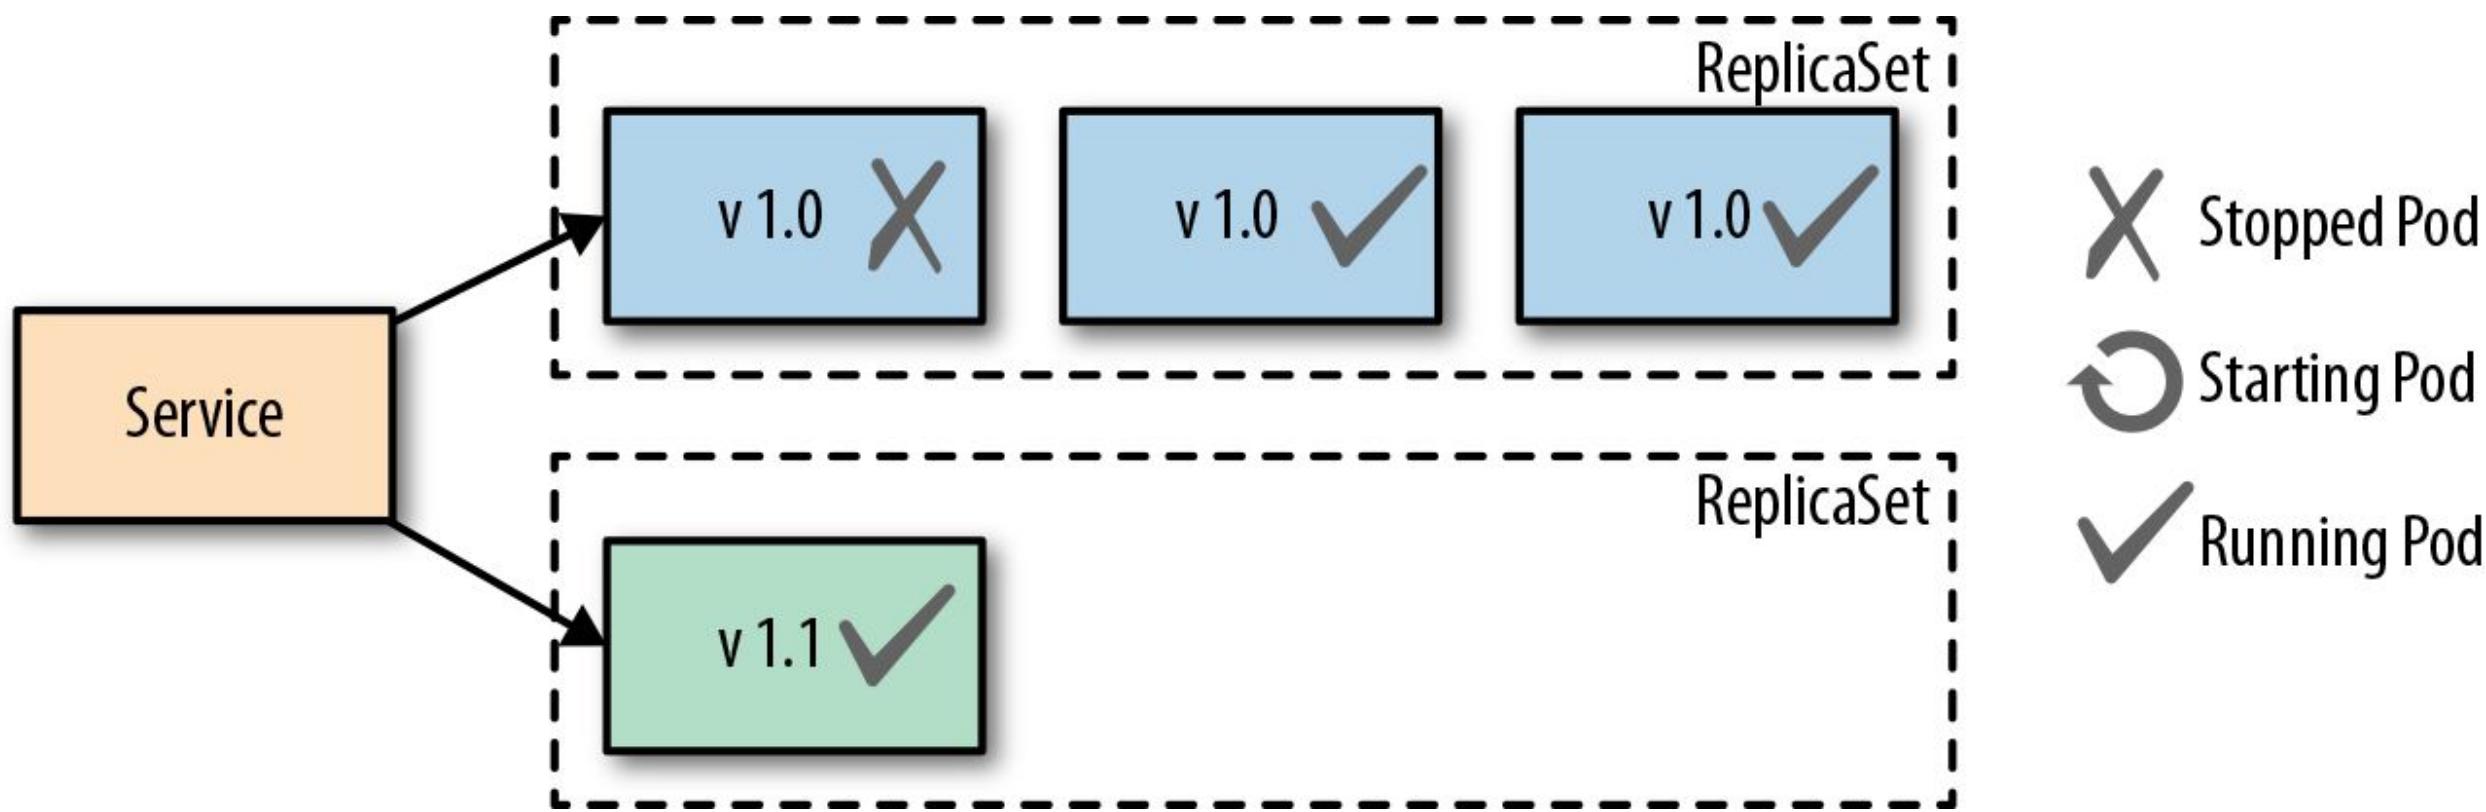
# Fixed Deployment



# Blue-Green Release

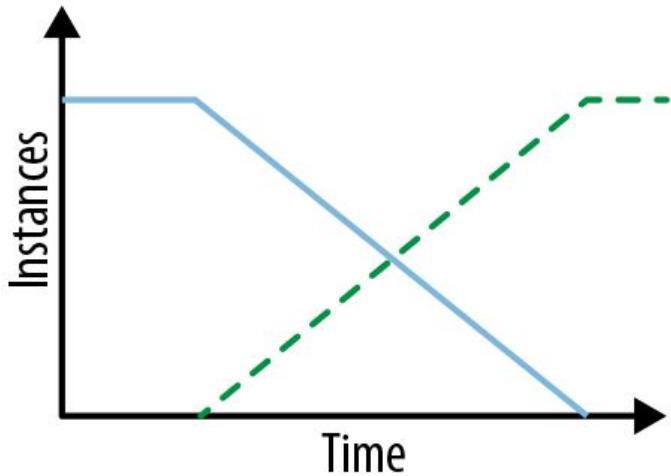


# Canary Release

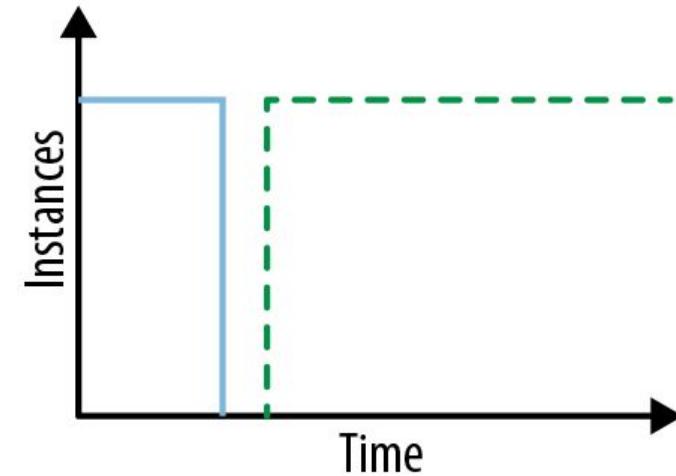


## Declarative Deployment

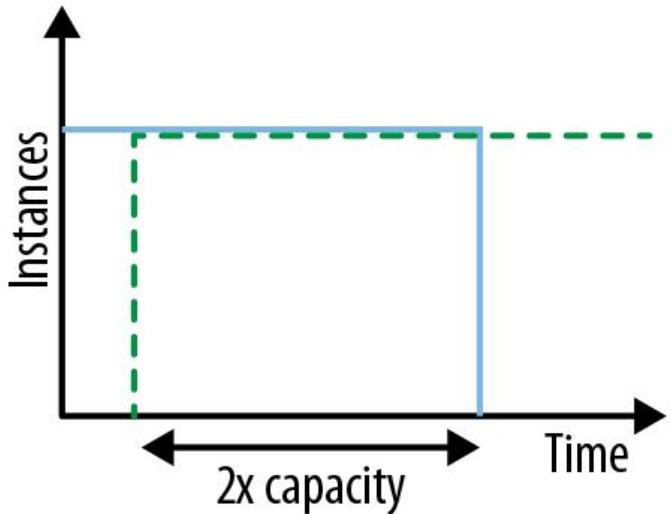
**Rolling deployment**



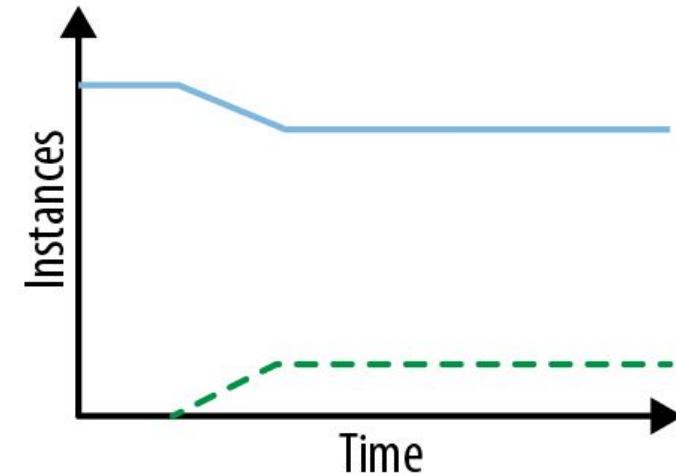
**Fixed deployment**

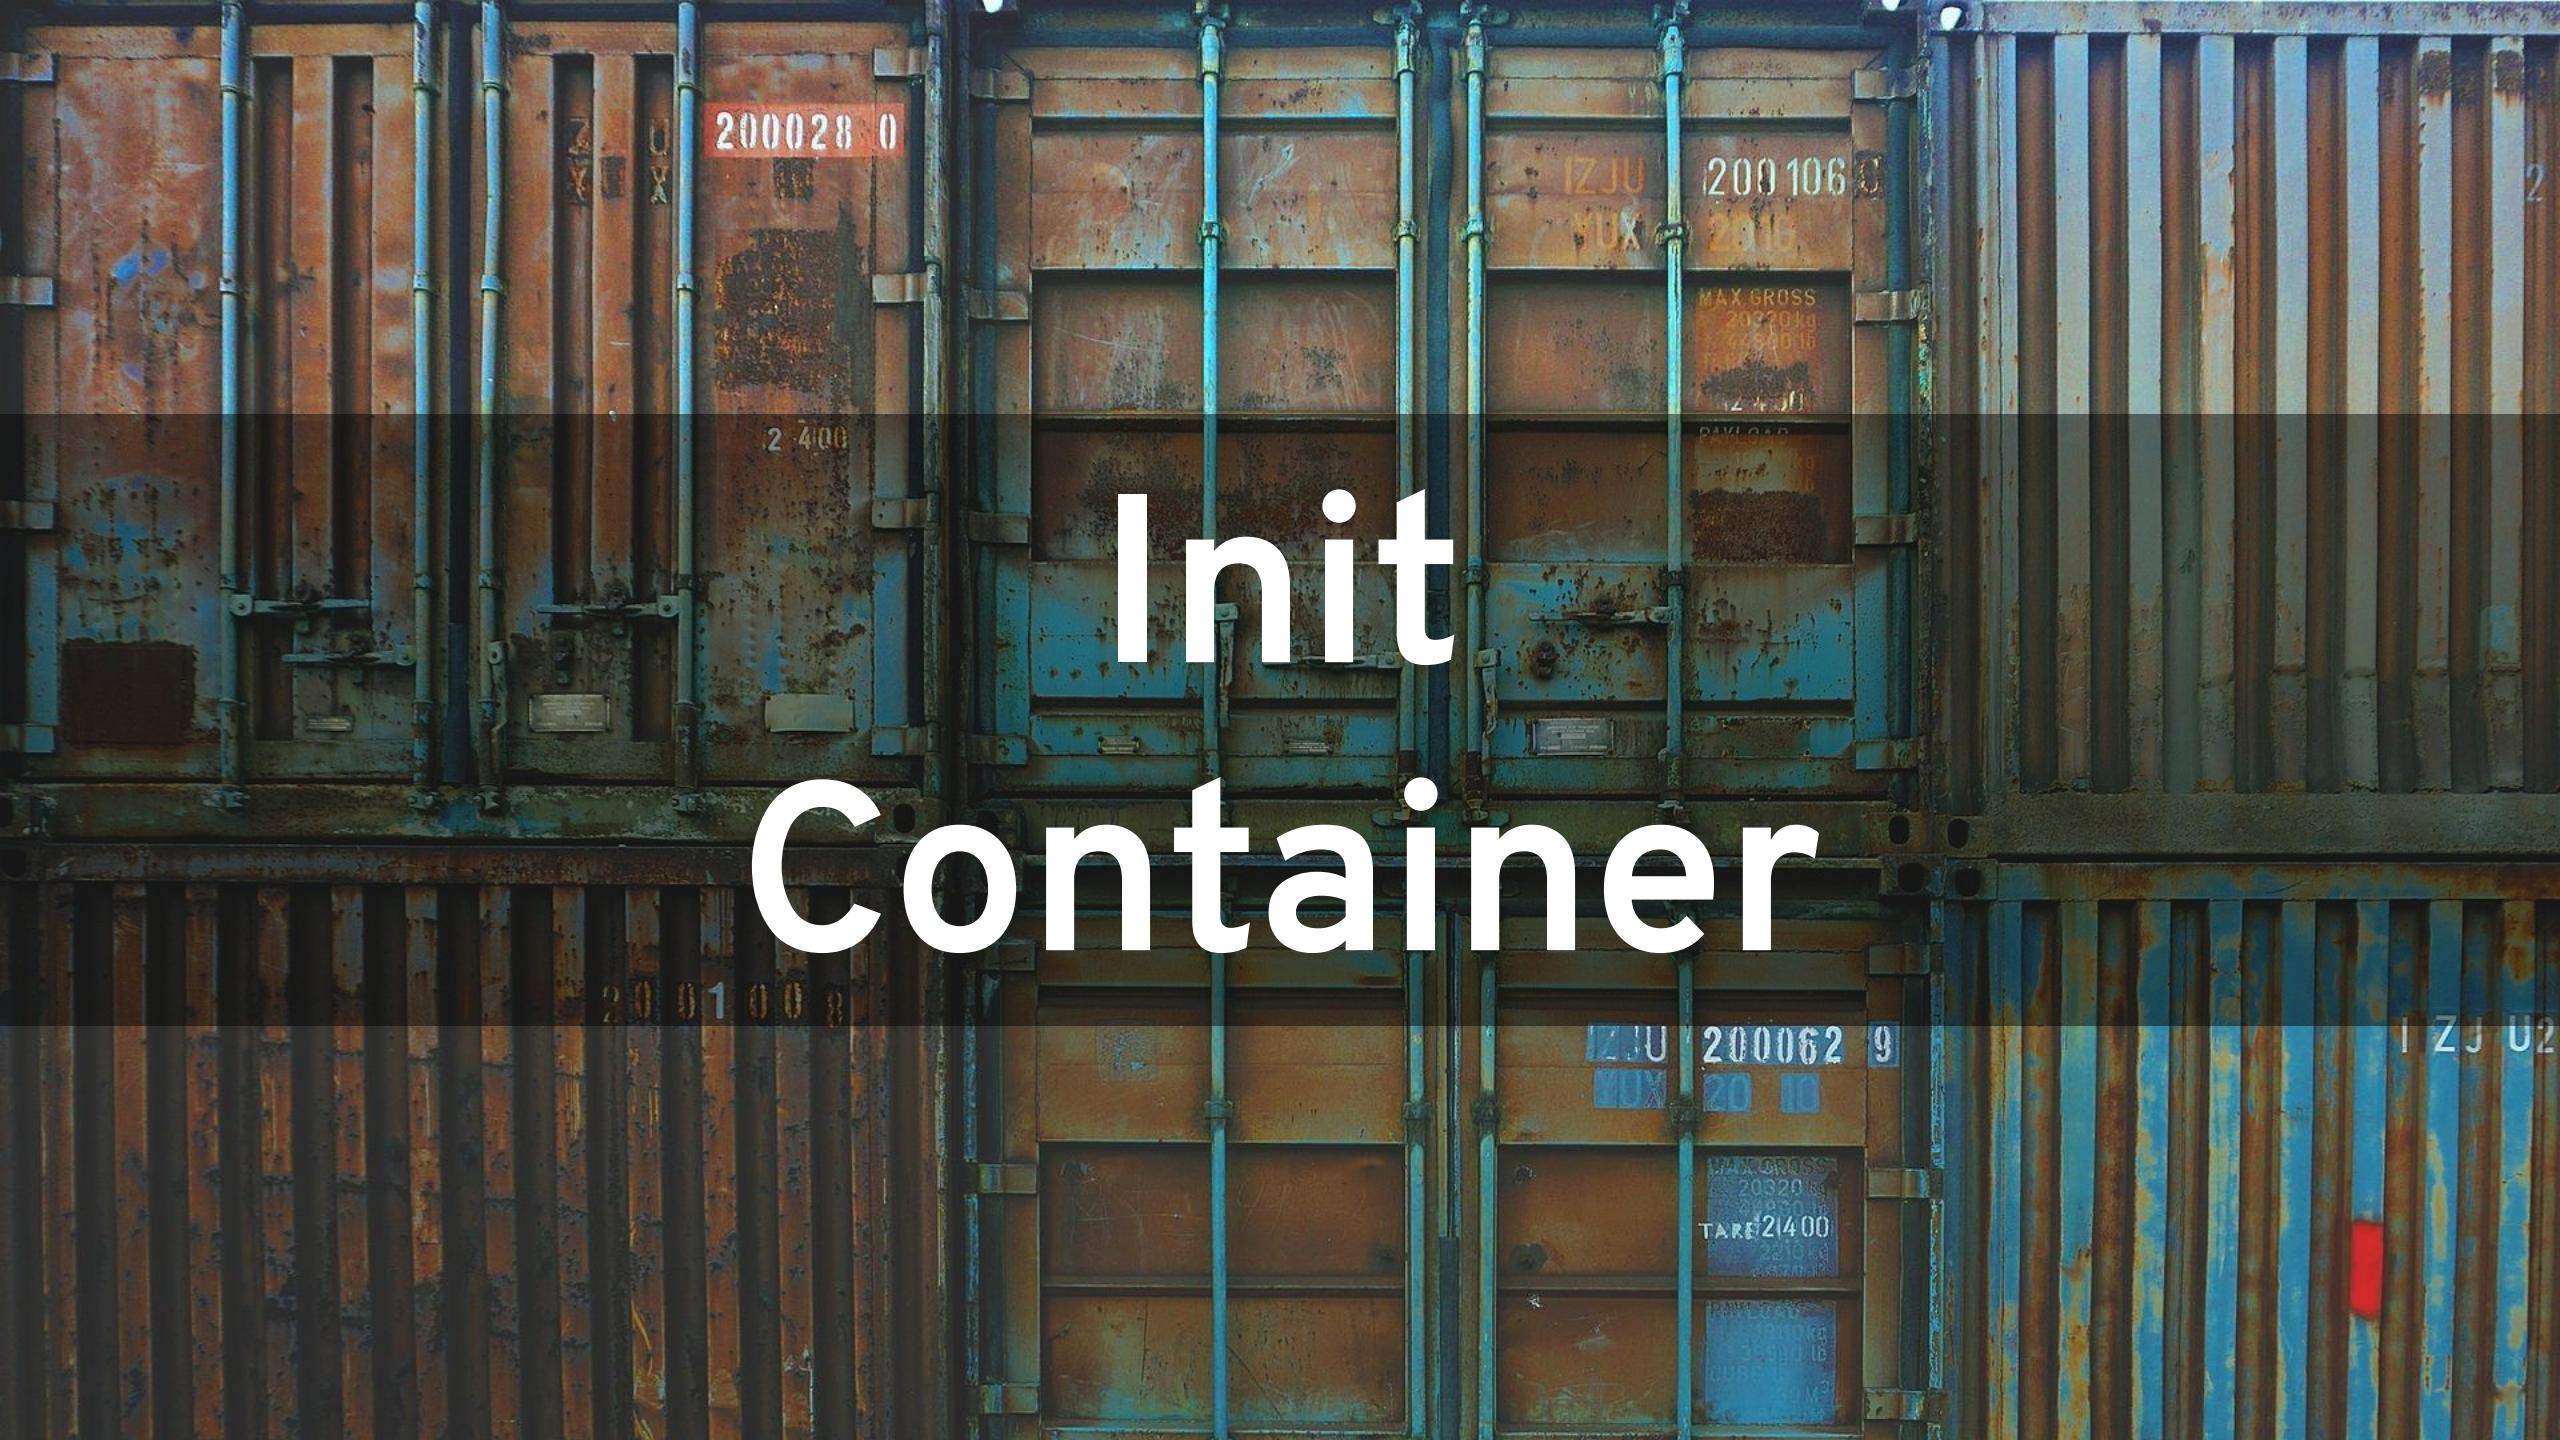


**Blue-green release**



**Canary release**





# Init Container

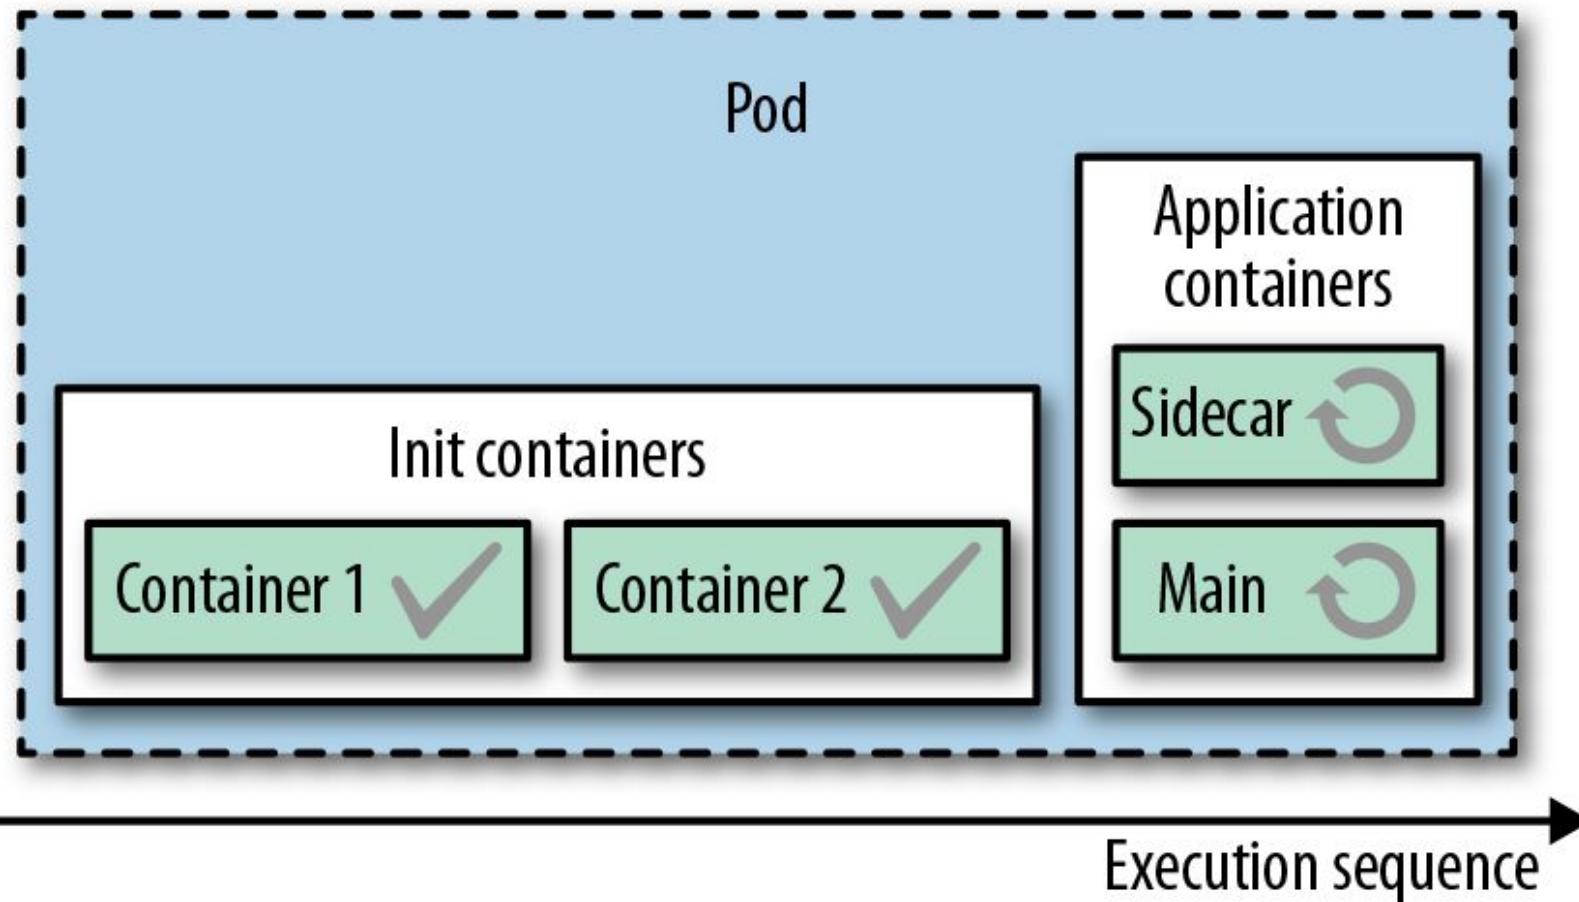
---

How to separate  
initialization steps from  
the main application.

# Init Container

- Part of a Pod
- One shot actions before application starts
- Needs to be idempotent
- Has own resource requirements

## Init Container



```
apiVersion: v1
kind: Pod
.....
spec:
  initContainers:
  - name: download
    image: axeclbr/git
    command: [ "git", "clone", "https://github.com/myrepo", "/data" ]
    volumeMounts:
      - mountPath: /var/lib/data
        name: source
  containers:
  - name: run
    image: docker.io/centos/httpd
    volumeMounts:
      - mountPath: /var/www/html
        name: source
  volumes:
  - emptyDir: {}
    name: source
```

The image is a collage of three photographs related to sidecar racing. The top photograph shows a rider in a purple and white patterned suit and helmet riding a sidecar through mud, with another rider visible behind them. The middle photograph is a close-up of a sidecar, heavily covered in mud, with its number plate partially visible. The bottom photograph shows a rider in a dark suit and helmet leaning into a turn on a muddy track.

Sidecar

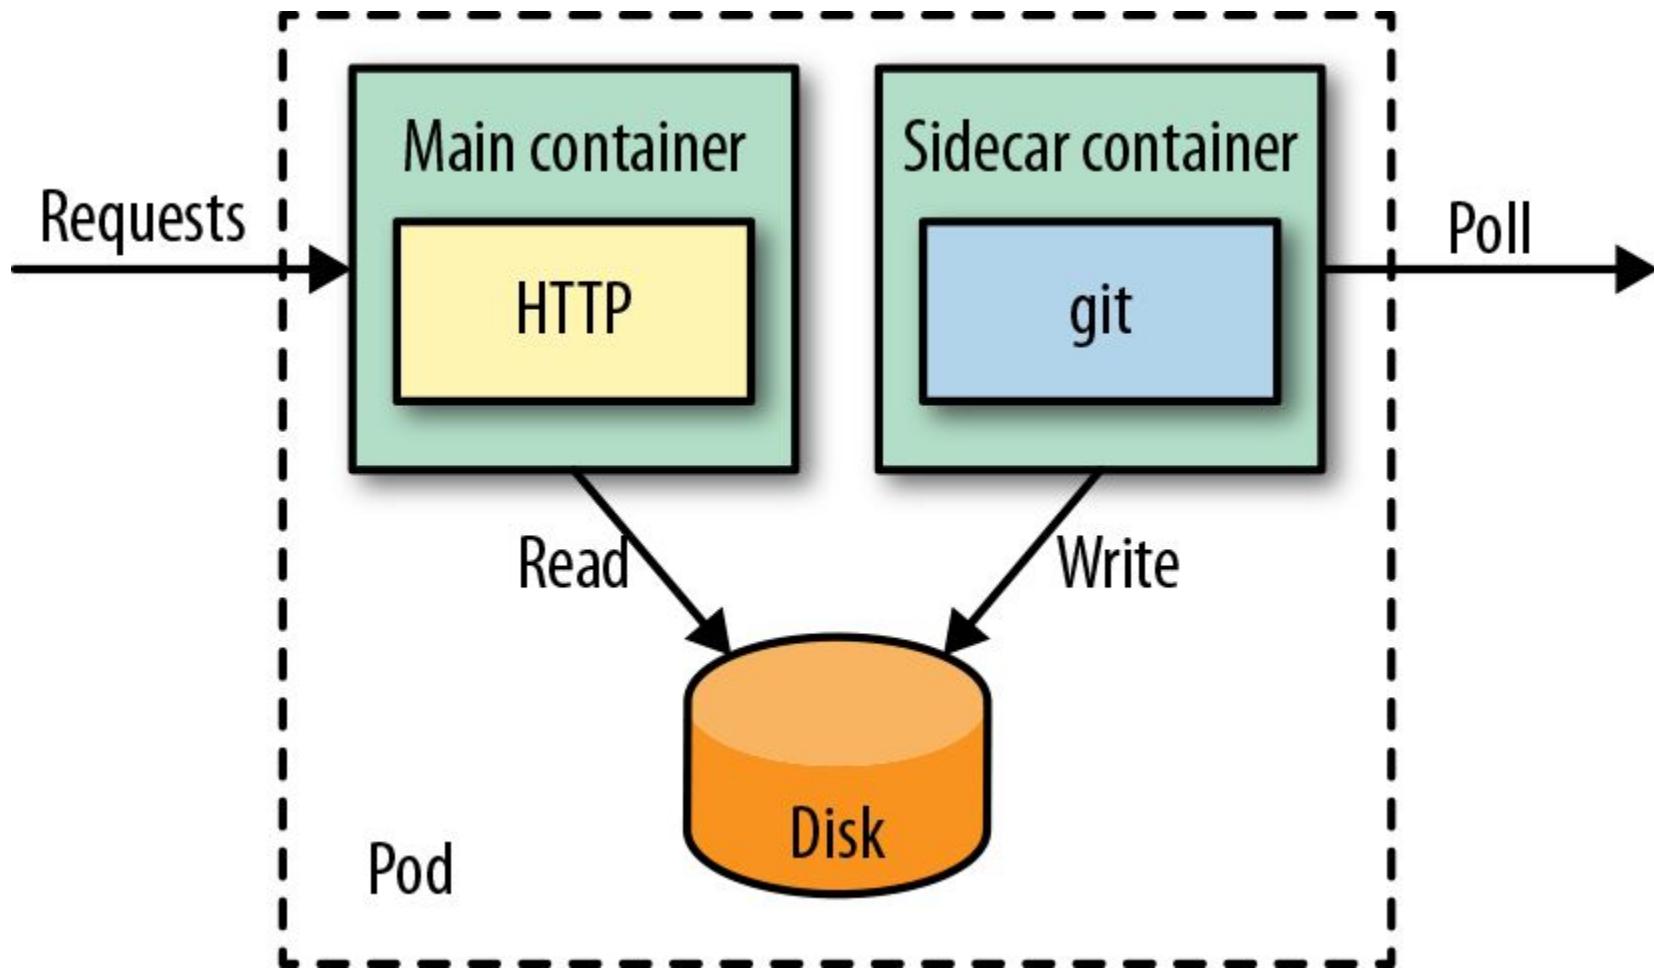
---

How to enhance the  
functionality of an application  
without changing it.

# Sidecar

- Runtime collaboration of containers
- Connected via shared resources:
  - Network
  - Volumes
- Similar what AOP is for programming
- Separation of concerns

# Sidecar



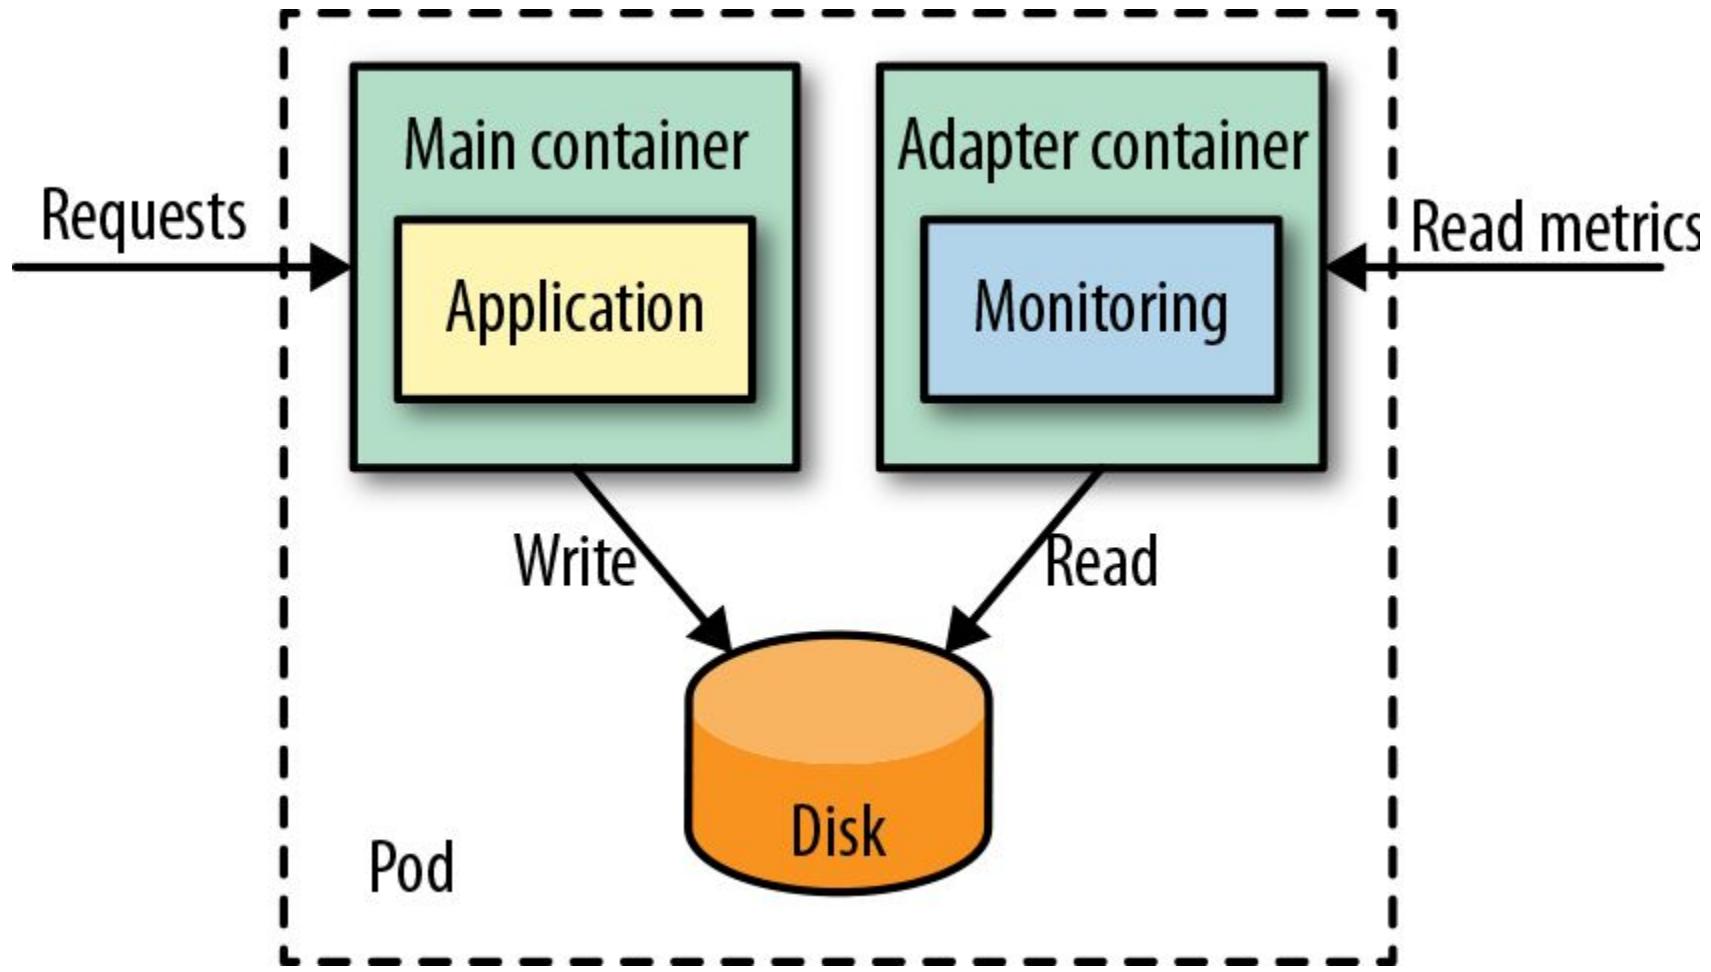
# Adapter



---

How to decouple access to a  
container **from** the outside  
world.

# Adapter



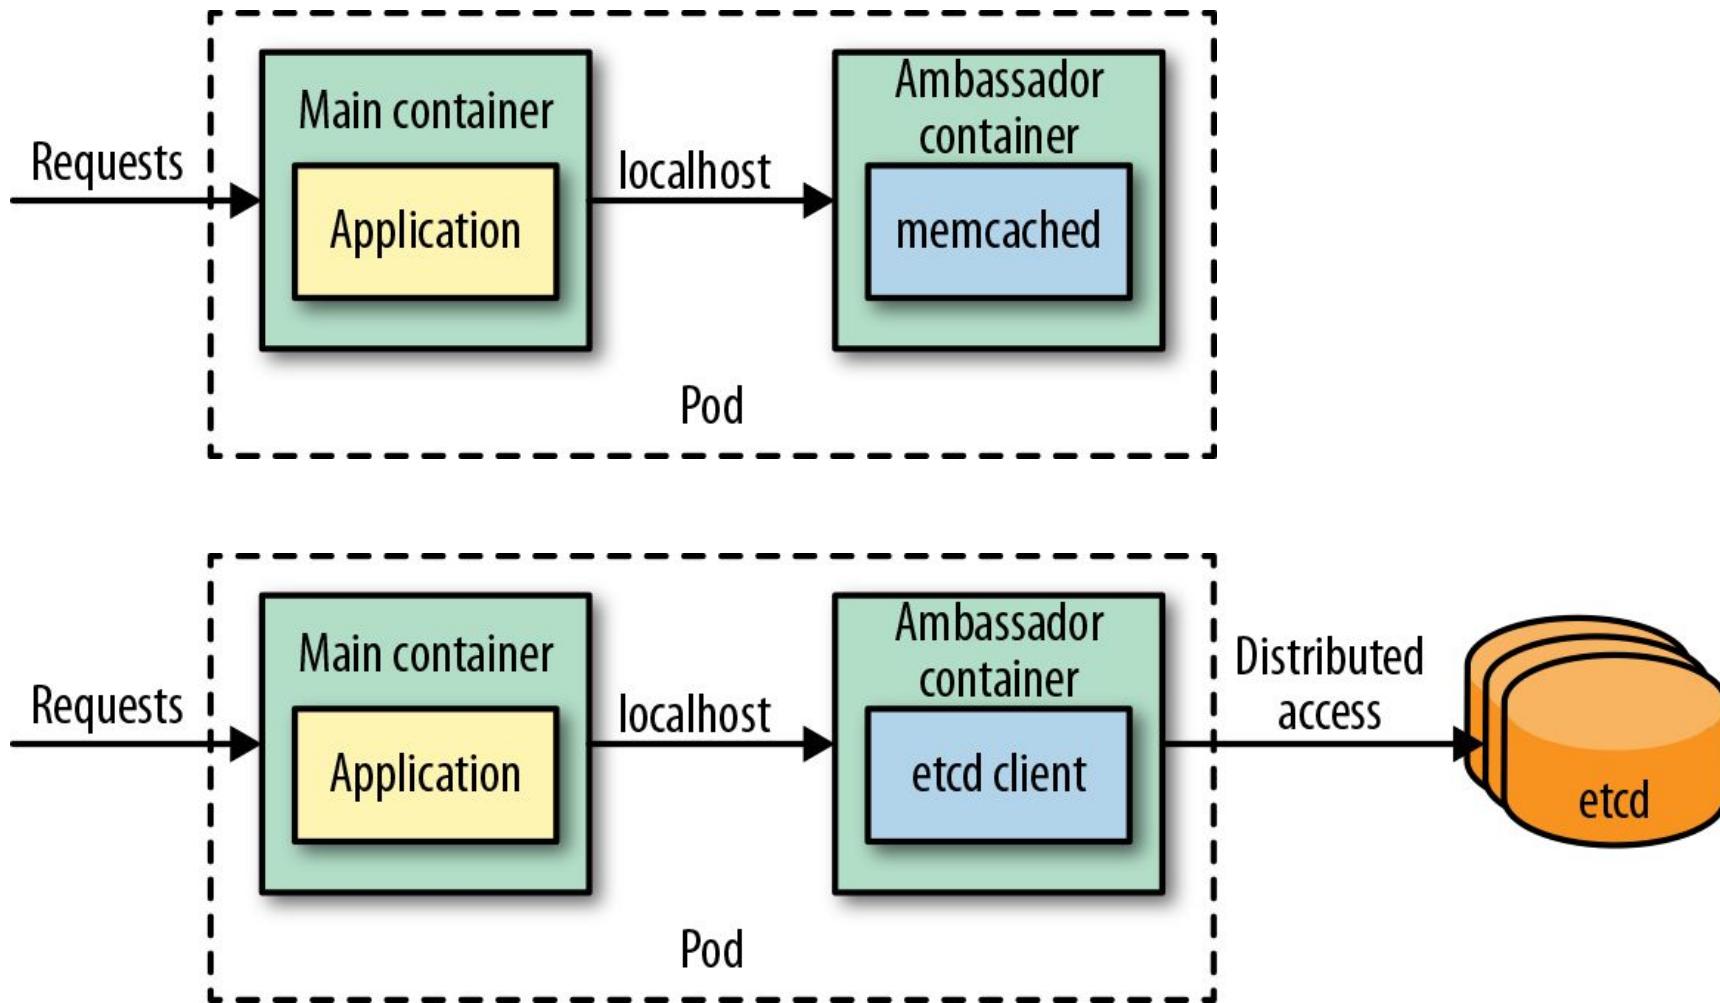


# Ambassador

---

# How to decouple a container's access **to** the outside world

# Ambassador



# Controller



---

How to get from the  
current state to a declared  
target state.

# State Reconciliation

- Kubernetes as distributed state manager
- Make the **actual** state more like the declared **target** state.



**Observe** - Discover the actual state

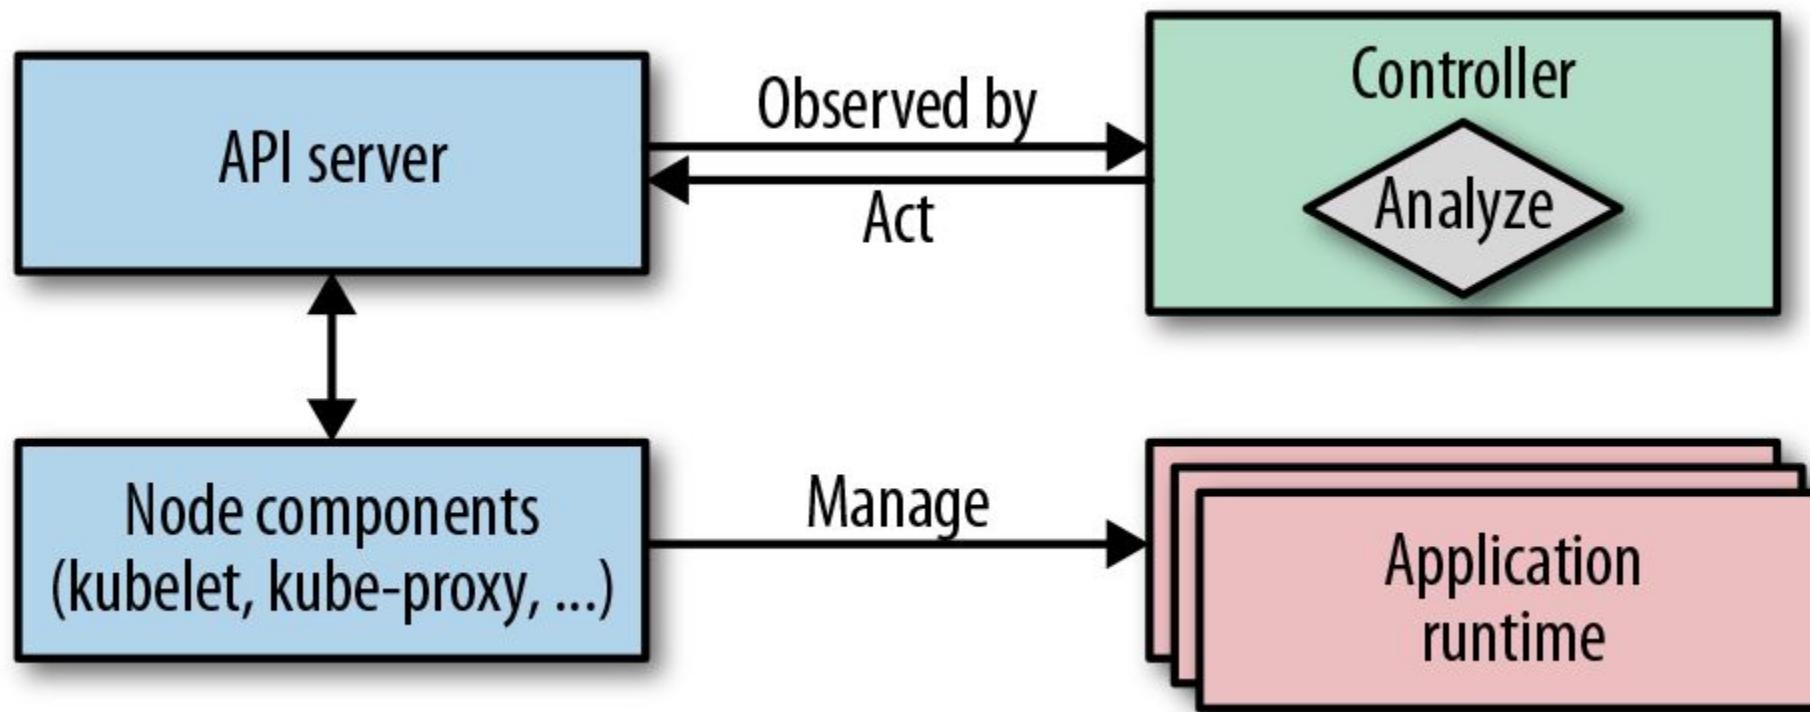


**Analyze** - Determine difference to target state



**Act** - Perform actions to drive the actual to the desired state

# Observe - Analyze - Act



# Common Triggers

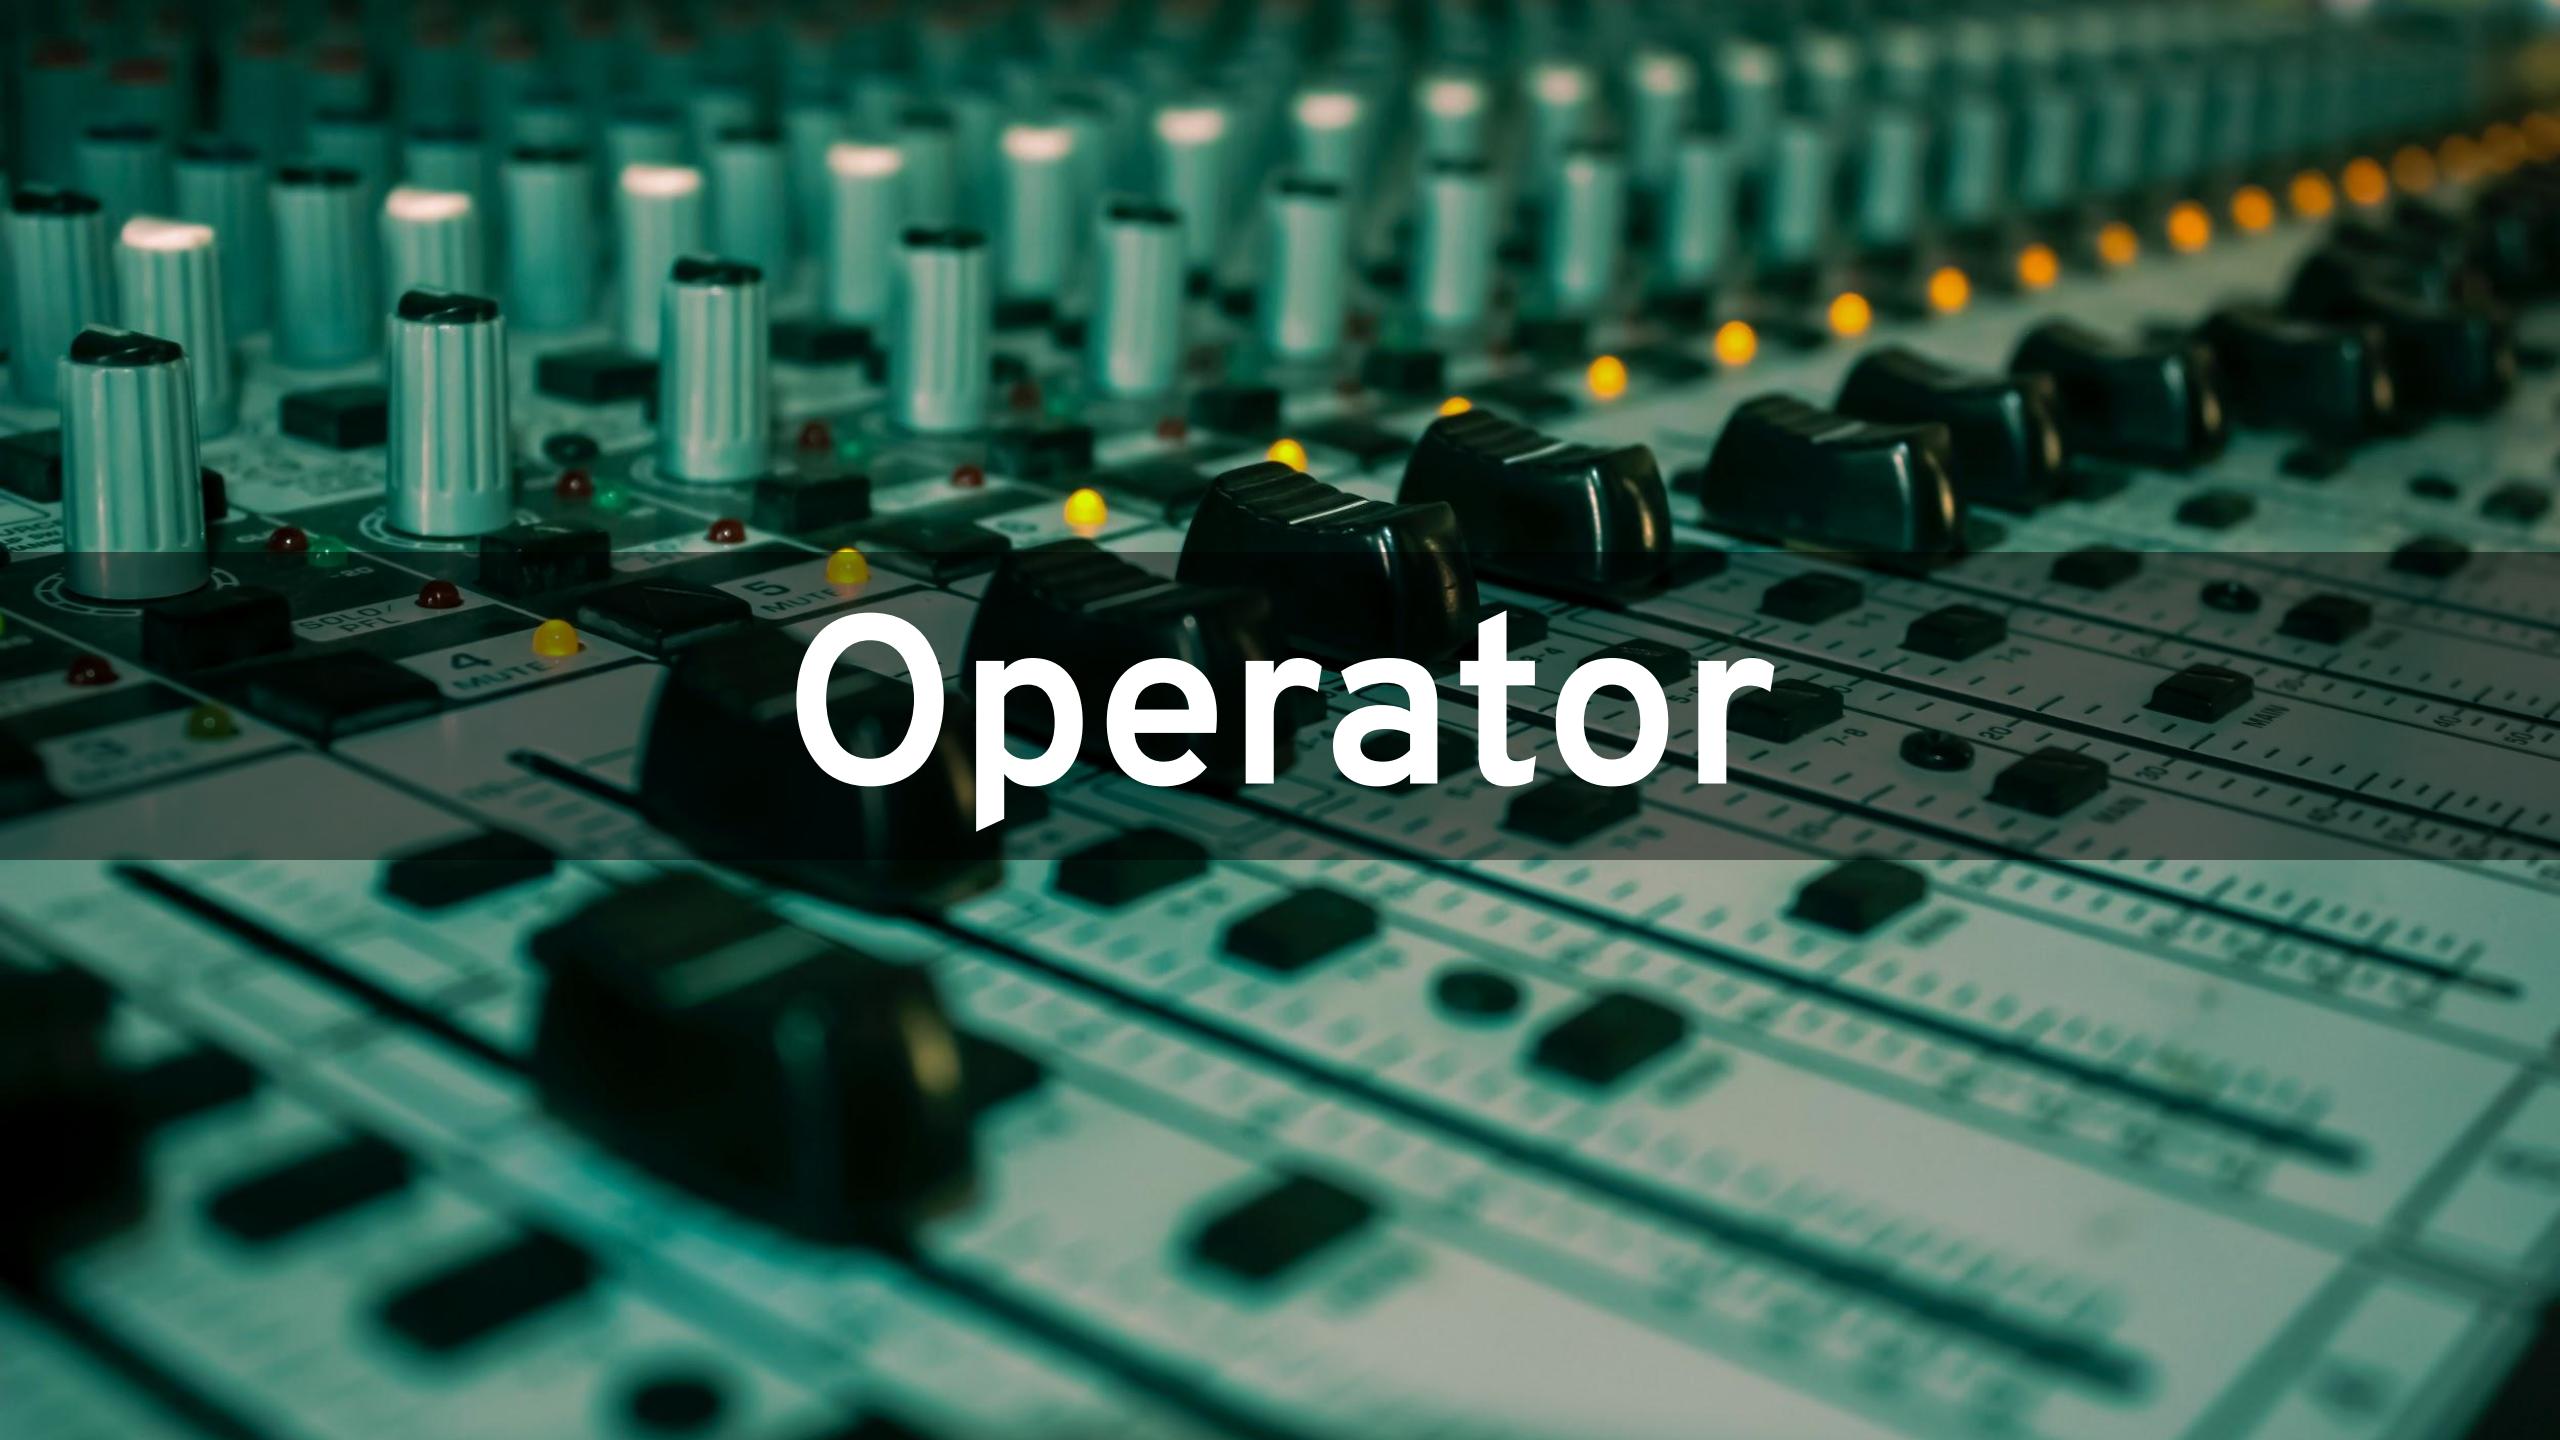
- Labels
  - Indexed by backend
  - Suitable for selector-like functionality
  - Limitation on charset for names and values
- Annotations
  - No syntax restrictions
  - Not indexed
- ConfigMaps
  - Good for complex structured state declarations
  - Simple alternative to CustomResourceDefinitions

# ConfigMap Watch Controller

```
namespace=${WATCH_NAMESPACE:-default}
base=http://localhost:8001
ns=namespaces/$namespace

curl -N -s $base/api/v1/${ns}/configmaps?watch=true | \
while read -r event
do
    type=$(echo "$event" | jq -r '.type')
    config_map=$(echo "$event" | jq -r '.object.metadata.name')
    annotations=$(echo "$event" | jq -r '.object.metadata.annotations')

    if [ $type = "MODIFIED" ]; then
        # Restart Pods using this ConfigMap
        # ...
    fi
done
```



# Operator

How to encapsulate  
operational knowledge  
into executable software.

# Definition

“” An **operator** is a Kubernetes **controller** that understands two domains: Kubernetes and *something else*. By combining knowledge of both areas, it can **automate** tasks that usually require a human operator that understands both domains.

---

Jimmy Zelinskie

<http://bit.ly/2Fjlx1h>

Technical:

**Operator = Controller + CustomResourceDefinition**

# CustomResourceDefinition

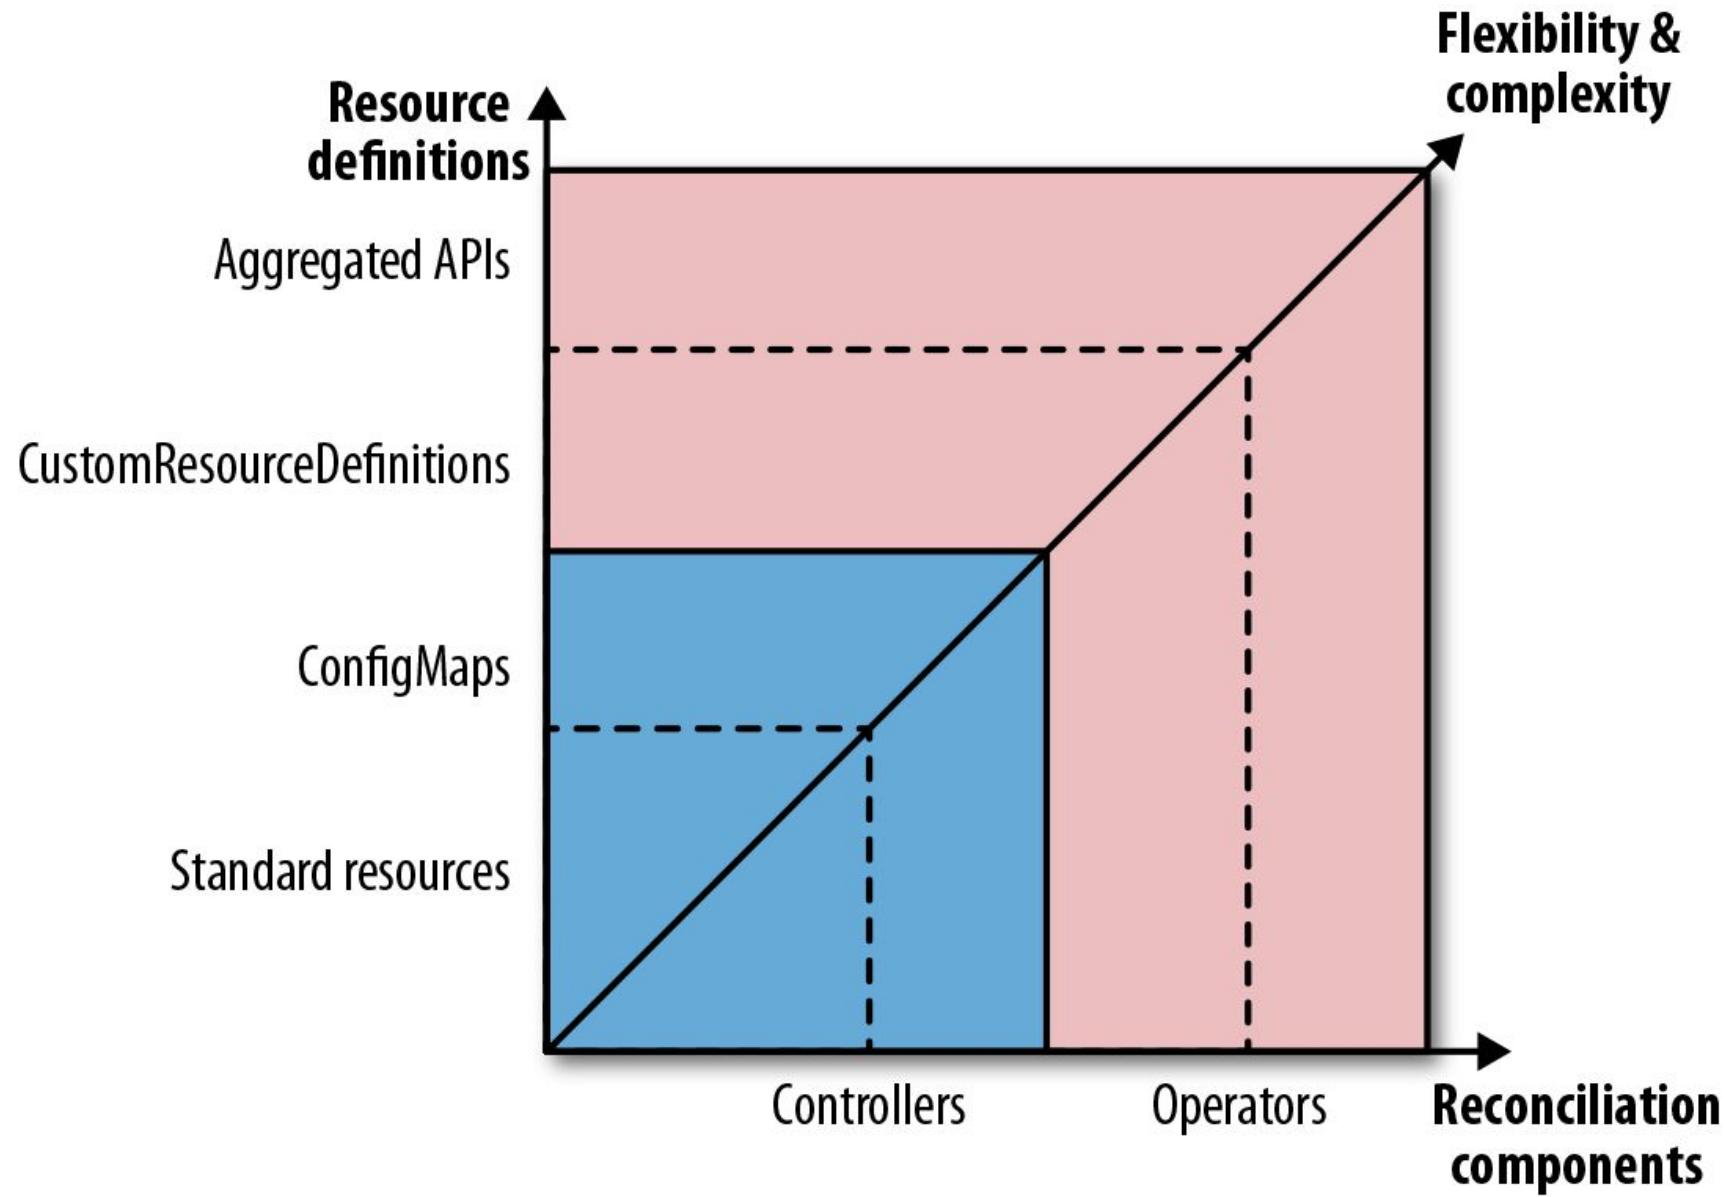
Custom resource is modelling a custom domain and managed through the Kubernetes API

```
apiVersion: apiextensions.k8s.io/v1beta1
kind: CustomResourceDefinition
metadata:
  name: configwatchers.k8spatterns.io
spec:
  scope: Namespaced
  group: k8spatterns.io
  version: v1
  names:
    kind: ConfigWatcher
    plural: configwatchers
  validation:
    openAPIV3Schema:
      ...

```

# Custom Resource

```
kind: ConfigWatcher
apiVersion: k8spatterns.io/v1
metadata:
  name: webapp-config-watcher
spec:
  configMap: webapp-config
  podSelector:
    app: webapp
```



# CRD Classification

- Installation CRDs
  - Installing and operating applications
  - Backup and Restore
  - Monitoring and self-healing
  - Example: Prometheus for installing Prometheus & components
- Application CRDs
  - Application specific domain concepts
  - Example: ServiceMonitor for registering Kubernetes service to be scraped by Prometheus

# Operator Hub

The screenshot shows the OperatorHub.io homepage with a search bar and a provider filter set to "Red+Hat". The results page displays five operator cards:

- Jaeger Tracing** (provided by Jaeger): Provides tracing, monitoring and troubleshooting for microservices-based applications.
- Kubernetes Federation** (provided by Red Hat): Gain Hybrid Cloud capabilities between your clusters with Kubernetes Federation.
- MongoDB** (provided by MongoDB, Inc): The MongoDB Enterprise Kubernetes Operator enables easy deploys of MongoDB.
- Prometheus Operator** (provided by Red Hat): The Prometheus Operator for Kubernetes provides easy monitoring definitions for your application metrics.
- Strimzi Kafka** (provided by Red Hat): Run an Apache Kafka cluster, including Kafka Connect, ZooKeeper and more.

# Operator Development

- Operator can be implemented in any language
- Frameworks:
  - Operator Framework (Golang, Helm, Ansible)  
<https://github.com/operator-framework>
  - Kubebuilder (Golang)  
<https://github.com/kubernetes-sigs/kubebuilder>
  - Metacontroller (Language agnostic)  
<https://metacontroller.app/>
  - jvm-operators (Java, Groovy, Kotlin, ...)  
<https://github.com/jvm-operators>

# Thank you



<https://k8spatterns.io>



<https://twitter.com/ro14nd>



<https://twitter.com/bibryam>



<https://twitter.com/k8spatterns>

Picture credits: <http://bit.ly/k8spatterns-picture-credits>

# Picture Credits

<https://www.pexels.com/photo/brown-and-black-pattern-2158386/>

<https://pixabay.com/photos/ship-helm-sunset-cutter-coast-guard-759954/>

<https://unsplash.com/photos/yo01Z-9HQAw>

<https://www.pexels.com/photo/turned-on-light-crane-1117211/>

<https://www.pexels.com/photo/shallow-focus-photography-of-black-ship-1095814/> <https://unsplash.com/photos/UHDx3BHIFvY>

<https://pixabay.com/photos/containers-storage-rusted-rusty-1209079/>

<https://pixabay.com/photos/motocross-sidecar-race-motorsport-1045661/>

<https://www.pexels.com/photo/golden-gate-bridge-san-francisco-california-1141> [https://unsplash.com/photos/ymf4\\_9Y9S\\_A](https://unsplash.com/photos/ymf4_9Y9S_A)

<https://www.pexels.com/photo/reflection-playstation-pad-gaming-18174/>

[https://unsplash.com/photos/UJP\\_QpCKj7M](https://unsplash.com/photos/UJP_QpCKj7M)

<https://www.pexels.com/photo/grayscale-photo-of-person-holding-chess-piece-1498958/>

<https://pixabay.com/photos/cans-manufacturing-business-2888650/>

<https://unsplash.com/photos/IRoX0shwjUQ>

<https://www.freeimages.com/photo/poppy-in-wheat-1344010>

<https://pixabay.com/photos/telescope-field-glass-spyglass-1966366/>

<https://unsplash.com/photos/UHDx3BHIFvY>

<https://pixabay.com/photos/bake-advent-christmas-cookie-1786926/>

<https://pixabay.com/photos/balloons-colors-party-celebration-1869790/>