

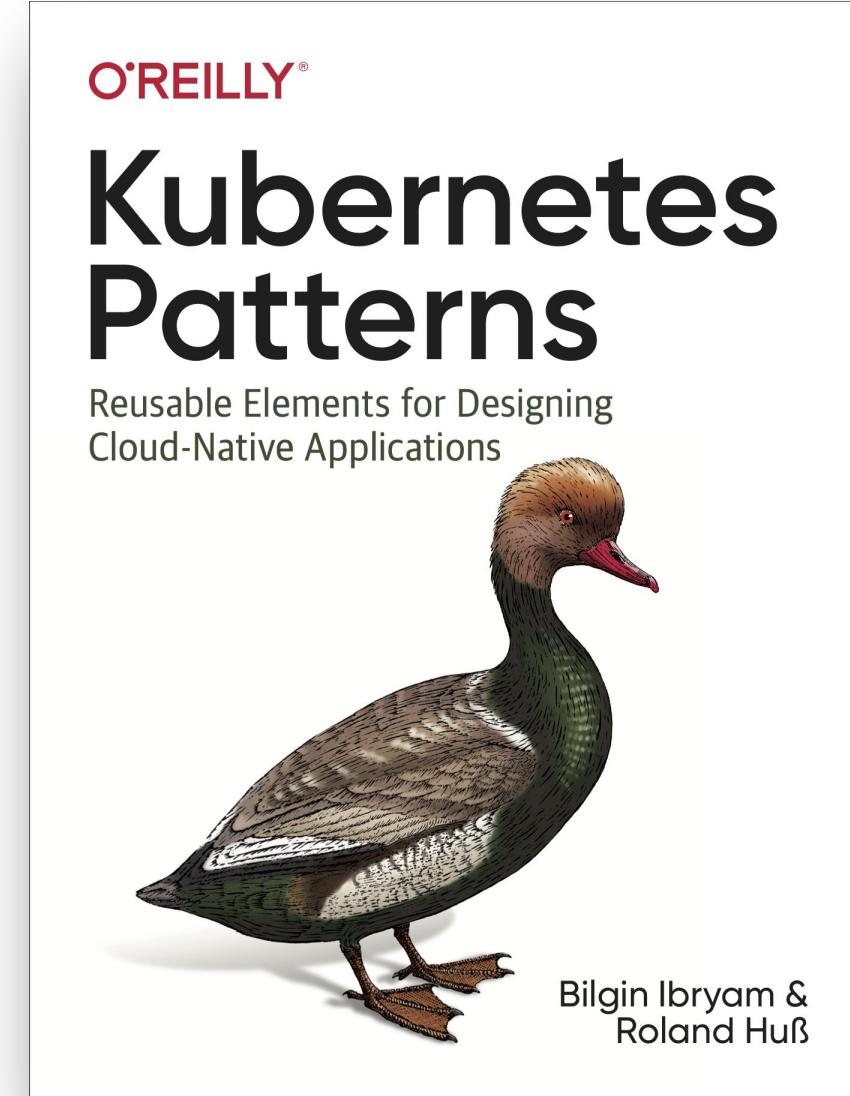
Kubernetes Patterns

Reusable Elements for Designing
Cloud-Native Applications

Dr. Roland Huß
Principal Software Engineer
@ro14nd
<https://k8spatterns.io>

Bilgin Ibryam
Principal Product Manager
@bibryam
<https://k8spatterns.io>





<https://k8spatterns.io>

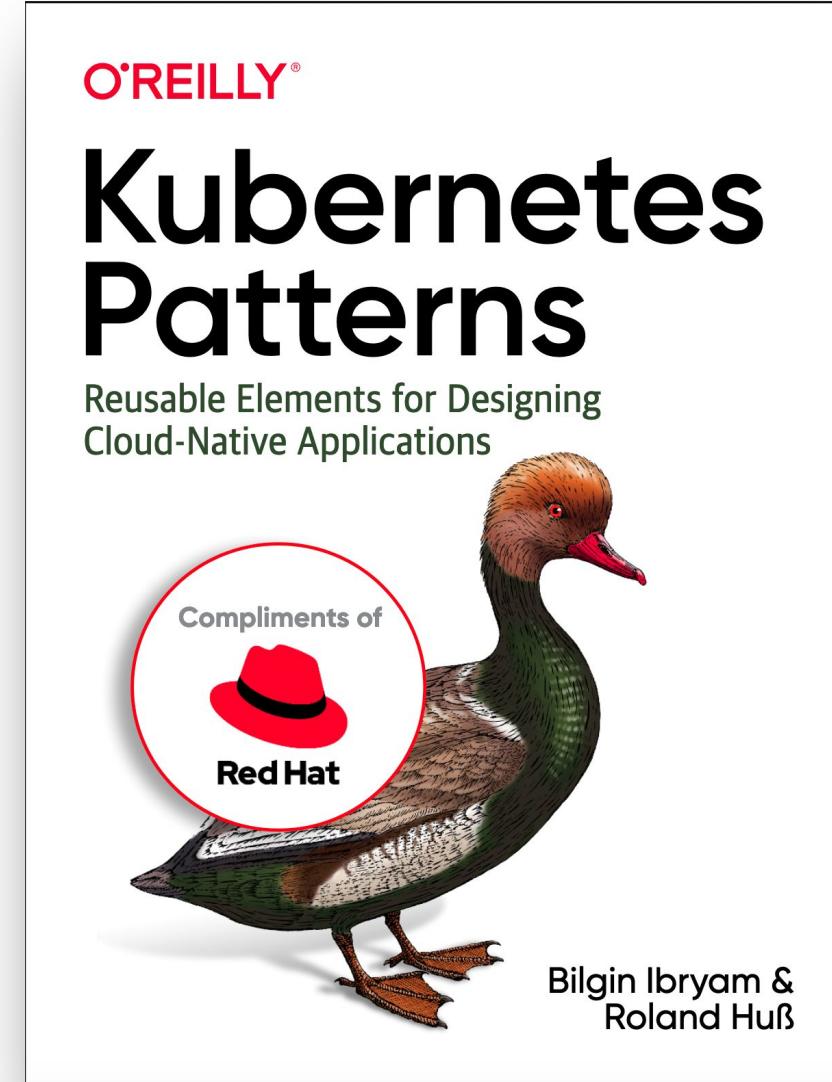


Grady Booch 
@Grady_Booch

Shout out to [@bibryam](#) and [@ro14nd](#):

You two have written a wonderful, incredibly informative, and intensely useful book. Thank you!

https://twitter.com/Grady_Booch/status/1252674791376449536



https://github.com/k8spatterns/examples Example X +

README.adoc

Kubernetes Patterns - Examples



This GitHub project contains the examples from *Kubernetes Patterns - Reusable Elements for Designing Cloud-Native Applications* book by Bilgin Ibryam and Roland Huß

Installation instructions for the example prerequisites are summarised in [INSTALL](#). By default, you need access to a vanilla Kubernetes installation, like Minikube. If addons are required, the example description explains this additional requirement.

For feedback, issues or questions in general, please use the [issue tracker](#) to open issues. Also, we love contributions like spelling fixes, bug fixes, improvements, ... Please open Pull Requests, we are happy to review them!

Patterns

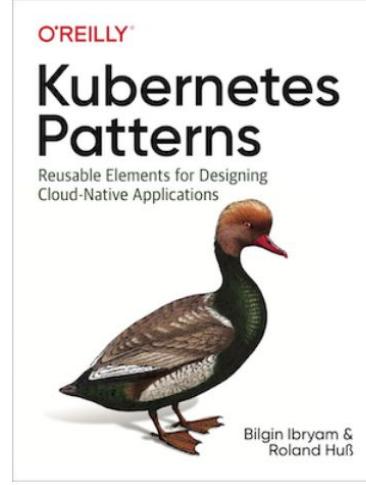
All example are categorised according to the Book's patterns category. Each of the examples is contained in an extra directory per pattern and is self-contained. [1]

Foundational Patterns

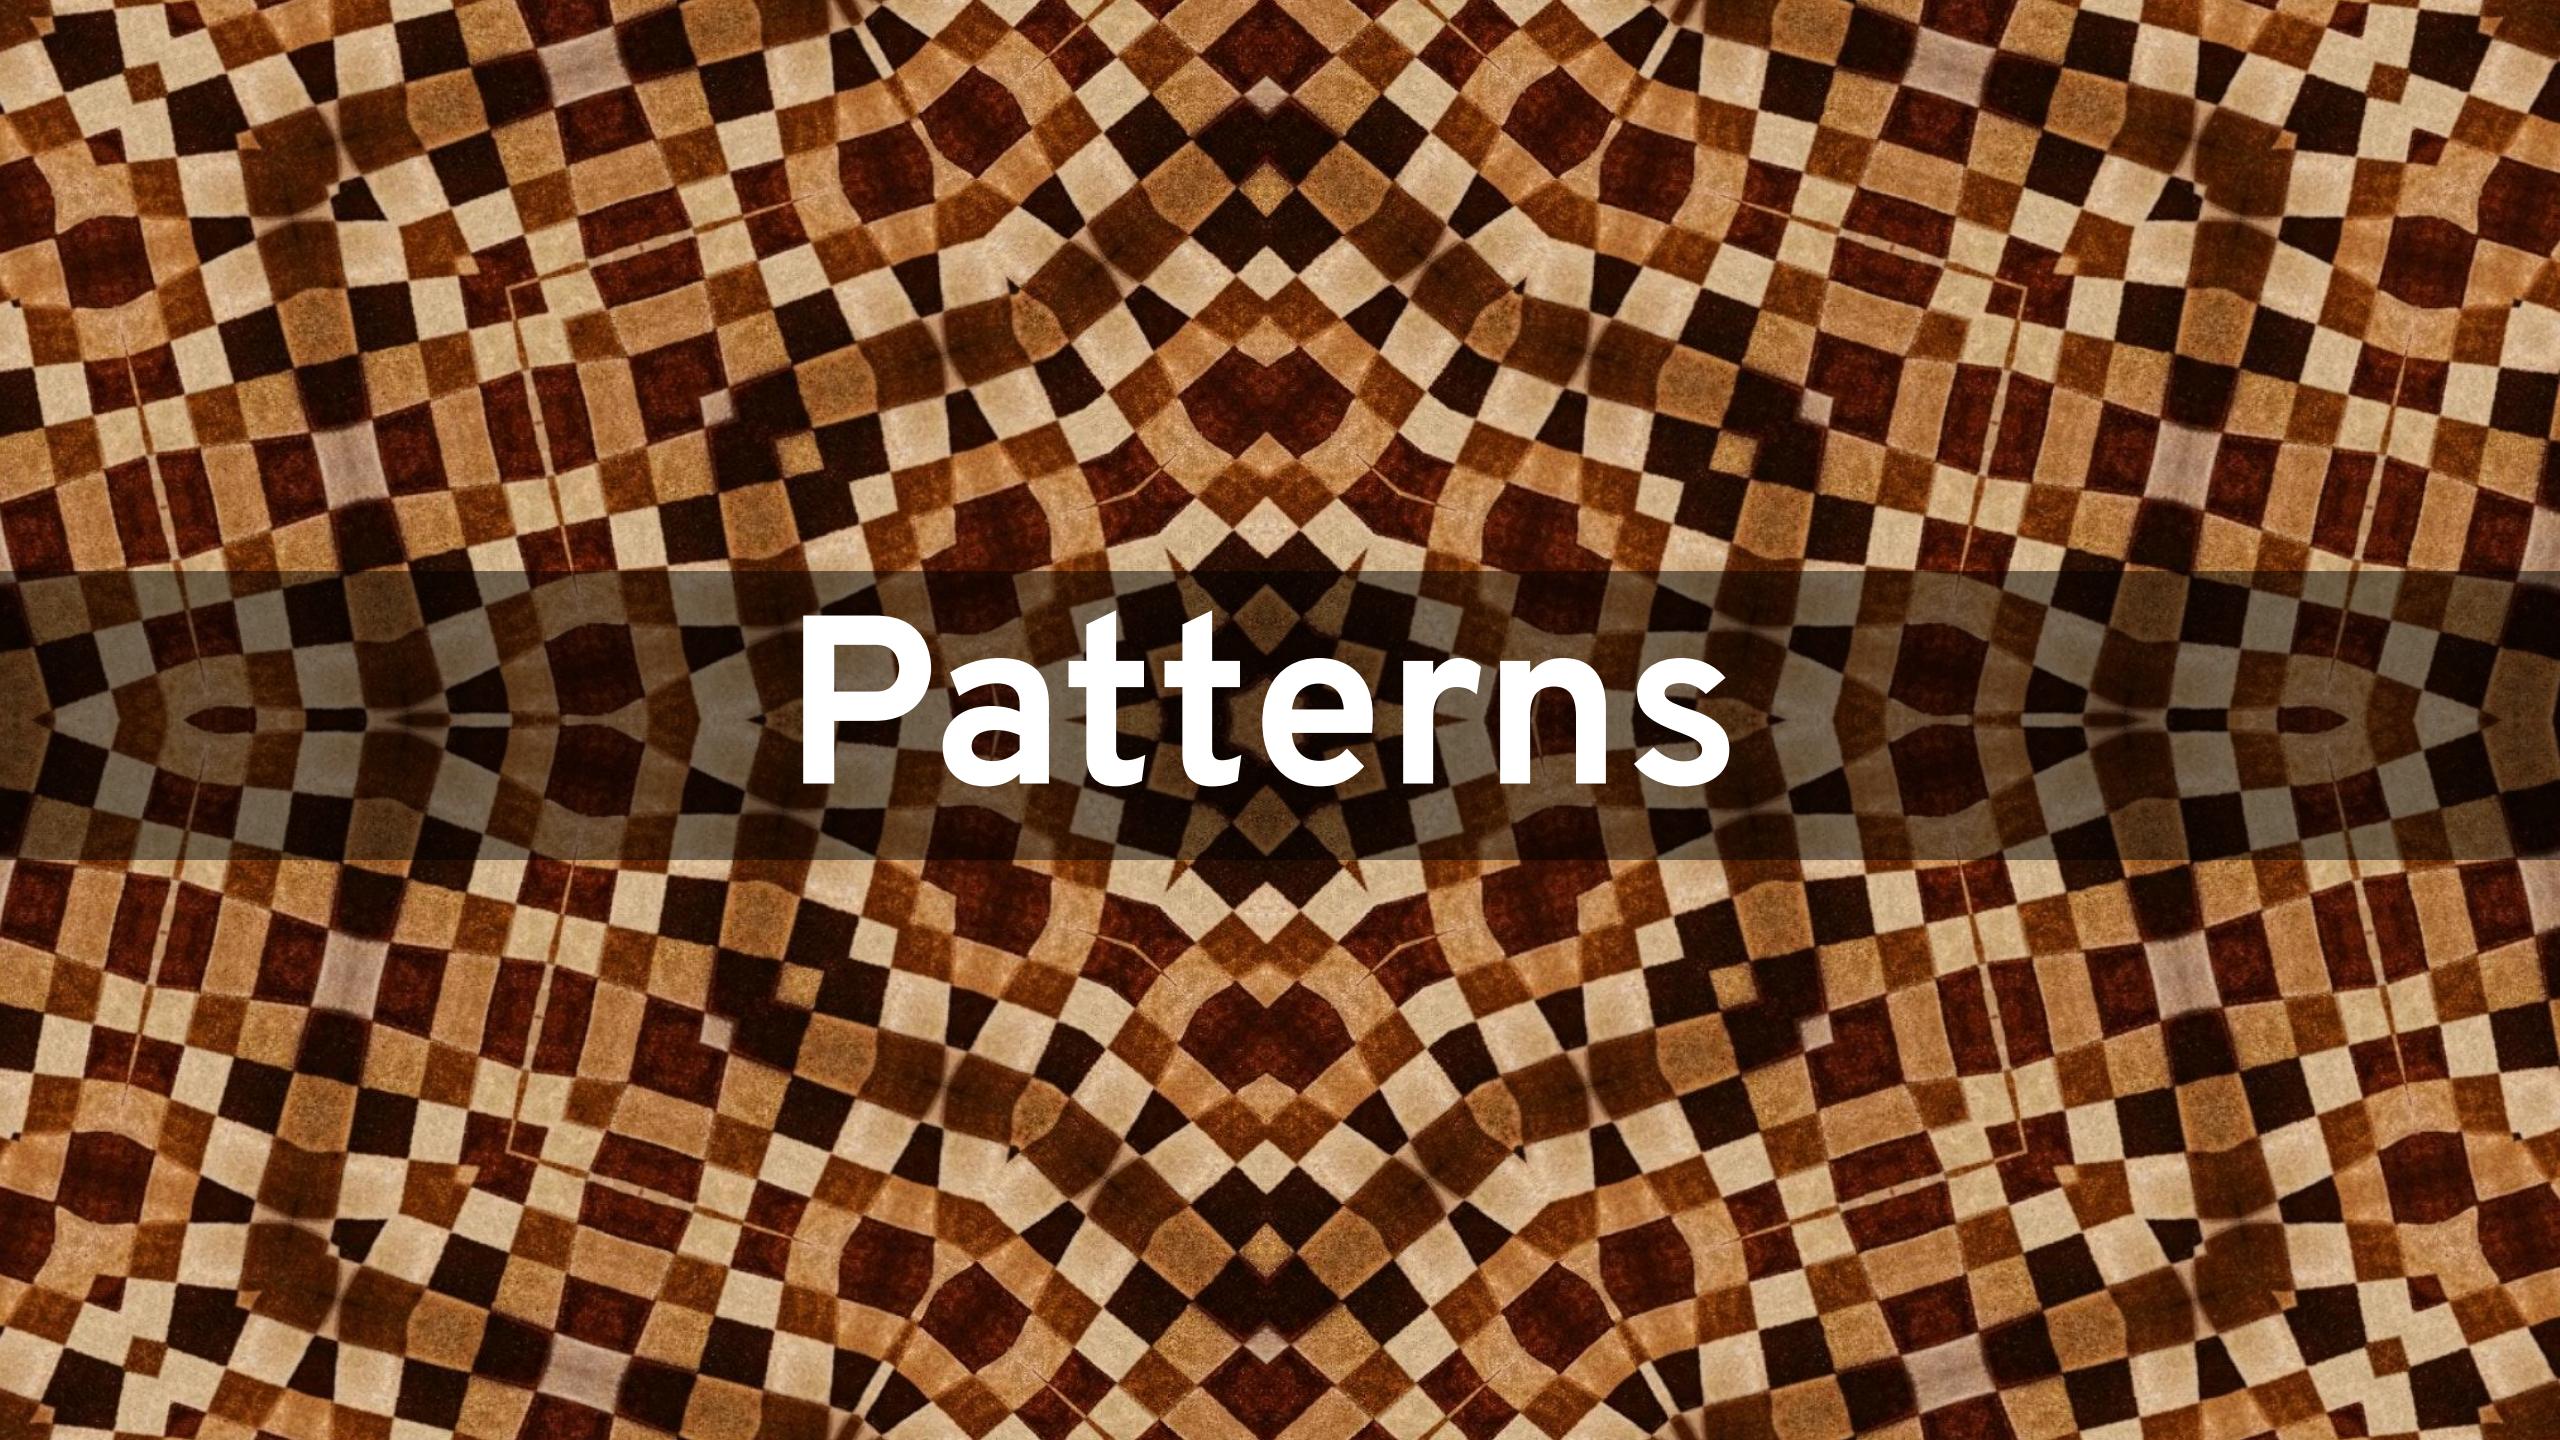
Predictable Demands
Our sample random generator dealing with hard requirements on ConfigMap and PersistentVolumeClaims as well as with resource limits.

Declarative Deployment
Rolling and fixed update of the random generator Deployment from version 1.0 to 2.0.

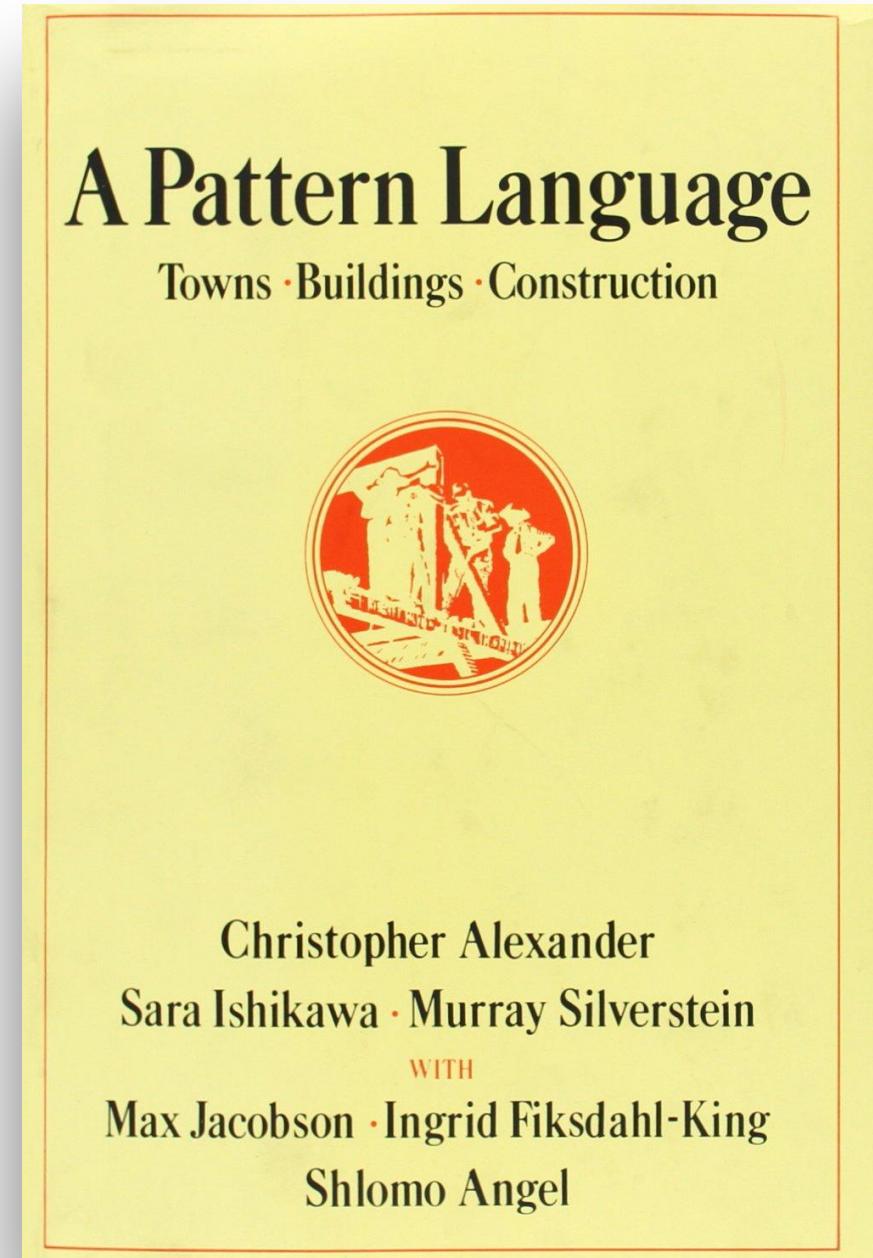
Health Probe
Liveness and Readiness probes for the random generator.



O'REILLY
Kubernetes Patterns
Reusable Elements for Designing Cloud-Native Applications
Bilgin Ibryam & Roland Huß

The background of the image is a intricate geometric pattern. It features a repeating motif of overlapping diamond shapes, creating a sense of depth and movement. The colors used are various shades of brown, tan, and black, which are arranged in a way that suggests a three-dimensional perspective. The overall effect is one of a sophisticated and artistic design.

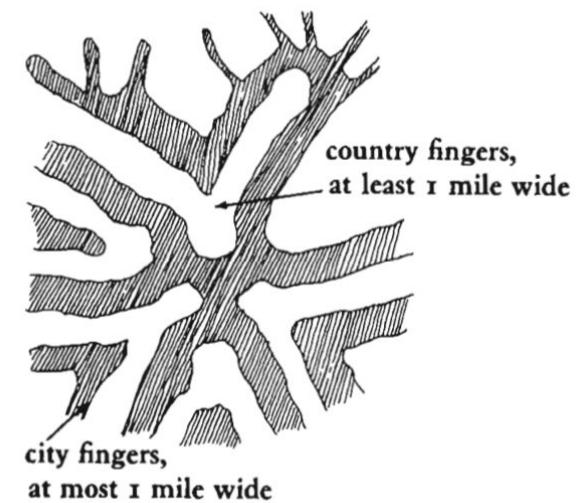
Patterns

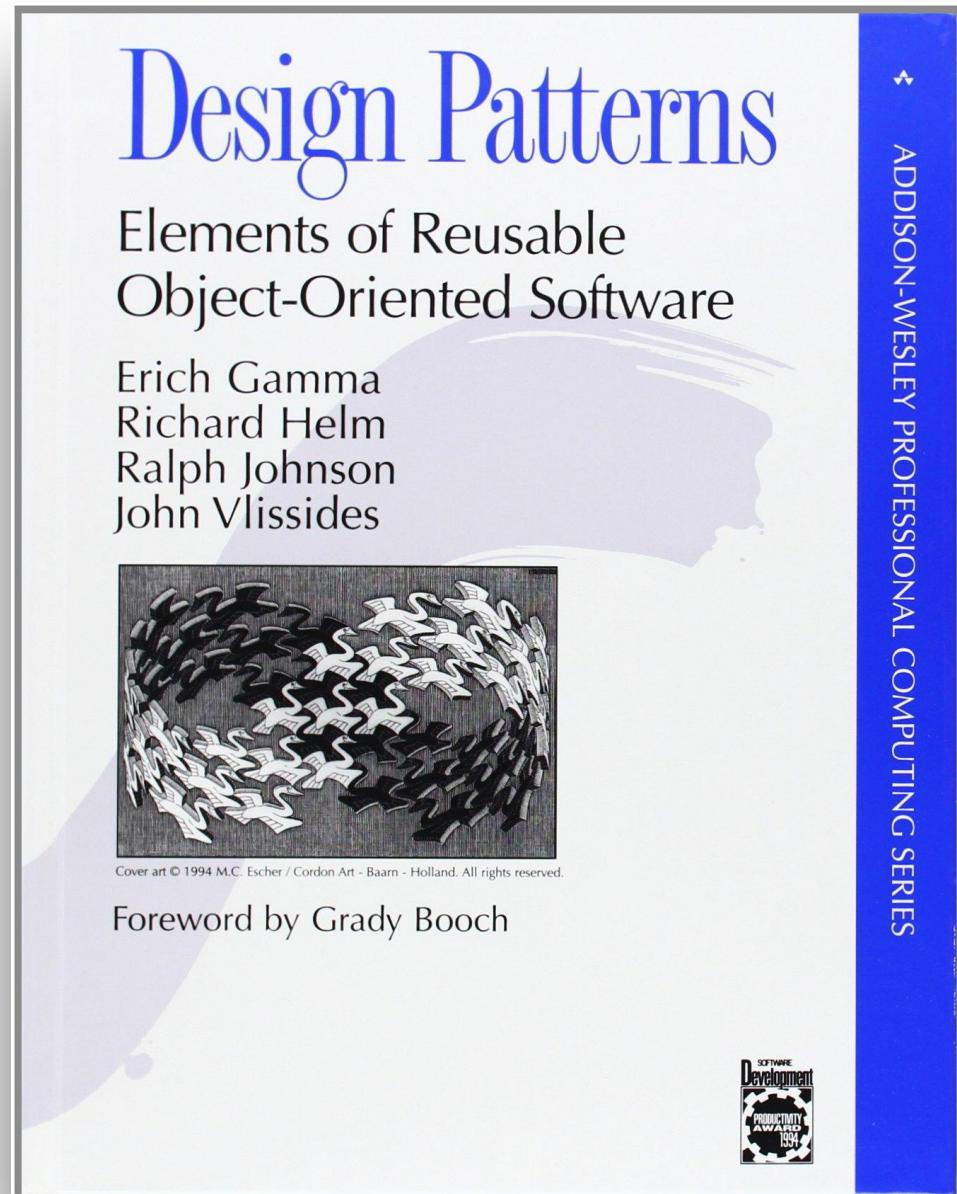


3 CITY COUNTRY FINGERS**



*When the countryside is far away
the city becomes a prison.*





Patterns

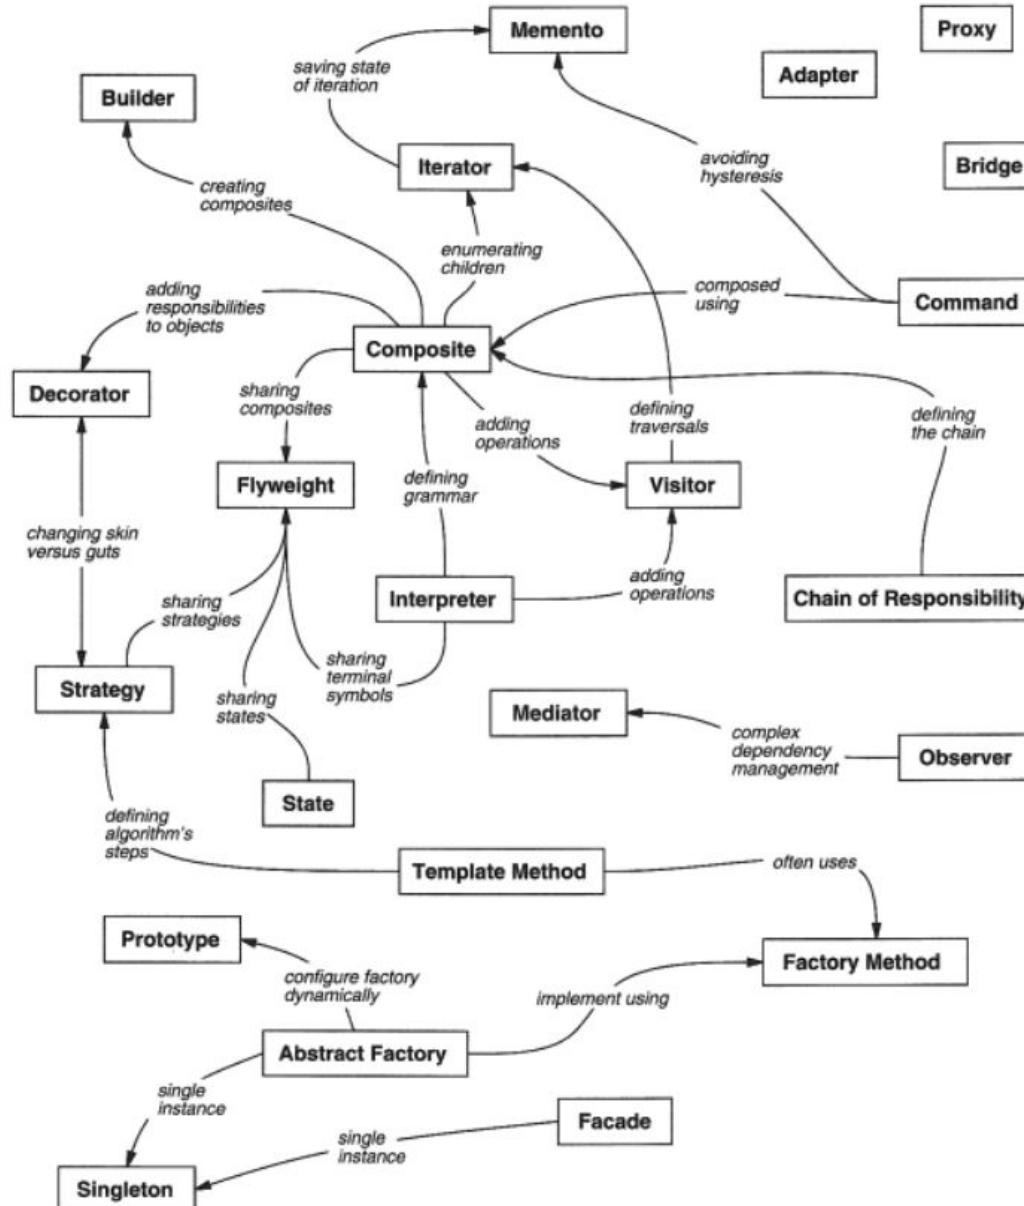


Figure 1.1: Design pattern relationships

A large, dark wooden ship's wheel is positioned in the foreground, angled towards the left. A rectangular metal plaque is attached to the right side of the wheel, featuring the words "U.S. COAST GUARD" in both English and Spanish. The background is a blurred photograph of a sunset or sunrise over the ocean, with warm orange and yellow hues reflecting on the water.

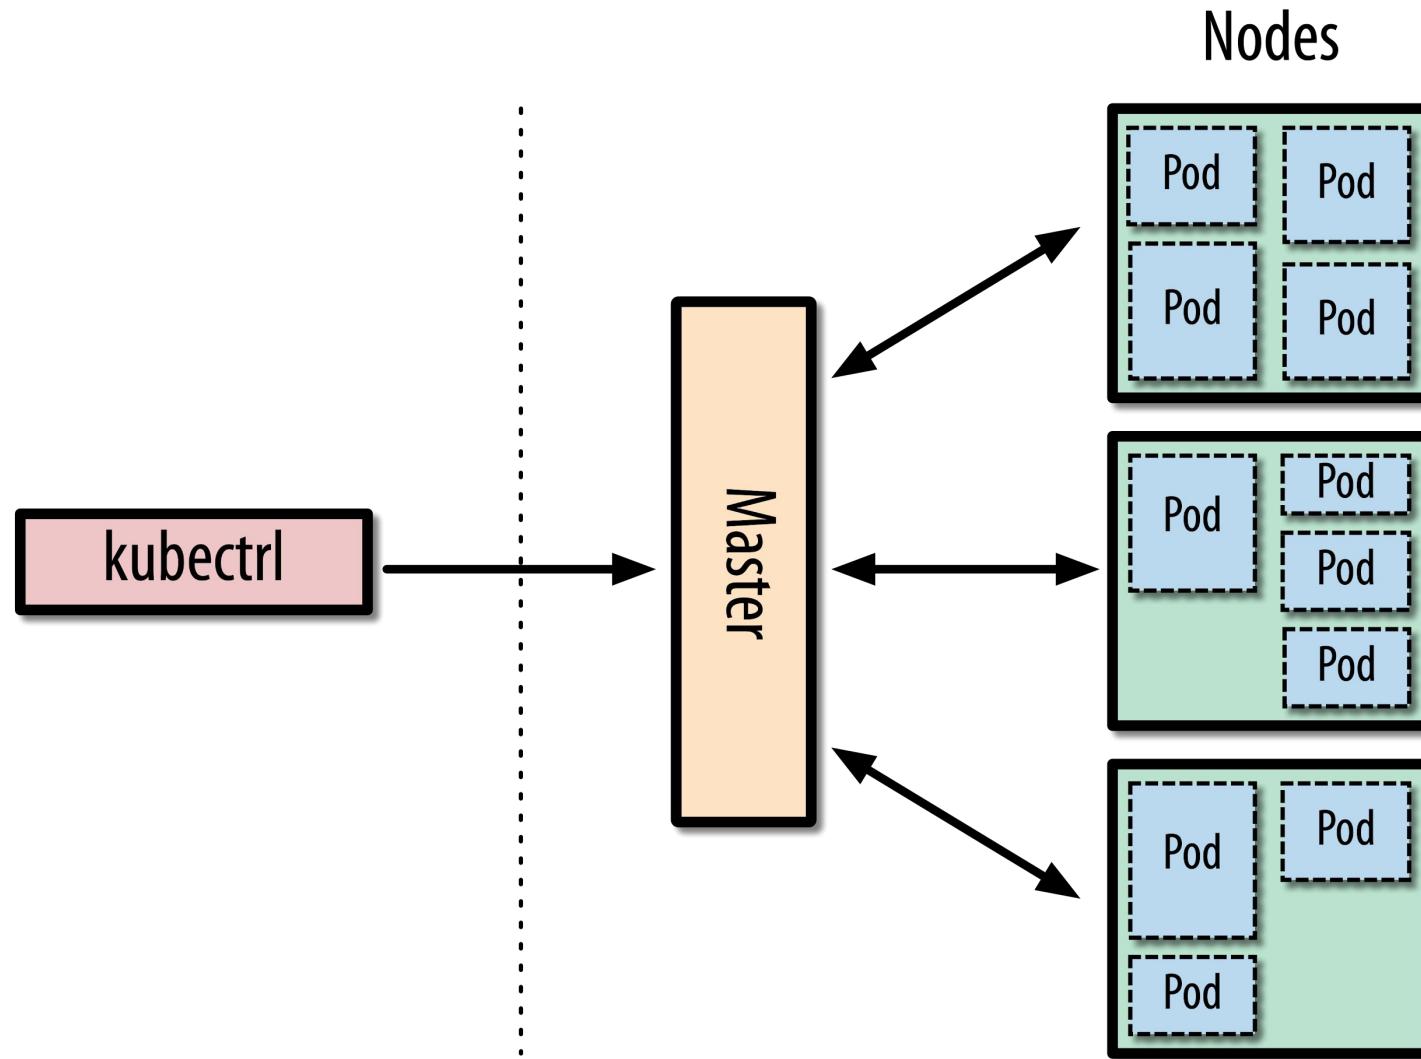
Kubernetes

Kubernetes

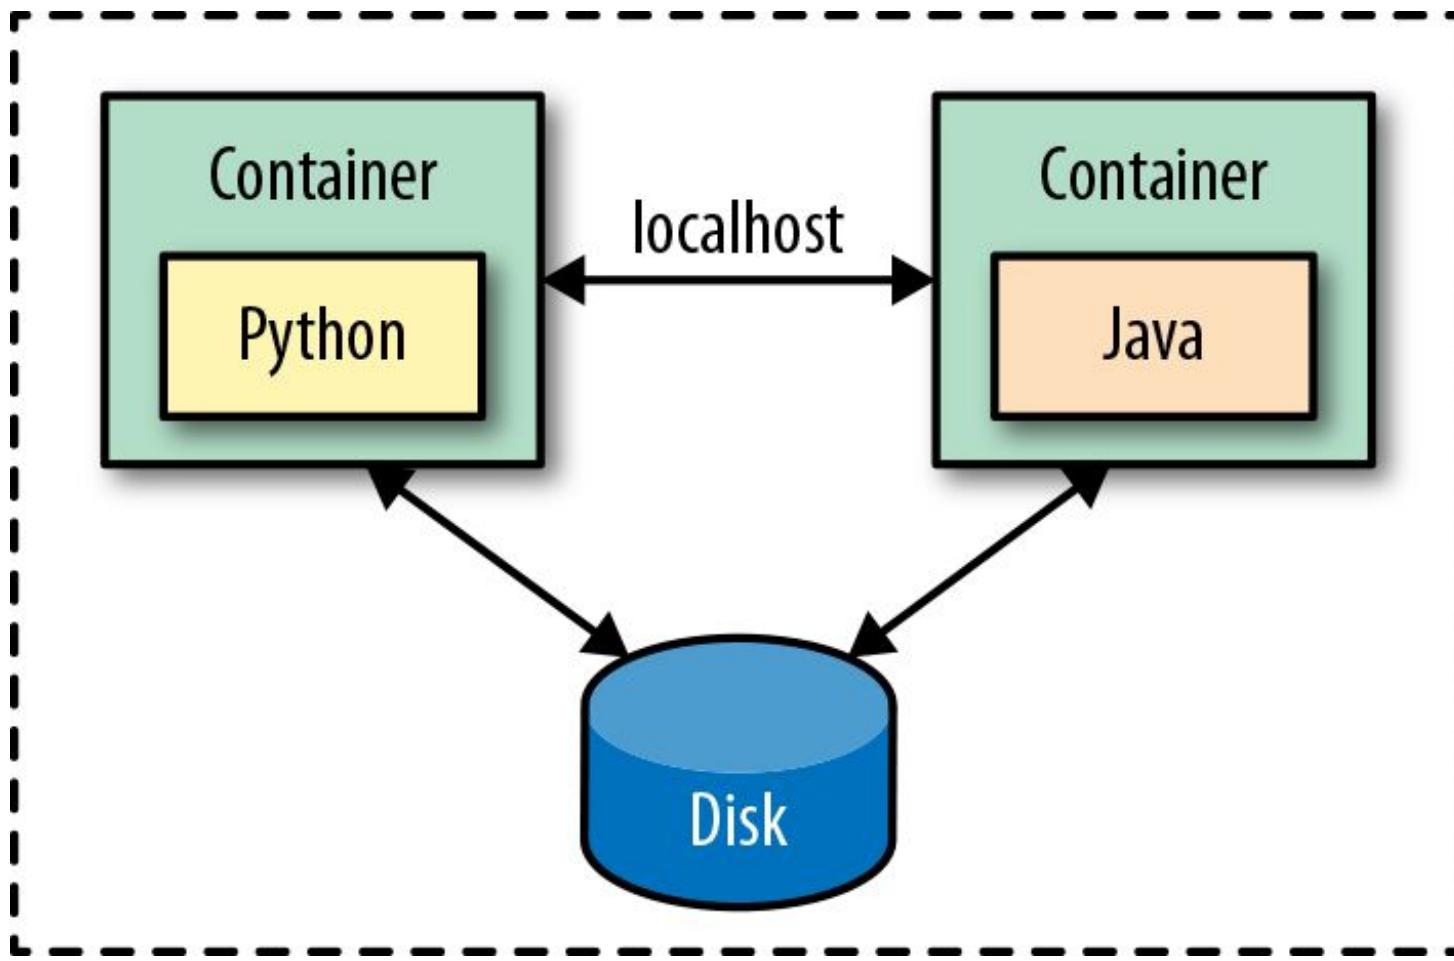


- Open Source container orchestration system started by Google in 2014
 - ※ Scheduling
 - ※ Self-healing
 - ※ Horizontal and vertical scaling
 - ※ Service discovery
 - ※ Rollout and Rollbacks
- Declarative resource-centric REST API

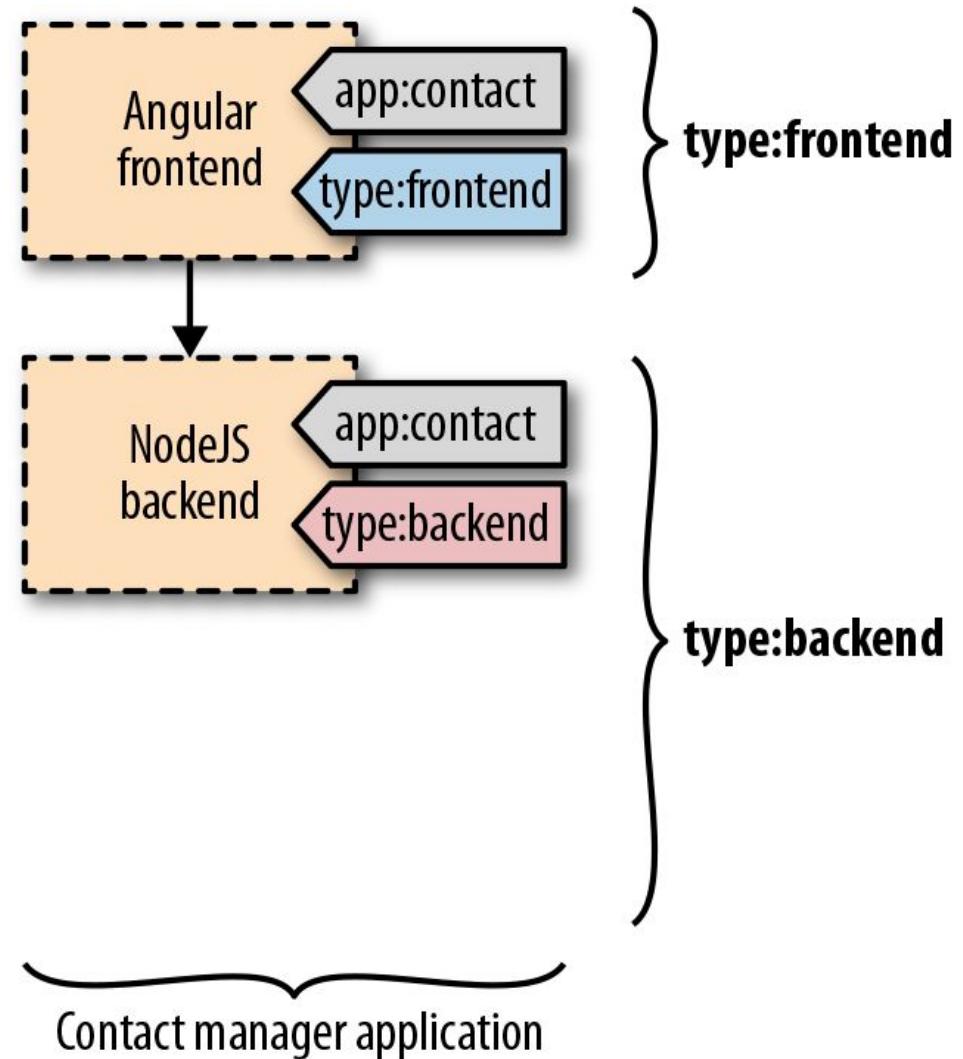
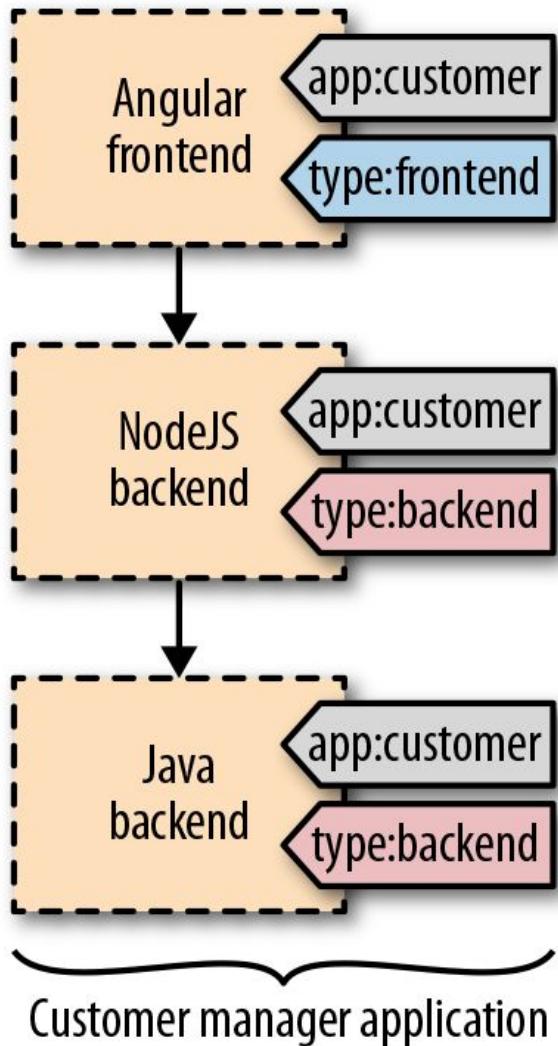
Architecture

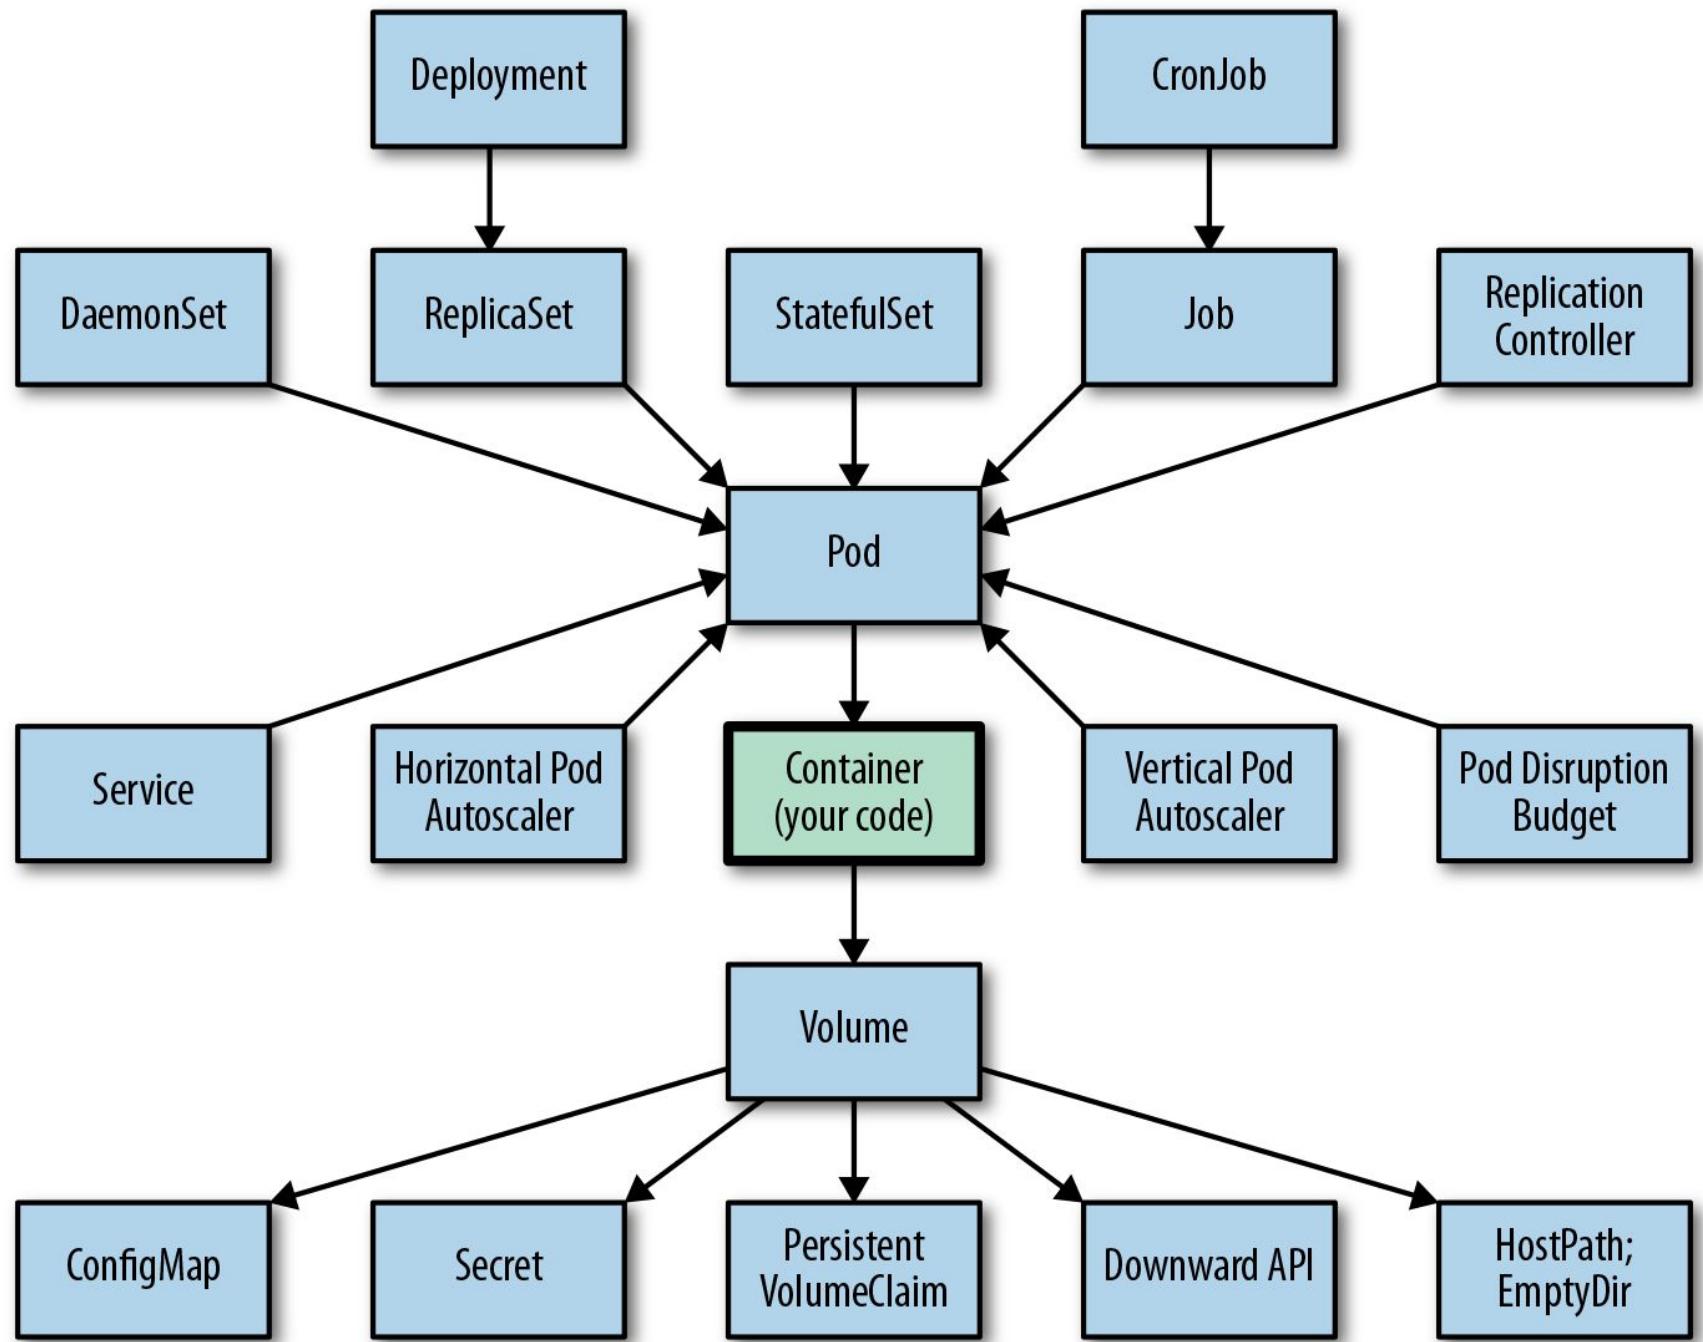


Pod



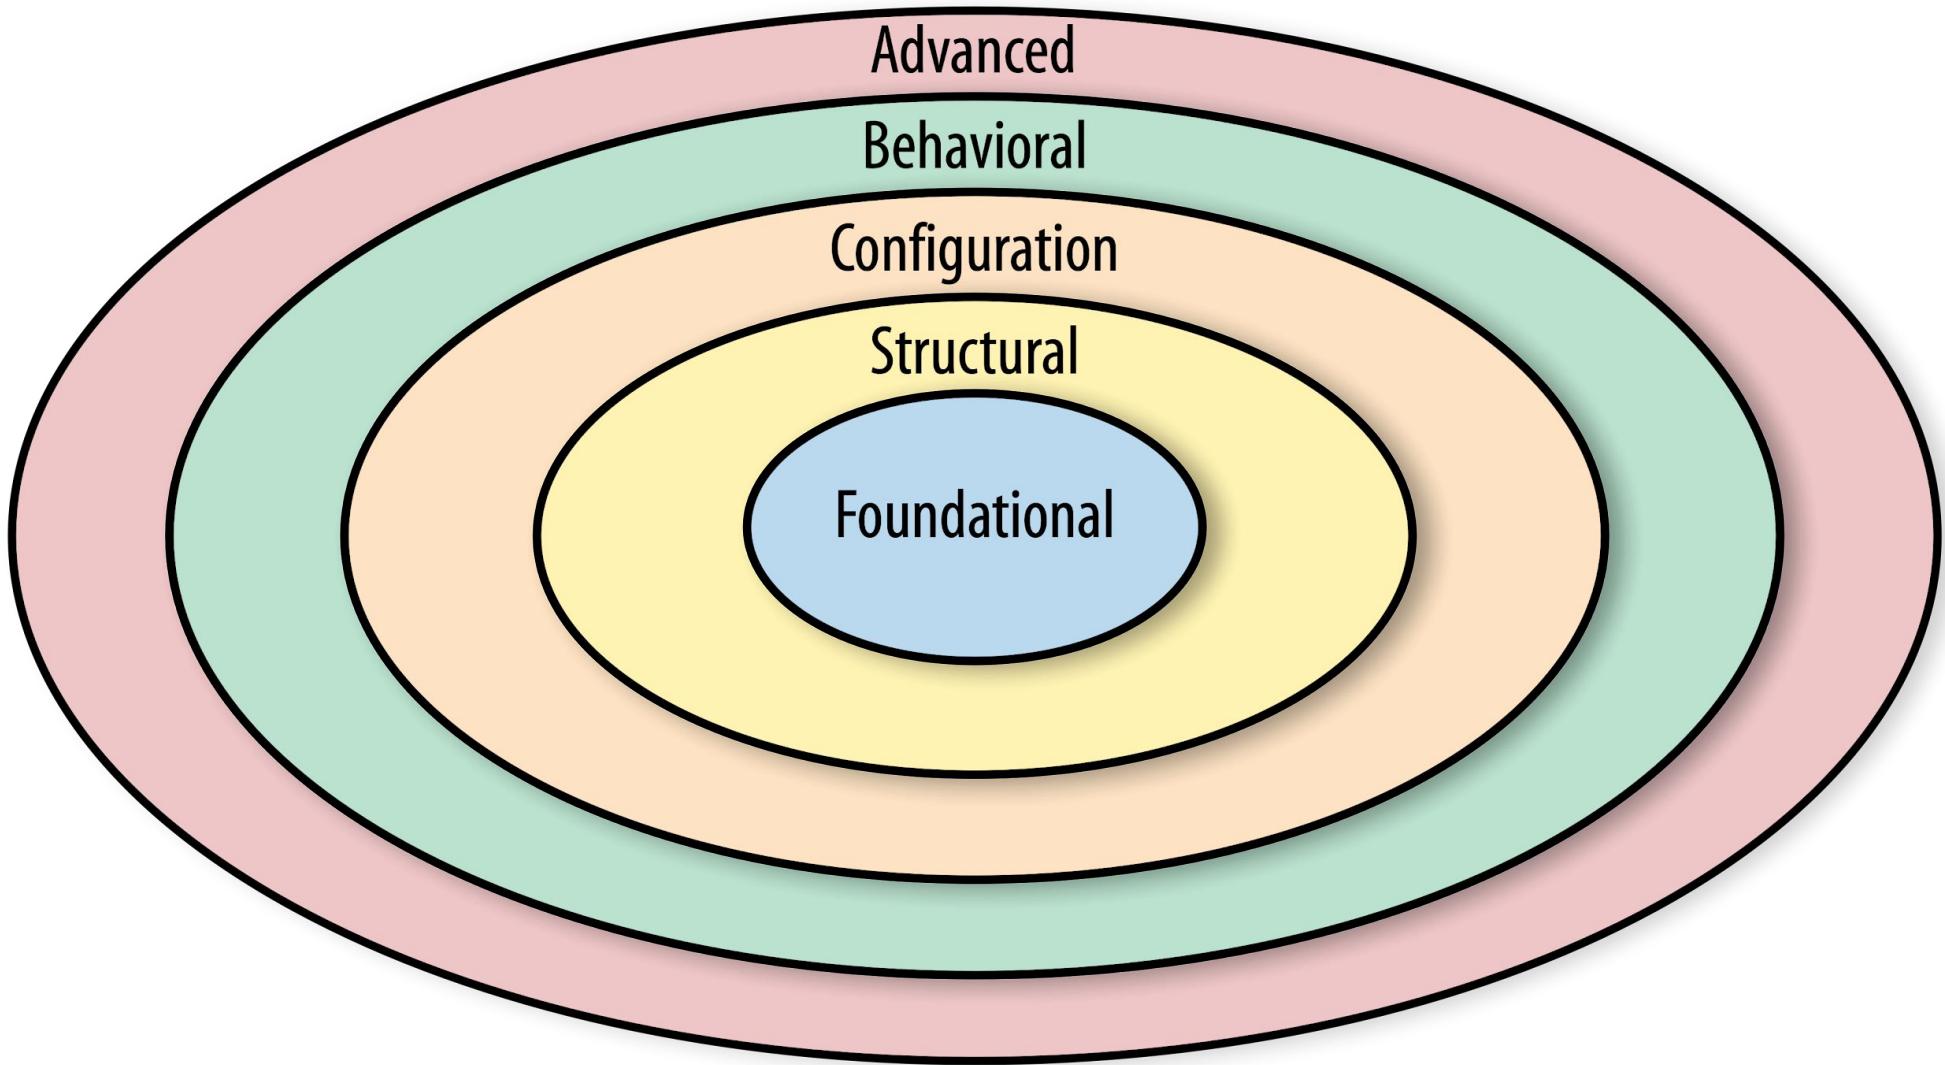
Labels

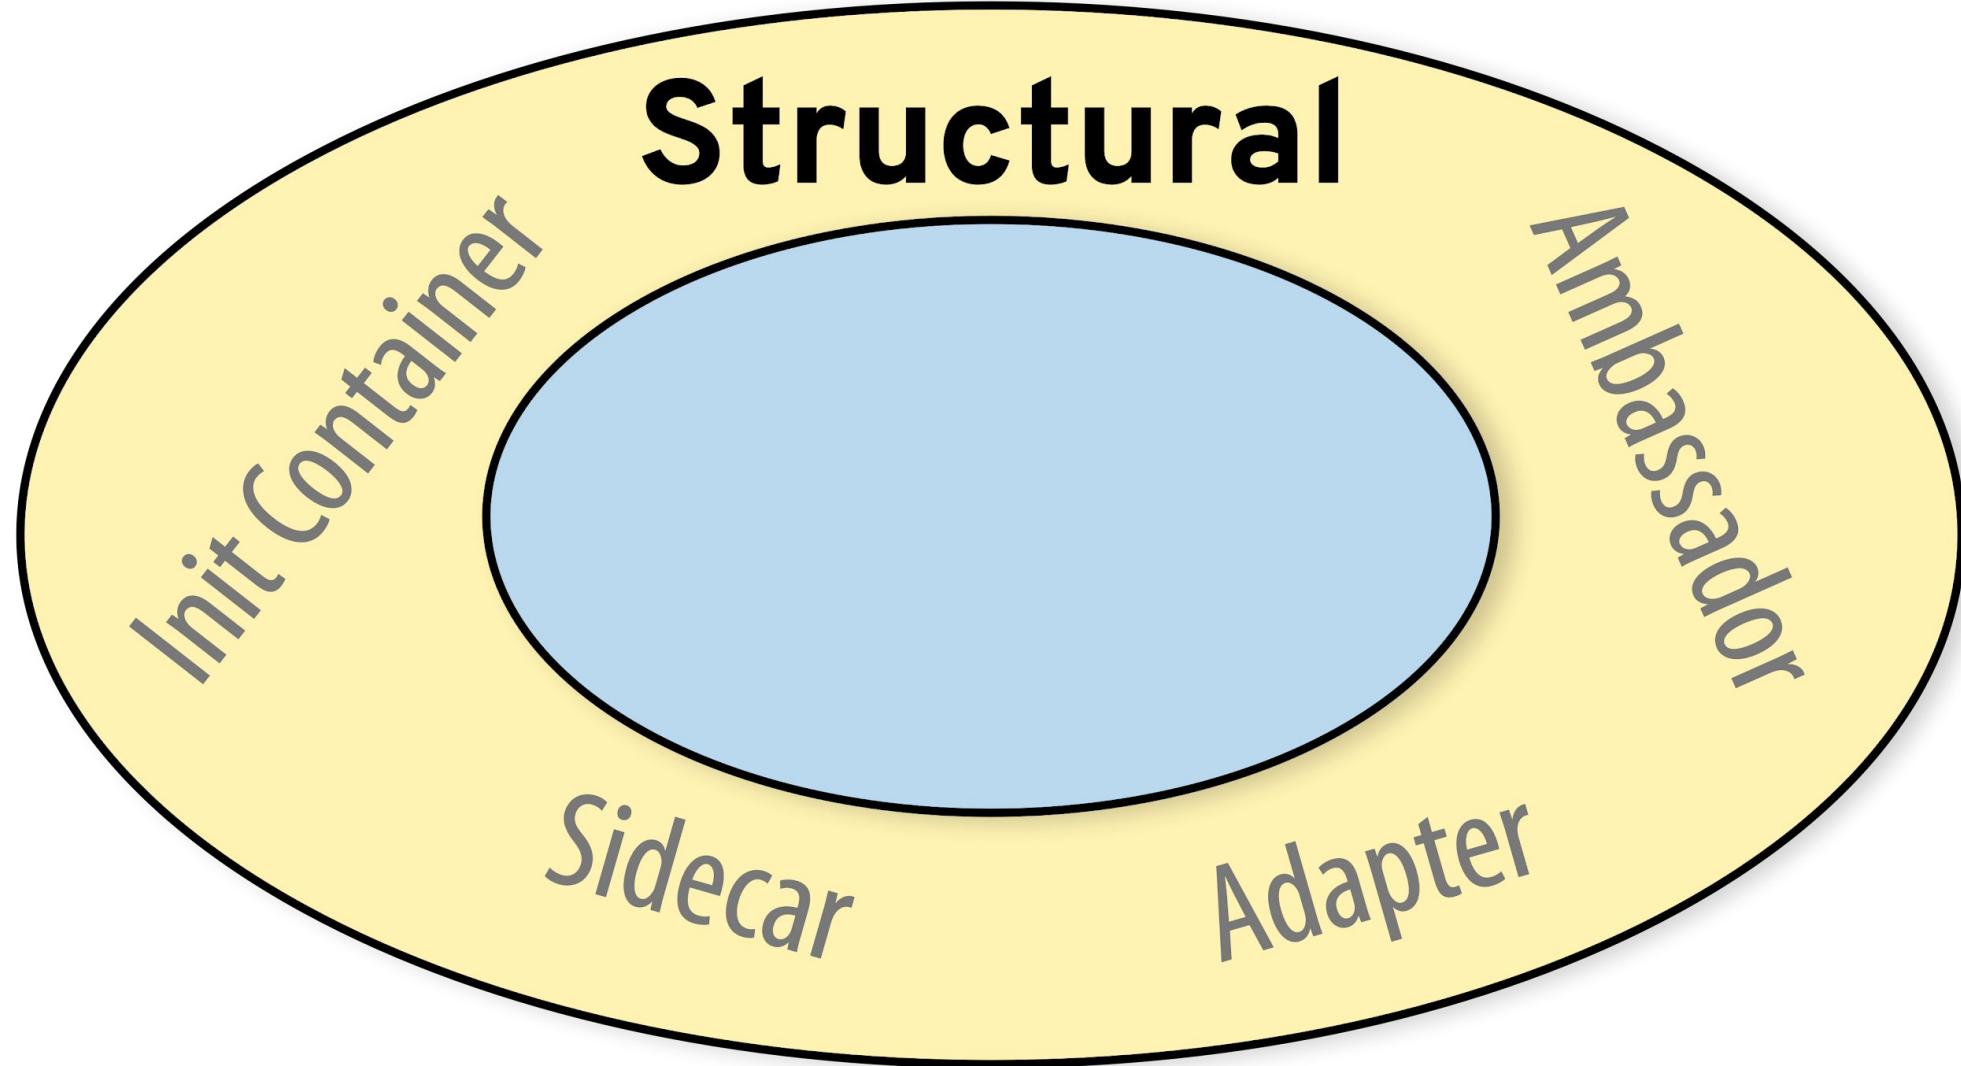


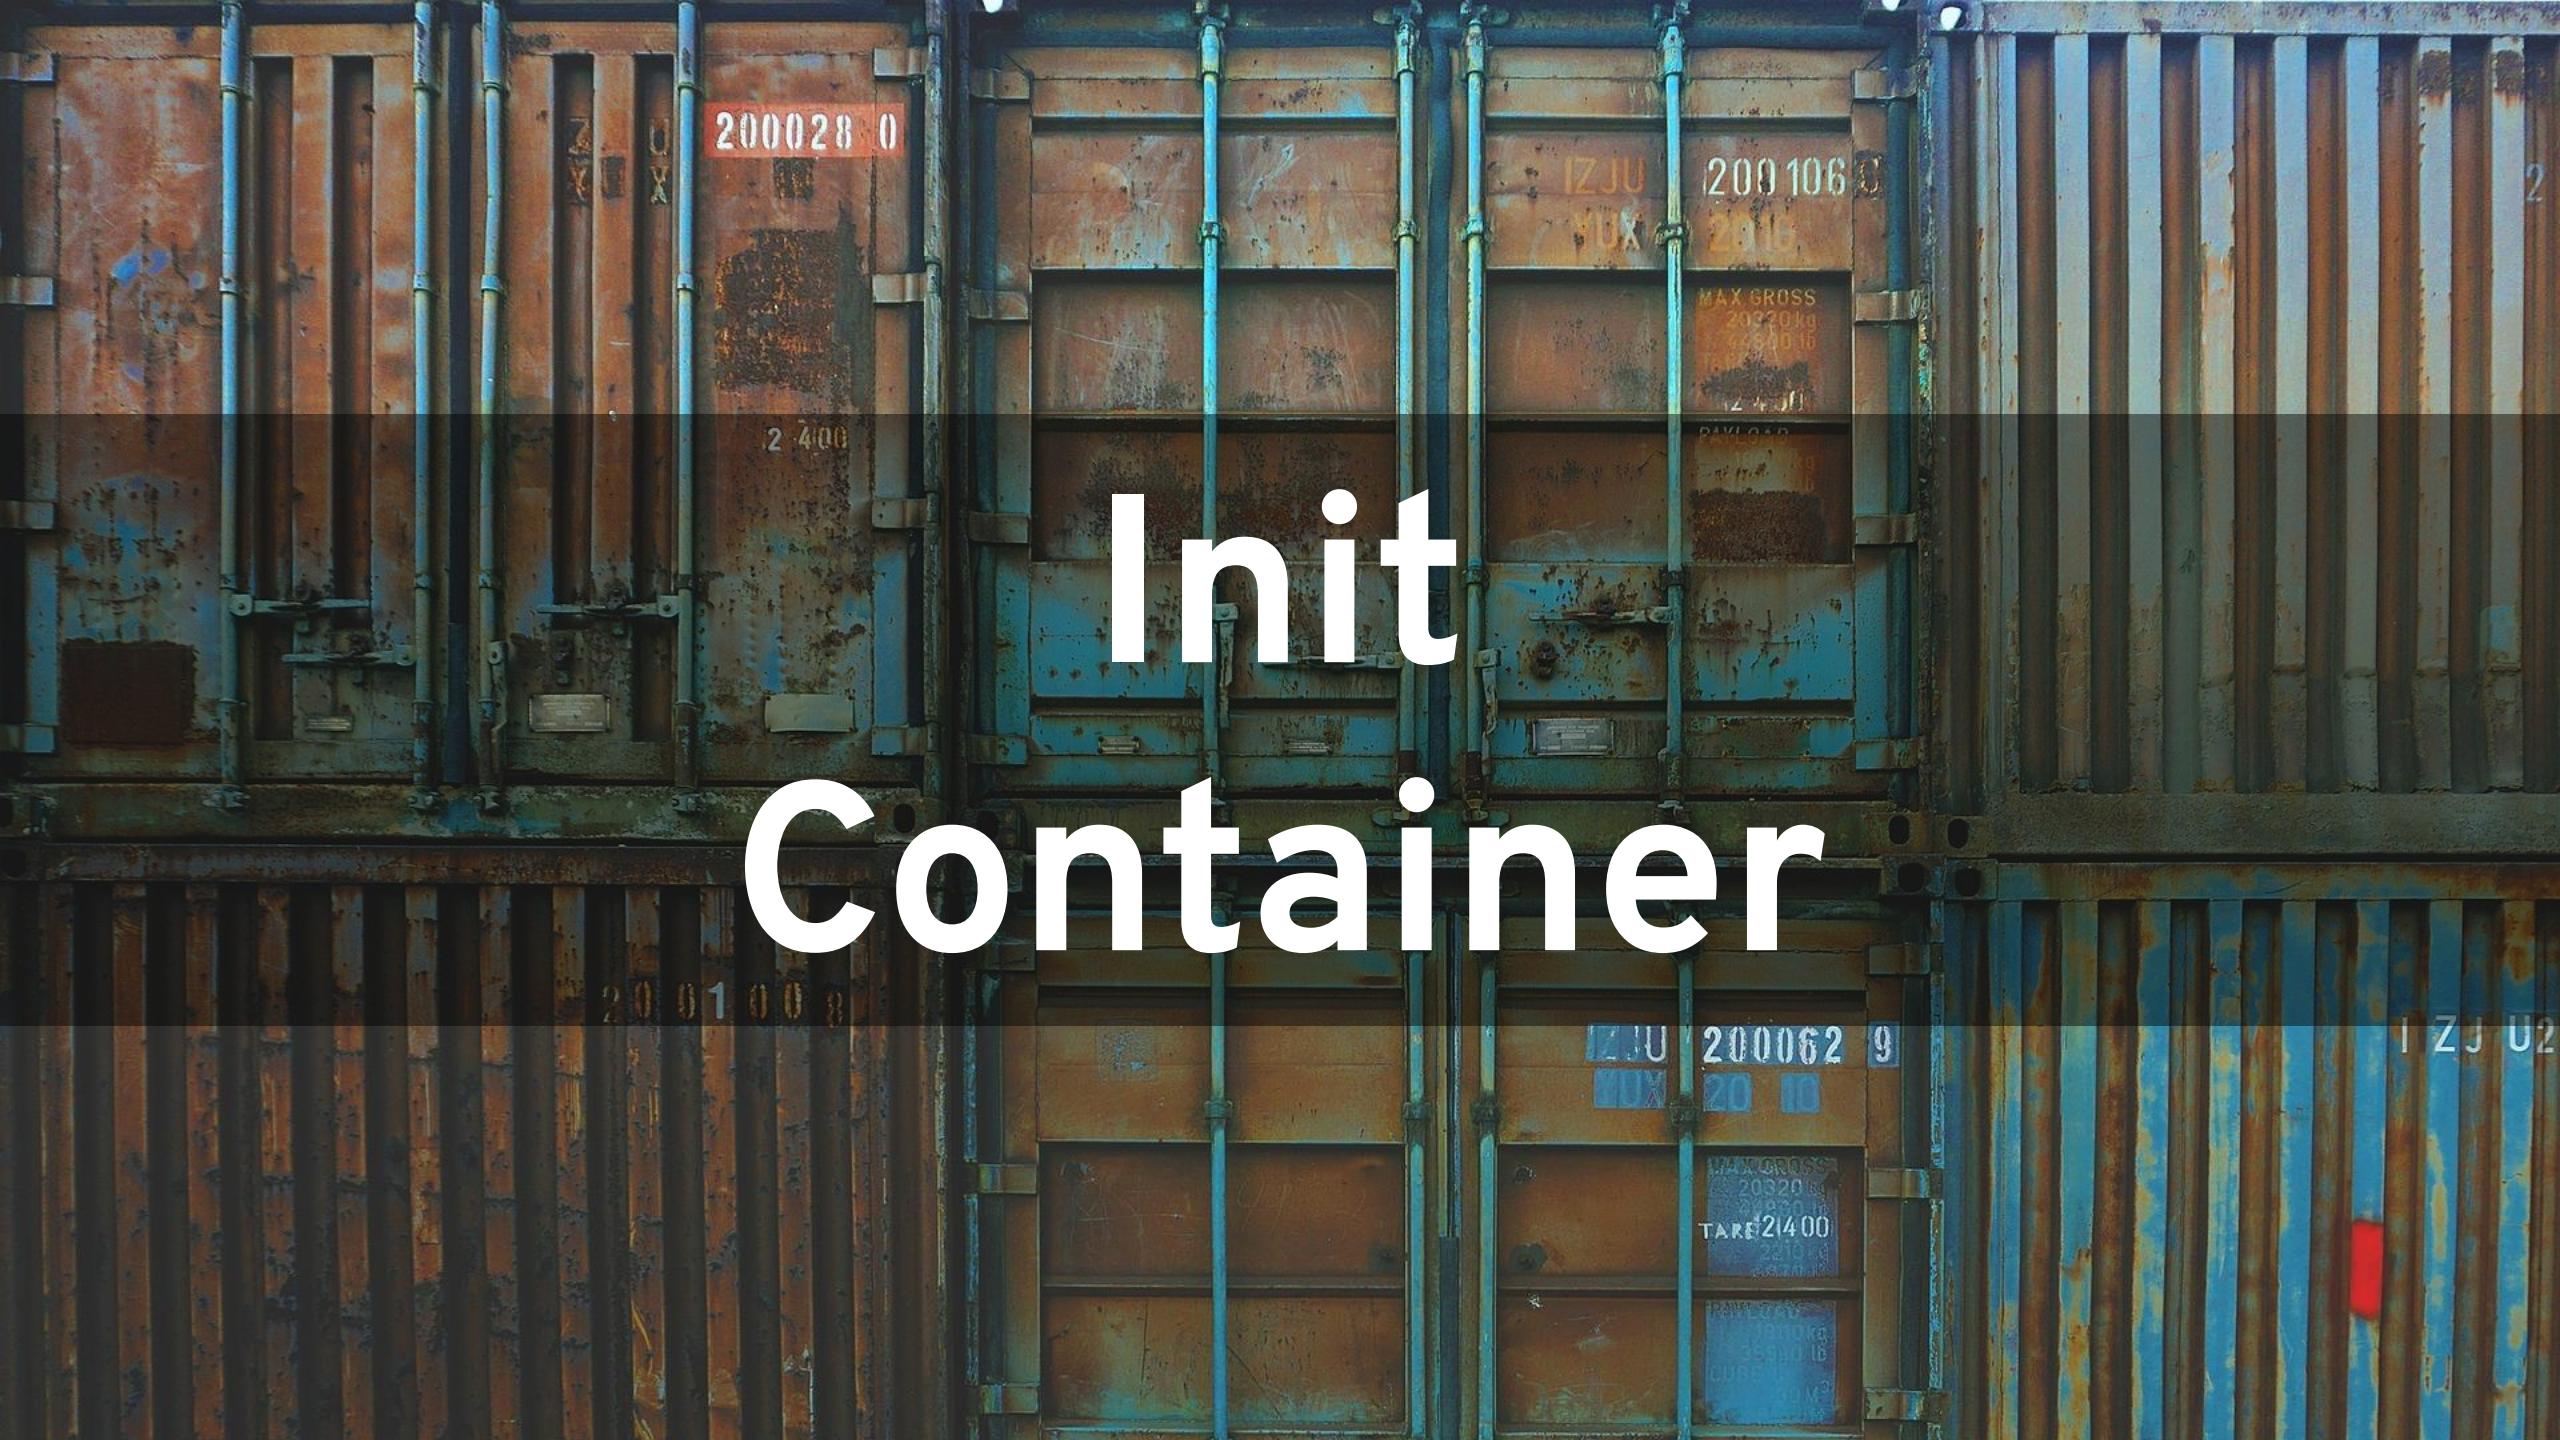


Local vs Distributed Primitives

Concept	Local primitive	Distributed primitive
Behavior encapsulation	Class	Container image
Behavior instance	Object	Container
Unit of reuse	.jar	Container image
Composition	Class A contains Class B	Sidecar pattern
Inheritance	Class A extends Class B	A container's FROM parent image
Deployment unit	.jar/.war/.ear	Pod
Buildtime/Runtime isolation	Module, Package, Class	Namespace, Pod, container
Initialization preconditions	Constructor	Init container
Postinitialization trigger	Init-method	postStart
Predestroy trigger	Destroy-method	preStop
Cleanup procedure	finalize(), shutdown hook	Defer container ^a
Asynchronous & parallel execution	ThreadPoolExecutor, ForkJoinPool	Job
Periodic task	Timer, ScheduledExecutorService	CronJob
Background task	Daemon thread	DaemonSet
Configuration management	System.getenv(), Properties	ConfigMap, Secret







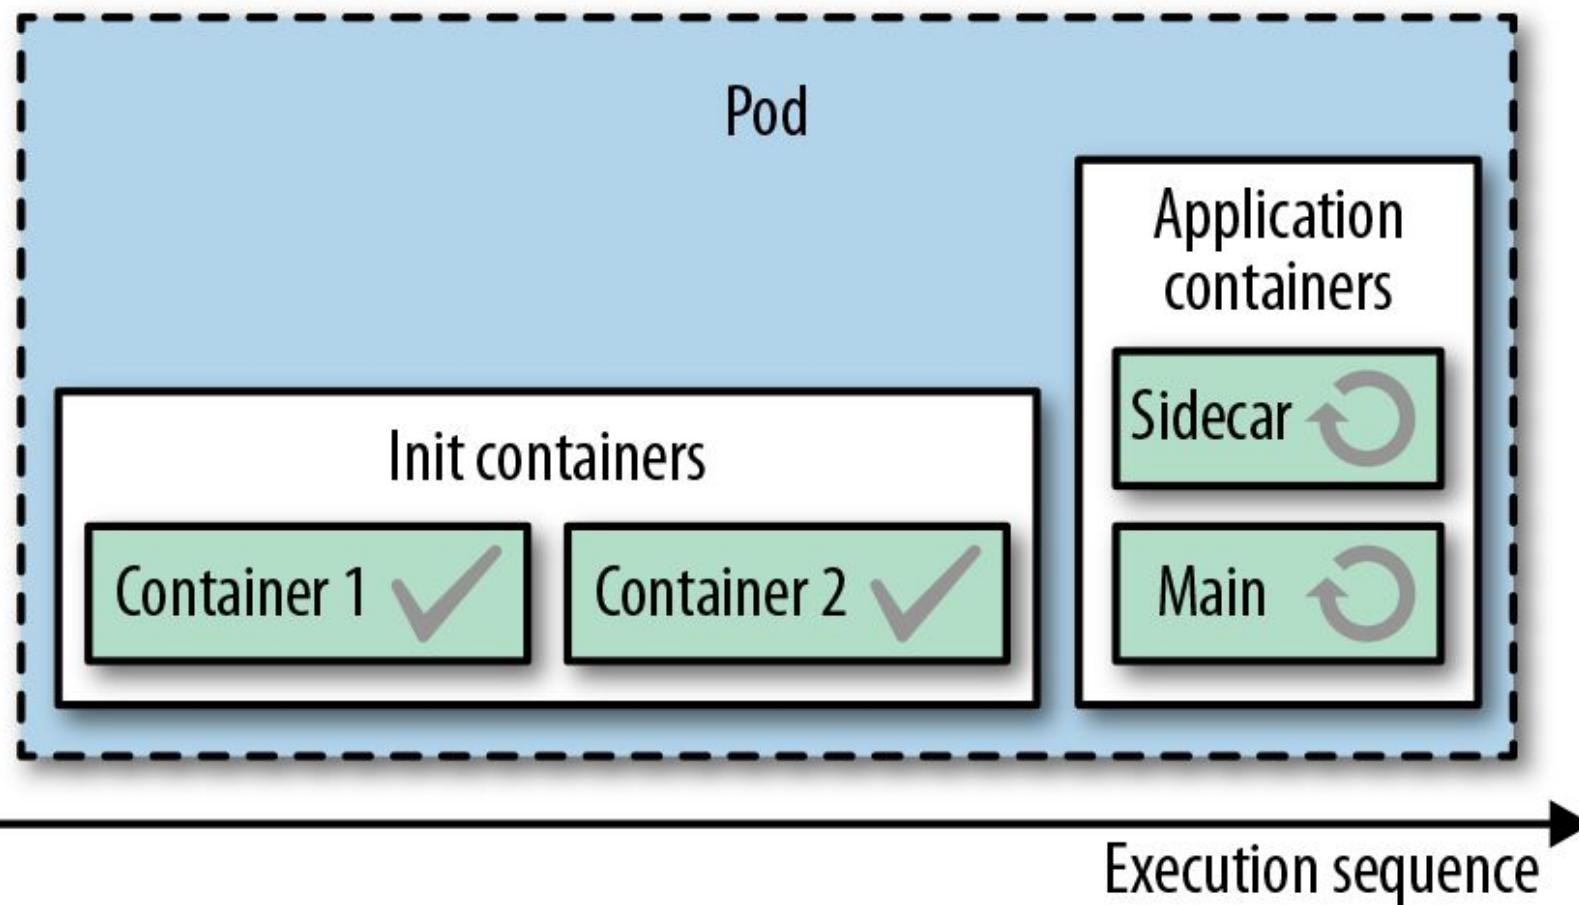
Init Container

How to separate initialization steps from the main application

Init Container

- Part of a Pod
- One shot actions before application starts
- Needs to be idempotent
- Has own resource requirements

Init Container



Init Container

```
apiVersion: v1
kind: Pod
.....
spec:
  initContainers:
  - name: download
    image: axeclbr/git
    command: [ "git", "clone", "https://github.com/myrepo", "/data" ]
    volumeMounts:
      - mountPath: /var/lib/data
        name: source
  containers:
  - name: run
    image: docker.io/centos/httpd
    volumeMounts:
      - mountPath: /var/www/html
        name: source
  volumes:
  - emptyDir: {}
    name: source
```

Demo



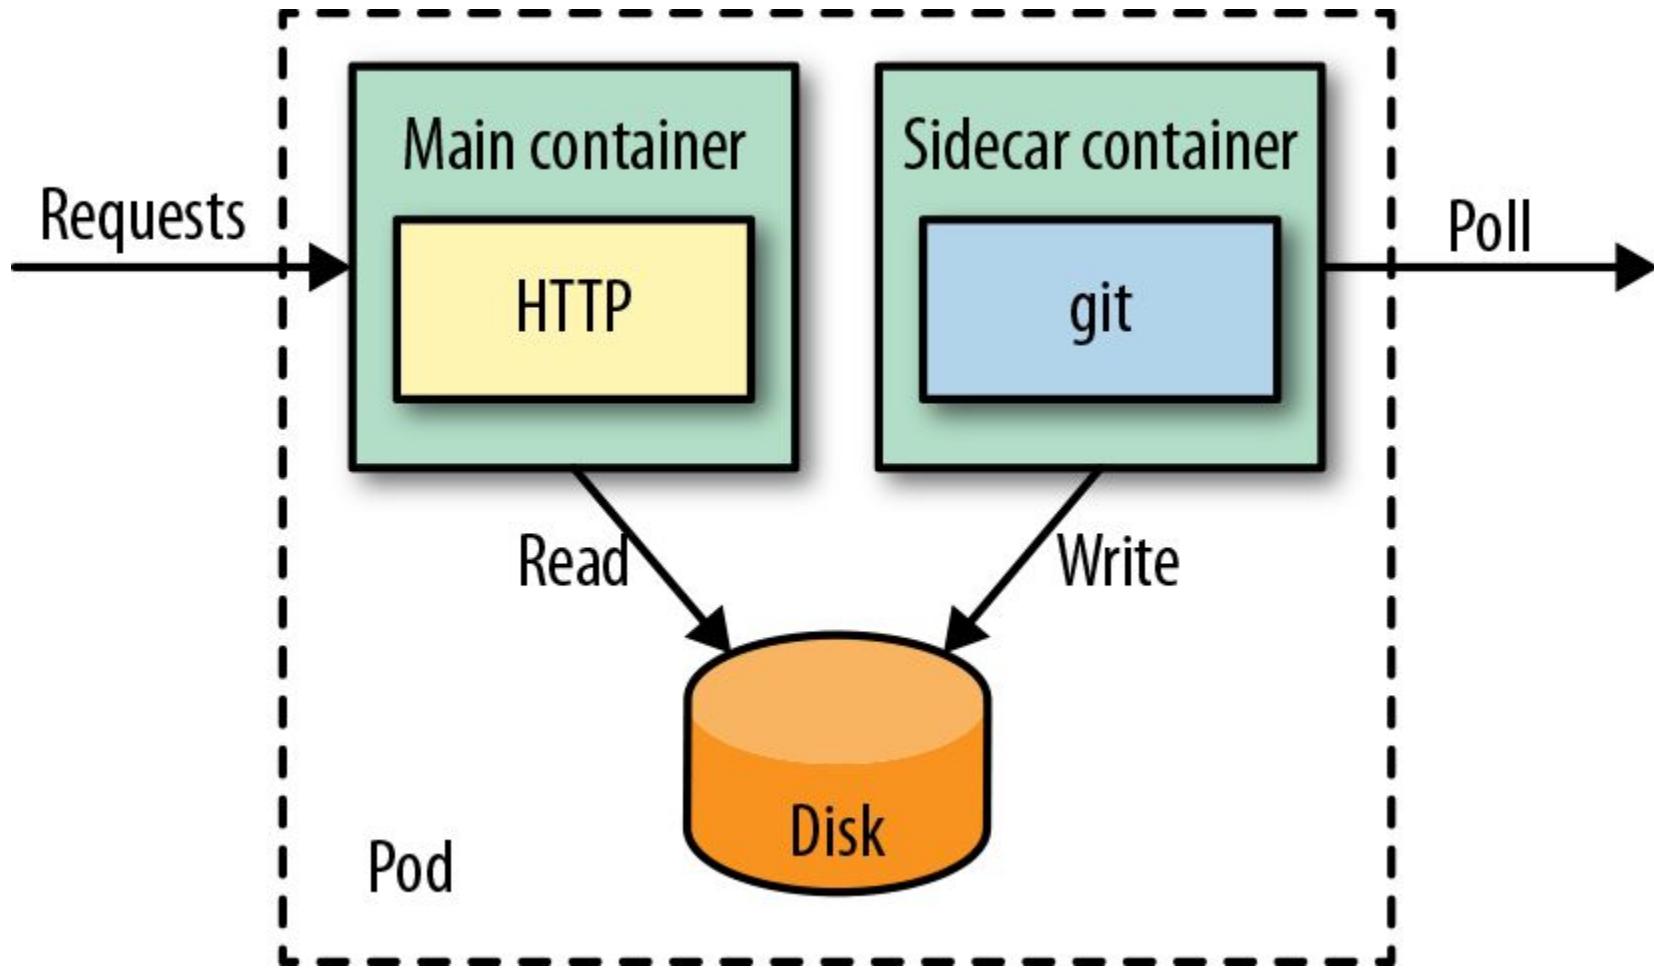
Sidecar

How to enhance the functionality of an application without changing it

Sidecar

- Runtime collaboration of containers
- Connected via shared resources:
 - Network
 - Volumes
- Similar what AOP is for programming
- Separation of concerns

Sidecar

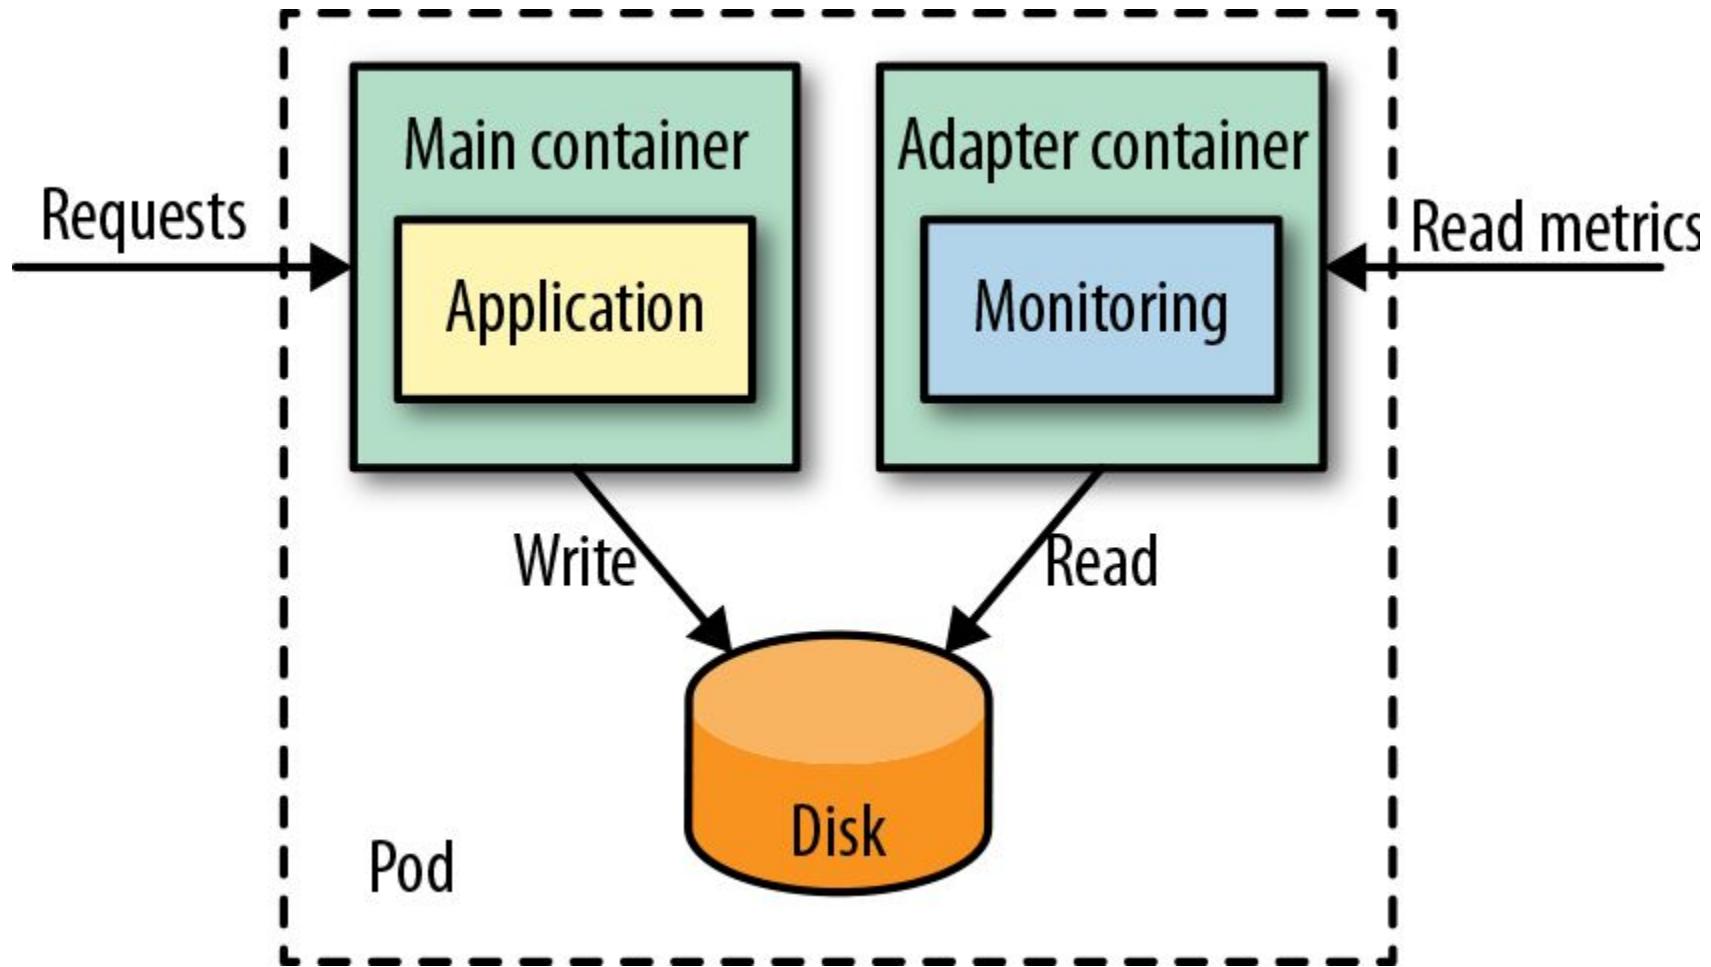


Adapter



How to decouple access to a container **from** the outside world

Adapter

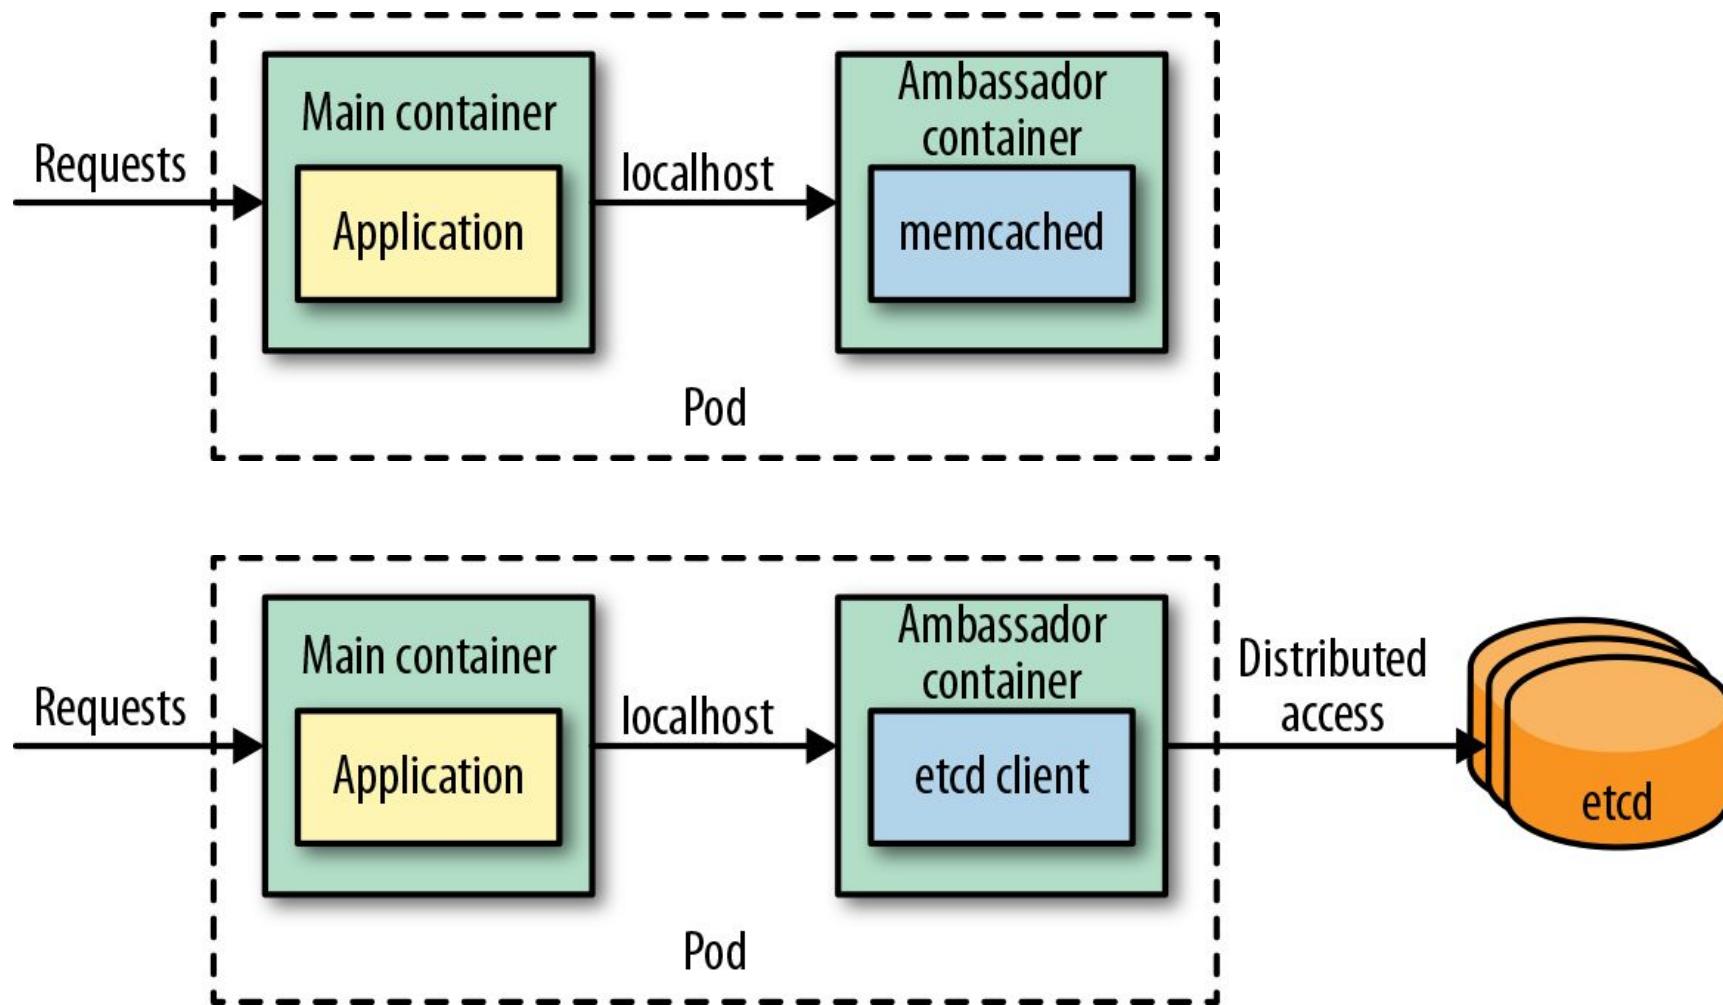




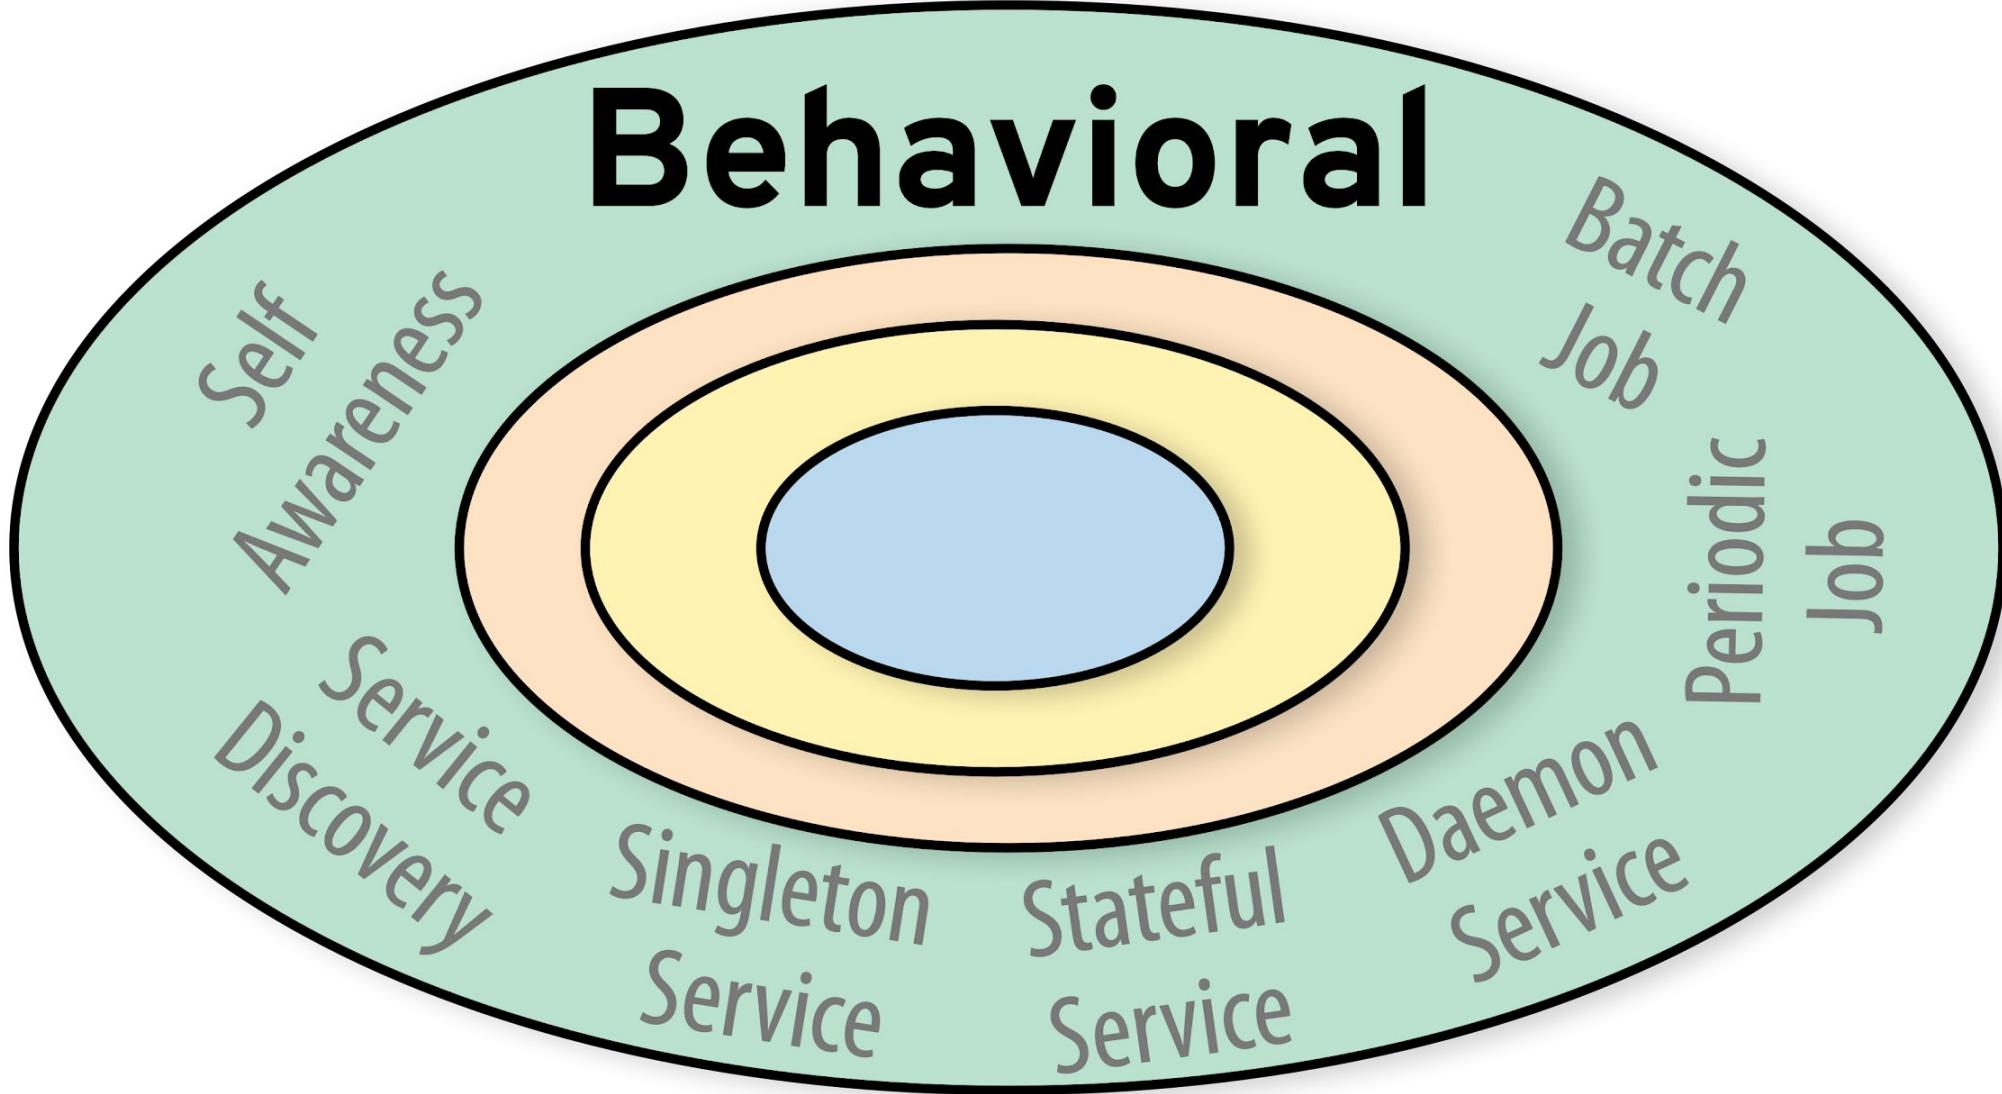
Ambassador

How to decouple a container's access **to** the outside world

Ambassador



Demo



Stateful Service



How to manage stateful workloads

Distributed Stateful Services

- Non-shared persistent Storage
- Unique and stable network address
- Unique identity
- Defined instance order
- Minimal availability

StatefulSet

- Similar to ReplicaSet
- Additional Elements
 - serviceName - reference to headless Service
 - volumeClaimTemplates - template for creating instance unique PVCs
- Assigned PVs are **not** automatically deleted
- Headless service for creating DNS entries for each instance's Pod

Headless Service

```
apiVersion: v1
kind: Service
metadata:
  name: random-generator
spec:
  clusterIP: None
  selector:
    app: random-generator
  ports:
  - name: http
    port: 8080
```

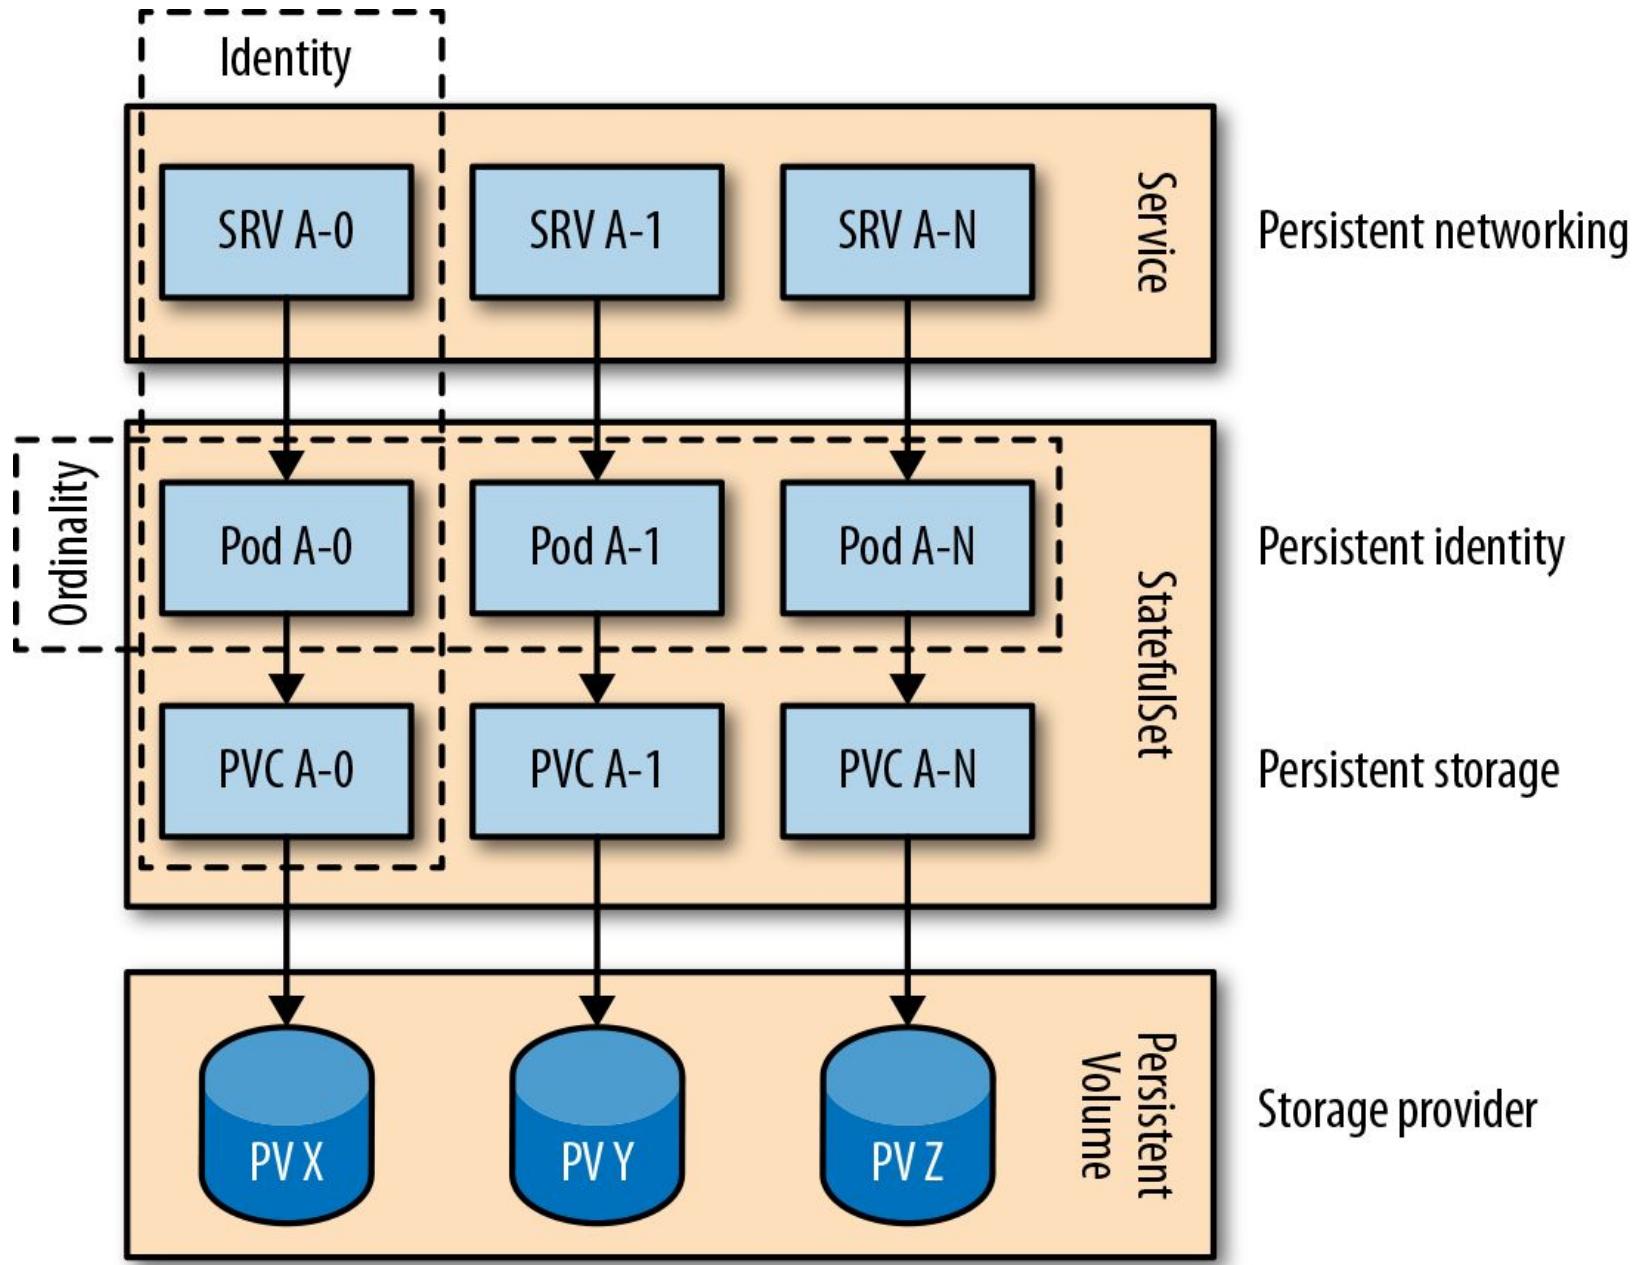
rg-0.random-generator.default.svc.cluster.local
rg-1.random-generator.default.svc.cluster.local

...

Stateful Service

```
apiVersion: apps/v1
kind: StatefulSet
metadata:
  name: rg
spec:
  serviceName: random-generator
  replicas: 2
  selector:
    matchLabels:
      app: random-generator
  template:
    metadata:
      labels:
        app: random-generator
    spec:
      containers:
        - image: k8spatterns/random-generator:1.0
          name: random-generator
          name: http
          volumeMounts:
            - name: logs
              mountPath: /logs
  volumeClaimTemplates:
    - metadata:
        name: logs
      spec:
        resources:
          requests:
            storage: 10Mi
```

Stateful Service



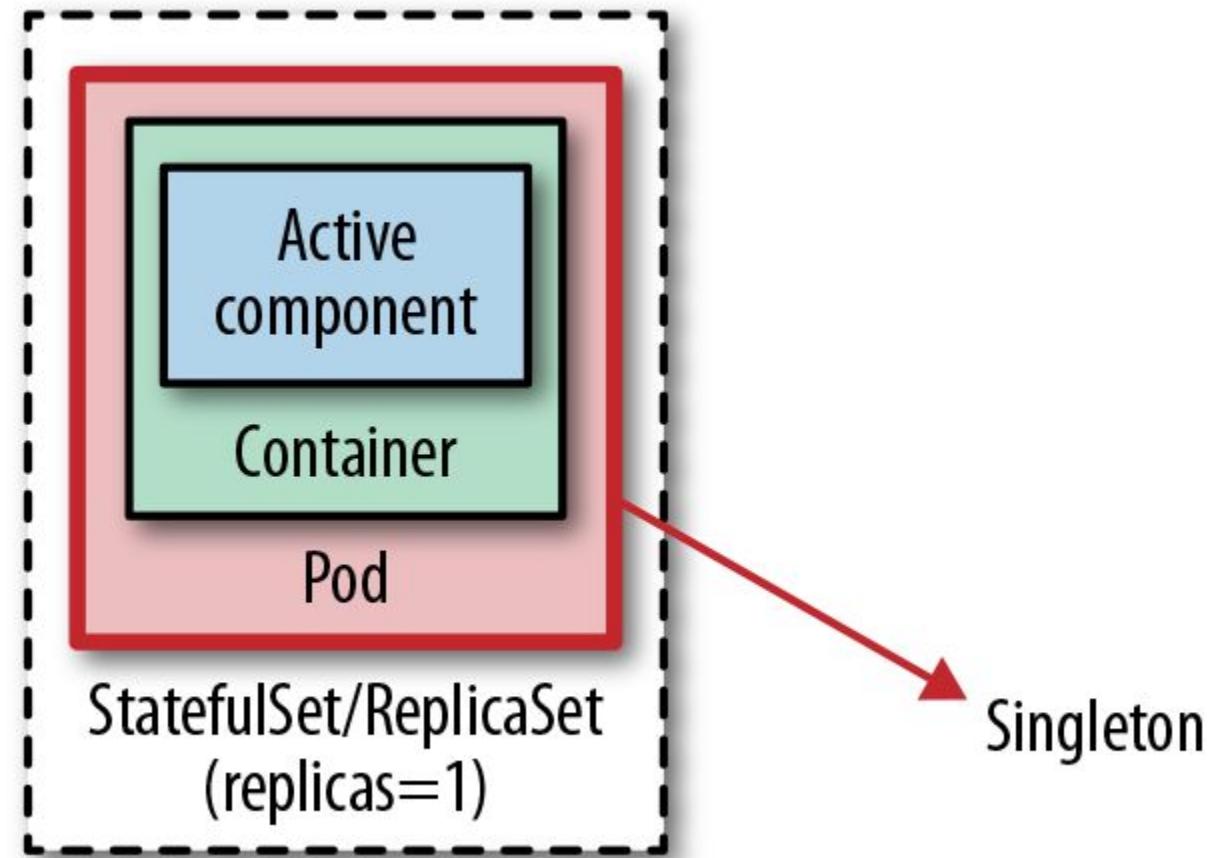
Demo

Singleton Service



How to ensure that only one application instance is active

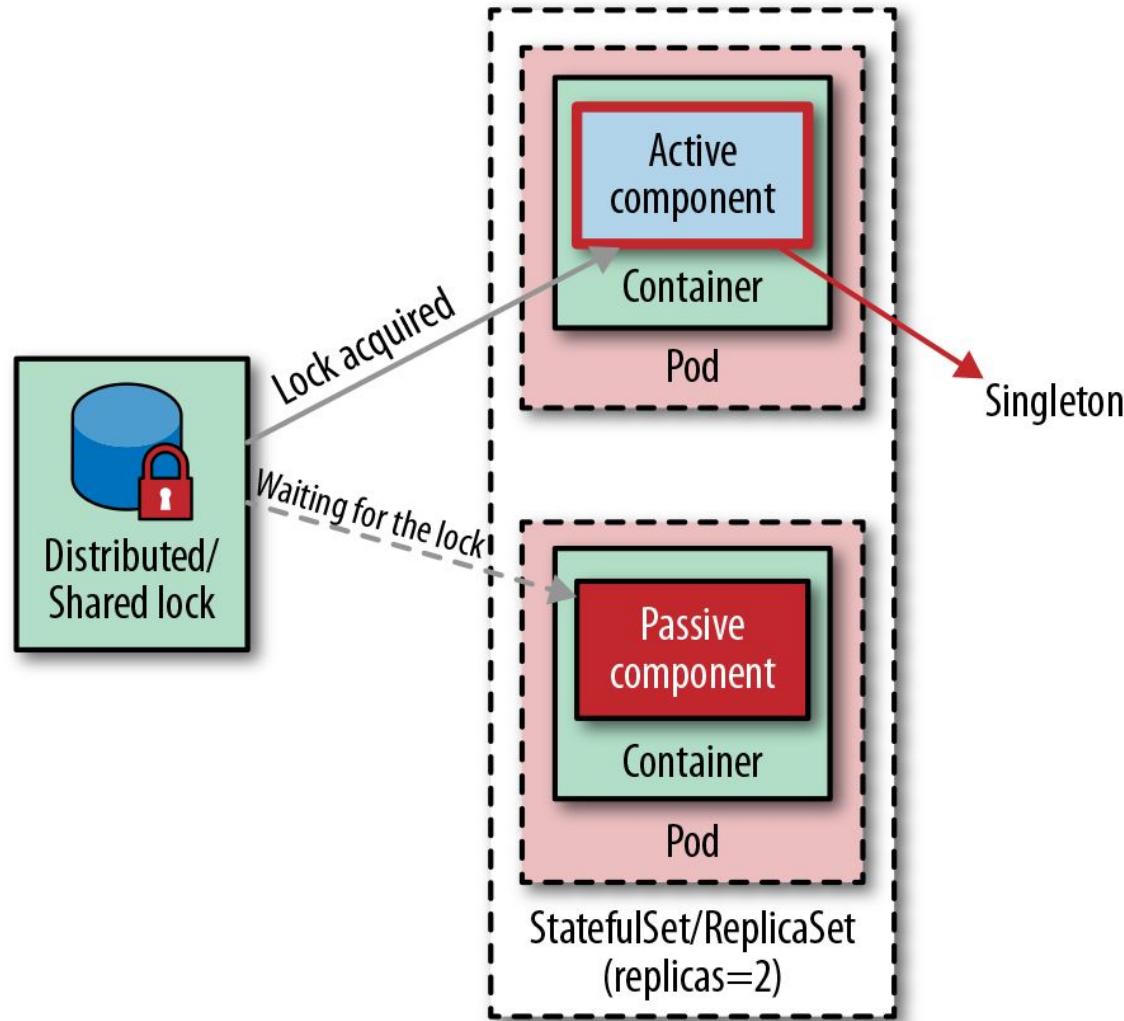
Out-of-Application Locking



Out-of-Application Locking

- ReplicaSet with 1 replica
- Highly available Pod which is monitored and restarted in case of failures
- ReplicaSet favors availability over consistency
 - more than one Pod can exist temporarily
- Alternative: StatefulSet with 1 replica

In-Application Locking



In-Application Locking

- Distributed lock shared by simultaneously running applications
- *Active-Passive* topology
- Distributed lock implementations e.g.
 - Zookeeper
 - Consul
 - Redis
 - etcd

Pod Disruption Budget

Ensures a certain number of Pods will not
voluntarily be evicted from a node

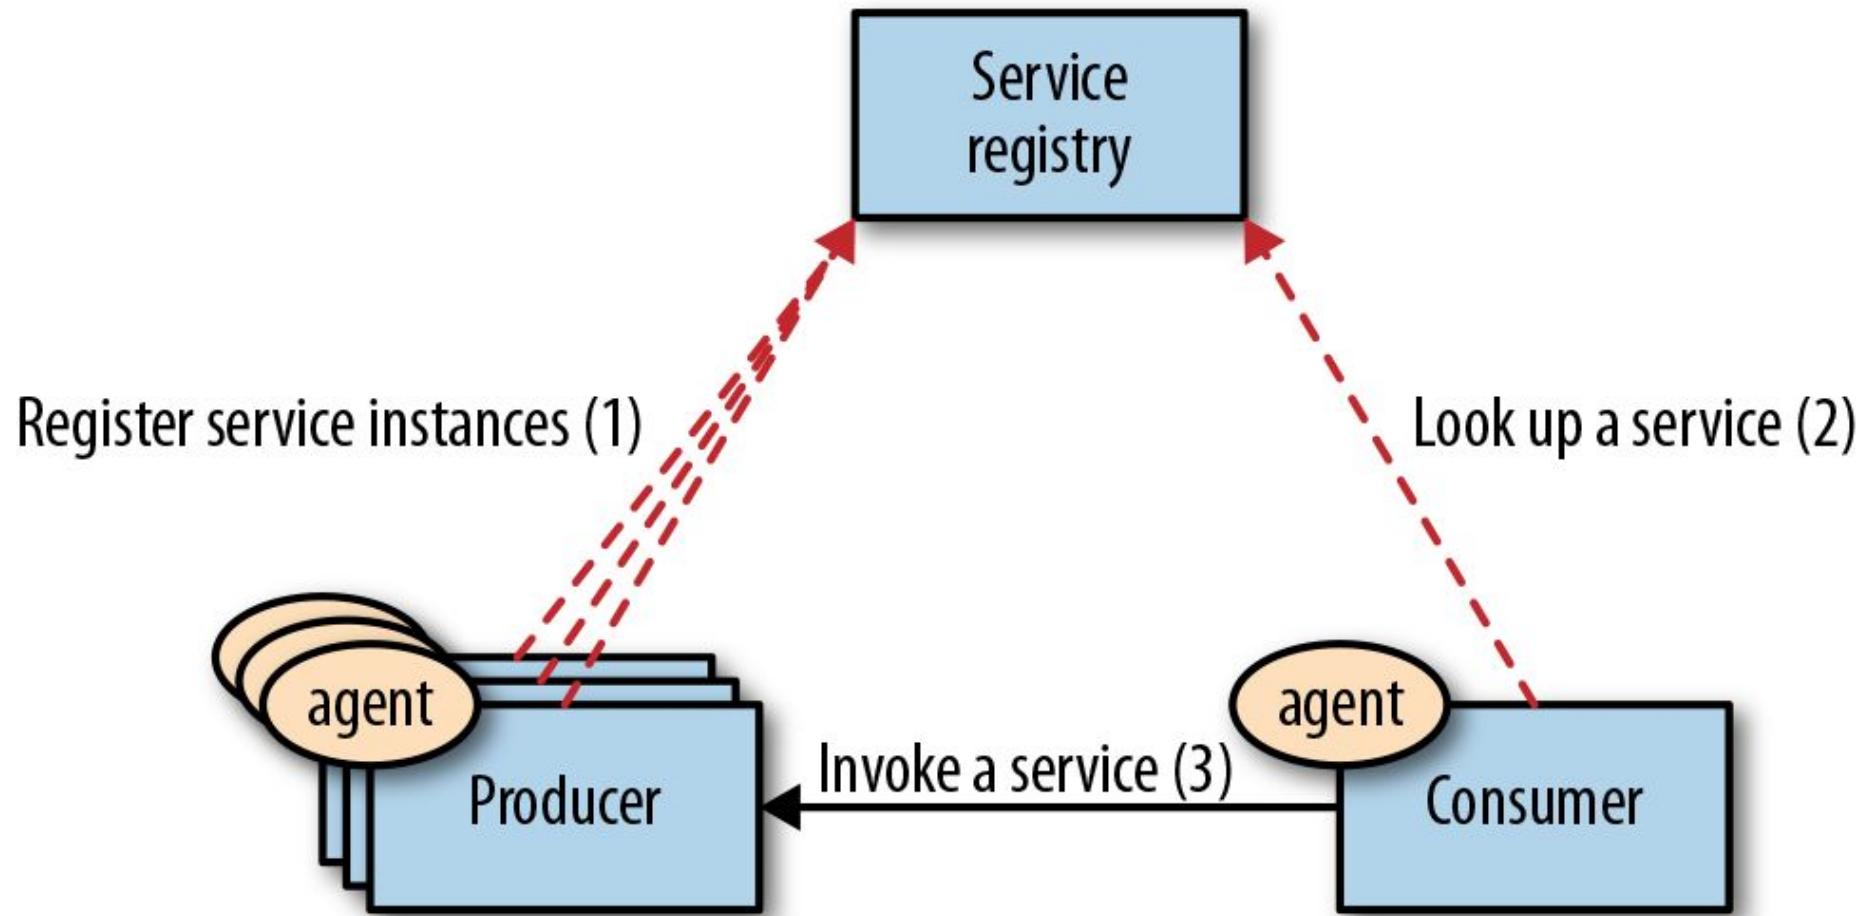
```
apiVersion: policy/v1beta1
kind: PodDisruptionBudget
metadata:
  name: random-generator-pdb
spec:
  selector:
    matchLabels:
      app: random-generator
  minAvailable: 2
```

A vintage-style telescope with a dark body and gold-colored eyepiece and objective lens, mounted on a silver-colored tripod. It is positioned diagonally across the frame, pointing from the top left towards the bottom right. The background is a scenic landscape with a blue sky, white clouds, and distant mountains.

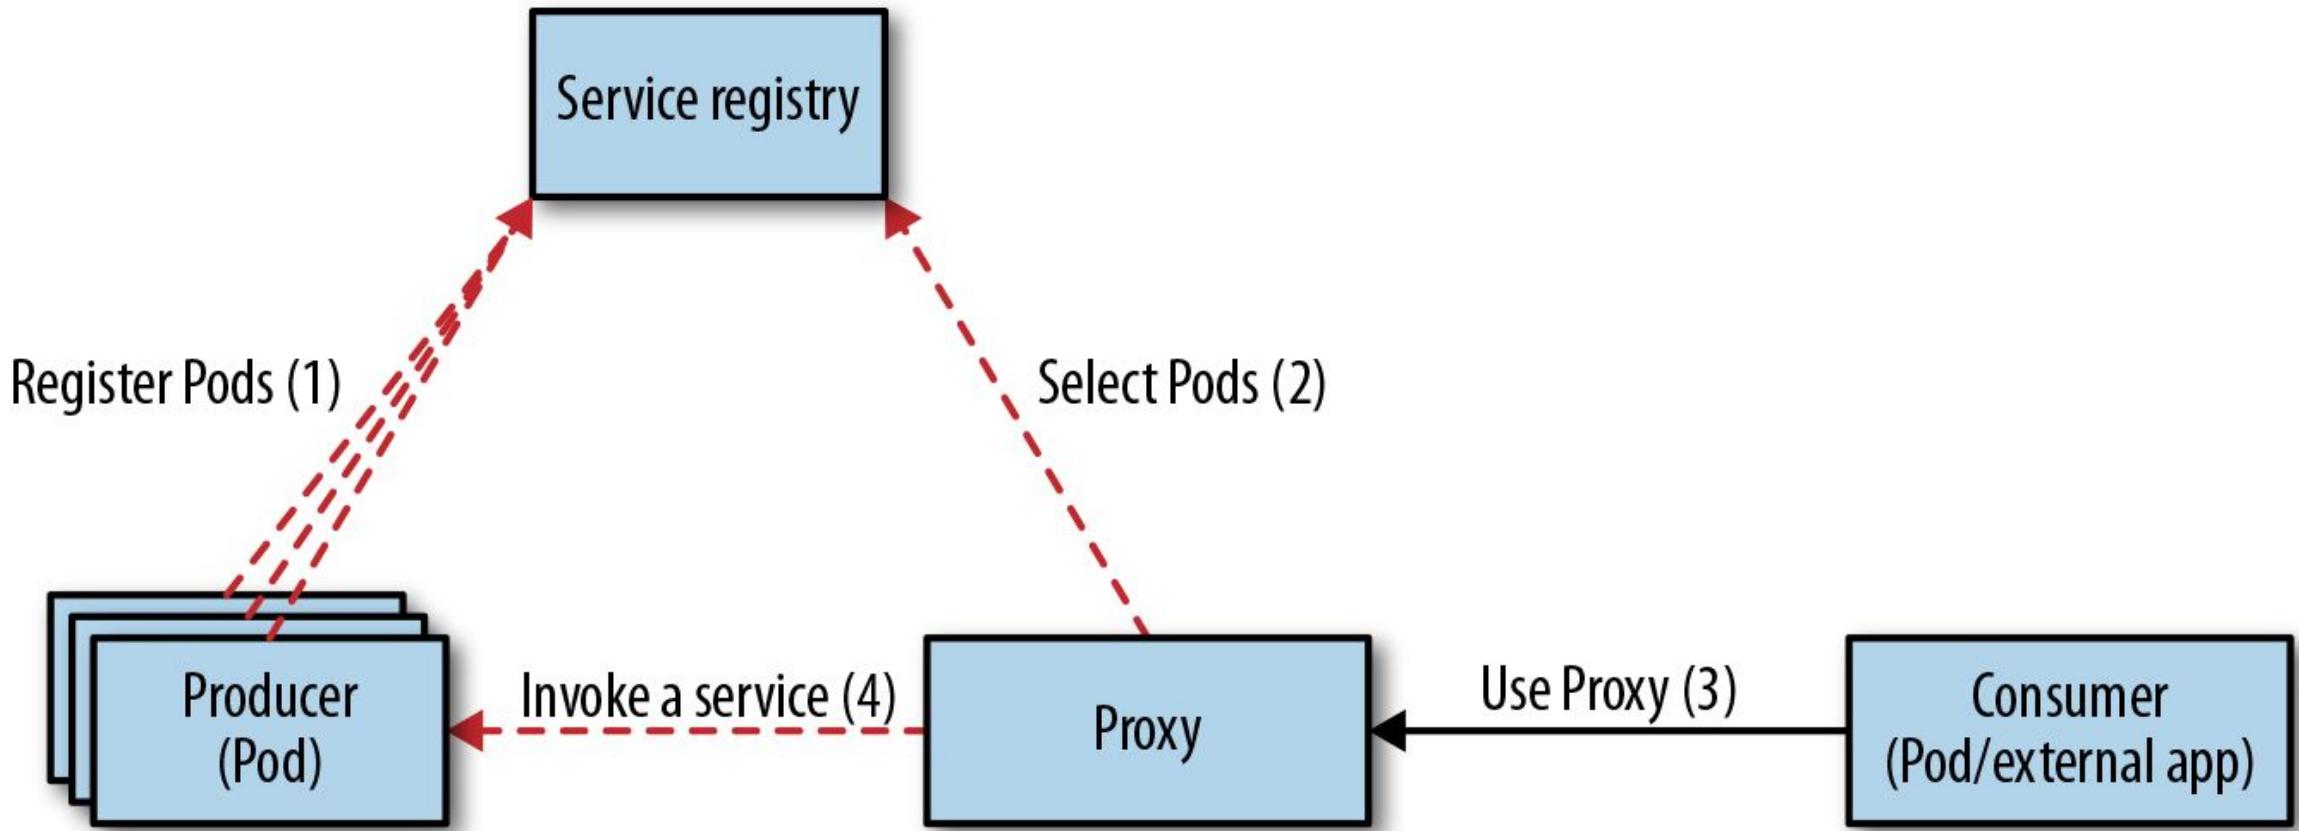
Service Discovery

How to discover and use services

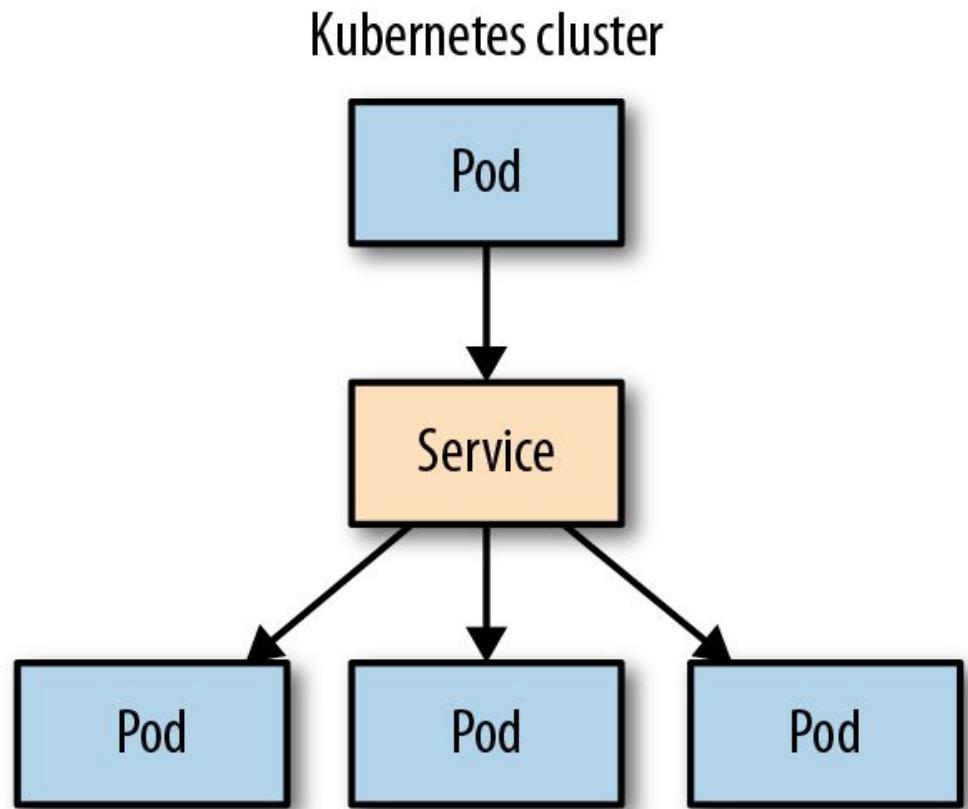
Client-side Service Discovery (non Kubernetes)



Server-side Service Discovery (Kubernetes)

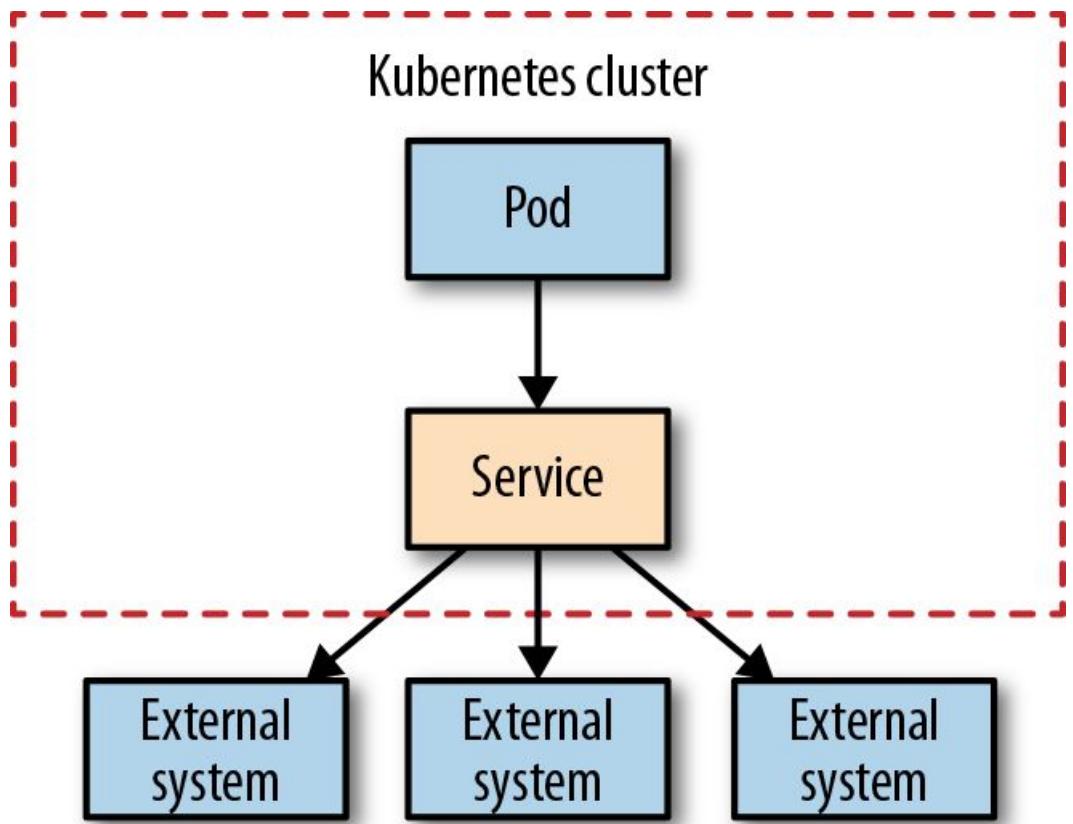


Internal Service Discovery



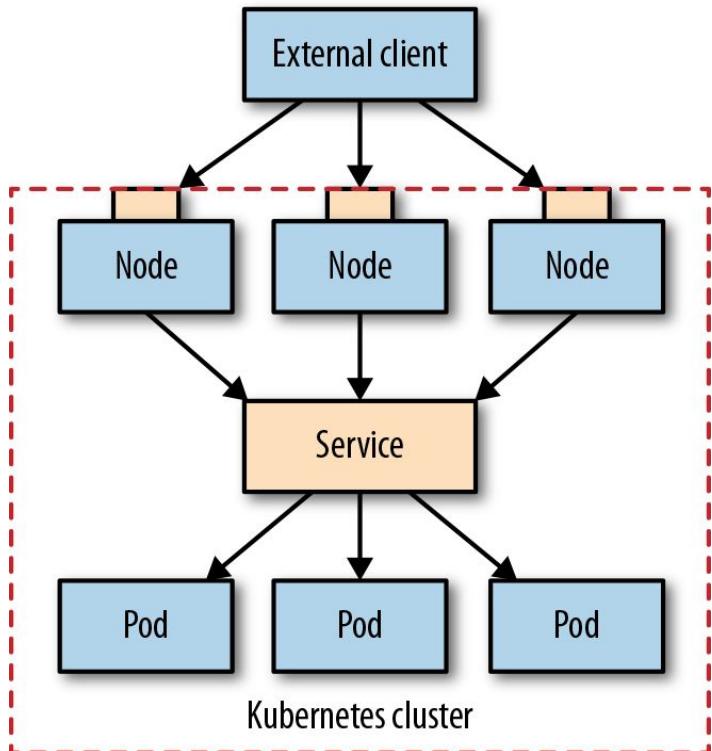
- Discovery through DNS lookups
- Pods picked by label selector
- Multiple ports per Service
- Session affinity on IP address
- Successful readiness probes required for routing
- Virtual IP address for each Service

Manual Service Discovery

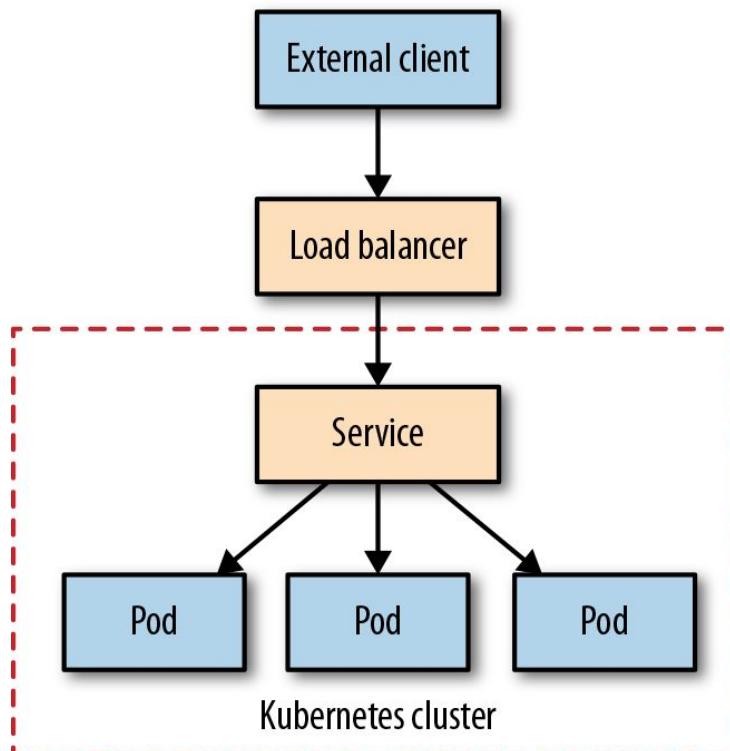


- Service without selector
- Manually creating Endpoint resource with the same name as the Service
- Service of type **ExternalName** map are registered as DNS CNAMEs

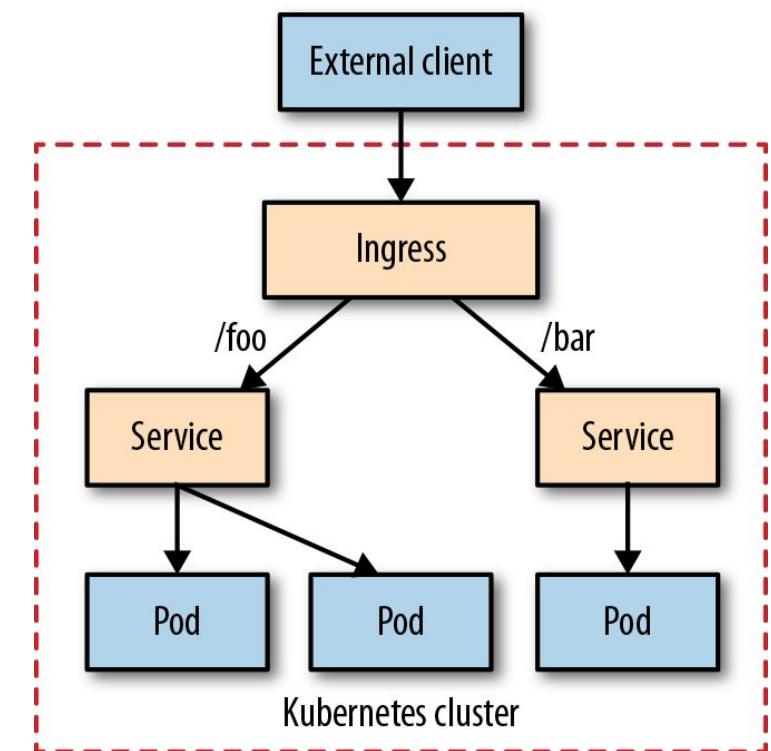
Node Port



Load Balancer



Ingress

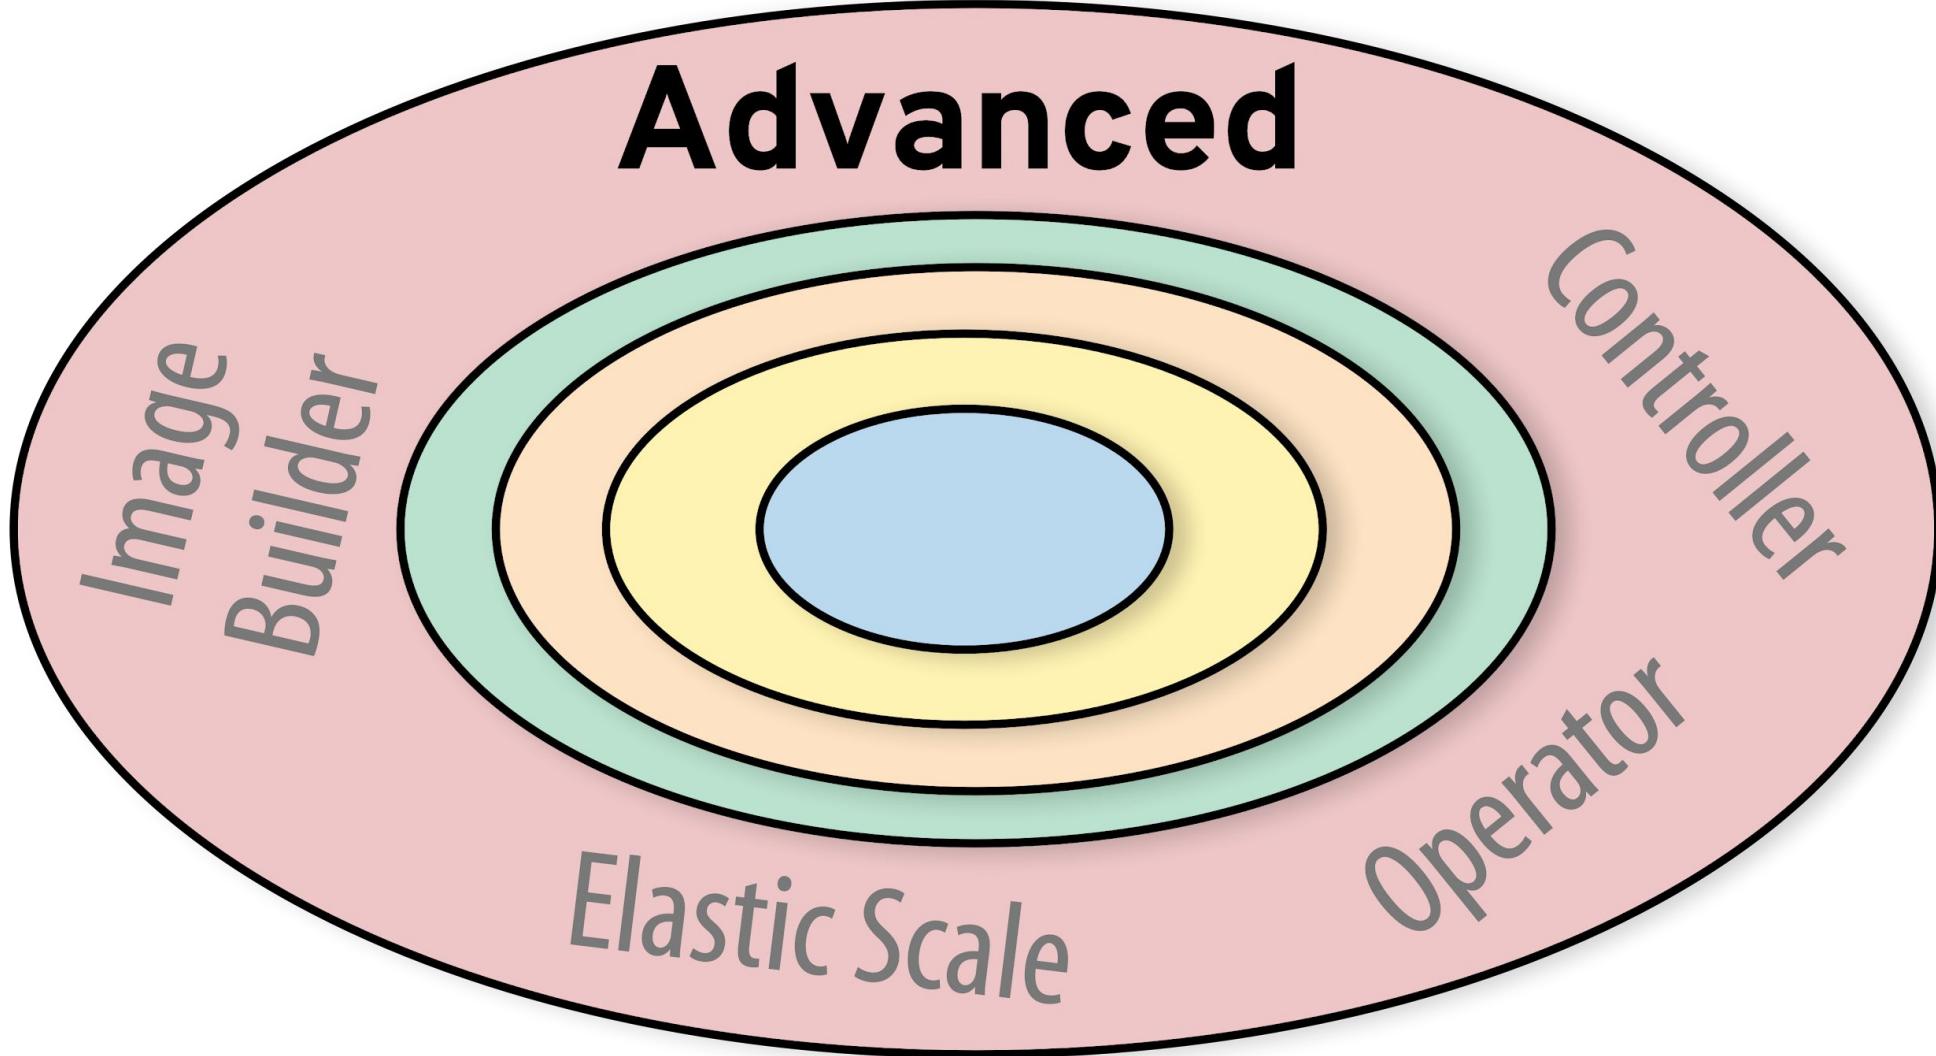


Ingress

```
apiVersion: extensions/v1beta1
kind: Ingress
metadata:
  name: random-generator
spec:
  rules:
  - http:
    paths:
    - path: /
      backend:
        serviceName: random-generator
        servicePort: 8080
    - path: /cluster-status
      backend:
        serviceName: cluster-status
        servicePort: 80
```

Service Discovery

Name	Configuration	Client type	Summary
ClusterIP	<code>type: ClusterIP .spec.selector</code>	Internal	The most common internal discovery mechanism
Manual IP	<code>type: ClusterIP kind: Endpoints</code>	Internal	External IP discovery
Manual FQDN	<code>type: ExternalName .spec.externalName</code>	Internal	External FQDN discovery
Headless Service	<code>type: ClusterIP .spec.clusterIP: None</code>	Internal	DNS-based discovery without a virtual IP
NodePort	<code>type: NodePort</code>	External	Preferred for non-HTTP traffic
LoadBalancer	<code>type: LoadBalancer</code>	External	Requires supporting cloud infrastructure
Ingress	<code>kind: Ingress</code>	External	L7/HTTP-based smart routing mechanism



Controller



How to get from the
current state to a
declared target state

State Reconciliation

- Kubernetes as distributed state manager
- Make the **actual** state more like the declared **target** state.



Observe - Discover the actual state

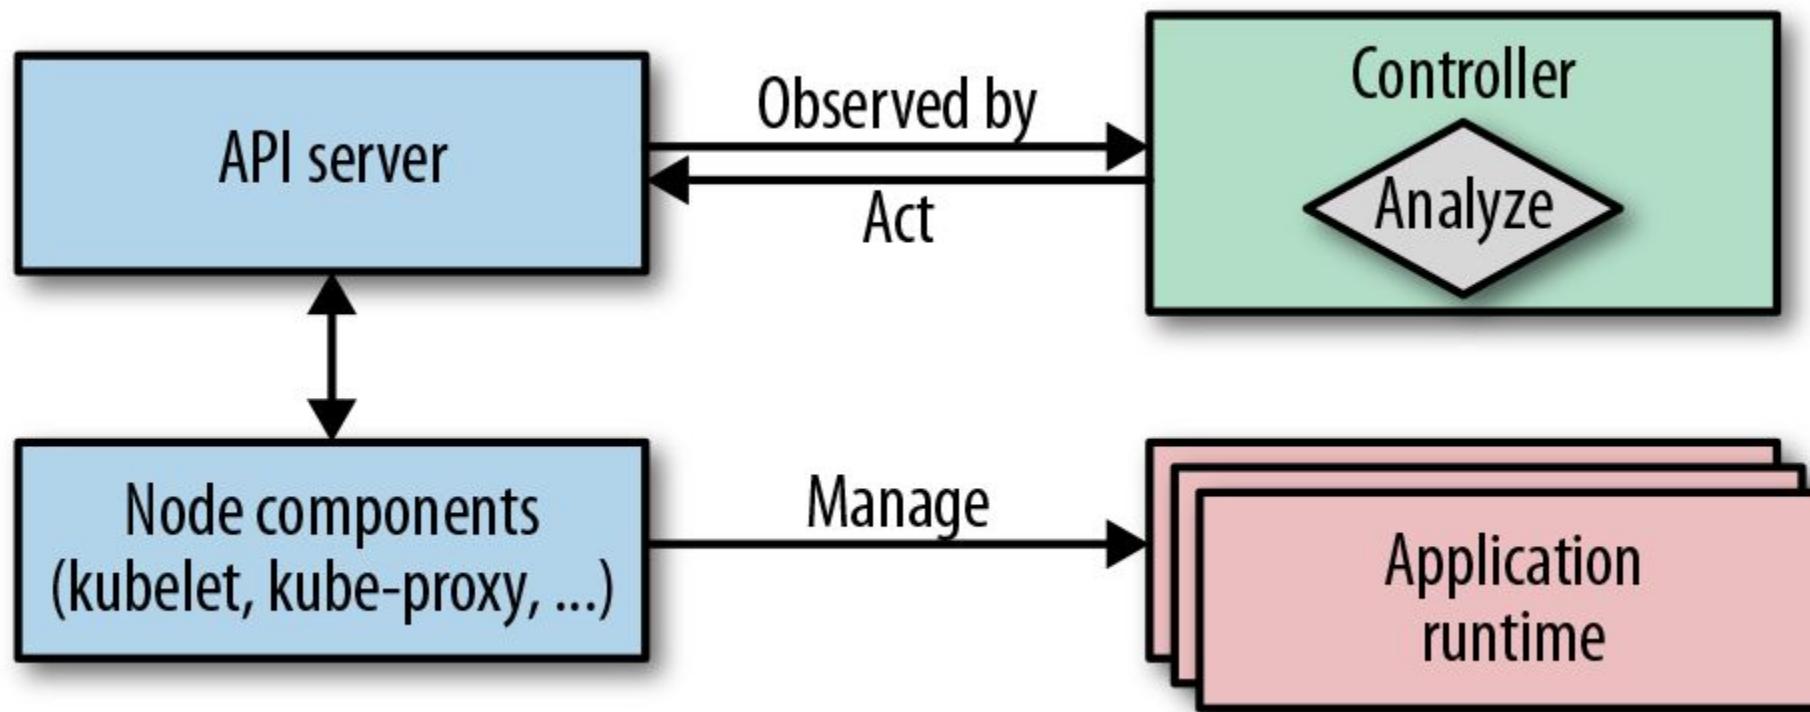


Analyze - Determine difference to target state



Act - Perform actions to drive the actual to the desired state

Observe - Analyze - Act



Common Triggers

- Labels
 - Indexed by backend
 - Suitable for selector-like functionality
 - Limitation on charset for names and values
- Annotations
 - No syntax restrictions
 - Not indexed
- ConfigMaps
 - Good for complex structured state declarations
 - Simple alternative to CustomResourceDefinitions

ConfigMap Watch Controller

```
namespace=${WATCH_NAMESPACE:-default}
base=http://localhost:8001
ns=namespaces/$namespace

curl -N -s $base/api/v1/$ns/configmaps?watch=true | \
while read -r event
do
    type=$(echo "$event" | jq -r '.type')
    config_map=$(echo "$event" | jq -r '.object.metadata.name' )
    annotations=$(echo "$event" | jq -r '.object.metadata.annotations' )

    if [ $type = "MODIFIED" ]; then
        # Restart Pods using this ConfigMap
        # ...
    fi
done
```



Operator

How to encapsulate operational knowledge into executable software

Definition

“” An **operator** is a Kubernetes **controller** that understands two domains: Kubernetes and *something else*. By combining knowledge of both areas, it can **automate** tasks that usually require a human operator that understands both domains.

Jimmy Zelinskie

<http://bit.ly/2Fjlx1h>

Technical:

Operator = Controller + CustomResourceDefinition

CustomResourceDefinition

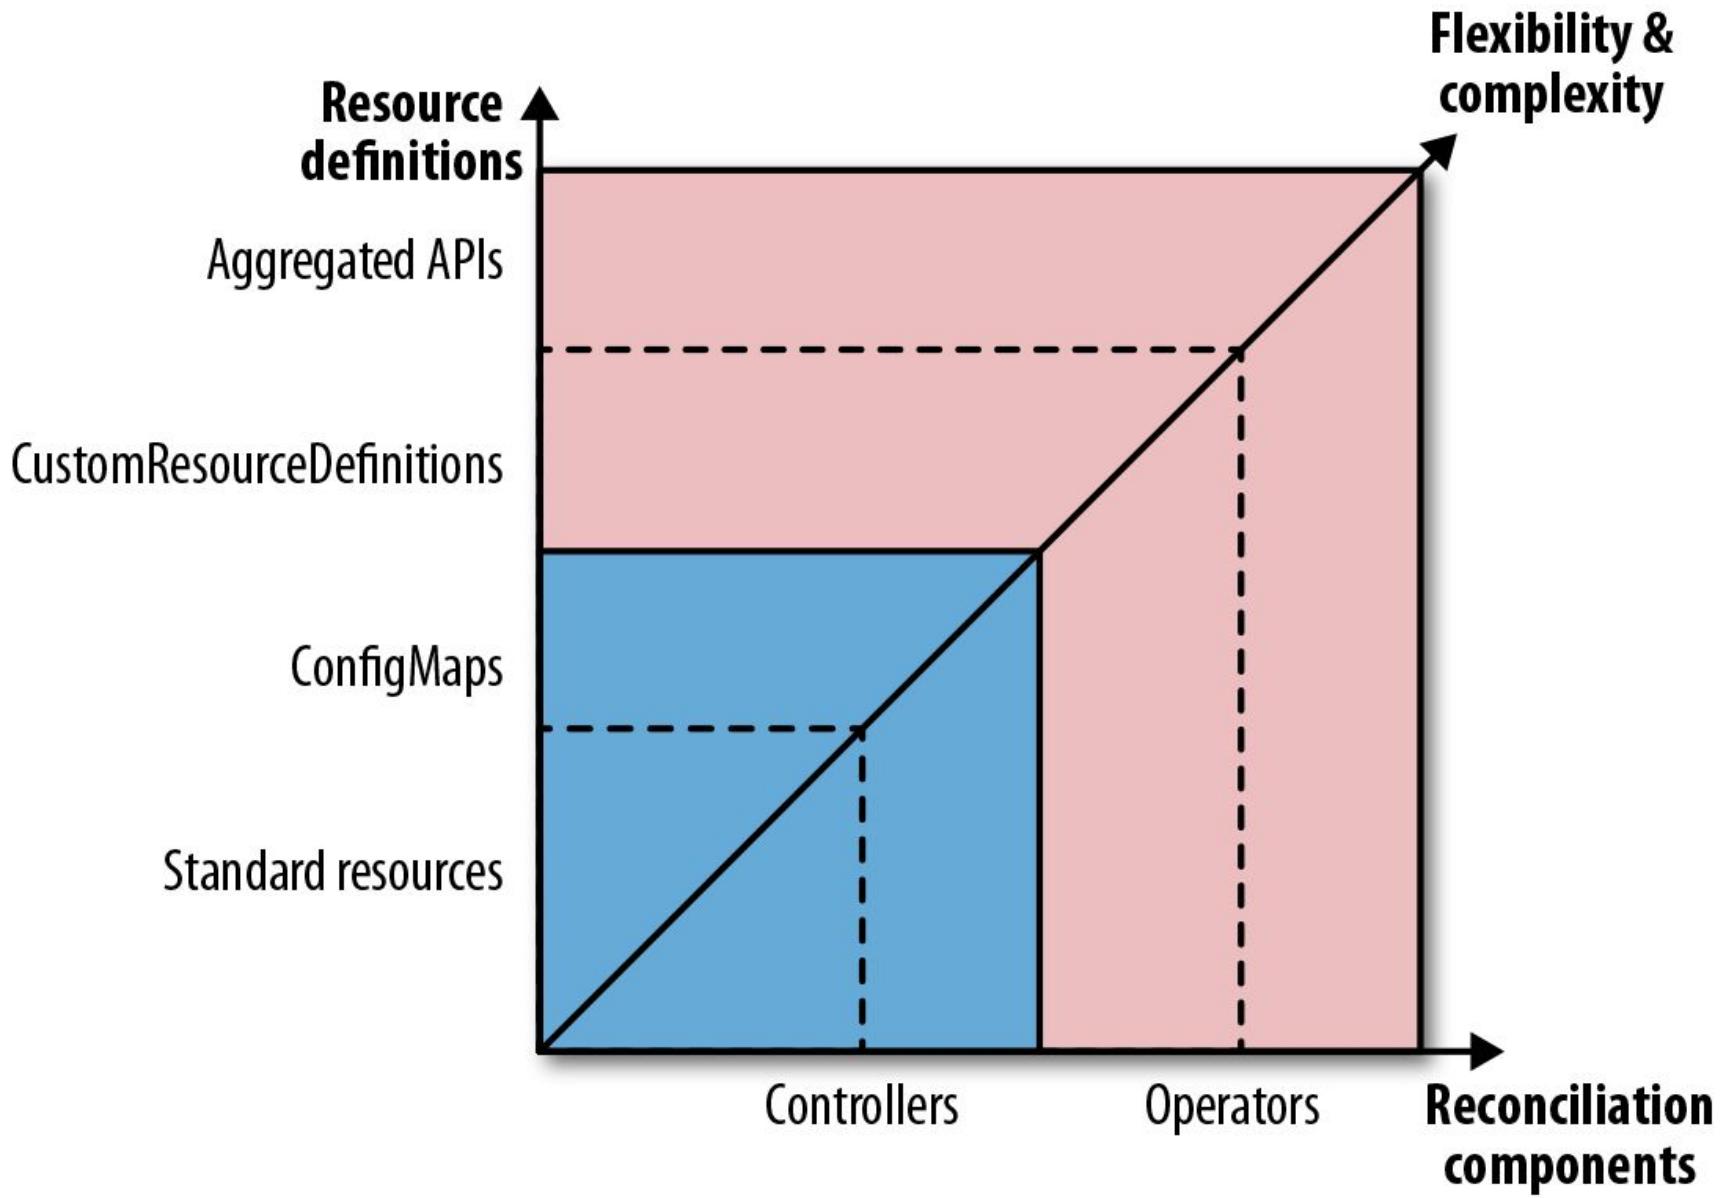
Custom resource is modelling a custom domain and managed through the Kubernetes API

```
apiVersion: apiextensions.k8s.io/v1beta1
kind: CustomResourceDefinition
metadata:
  name: configwatchers.k8spatterns.io
spec:
  scope: Namespaced
  group: k8spatterns.io
  version: v1
  names:
    kind: ConfigWatcher
    plural: configwatchers
  validation:
    openAPIV3Schema:
      ...

```

Custom Resource

```
kind: ConfigWatcher
apiVersion: k8spatterns.io/v1
metadata:
  name: webapp-config-watcher
spec:
  configMap: webapp-config
  podSelector:
    app: webapp
```



CRD Classification

- Installation CRDs
 - Installing and operating applications
 - Backup and Restore
 - Monitoring and self-healing
 - Example: Prometheus for installing Prometheus & components
- Application CRDs
 - Application specific domain concepts
 - Example: ServiceMonitor for registering Kubernetes service to be scraped by Prometheus

Operator Hub

The screenshot shows the OperatorHub.io website interface. At the top, there is a navigation bar with a search bar labeled "Search OperatorHub..." and a "Contribute" button. Below the header, a main banner reads "Welcome to OperatorHub.io" and "OperatorHub.io is a new home for the Kubernetes community to share Operators. Find an existing Operator or list your own today." On the left side, there is a sidebar with a "PROVIDER" section containing a list of providers with checkboxes, where "Jaeger" and "Red Hat" are checked. Below this is a "CAPABILITY LEVEL" section with options like "Basic Install", "Seamless Upgrades", and "Full Lifecycle". The main content area displays "5 ITEMS" found, sorted by provider. Each item has a thumbnail icon, the operator name, the provider, and a brief description. The operators shown are Jaeger Tracing, Kubernetes Federation, MongoDB, Prometheus Operator, and Strimzi Kafka.

Provider	Operator	Description
Jaeger	Jaeger Tracing	Provides tracing, monitoring and troubleshooting microservices-based
Red Hat	Kubernetes Federation	Gain Hybrid Cloud capabilities between your clusters with Kubernetes Federation.
Red Hat	MongoDB	The MongoDB Enterprise Kubernetes Operator enables easy deploys of MongoDB.
Red Hat	Prometheus Operator	The Prometheus Operator for Kubernetes provides easy monitoring definitions for
Red Hat	Strimzi Kafka	Run an Apache Kafka cluster, including Kafka Connect, ZooKeeper and more.

Operator Development

- Operator can be implemented in any language
- Frameworks:
 - Operator Framework (Golang, Helm, Ansible)
<https://github.com/operator-framework>
 - Kubebuilder (Golang)
<https://github.com/kubernetes-sigs/kubebuilder>
 - Metacontroller (Language agnostic)
<https://metacontroller.app/>
 - jvm-operators (Java, Groovy, Kotlin, ...)
<https://github.com/jvm-operators>

Demo

https://github.com/k8spatterns/examples Example X +

README.adoc

Kubernetes Patterns - Examples



This GitHub project contains the examples from *Kubernetes Patterns - Reusable Elements for Designing Cloud-Native Applications* book by Bilgin Ibryam and Roland Huß

Installation instructions for the example prerequisites are summarised in [INSTALL](#). By default, you need access to a vanilla Kubernetes installation, like Minikube. If addons are required, the example description explains this additional requirement.

For feedback, issues or questions in general, please use the [issue tracker](#) to open issues. Also, we love contributions like spelling fixes, bug fixes, improvements, ... Please open Pull Requests, we are happy to review them!

Patterns

All example are categorised according to the Book's patterns category. Each of the examples is contained in an extra directory per pattern and is self-contained. [1]

Foundational Patterns

Predictable Demands

Our sample random generator dealing with hard requirements on ConfigMap and PersistentVolumeClaims as well as with resource limits.

Declarative Deployment

Rolling and fixed update of the random generator Deployment from version 1.0 to 2.0.

Health Probe

Liveness and Readiness probes for the random generator.

O'REILLY
Kubernetes Patterns
Reusable Elements for Designing Cloud-Native Applications

Bilgin Ibryam & Roland Huß

<https://github.com/k8spatterns/examples>



Thank you



<https://k8spatterns.io>



@ro14nd



@bibryam



@k8spatterns

Picture Credits

<https://www.pexels.com/photo/brown-and-black-pattern-2158386/>

<https://pixabay.com/photos/ship-helm-sunset-cutter-coast-guard-759954/>

<https://unsplash.com/photos/yo01Z-9HQAw>

<https://www.pexels.com/photo/turned-on-light-crane-1117211/>

<https://www.pexels.com/photo/shallow-focus-photography-of-black-ship-1095814/>

<https://pixabay.com/photos/containers-storage-rusted-rusty-1209079/>

<https://pixabay.com/photos/motocross-sidecar-race-motorsport-1045661/>

<https://www.pexels.com/photo/golden-gate-bridge-san-francisco-california-1141853/>

https://unsplash.com/photos/M_I-crkDO-k

<https://www.pexels.com/photo/reflection-playstation-pad-gaming-18174/>

https://unsplash.com/photos/UJP_QpCKj7M

<https://www.pexels.com/photo/grayscale-photo-of-person-holding-chess-piece-1498958/>

<https://pixabay.com/photos/cans-manufacturing-business-2888650/>

<https://unsplash.com/photos/lRoX0shwjUQ>

<https://www.freeimages.com/photo/poppy-in-wheat-1344010>

<https://pixabay.com/photos/telescope-field-glass-spyglass-1966366/>

<https://unsplash.com/photos/UHDx3BHIFvY>

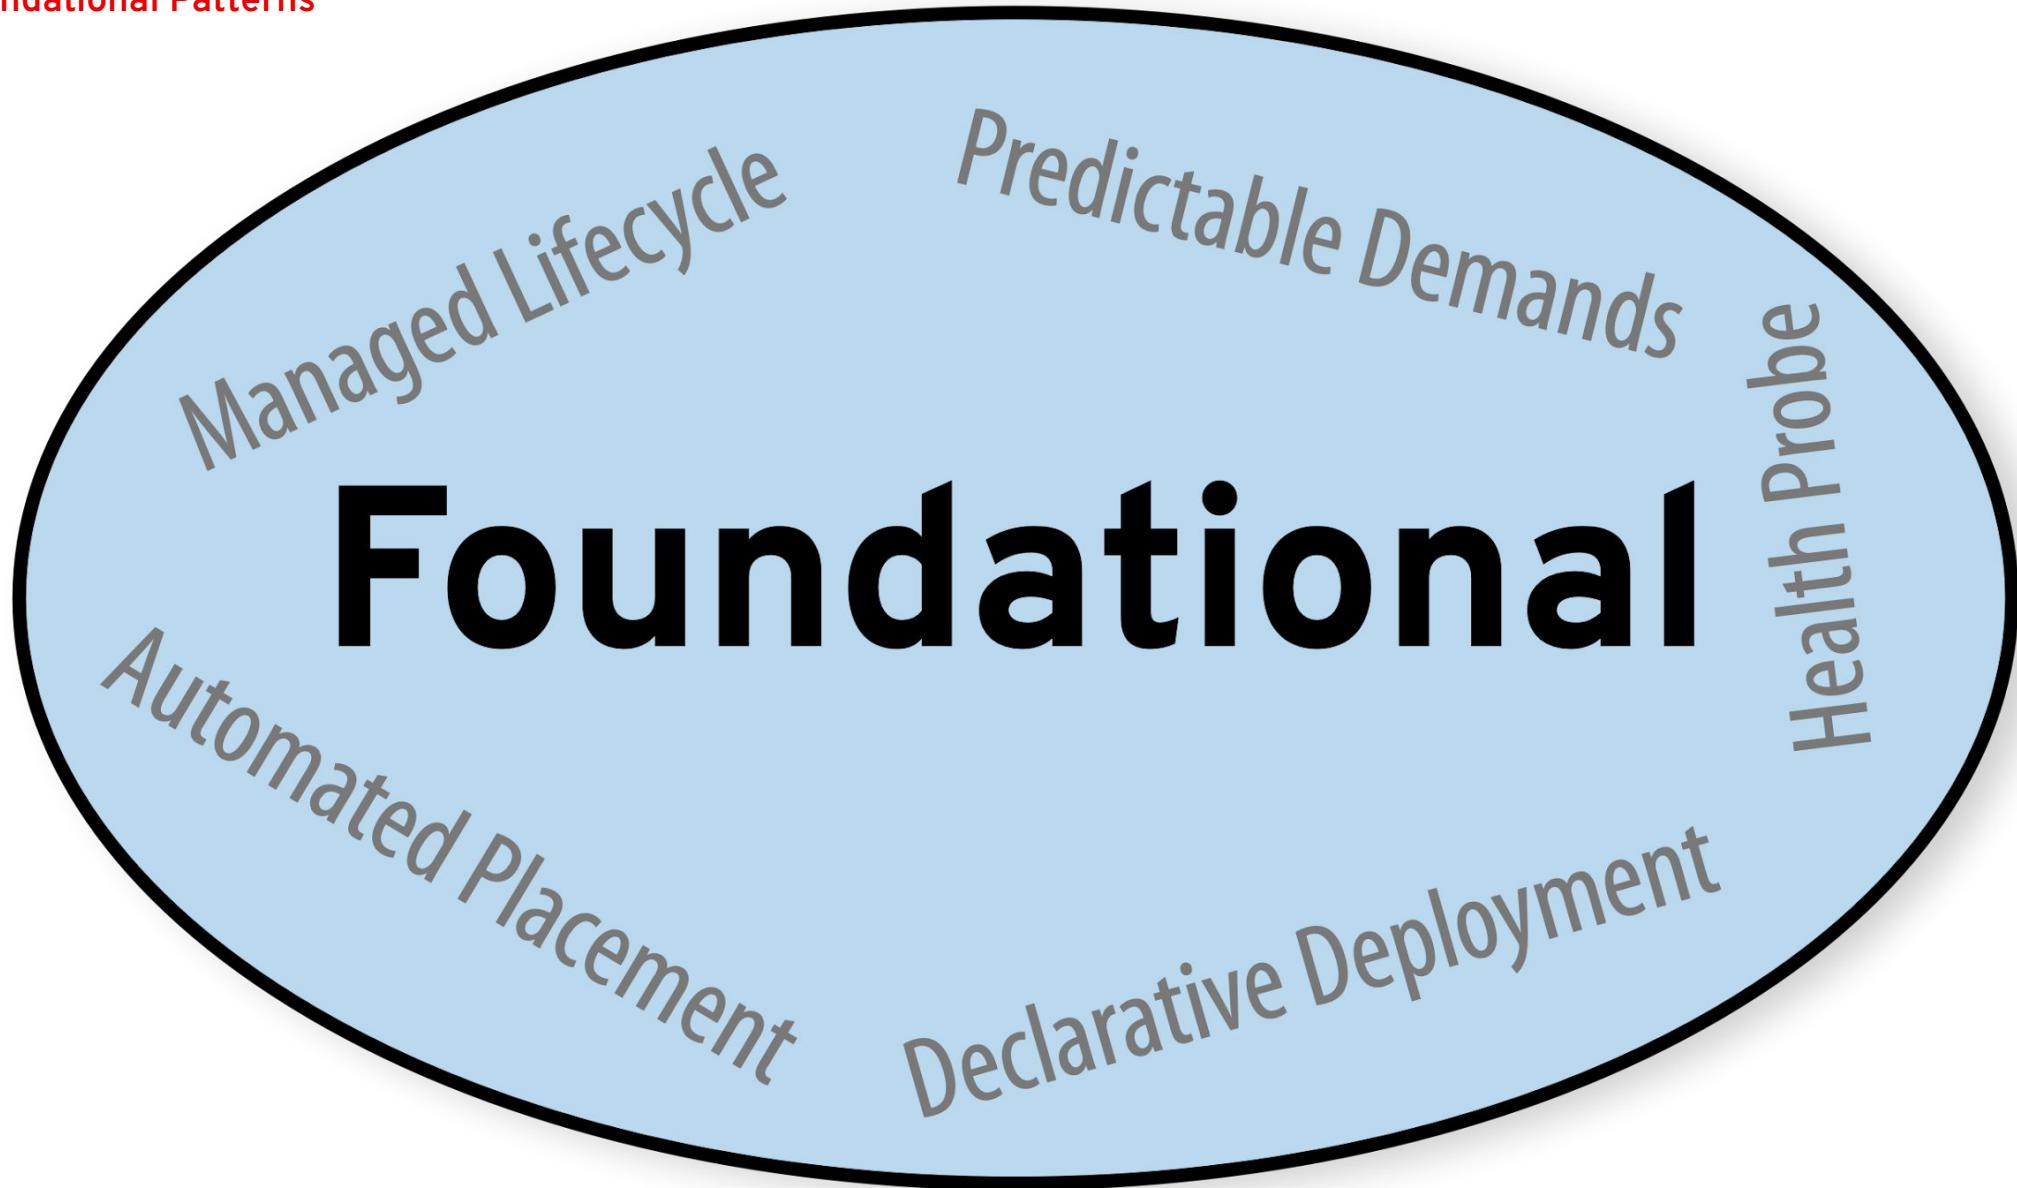
<https://pixabay.com/photos/bake-advent-christmas-cookie-1786926/>

<https://pixabay.com/photos/balloons-colors-party-celebration-1869790/>

https://unsplash.com/photos/ymf4_9Y9S_A

<https://unsplash.com/photos/ZUabNmumOcA>

<https://pixabay.com/photos/lost-places-machines-old-factory-3991951/>





Predictable Demands

How to declare application requirements

Application Requirements

- Declared requirements help in
 - Scheduling decisions
 - Capacity planning
 - Matching infrastructure services
- Hard runtime dependencies
 - Persistent Volumes
 - Host ports
 - Dependencies on ConfigMaps and Secrets

Resource Profile

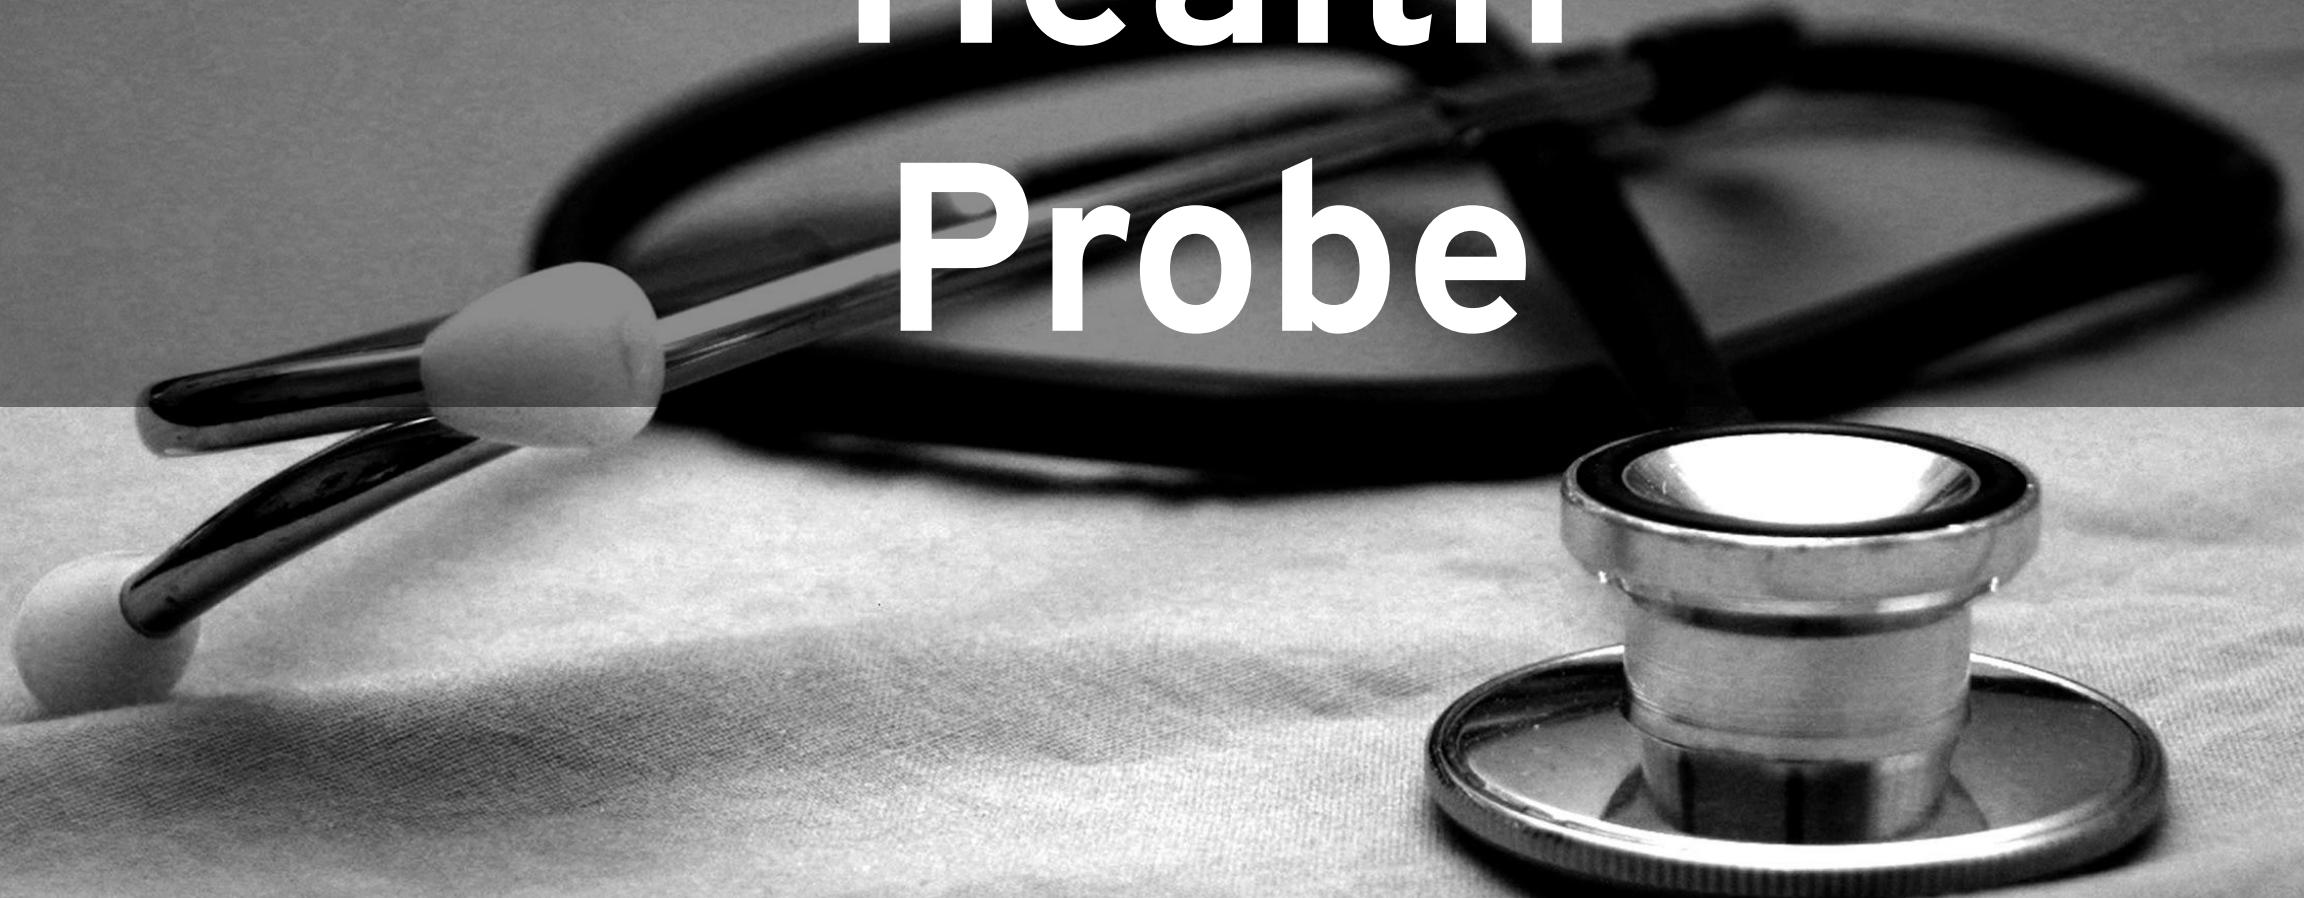
```
apiVersion: v1
kind: Pod
metadata:
  name: http-server
spec:
  containers:
  - image: nginx
    name: nginx
  resources:
    requests:
      cpu: 200m
      memory: 100Mi
    limits:
      cpu: 300m
      memory: 200Mi
```

Quality-of-Service Classes

- **Best Effort**
 - No requests or limits
- **Burstable**
 - requests < limits
- **Guaranteed**
 - requests == limits

Exercise

Health Probe



How to communicate an application's health state to Kubernetes

Monitoring Health

- Process health checks
 - Checks running process
 - Restarts container if container process died
- Application provided Health Probes
 - **Liveness Probe:** Check application health
 - **Readiness Probe:** Check readiness to process requests

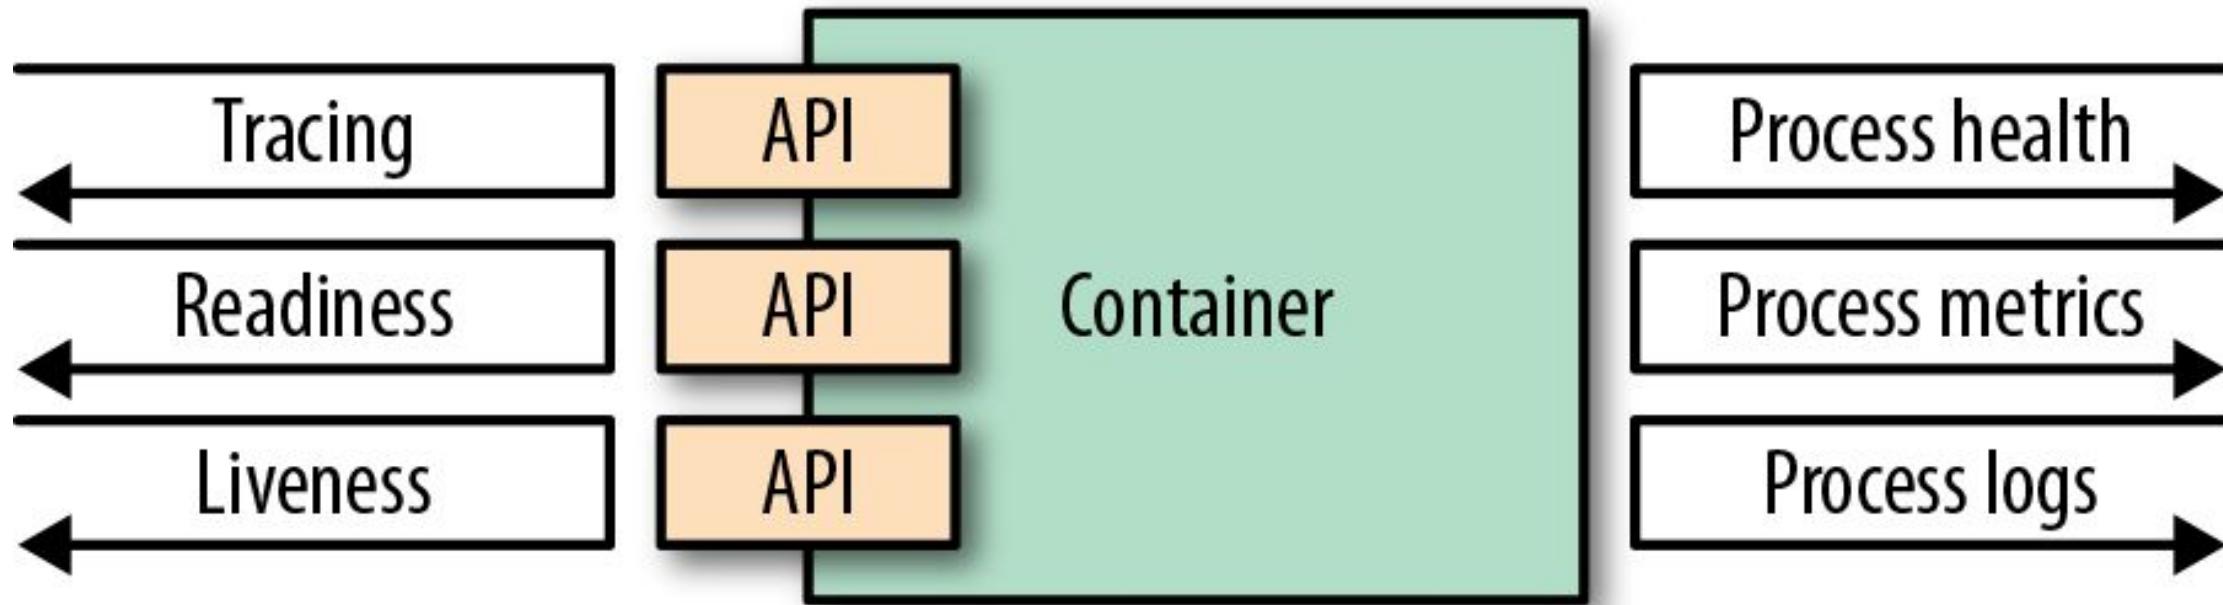
Liveness & Readiness

- Liveness Probe
 - Restarting containers if liveness probes fail
- Readiness Probe
 - Removing from service endpoint if readiness probe fails
- Probe methods
 - HTTP endpoint
 - TCP socket endpoint
 - Unix command's return value

Health Probe

```
apiVersion: v1
kind: Pod
metadata:
  name: pod-with-readiness-check
spec:
  containers:
  - image: k8spatterns/random-generator:1.0
    name: random-generator
    livenessProbe:
      httpGet:
        path: /actuator/health
        port: 8080
      initialDelaySeconds: 30
    readinessProbe:
      exec:
        command: [ "stat", "/var/run/random-generator-ready" ]
```

Container Observability Options



Exercise



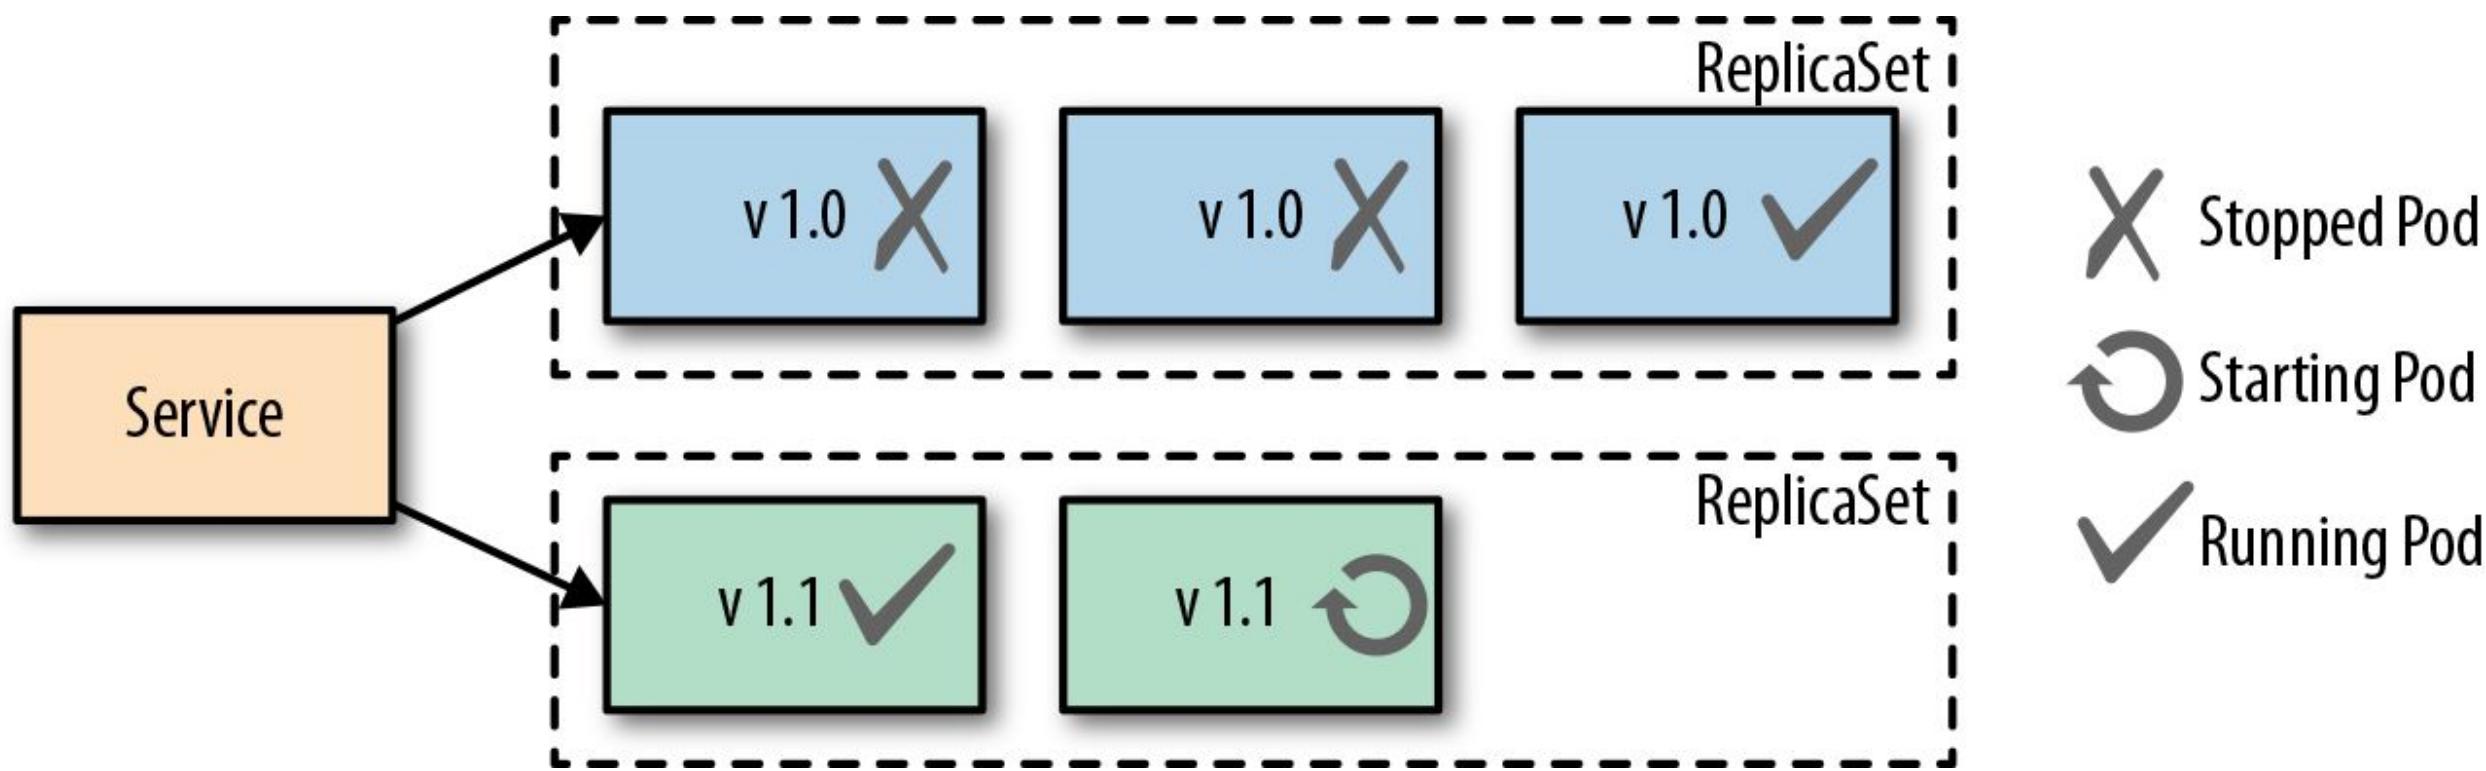
Declarative Deployment

How to perform
application installations
and upgrades by
configuration.

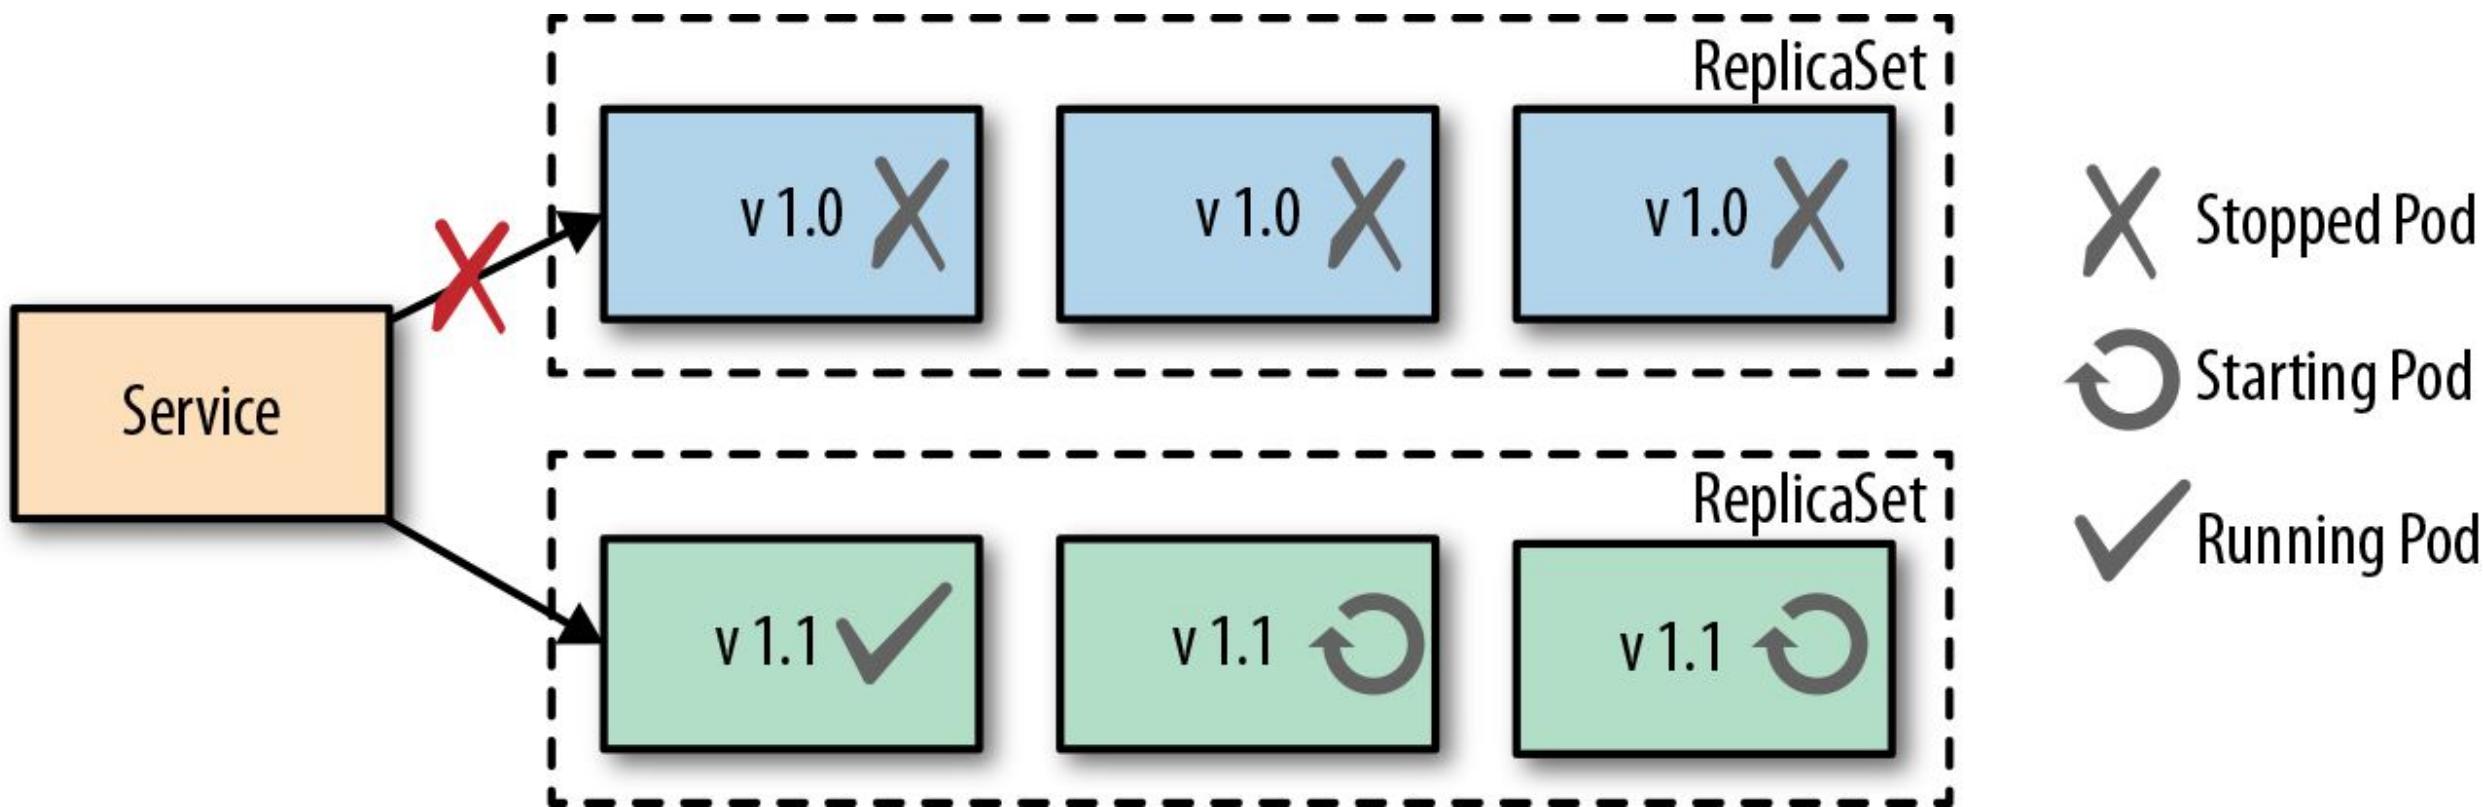
Deployment

- **Declarative** vs. **Imperative** deployment
- Deployment Kubernetes Resource:
 - Holds template for Pod
 - Creates ReplicaSet on the fly
 - Allows rollback
 - Update strategies are declarable

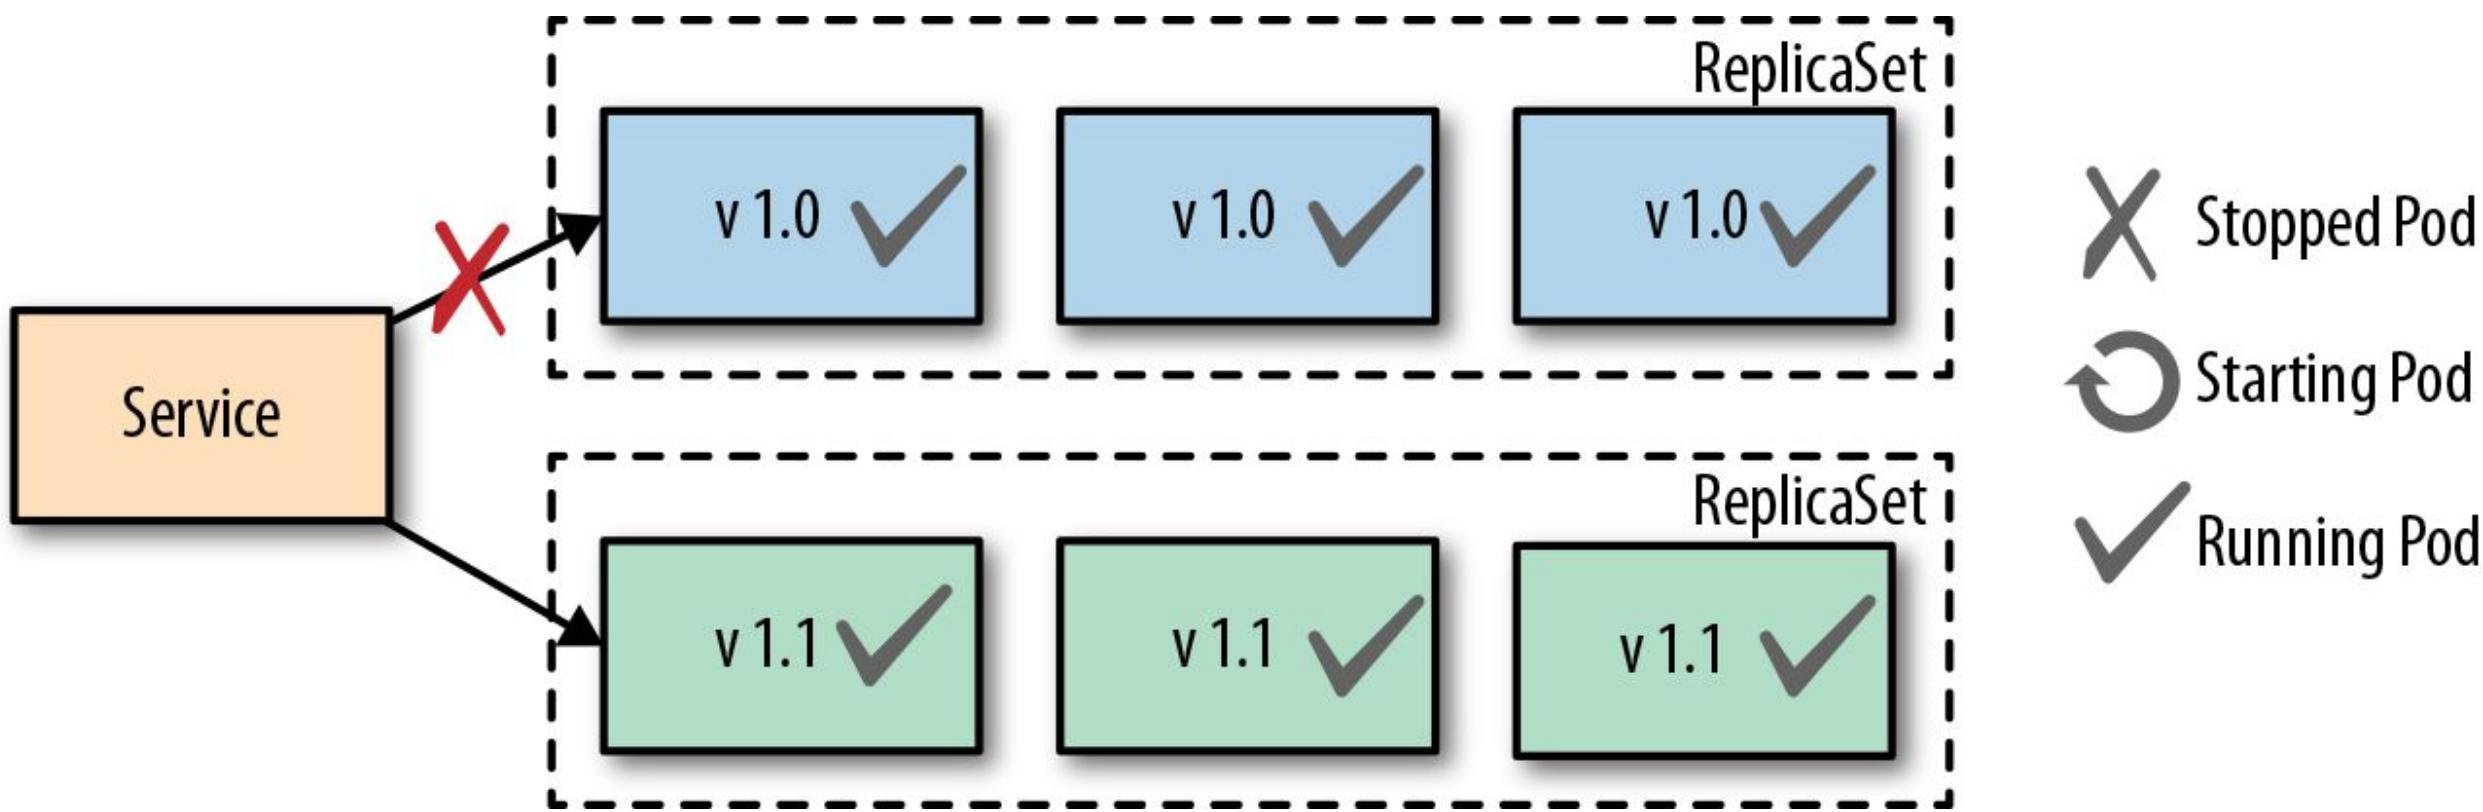
Rolling Deployment



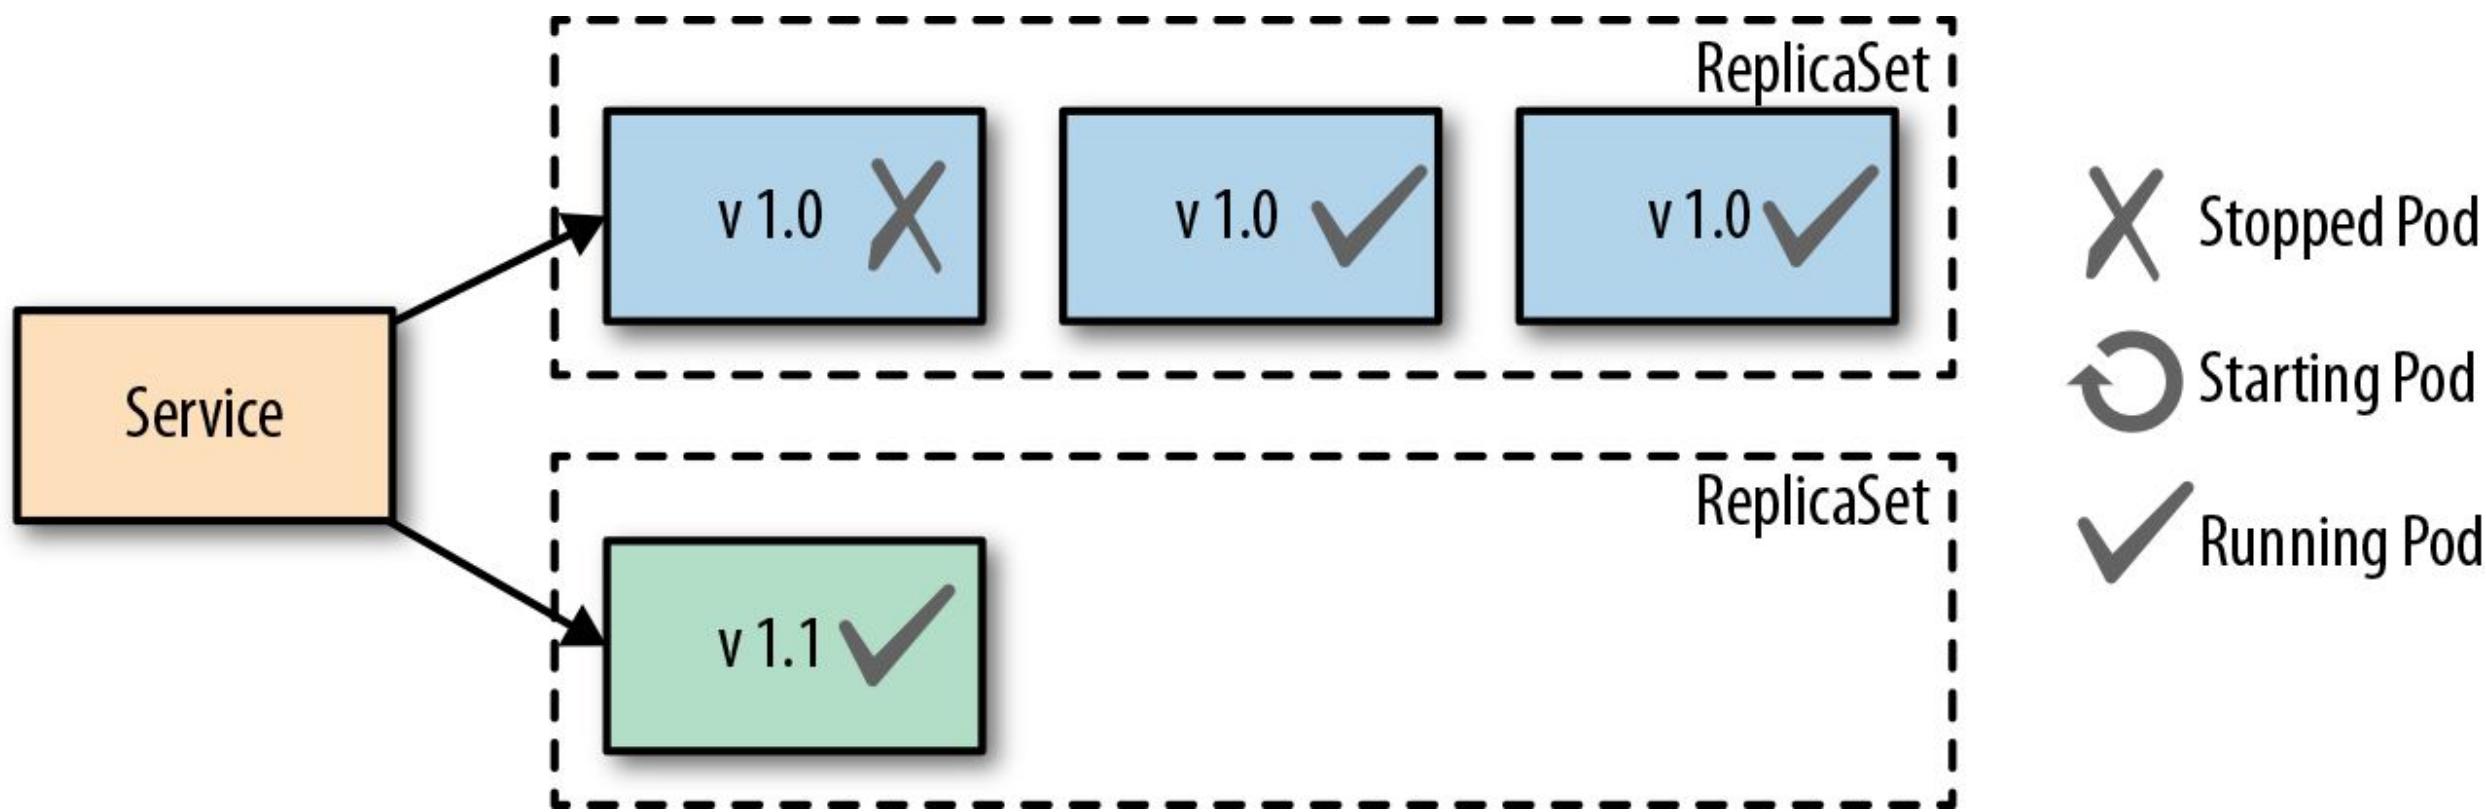
Fixed Deployment



Blue-Green Release

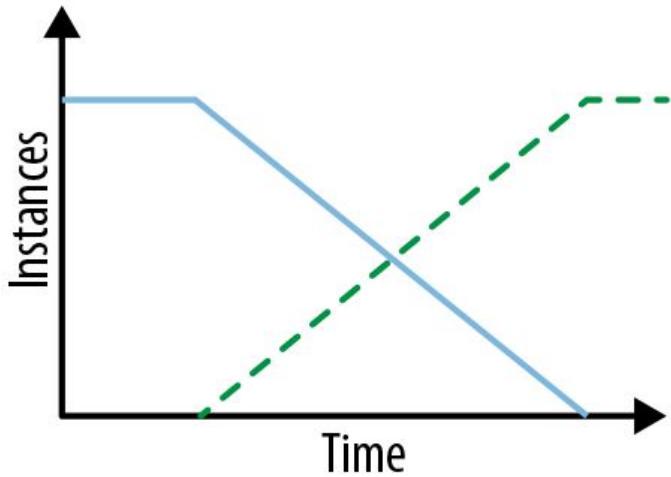


Canary Release

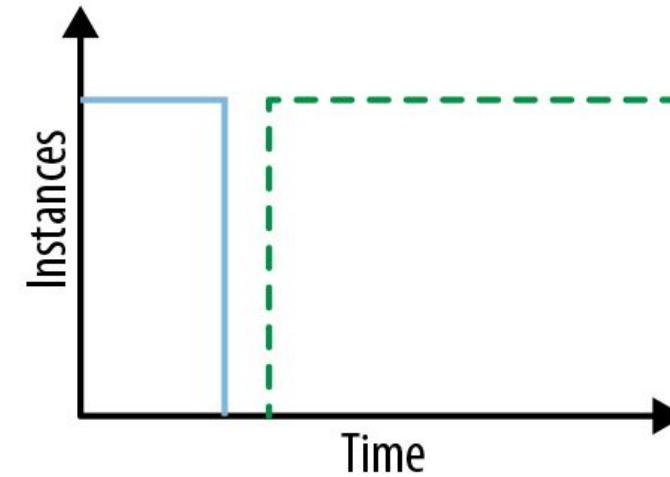


Declarative Deployment

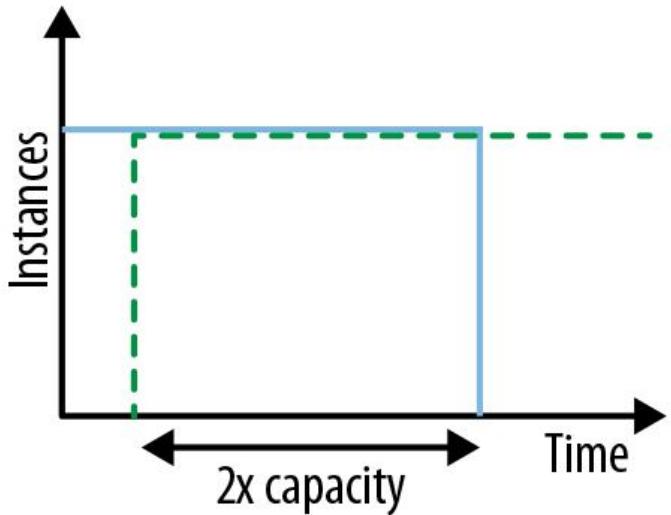
Rolling deployment



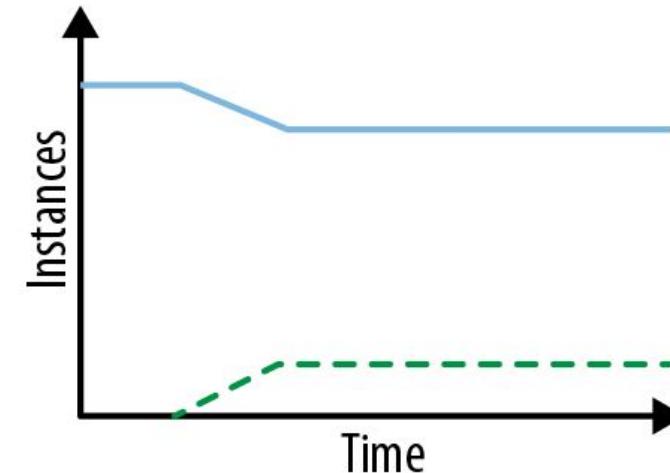
Fixed deployment



Blue-green release



Canary release



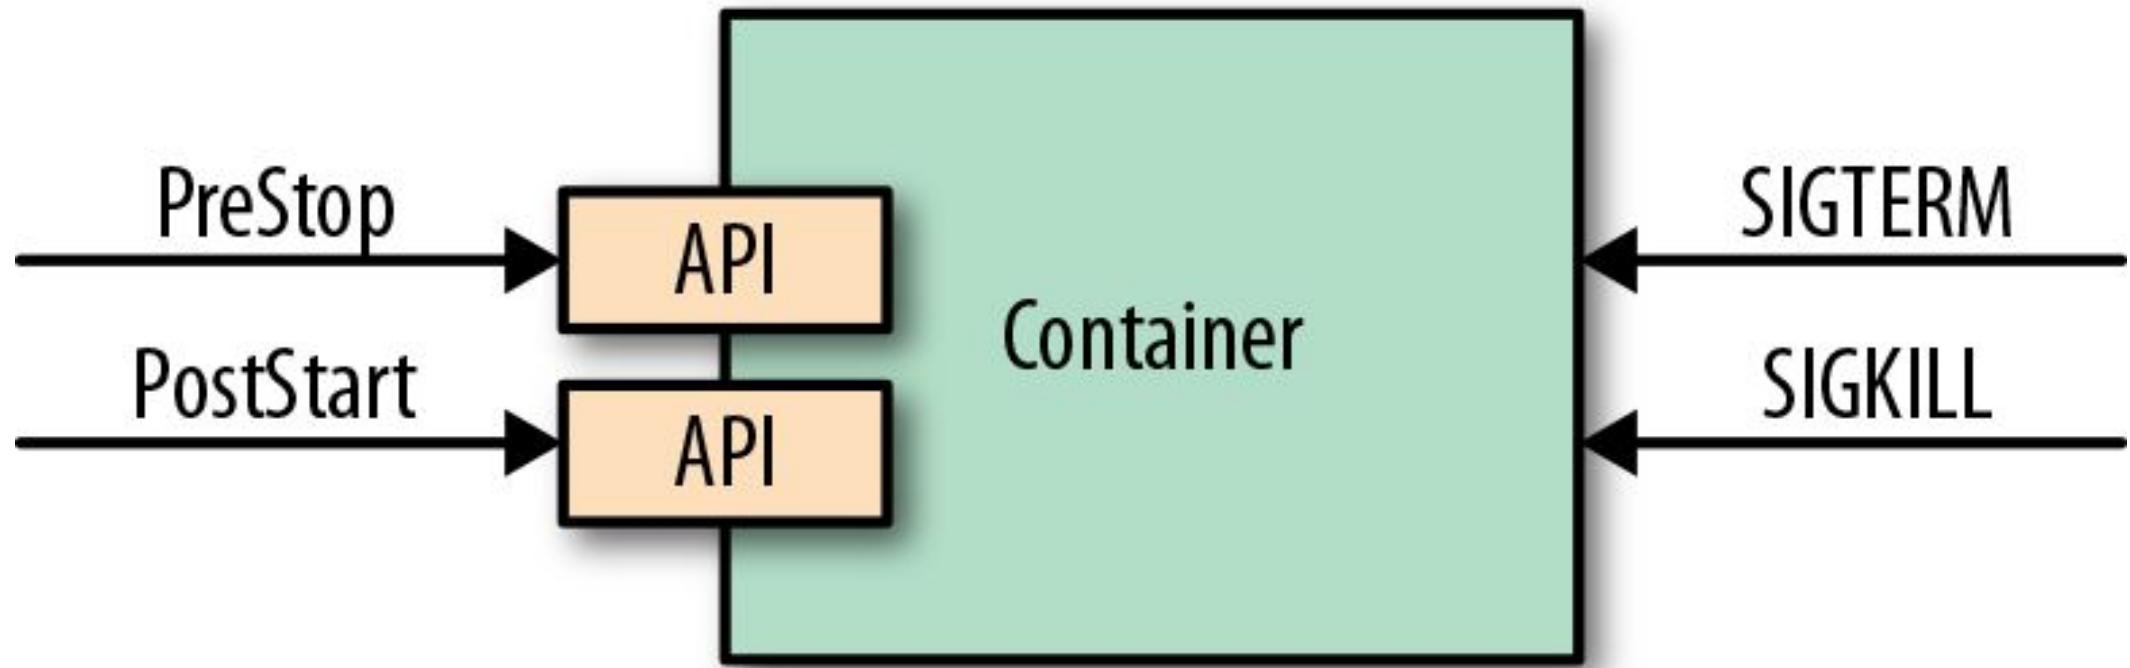
Exercise

Managed Lifecycle



How applications should react on lifecycle events

Managed Container Lifecycle



Lifecycle Events

- SIGTERM
 - Initial event issued when a container is going to shutdown
 - Application should listen to this event to cleanup properly and then exit
- SIGKILL
 - Final signal sent after a grace period which can't be catched
 - terminationGracePeriodSeconds: Time to wait after SIGTERM (default: 30s)

Lifecycle Hooks

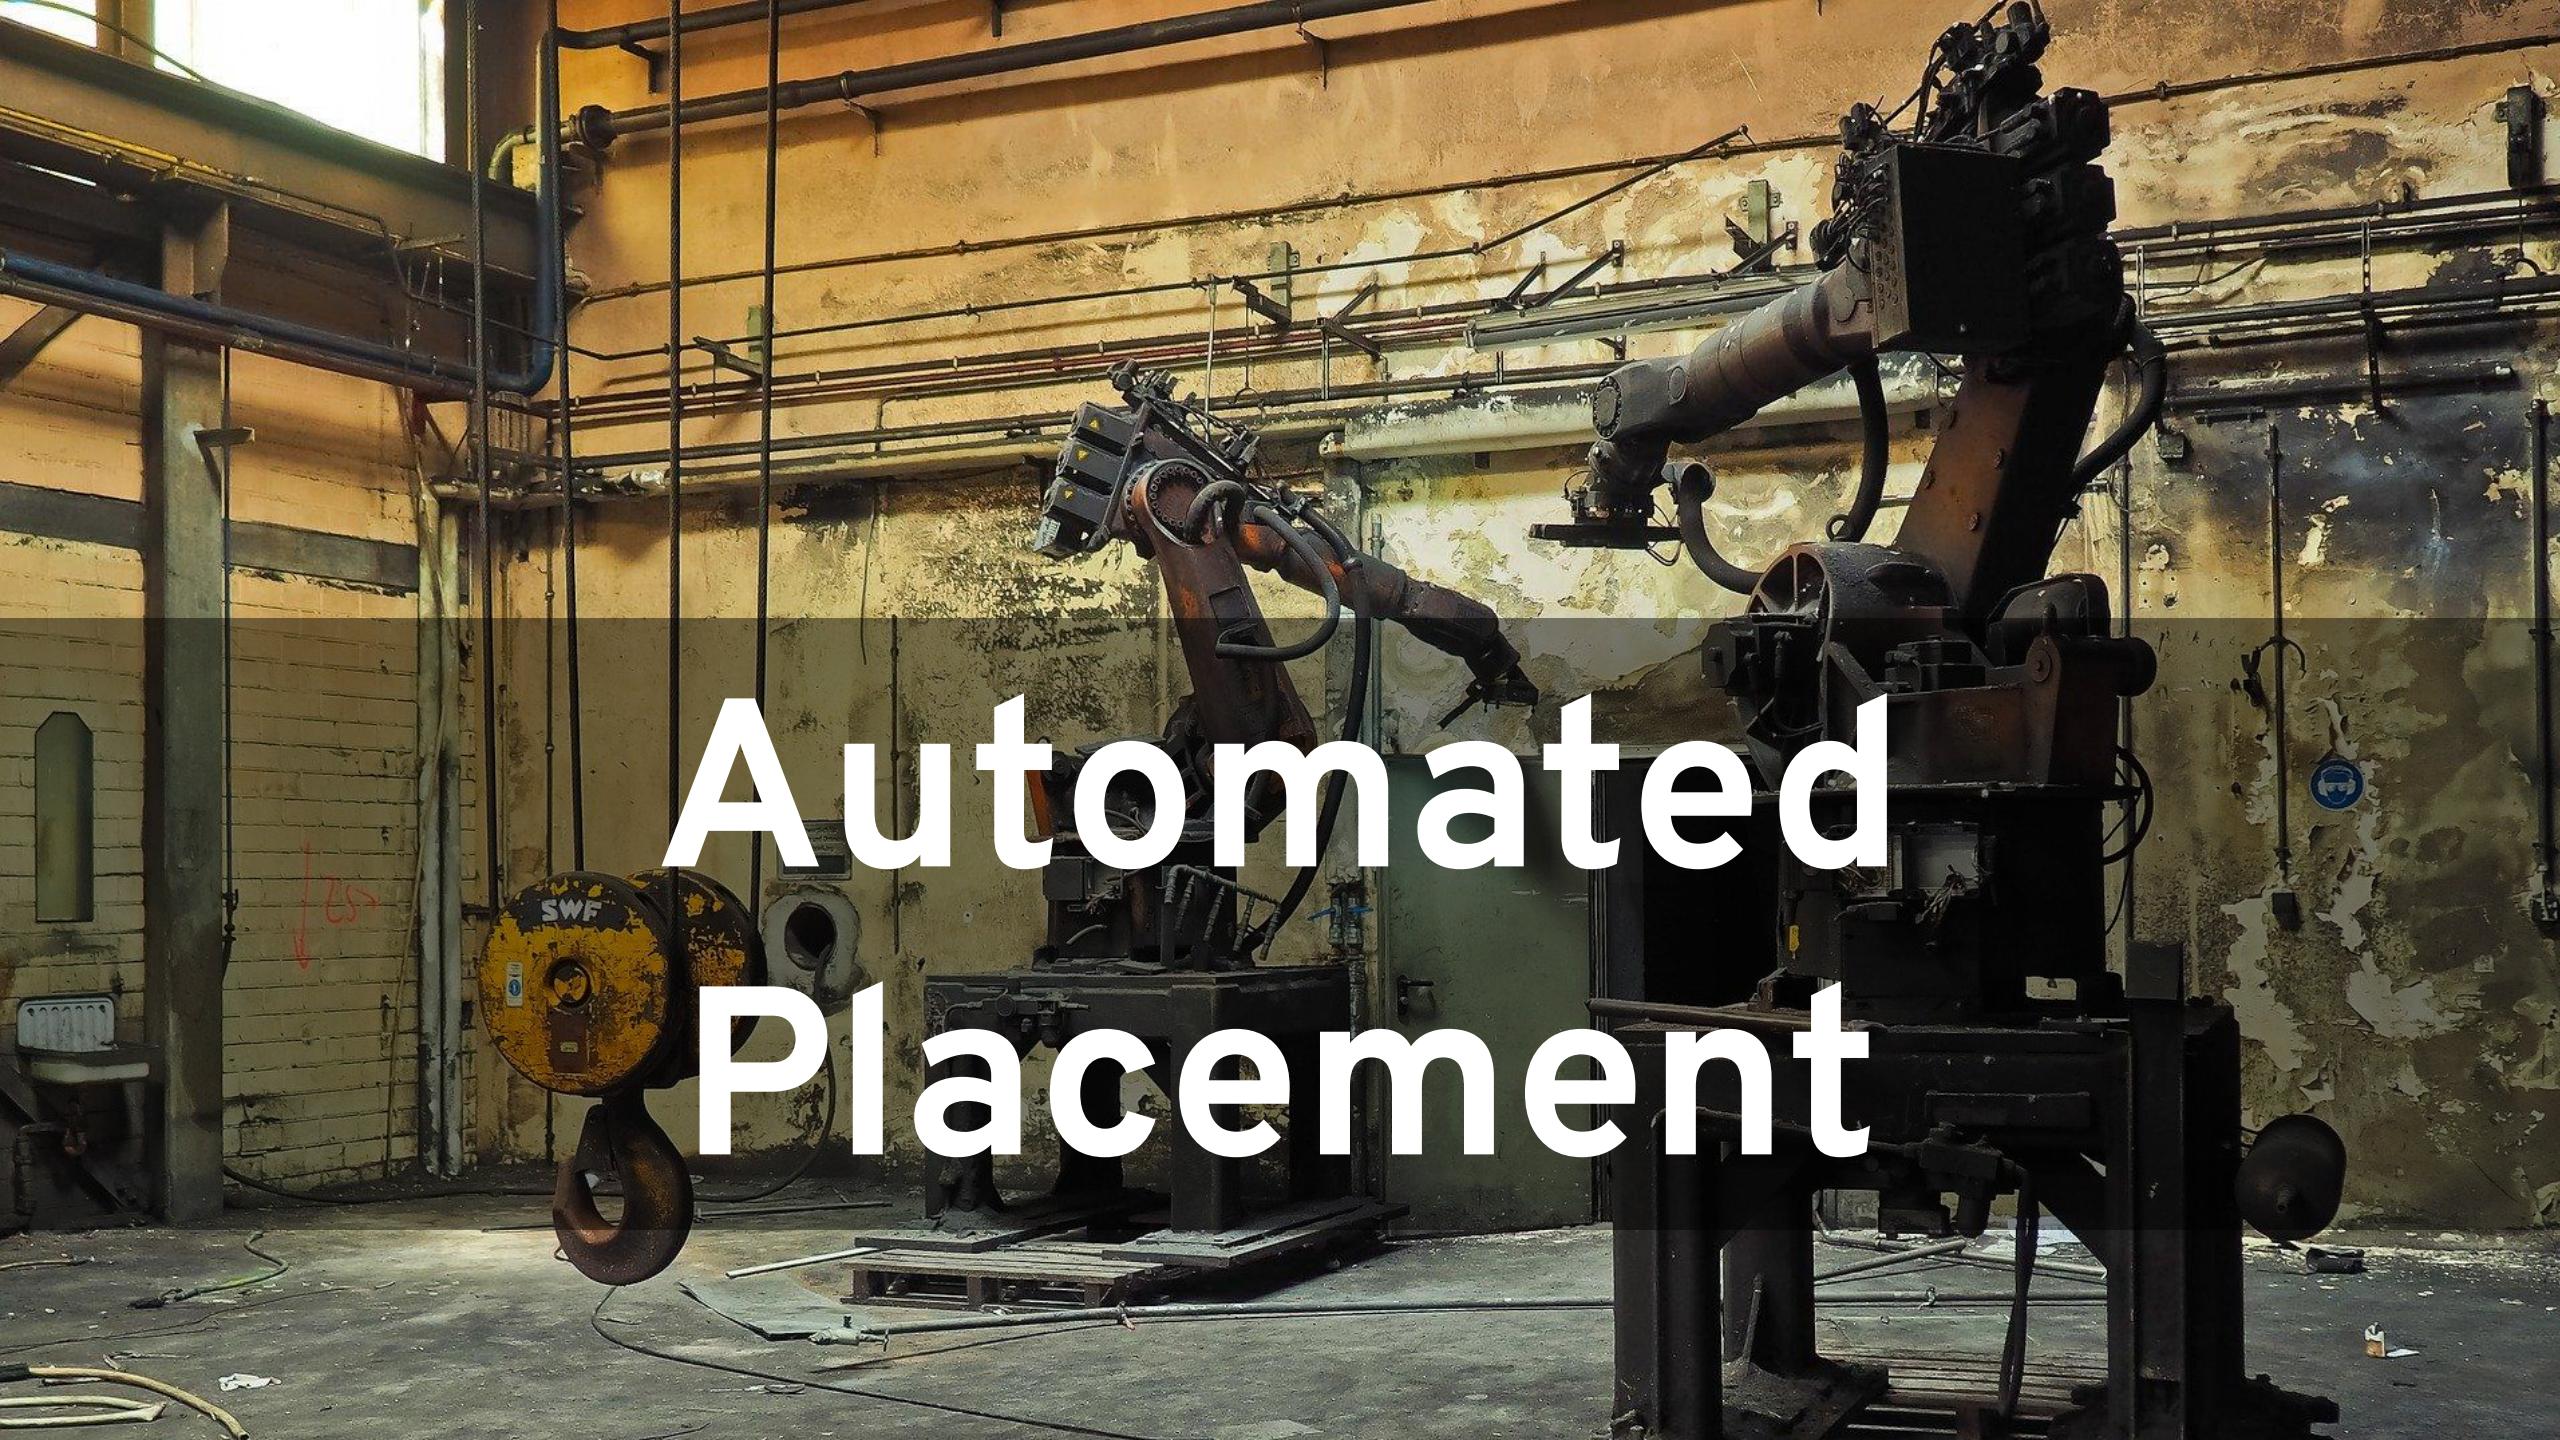
- **postStart**
 - Called after container is created
 - Runs in parallel to the main container
 - Keeps Pod in status *Pending* until exited successfully
 - **exec** or **httpGet** handler types (like *Health Probe*)
- **preStop**
 - Called before container is stopped
 - Same purpose & semantics as for SIGTERM

Managed Lifecycle

```
apiVersion: v1
kind: Pod
metadata:
  name: pre-stop-hook
spec:
  containers:
  - image: k8spatterns/random-generator:1.0
    name: random-generator
    lifecycle:
      postStart:
        exec:
          command:
          - sh
          - -c
          - sleep 30 && echo "Wake up!" > /tmp/postStart_done
      preStop:
        httpGet:
          port: 8080
          path: /shutdown
```

Exercise

Automated Placement



How Kubernetes schedules containers

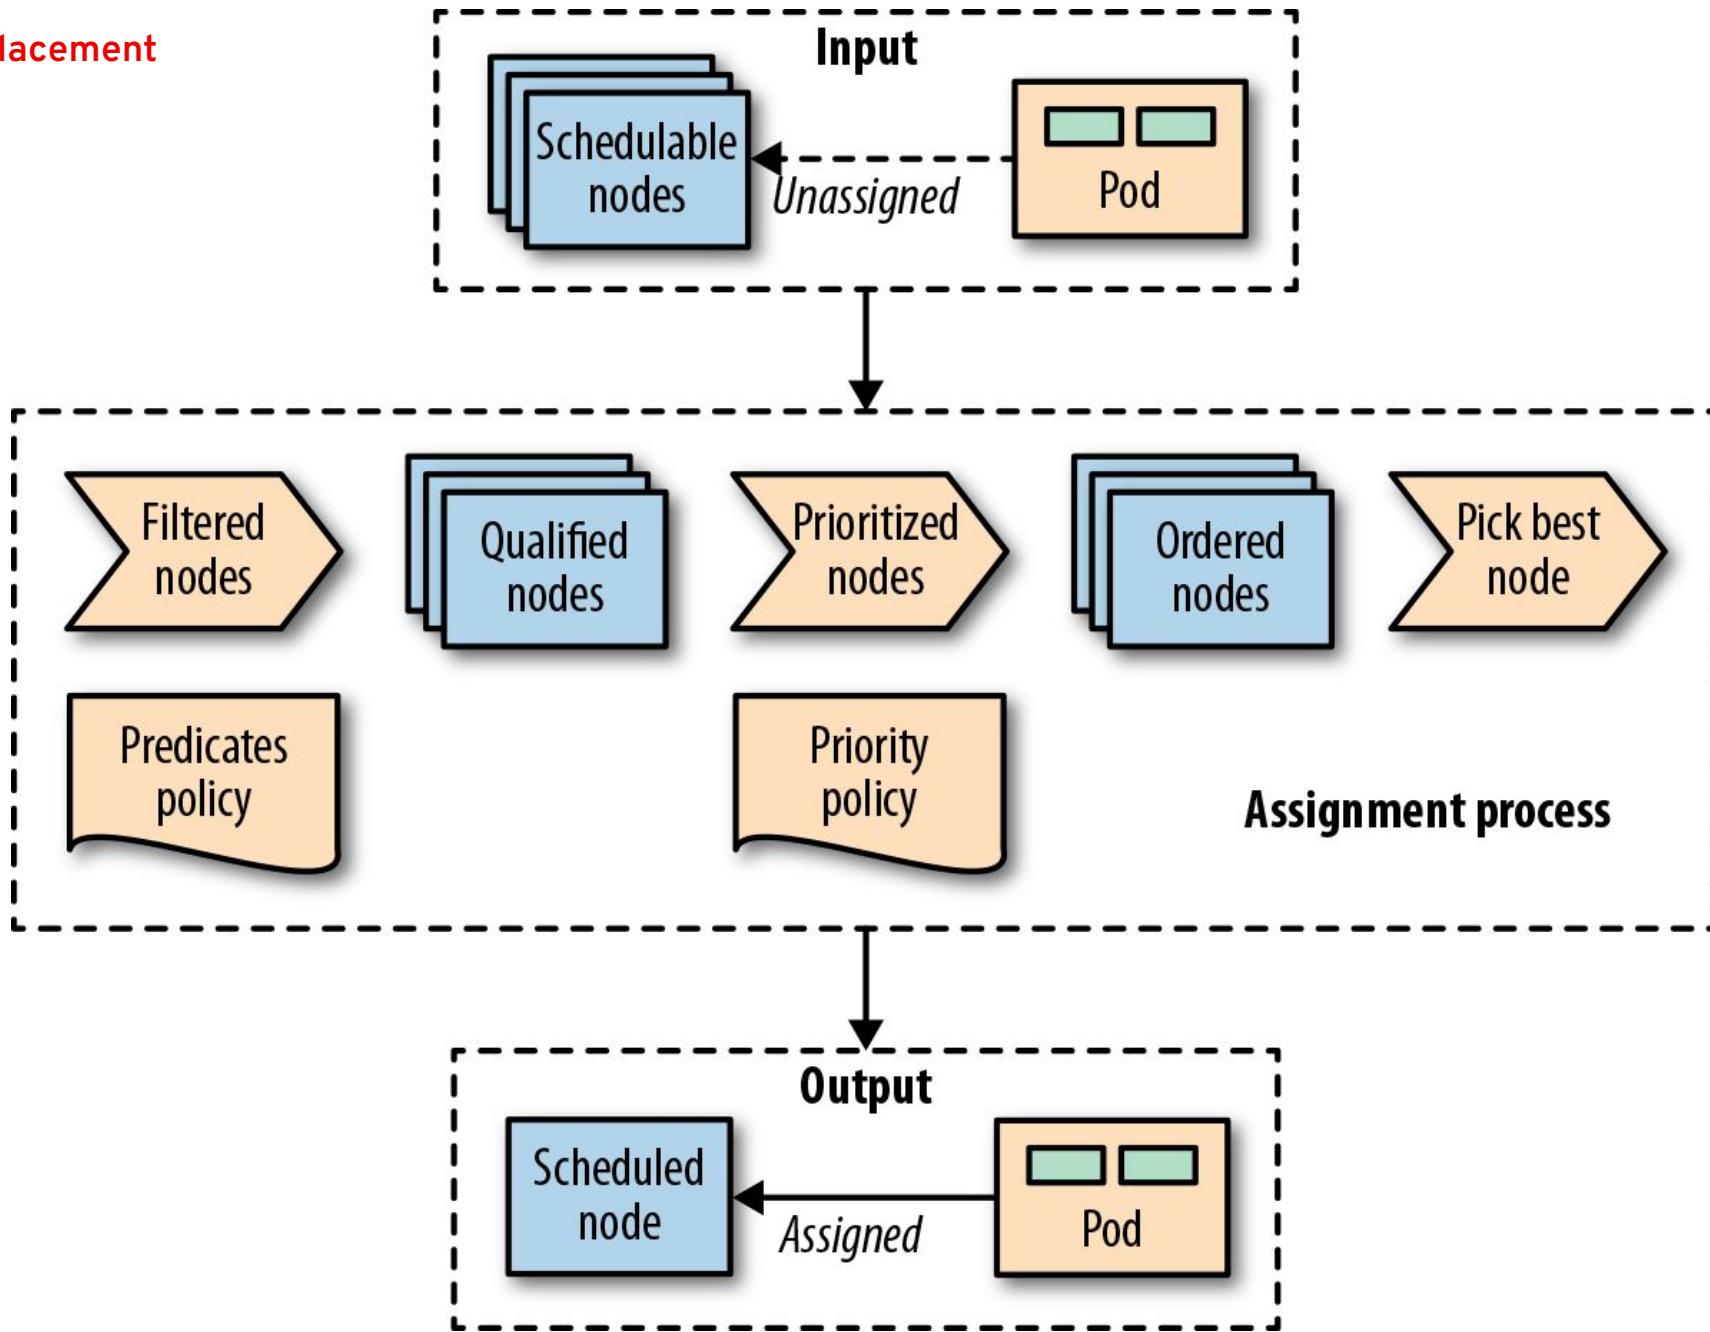
Scheduler

- Responsible for assigning Pods to nodes
- **Predicates:** Filter of suitable nodes
- **Priorities:** Ordering of nodes according to preference
- Predicates + Priorities = Scheduler Policy
- Multiple schedulers with different policies can be defined
- Pod can overwrite with `.spec.schedulerName`

Scheduler Policy

```
{  
  "kind": "Policy",  
  "apiVersion": "v1",  
  "predicates": [  
    {  
      "name": "PodFitsHostPorts"  
    },  
    {  
      "name": "PodFitsResources"  
    },  
    ...  
  ],  
  "priorities": [  
    {  
      "name": "LeastRequestedPriority",  
      "weight": 2  
    },  
    ...  
  ]  
}
```

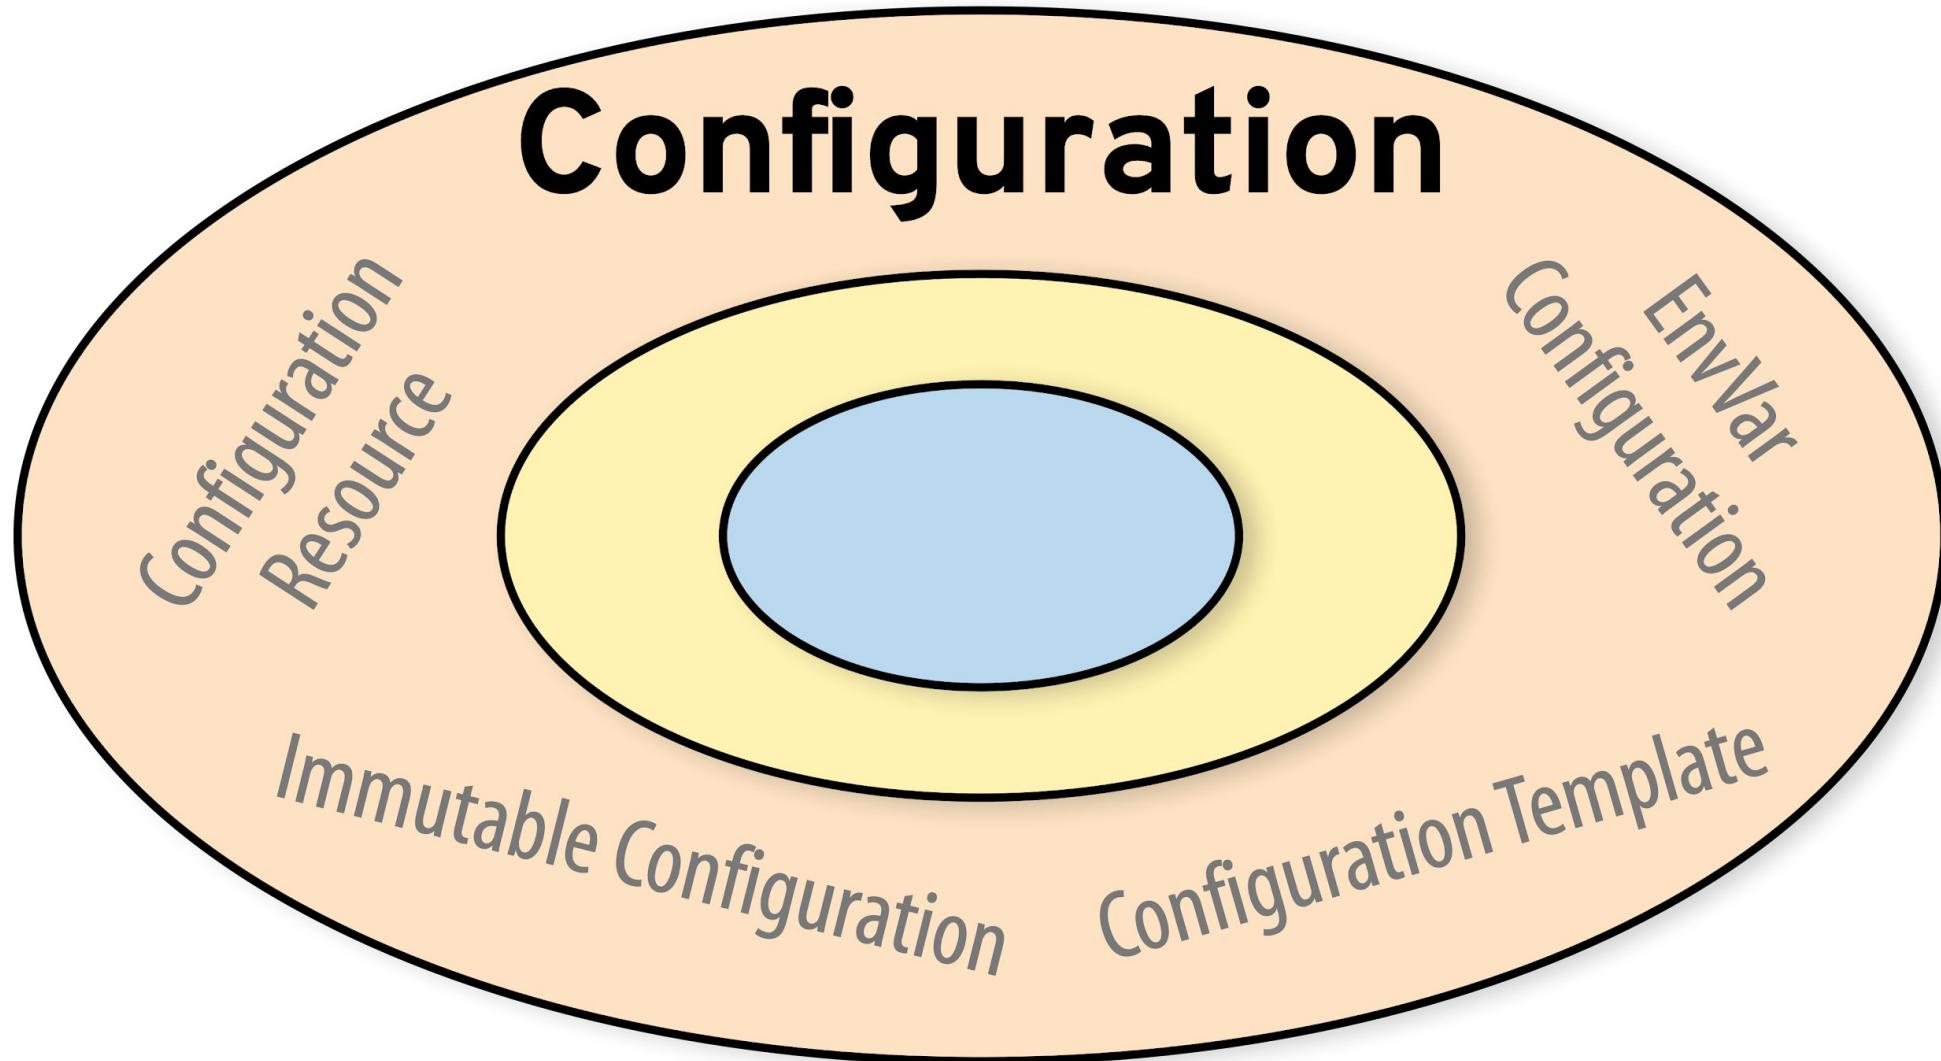
Automated Placement



How to influence the Scheduler

- General rule: Let the schedule do its job
- **Node Selector:** Label selector to pick a node
- **Node Affinity:** Required and preferred scheduling rules on nodes (label and field matching)
- **Pod Affinity and Antiaffinity:** Specify rules how pods are placed with respect to other Pods
- **Taints and Tolerations:** Lets node specify which pods might be scheduled on them (opposite of node selector)

Exercise



EnvVar Configuration

How to use environment variables to configure your applications

Environment Variables

- Simplest way to externalize configuration for containerized applications
- 12factor.net: Rule 3 - “Store config in the environment”
- When specified in a Pod:
 - Literally in the container spec
 - Imported from a ConfigMap or Secret
- Limitations
 - Immutable for a running process
 - Not suitable for a large configuration set

Exercise

A close-up, slightly blurred photograph of a foosball table. The table has a green playing surface with white lines. There are two sets of players: one set in red shirts and another in blue shirts. The blue team is positioned at the bottom of the frame, while the red team is at the top. The foosball rods are visible across the middle of the table.

Configuration Resource

How to manage configuration with Kubernetes resources

ConfigMap

- Key-Value Map
- Use in Pods as:
 - environment variables
 - volumes with keys as file names and values as file content

```
kubectl create cm spring-boot-config \
--from-literal=JAVA_OPTIONS=-Djava.security.egd=file:/dev/urandom \
--from-file=application.properties
```

Secret

- Like ConfigMap but content Base64 encoded
- Secrets are ...
 - ... only distributed to nodes running Pods that need it
 - ... only stored in memory in a tmpfs and never written to physical storage
 - ... stored encrypted in the backend store (etcd)
- Access can be restricted with RBAC rules
- But: For high security requirements application based encryption is needed

Exercise

A large, rugged, reddish-brown rock formation rises from the ocean under a clear blue sky. The rock has a distinct, rounded peak on the left and a more jagged, craggy base on the right. The water in the foreground is a deep blue.

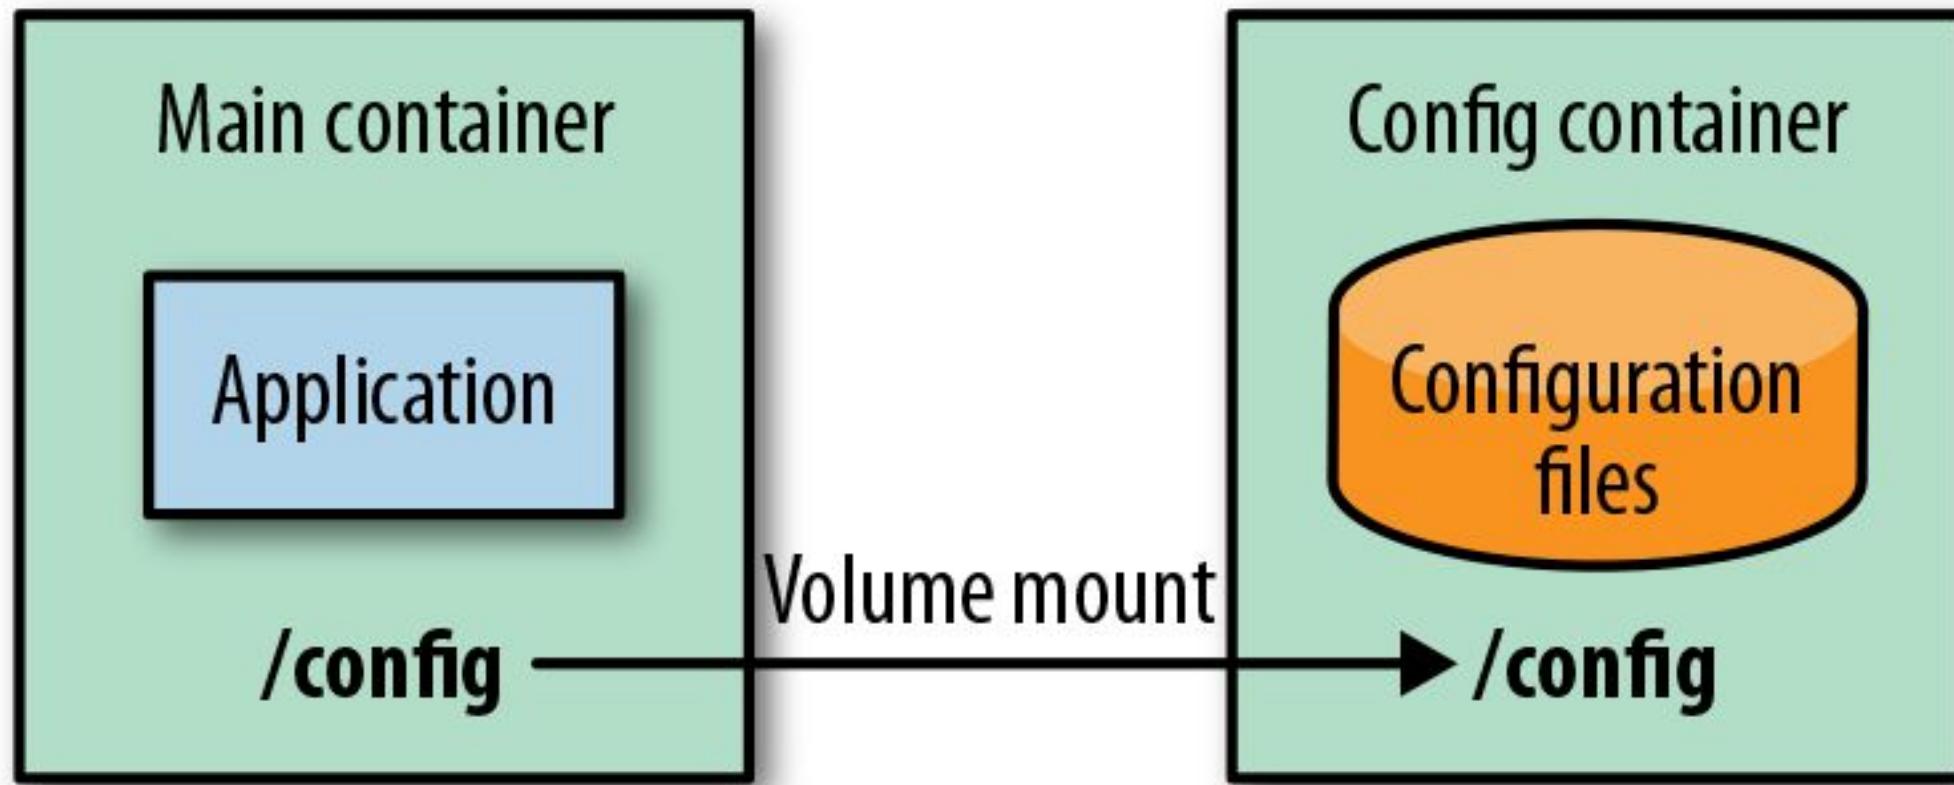
Immutable Configuration

How to configure your application with immutable container images

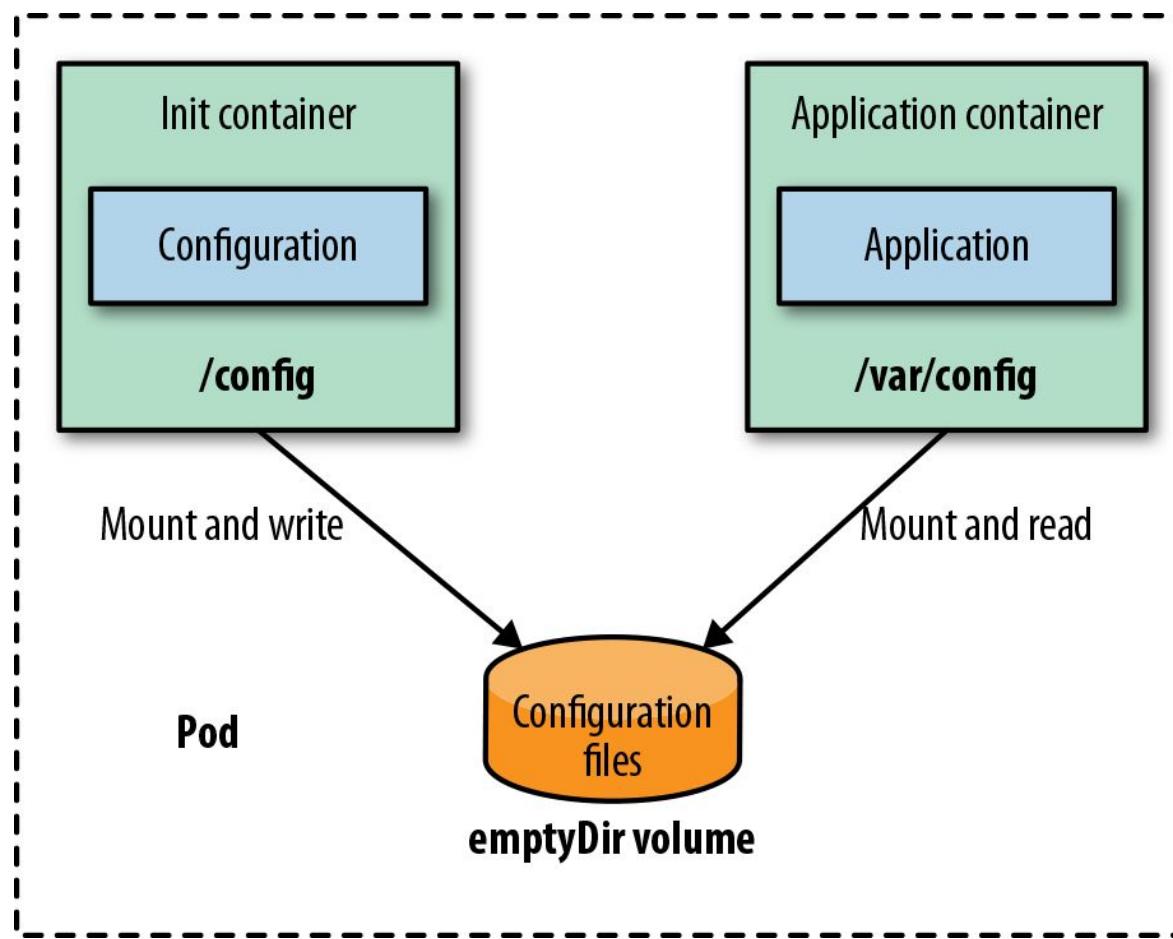
Immutable Configuration

- Configuration that can not be changed
- Every configuration change results in a new *Immutable Configuration* set
- Different sets of configuration for different environments (dev/prod)
- Revisioning and auditing for configuration
- Solution: Configuration stored in container images
 - Versioned via tags and digests
 - Distributed via image registry

Docker volume mount



Configuration provided by Init Container



Stateful Service

```
initContainers:  
- image: k8spatterns/config-dev:1  
  name: init  
  args:  
  - "/config"  
volumeMounts:  
- mountPath: "/config"  
  name: config-directory  
containers:  
- image: k8spatterns/demo:1  
  name: demo  
  ports:  
  - containerPort: 8080  
name: http  
  protocol: TCP  
volumeMounts:  
- mountPath: "/config"  
  name: config-directory  
volumes:  
- name: config-directory  
  emptyDir: {}
```

A large red arrow points from the left side of the diagram to the right, indicating a transformation or comparison between the two sections.

```
FROM busybox  
ADD dev.properties /config-src/demo.properties  
ENTRYPOINT \  
[ "sh", "-c", "cp /config-src/* $1", "--" ]
```

Init Containers to the rescue

- Kubernetes: No container image volume mounts
- Instead: Copy over configuration from image to a shared volume in an *Init Container*
- Exchanging the configuration image for different environments is challenging
- Templating of Deployments to the rescue
 - OpenShift Templates, Helm Charts, Operator with CRD ...
- Good for large to huge configuration data (like a Wildfly server configuration)

OpenShift Template

```
apiVersion: v1
kind: Template
metadata:
  name: demo
parameters:
  - name: CONFIG_IMAGE
    description: Name of configuration image
    value: k8spatterns/config-dev:1
objects:
- apiVersion: v1
  kind: DeploymentConfig
  // ....
  spec:
    template:
      metadata:
        // ....
      spec:
        initContainers:
        - name: init
          image:  ${CONFIG_IMAGE} 
        ...
```

Exercise

A photograph of a light-colored wooden table used for baking. In the top left corner, there are three whole oranges. Next to them are two dark brown, round objects, possibly dates or chocolate chips. In the center, a wooden rolling pin lies next to a rectangular block of yellowish dough. To the right of the dough is a white wooden tray with a scalloped edge, containing several metal cookie cutters in various shapes: a star, a snowflake, a Christmas tree, and a reindeer. The entire scene is set against a dark, slightly out-of-focus background.

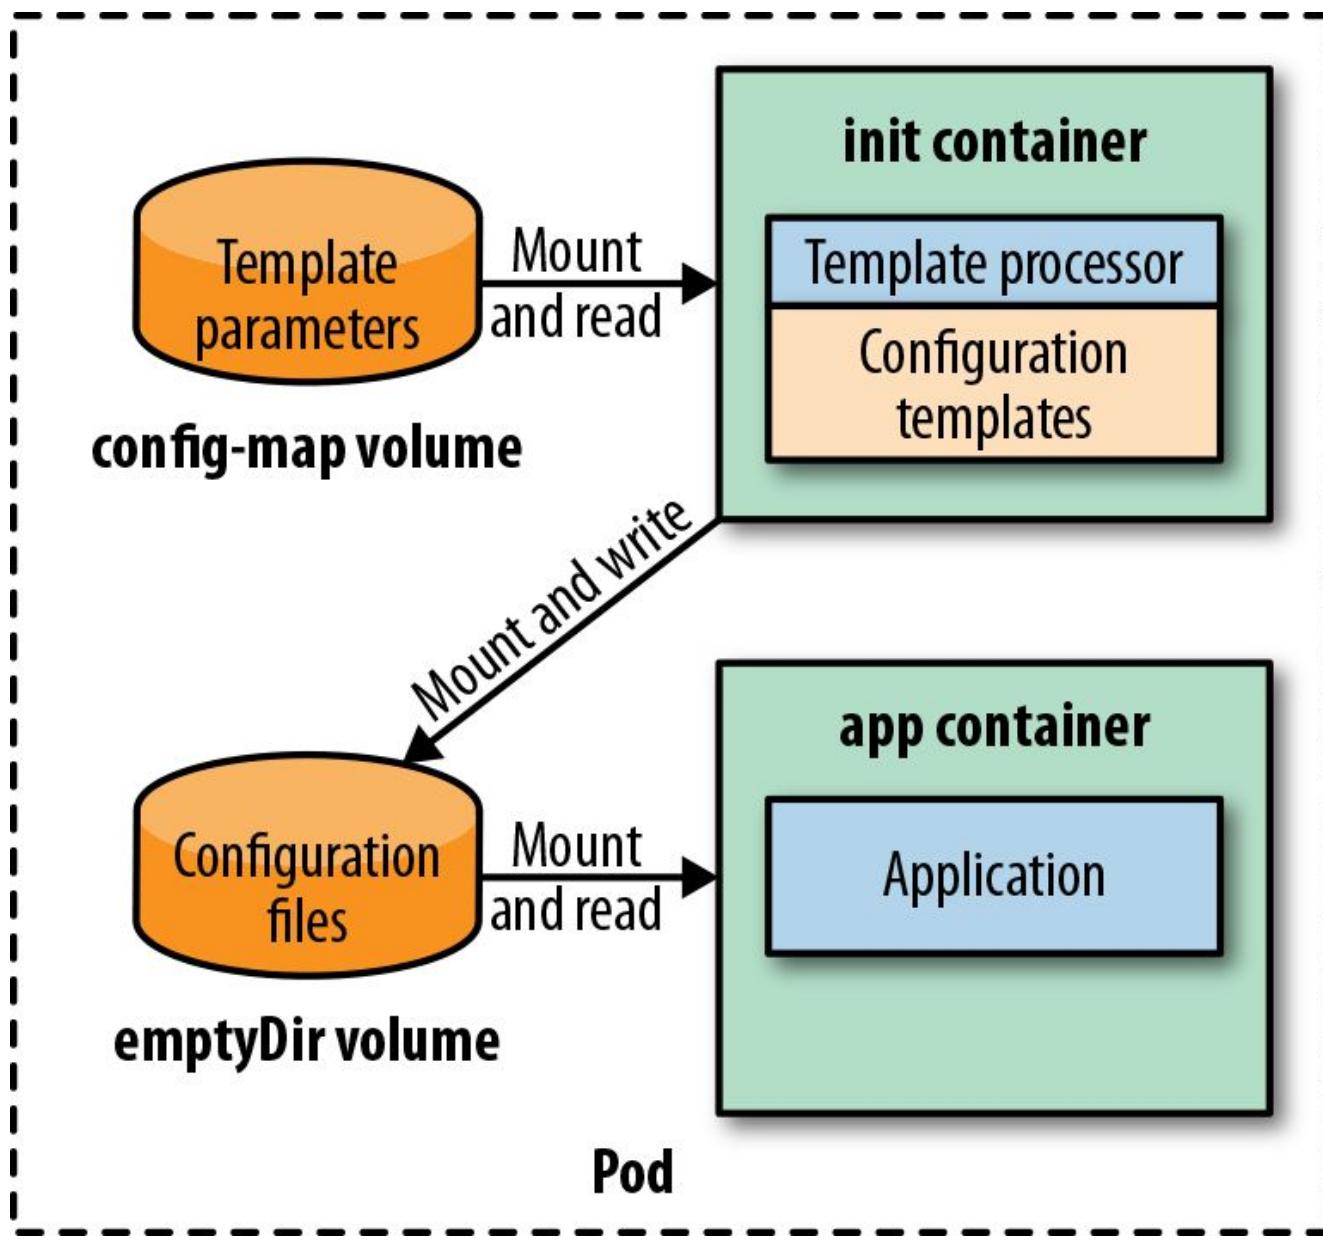
Configuration Template

How to manage large and complex similar configuration data

Preparing Configuration during Startup

- Init Container ...
 - ... contains a template processor
 - ... holds the configuration template
 - ... picks up template parameter from a ConfigMap
 - ... stores final configuration on a shared volume
- Main Container
 - ... accesses created configuration from shared volume

Configuration
Template



Configuration Template

```
apiVersion: extensions/v1beta1
kind: Deployment
metadata:
  name: wildfly-cm-template
spec:
  replicas: 1
  template:
    spec:
      initContainers:
        - image: k8spatterns/config-init
          name: init
          volumeMounts:
            - mountPath: "/params"
              name: wildfly-parameters
            - mountPath: "/out"
              name: wildfly-config
```

```
containers:
- image: jboss/wildfly:10.1.0.Final
  name: server
  volumeMounts:
    - mountPath: "/config"
      name: wildfly-config
  volumes:
    - name: wildfly-parameters
      configMap:
        name: wildfly-params-cm
    - name: wildfly-config
      emptyDir: {}
```

Exercise

A silhouette of a person performing a yoga pose, specifically a standing variation of the Anjali Mudra or Namaste, with hands raised and joined near the head. The background is a vibrant sunset over a body of water, with colors transitioning from deep blue at the bottom to orange, yellow, and red at the top, with a bright sun visible in the upper center.

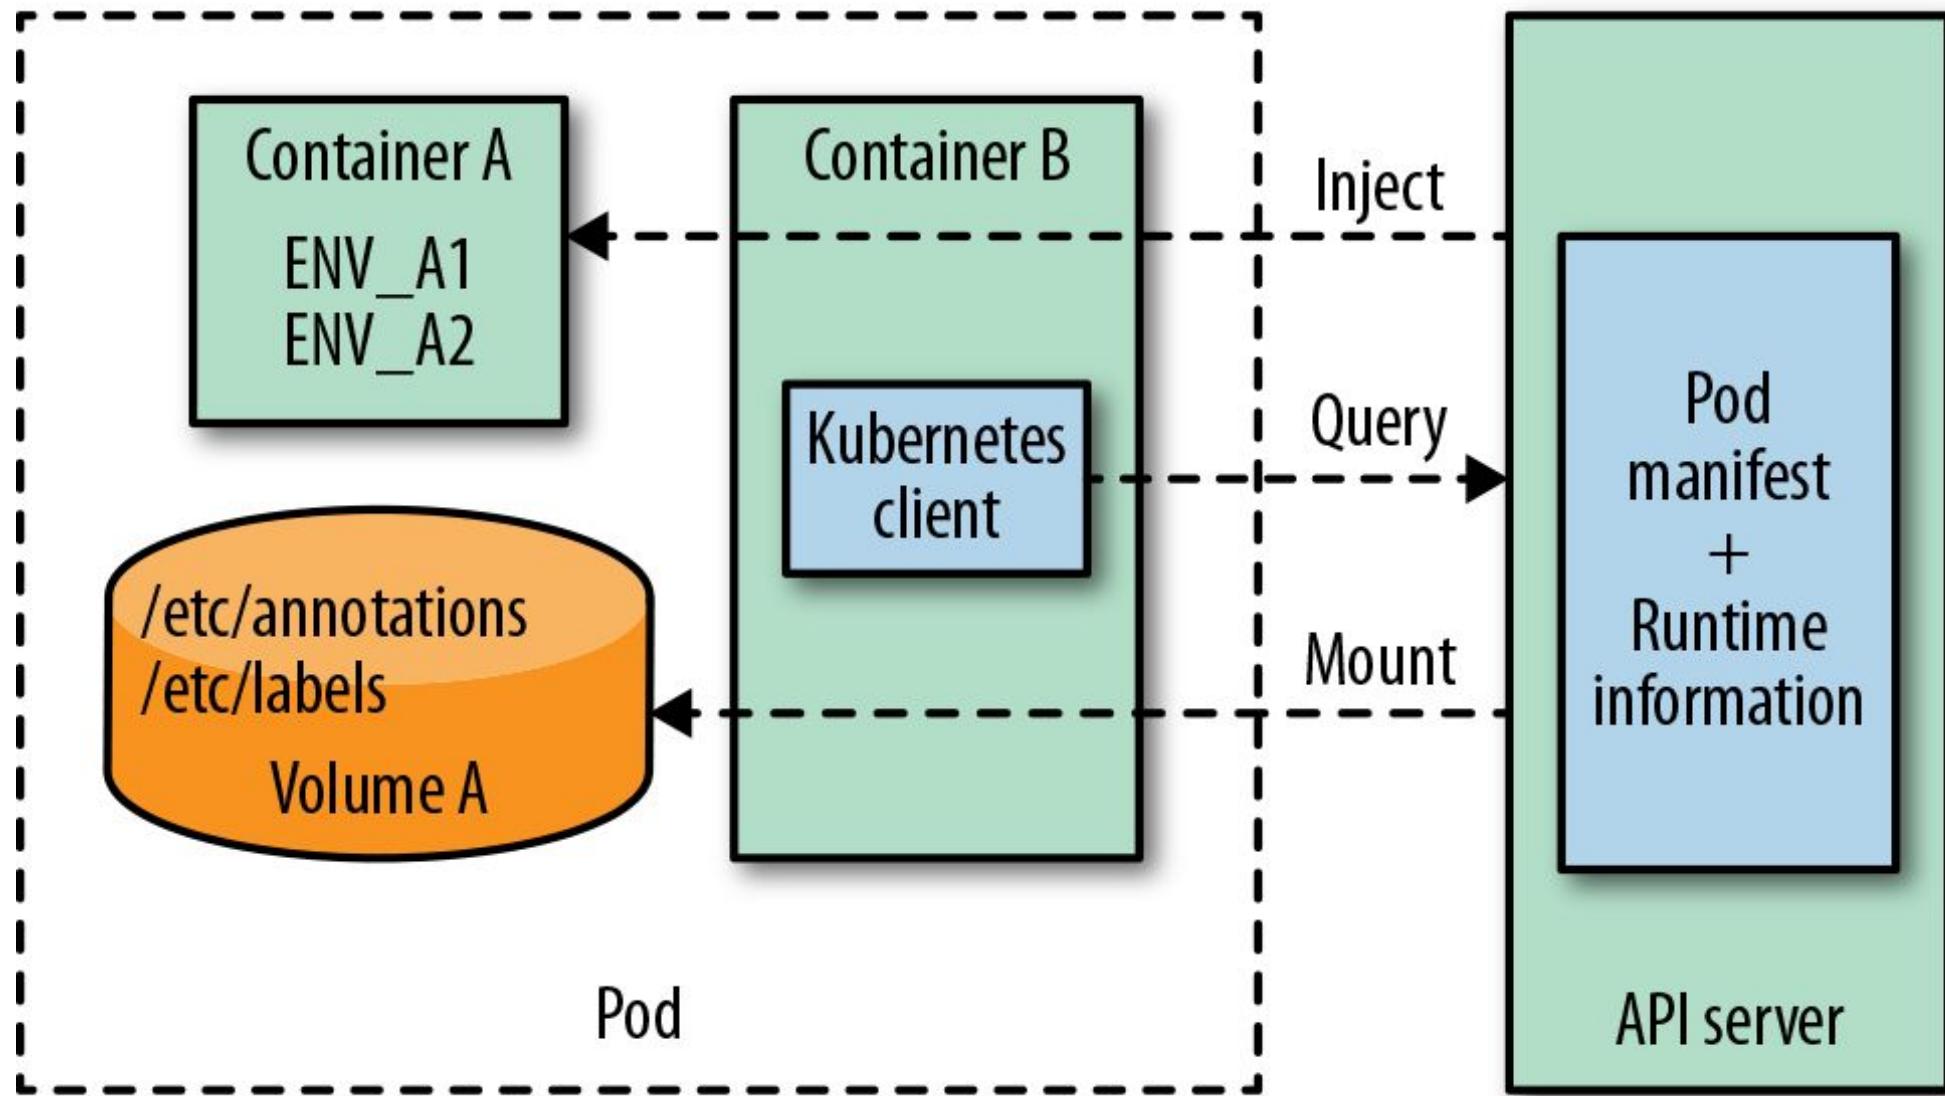
Self Awareness

How an application can access cluster context information

Downward API

- Passing Pod metadata to containers
 - ... via environment variables
 - ... via files
- “Injection” instead of active lookup
(Inversion-of-control)

Self Awareness



```
apiVersion: v1
kind: Pod
metadata:
  name: random-generator
spec:
  containers:
    - image: k8spatterns/random-generator:1.0
      name: random-generator
      env:
        - name: POD_IP
          valueFrom:
            fieldRef:
              fieldPath: status.podIP
        - name: MEMORY_LIMIT
          valueFrom:
            resourceFieldRef:
              container: random-generator
              resource: limits.memory
```

```
apiVersion: v1
kind: Pod
metadata:
  name: random-generator
spec:
  containers:
    - image: k8spatterns/random-generator:1.0
      name: random-generator
      volumeMounts:
        - name: pod-info
          mountPath: /pod-info
  volumes:
    - name: pod-info
      downwardAPI:
        items:
          - path: labels
            fieldRef:
              fieldPath: metadata.labels
```

fieldRef.fieldPath

Name	Description
spec.nodeName	Name of node hosting the Pod
status.hostIP	IP address of node hosting the Pod
metadata.name	Pod name
metadata.namespace	Namespace in which the Pod is running
status.podIP	Pod IP address
spec.serviceAccountName	ServiceAccount that is used for the Pod
metadata.uid	Unique ID of the Pod
metadata.labels['key']	Value of the Pod's label <i>key</i>
metadata.annotations['key']	Value of the Pod's annotation <i>key</i>

resourceFieldRef.resource

Name	Description
requests.cpu	A container's CPU request
limits.cpu	A container's CPU limit
limits.memory	A container's memory request
requests.memory	A container's memory limit

Exercise

The background of the image consists of a dense, repeating pattern of small, silver-colored cylindrical objects, which appear to be aerosol cans or similar containers. They are arranged in a grid-like fashion across the entire frame, creating a sense of depth and repetition.

Batch
Job

How to run short-lived Pods reliably until completion

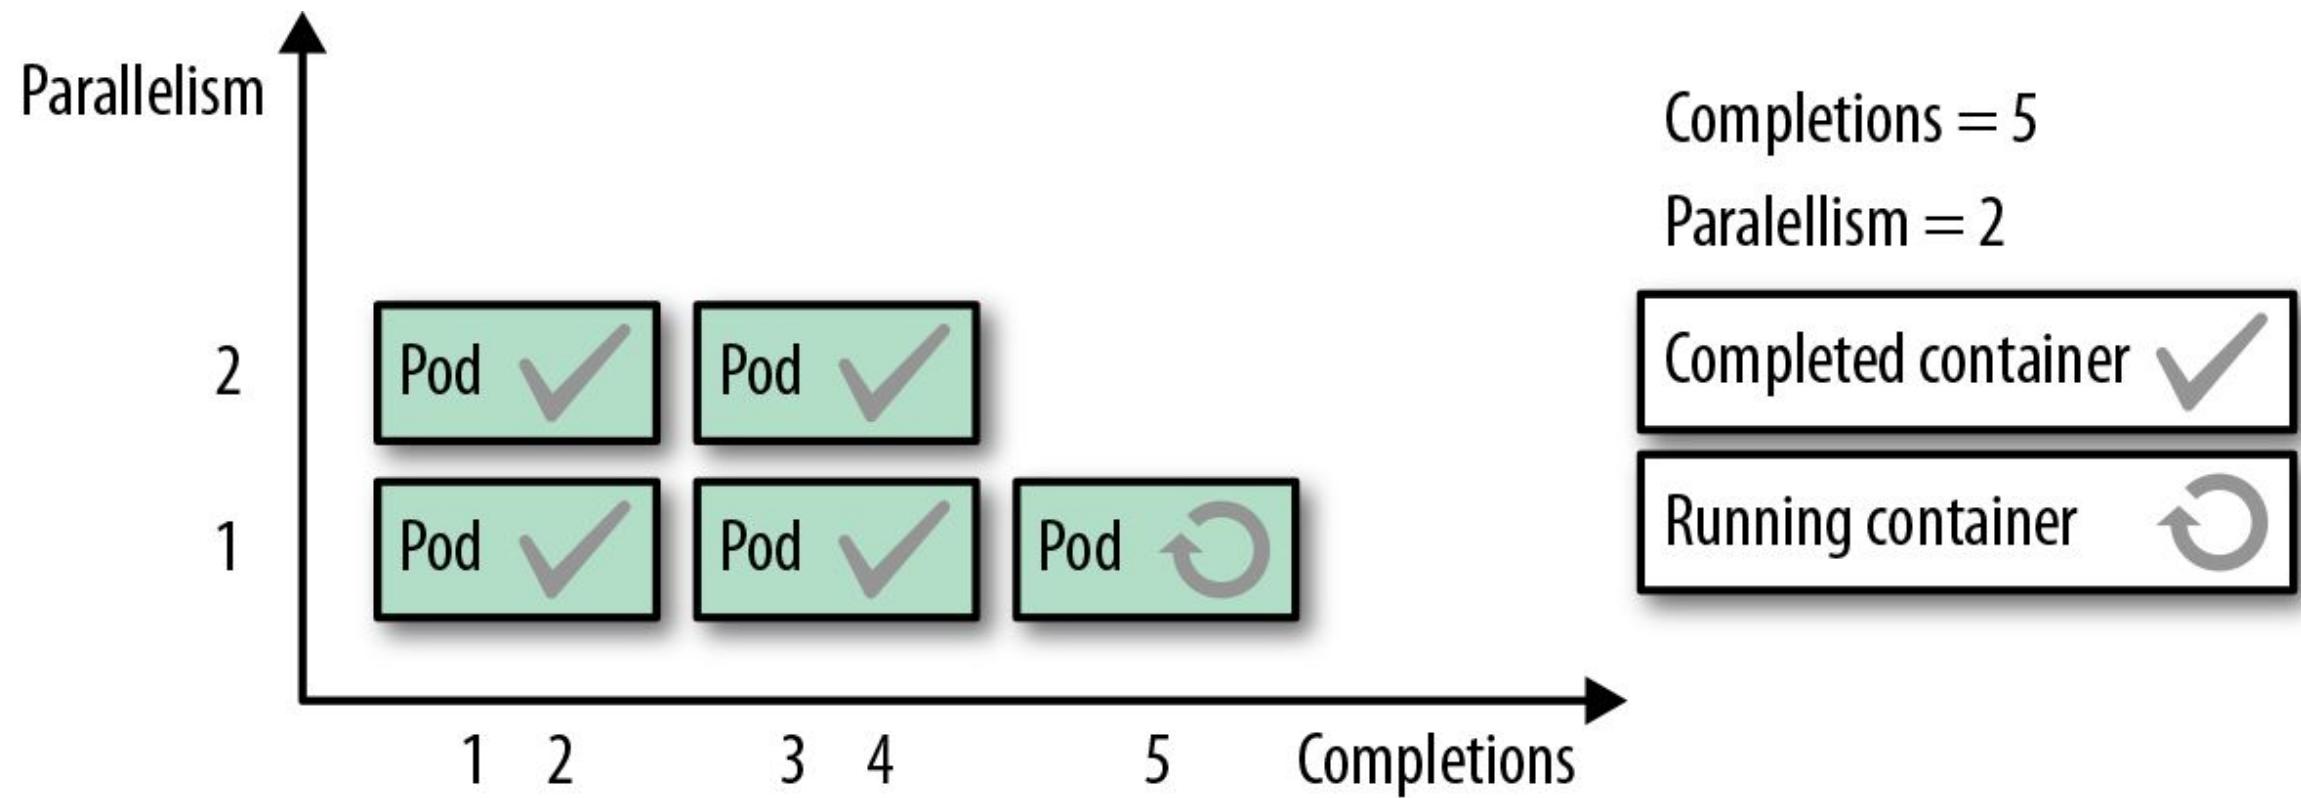
Job

- Resource for a predefined *finite* unit-of-work
- Survives cluster restarts and node failures
- Support for multiple Pod runs
- `.spec.completions`
 - How many Pods to run to complete Job
- `.spec.parallelism`
 - How many Pods to run in parallel

Job

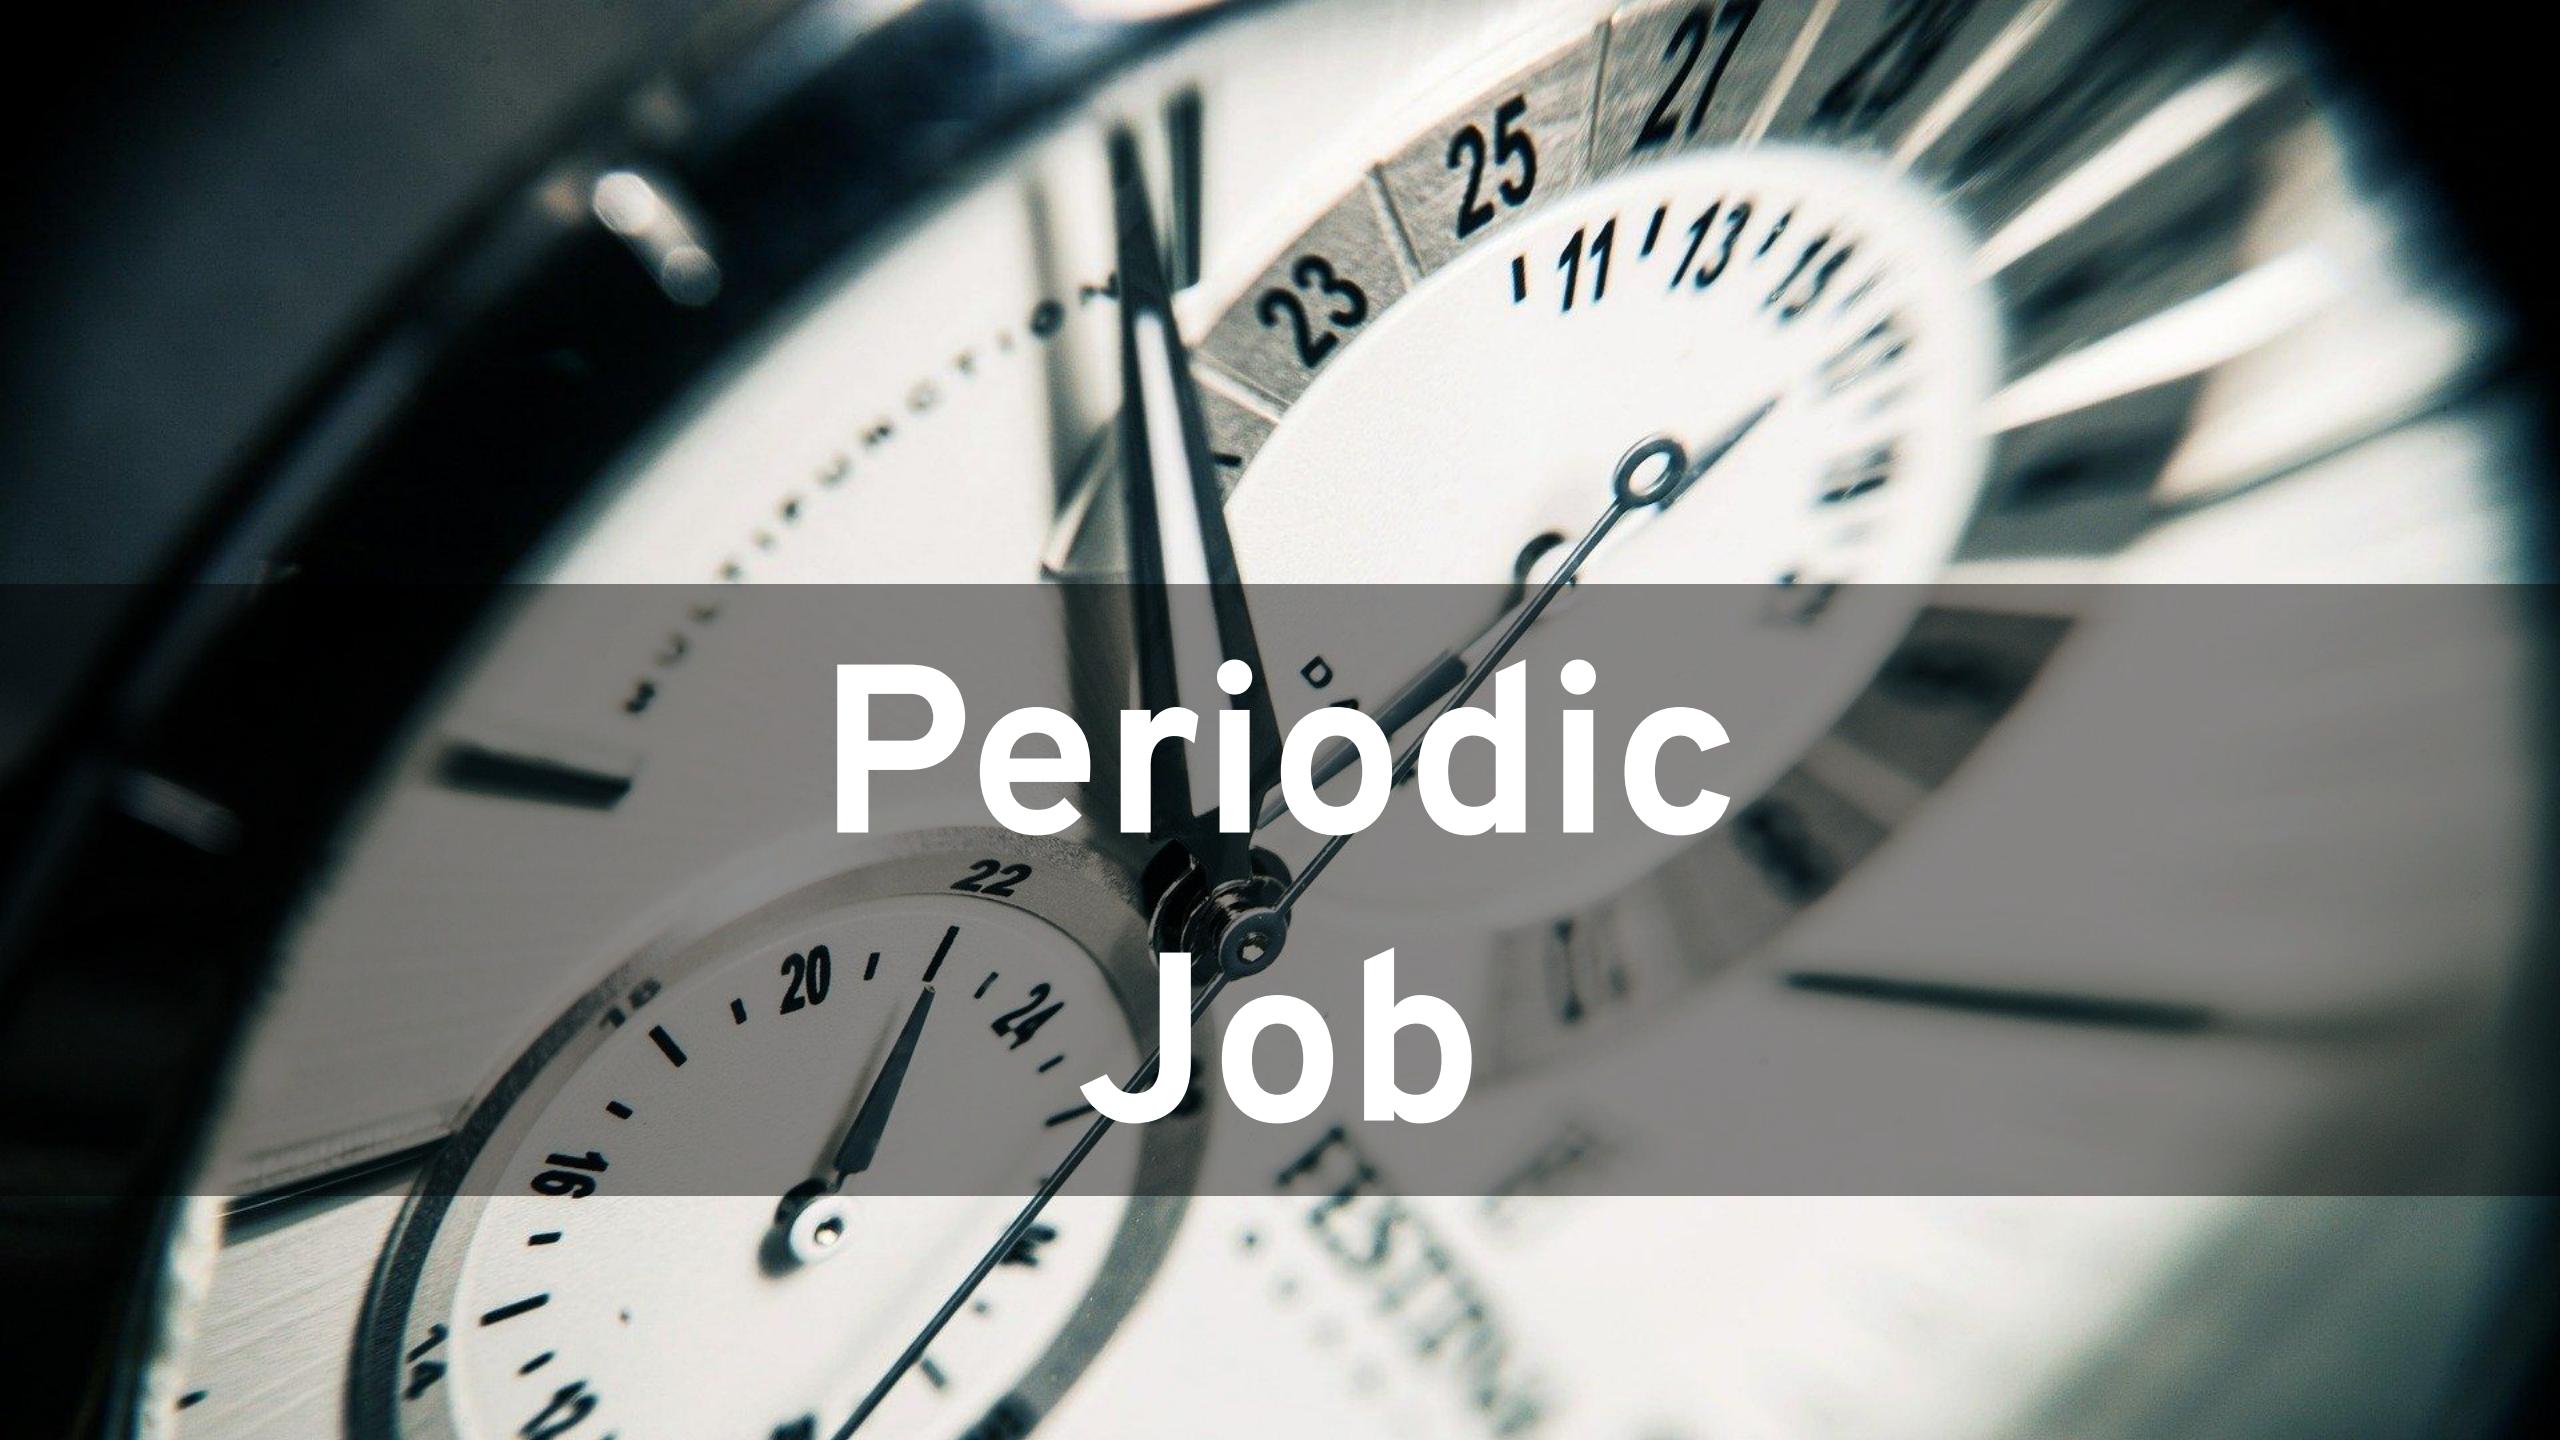
```
apiVersion: batch/v1
kind: Job
metadata:
  name: random-generator
spec:
  completions: 5
  parallelism: 2
  template:
    metadata:
      name: random-generator
    spec:
      restartPolicy: OnFailure
      containers:
        - image: k8spatterns/random-generator:1.0
          name: random-generator
          command: [ "java", "-cp", "/", "RandomRunner", "/numbers.txt", "10000" ]
```

Parallel Batch Job



Job Types

- Single Pod Job
 - completions = 1 and parallelism = 1
- Fixed completion count Jobs
 - completions > 1
 - One Pod per work item
- Work queue Jobs
 - completions = 1 and parallelism > 1
 - All Pods need to finish, with at least one Pod exiting successfully
 - Pods needs to coordinate to shutdown in a coordinate fashion



Periodic Job

How to run a *Batch Job* triggered by a temporal event

CronJob

- Executing a *Batch Job* periodically
- Fields:
 - **schedule** - Cron expression to specify when the Job should run
 - **template** - Job specification
 - **startingDeadlineSeconds** - Deadline for starting jobs that missed the schedule time
 - **concurrencyPolicy** - how to react if previous *Batch Job* is still running
 - **suspend** - set to true to pause a CronJob
 - **successfulJobsHistoryLimit** and **failedJobsHistoryLimit** - how many jobs to keep

CronJob

```
apiVersion: batch/v1beta1
kind: CronJob
metadata:
  name: random-generator
spec:
  # Every three minutes
  schedule: "*/3 * * * *"
  jobTemplate:
    spec:
      template:
        spec:
          containers:
            - image: k8spatterns/random-generator:1.0
              name: random-generator
              command: [ "java", "-cp", "/", "RandomRunner", "/numbers.txt", "10000" ]
  restartPolicy: OnFailure
```



A close-up photograph of a campfire with bright orange flames and glowing embers against a dark background. The fire is contained within a black metal fire pit. The text "Daemon Service" is overlaid in the center of the image.

Daemon
Service

How to run a Pod replica on every node

DaemonSet

- Enables to run one Pod on every cluster node
- Good for infrastructure level processes
 - Log collectors, Metrics exporters, ...
- **nodeSelector** for running only on a subset of nodes
- Not managed by the Kubernetes scheduler
- Treated specially by controllers like the descheduler or cluster-auto scaler
- Alternative: Static Pods started by kubelet from a directory

DaemonSet

```
apiVersion: extensions/v1beta1
kind: DaemonSet
metadata:
  name: random-refresher
spec:
  selector:
    matchLabels:
      app: random-refresher
  template:
    metadata:
      labels:
        app: random-refresher
    spec:
      nodeSelector:
        feature: hw-rng
      containers:
        - image: k8spatterns/random-generator:1.0
          volumeMounts:
            - mountPath: /host_dev
              name: devices
          volumes:
            - name: devices
      hostPath:
        path: /dev
```



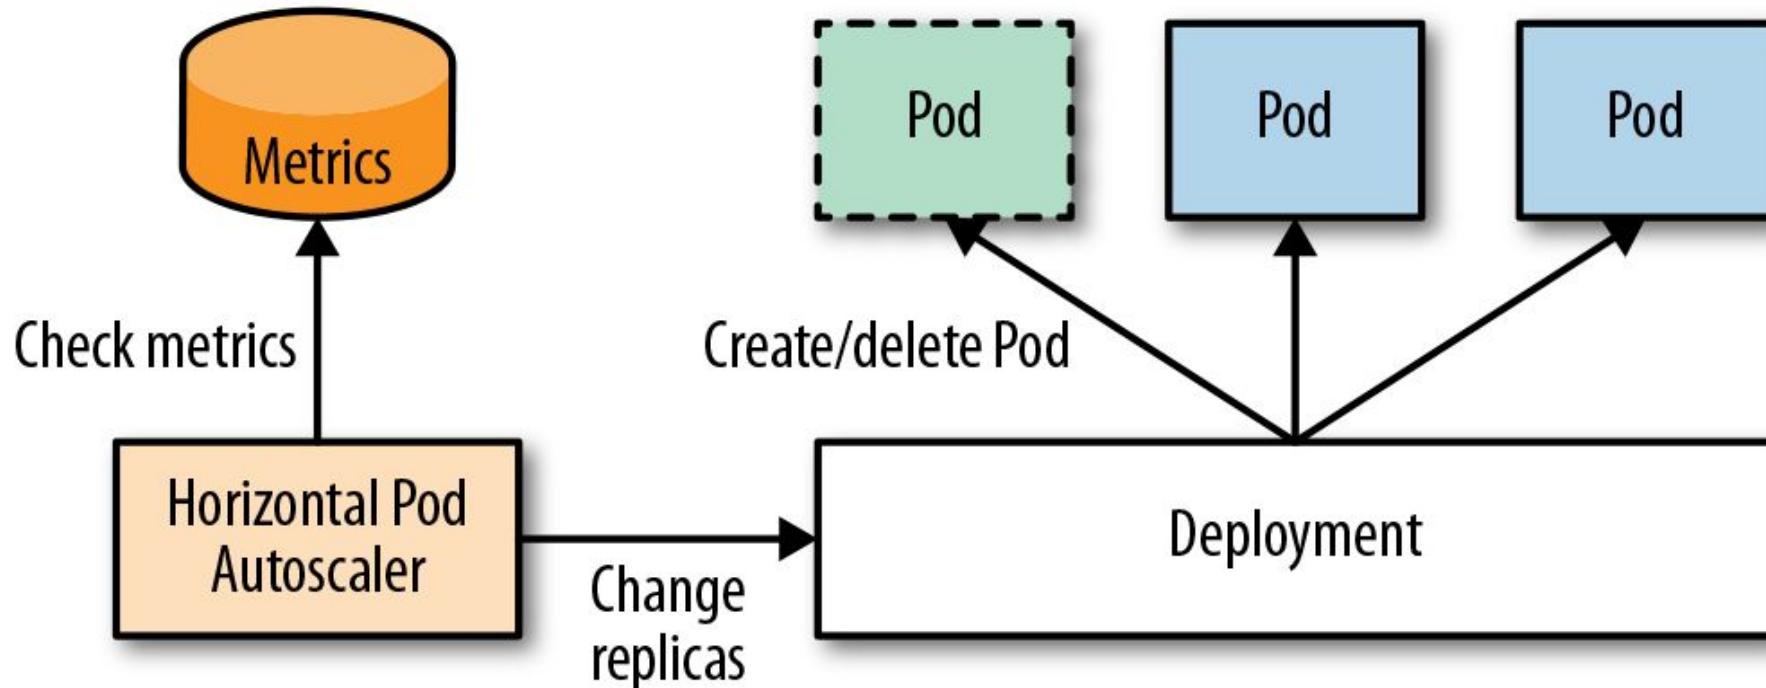
Elastic
Scale

How to automatically react to dynamically changing resource requirements

Scaling

- **Horizontal:** Changing replicas of a Pod
- **Vertical:** Changing resource constraints of containers in a single Pod
- **Cluster:** Adding new nodes to a cluster
- **Manual:** Changing scale parameters manually, imperatively or declaratively
- **Automatic:** Change scaling parameters based on observed metrics

Horizontal Pod Autoscaler (HPA)

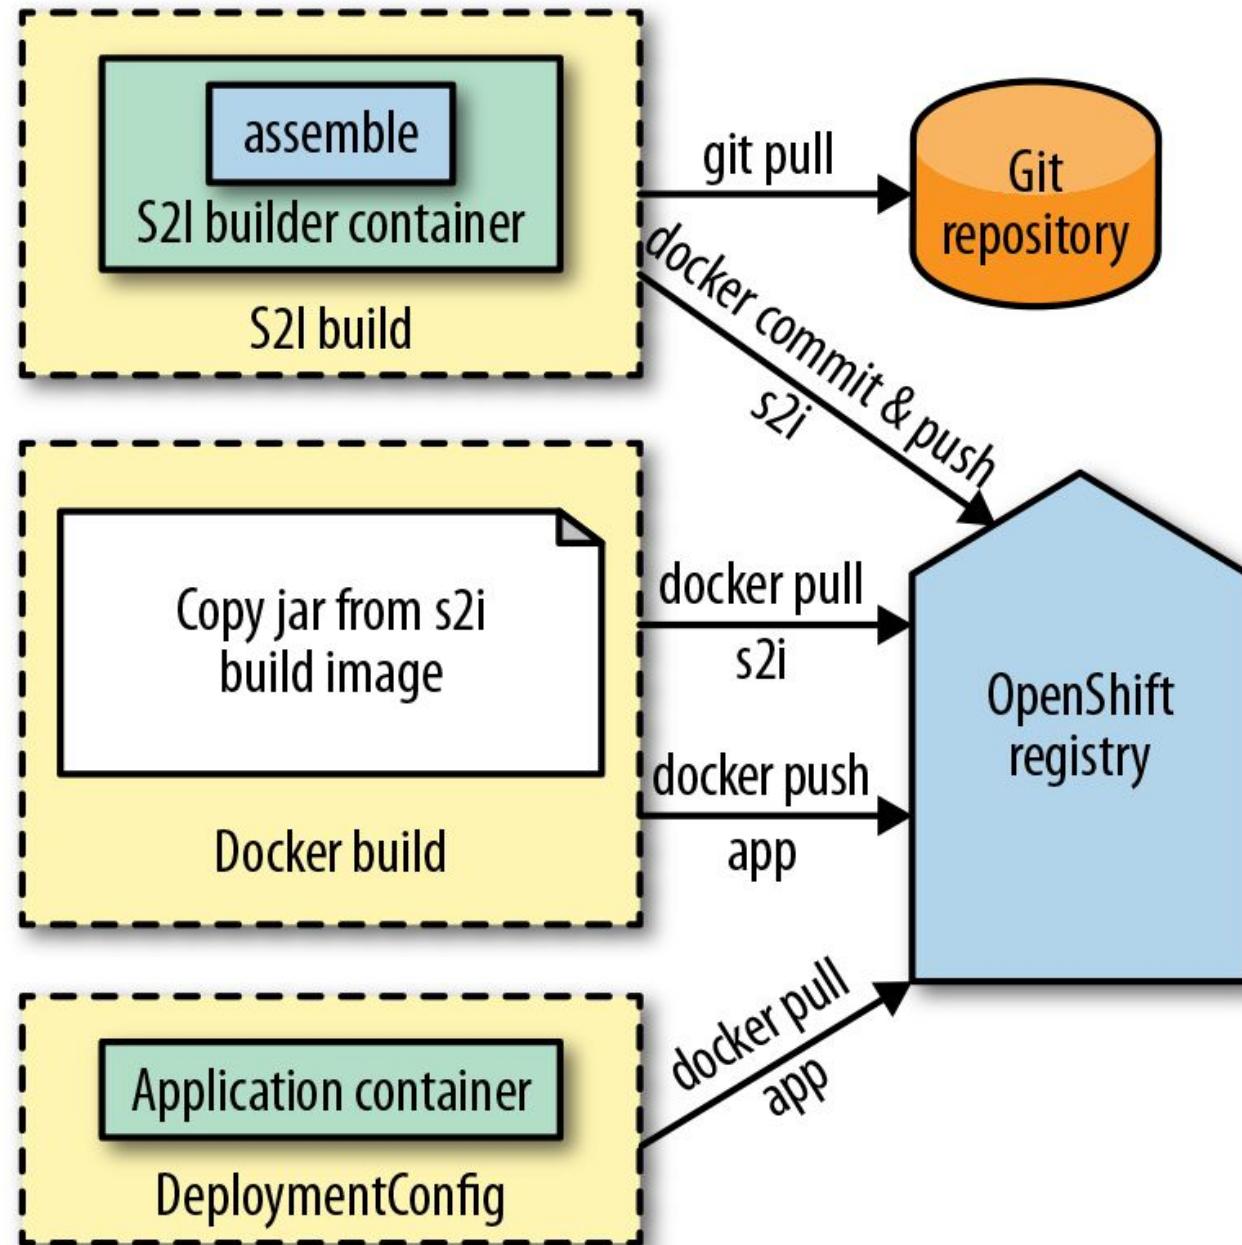


```
kubectl autoscale deployment random-generator --cpu-percent=50 --min=1 --max=5
```

BuildConfig

```
apiVersion: v1
kind: BuildConfig
metadata:
  name: random-generator-build
spec:
  source:
    git:
      uri: https://github.com/k8spatterns/random-generator
  strategy:
    sourceStrategy:
      from:
        kind: DockerImage
        name: fabric8/s2i-java
  output:
    to:
      kind: ImageStreamTag
      name: random-generator-build:latest
triggers:
- type: ImageChange
```

S2I Chained Build



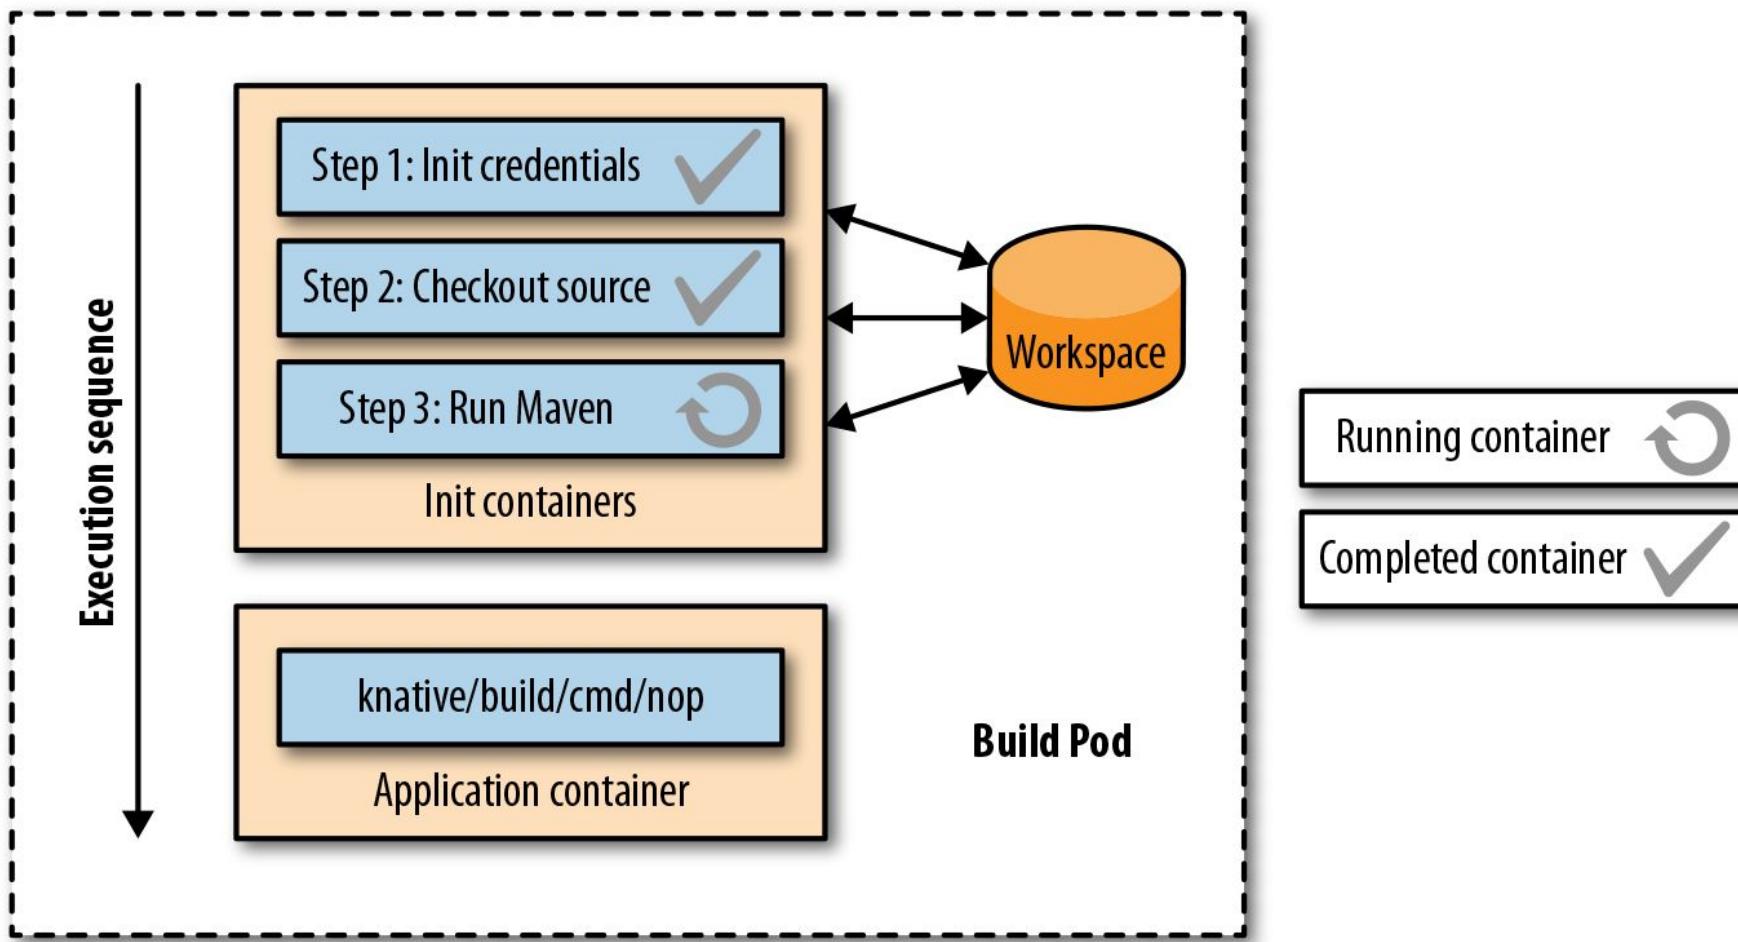
Knative

- **Knative serving**
 - Scale-to-Zero
- **Knative eventing**
 - Event infrastructure for triggering services
- **Knative build**
 - Transforming source to container image
- Build templates allows reusing build strategies

Build

```
apiVersion: build.knative.dev/v1alpha1
kind: Build
metadata:
  name: random-generator-build-jib
spec:
  source:
    git:
      url: https://github.com/k8spatterns/random-generator.git
      revision: master
  steps:
  - name: build-and-push
    image: gcr.io/cloud-builders/mvn
    args:
      - compile
      - com.google.cloud.tools:jib-maven-plugin:build
      - -Djib.to.image=registry/k8spatterns/random-generator
    workingDir: /workspace
```

Knative Build



Exercise