

# Containerful Serverless

Roland Huß @[ro14nd@hachyderm.io](mailto:ro14nd@hachyderm.io)

Senior Principal Software Engineer, Red Hat  
OpenShift Serverless Architect  
Ex-Knative TOC member

Matthias Weßendorf @[matzew@hachyderm.io](mailto:matzew@hachyderm.io)

Principal Software Engineer, Red Hat  
OpenShift Serverless Engineer  
Knative contributor



# Wait ... wat ?

---



# What is “Serverless” ?

“Serverless computing refers to the concept of building and running **applications** that **do not require server management**. It describes a finer-grained **deployment model** where applications, bundled as one or more functions are uploaded to a platform and then **executed, scaled, and billed** in response to the exact **demand** needed at the moment”

-- CNCF Definition, <https://www.cncf.io/blog/2018/02/14/cncf-takes-first-step-towards-serverless-computing/>

# Serverless vs. FaaS

**Serverless** is a **Deployment Model** that abstracts away the driving machine infrastructure.

- No server management required
- Executed, scaled and billed according to demand
- Defines a deployment packaging, but otherwise agnostic to the application

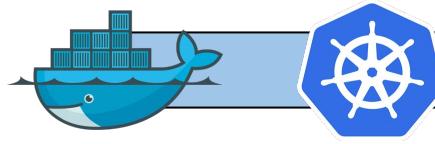
**FaaS** (Function-as-a-Service) is a **Programming Model** that mandates developing your application with fine grained function that match a given signature.

- Deployed as Serverless application
- Typically used as *glue code* to connect services

# Serverless Workloads

The ideal Serverless Workload is ...

- ... **stateless**
- ... **short-running**
- ... **HTTP**-based web services or **Event-Driven** Applications



*“From PaaS to FaaS”*

AWS  
Elastic Beanstalk

*“Time of the FaaS”*

AWS  
**Lambda**

Serverless  
Framework

Google Cloud  
Functions

2008

2010

2012

2014

2016

2018

2020

2022

Google  
App Engine

IBM Cloud  
Functions

Azure  
Functions

*“Containerful Serverless”*

OpenFaaS

Azure Container  
Instances

AWS  
Fargate

Knative  
0.1.0

IBM Cloud  
Code Engine

VMware Cloud  
Native Runtimes

**Knative** 1.0

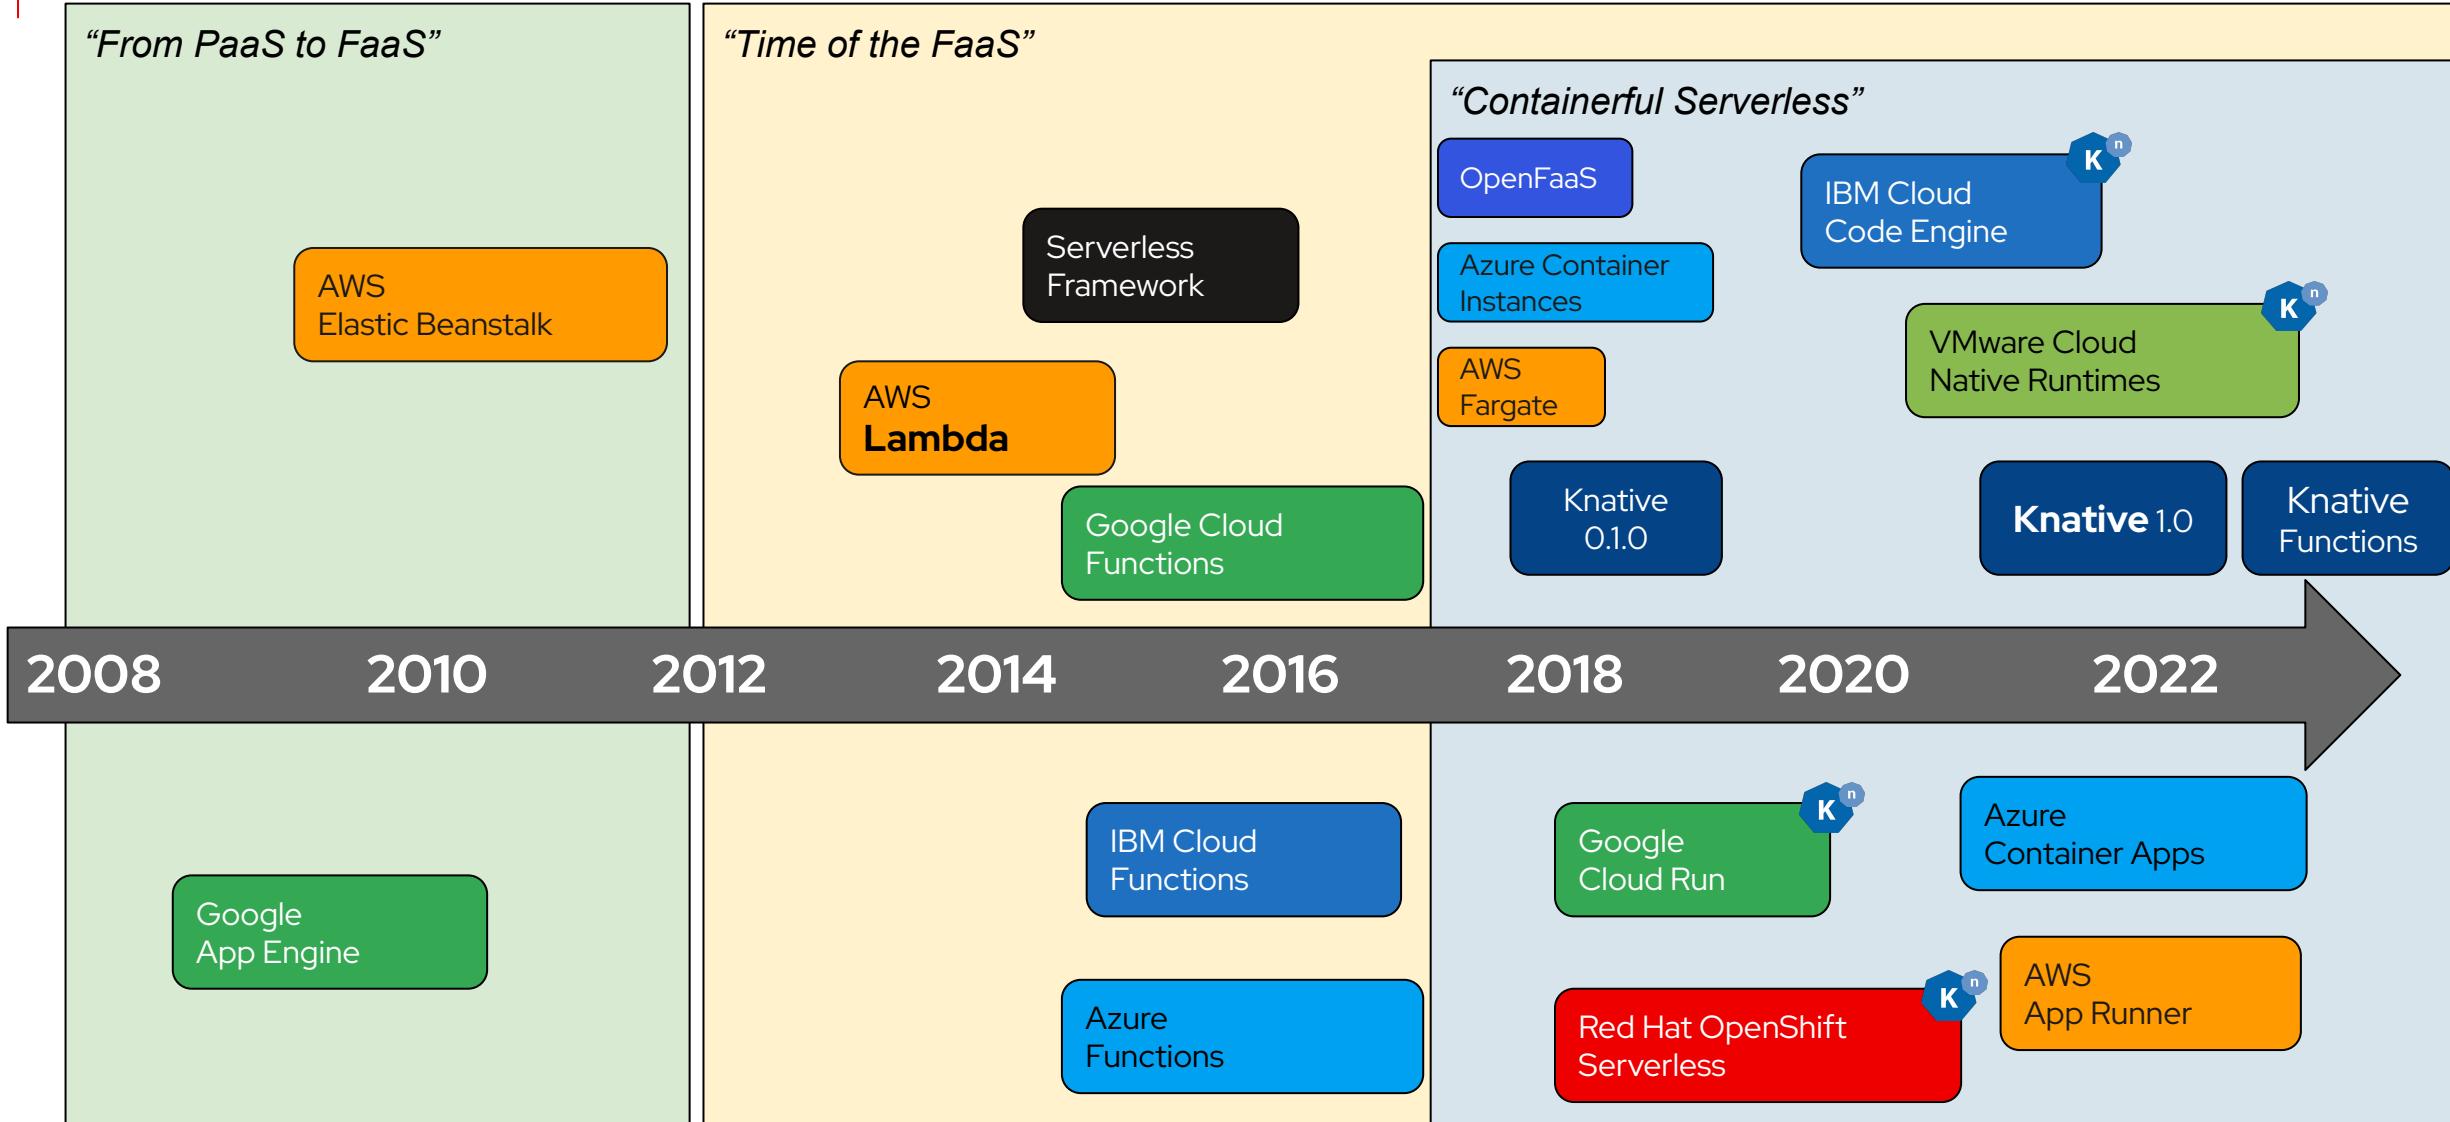
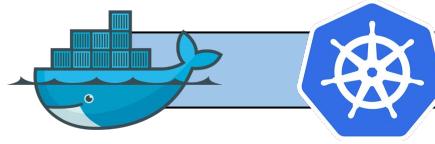
Knative  
Functions

Google  
Cloud Run

Azure  
Container Apps

Red Hat OpenShift  
Serverless

AWS  
App Runner



# Serverless Packaging

- **Custom** (zip, tar, ...)
  - All Cloud Providers
- **OCI Container Images**
  - Docker, Podman, Kubernetes, ...
- **Virtual Machines**
- **MicroVMs**
  - Firecracker, Kata containers, gVisor, ...
- **Web Assembly Modules (WASM)**
  - Krustlets, WasmEdge, Wasmtime

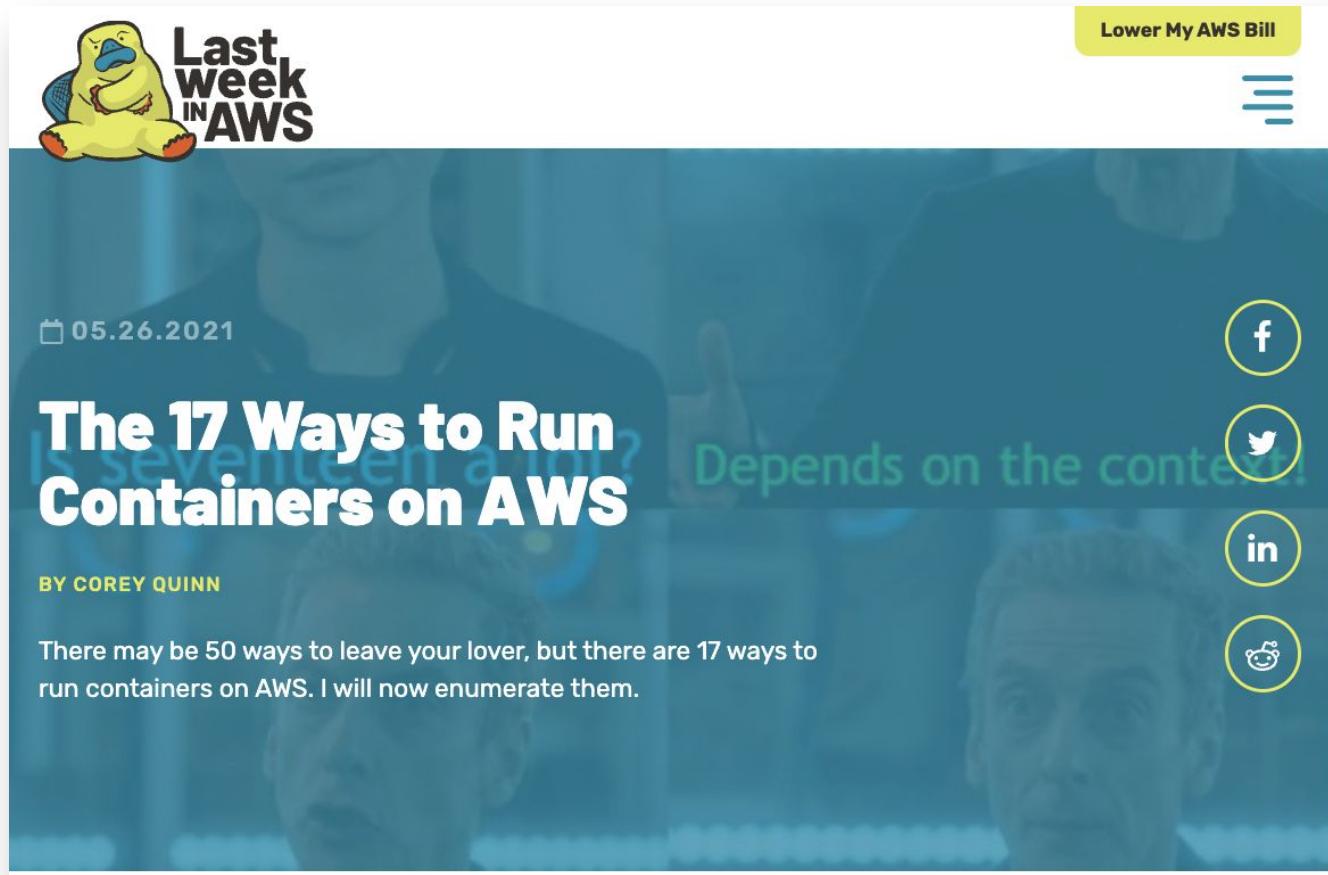
# Containerful Serverless

- **AWS:** Lambda, Fargate, App Runner
- **Microsoft Azure:** Container Instances, Container Apps
- **OpenFaaS**
- **Knative**
  - **Google** Cloud Run, **IBM** Code Engine, **VMware** Cloud Native Runtimes for Tanzu, **Red Hat** OpenShift Serverless

An aerial photograph showing a river flowing through a dense forest. The river curves through the landscape, creating a dark, winding path through the green trees. The forest appears to be a mix of coniferous and deciduous trees, with some yellow and orange foliage visible, suggesting autumn. The overall scene is a natural, undisturbed environment.

AWS

# Containers on AWS

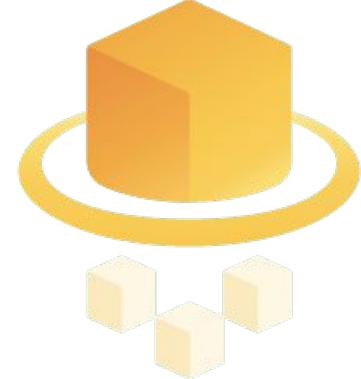


# AWS Lambda



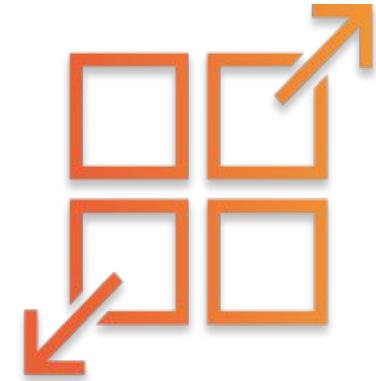
- Opinionated, **event-driven** computing platform
- Introduced **2014** with **Node.js** as runtime
- Main use case as '**glue code**' for connecting AWS services
- **Pay-as-you go** with 1 ms billing granularity
- **Container** packaging since 2019
  - up to 10 GB in size (instead of 50 MB zipped)
  - specific base images or clients needed to fulfill Lambda runtime contract

# AWS Fargate



- Introduced in 2017 as a computing **backend for ECS**
- Deploying containers without managing EC2 instances.
- Fargate can be used in EKS do deploy Kubernetes Pods on Fargate
  - “Infinite large Kubernetes Node”

# AWS App Runner



- Simplest and most opinionated way to deploy containers on AWS
- HTTP workloads
- Automatic Scaling based on traffic
- Automatic routing (public or private URL)
- No scale to zero

A wide-angle photograph of a tropical island. The foreground is filled with clear, turquoise-blue ocean water. In the middle ground, a dense forest of tall palm trees stretches across the island. A small, traditional wooden boat with a blue hull and white trim is anchored near the shore. The sky above is a bright, clear blue with a few wispy white clouds.

# Azure

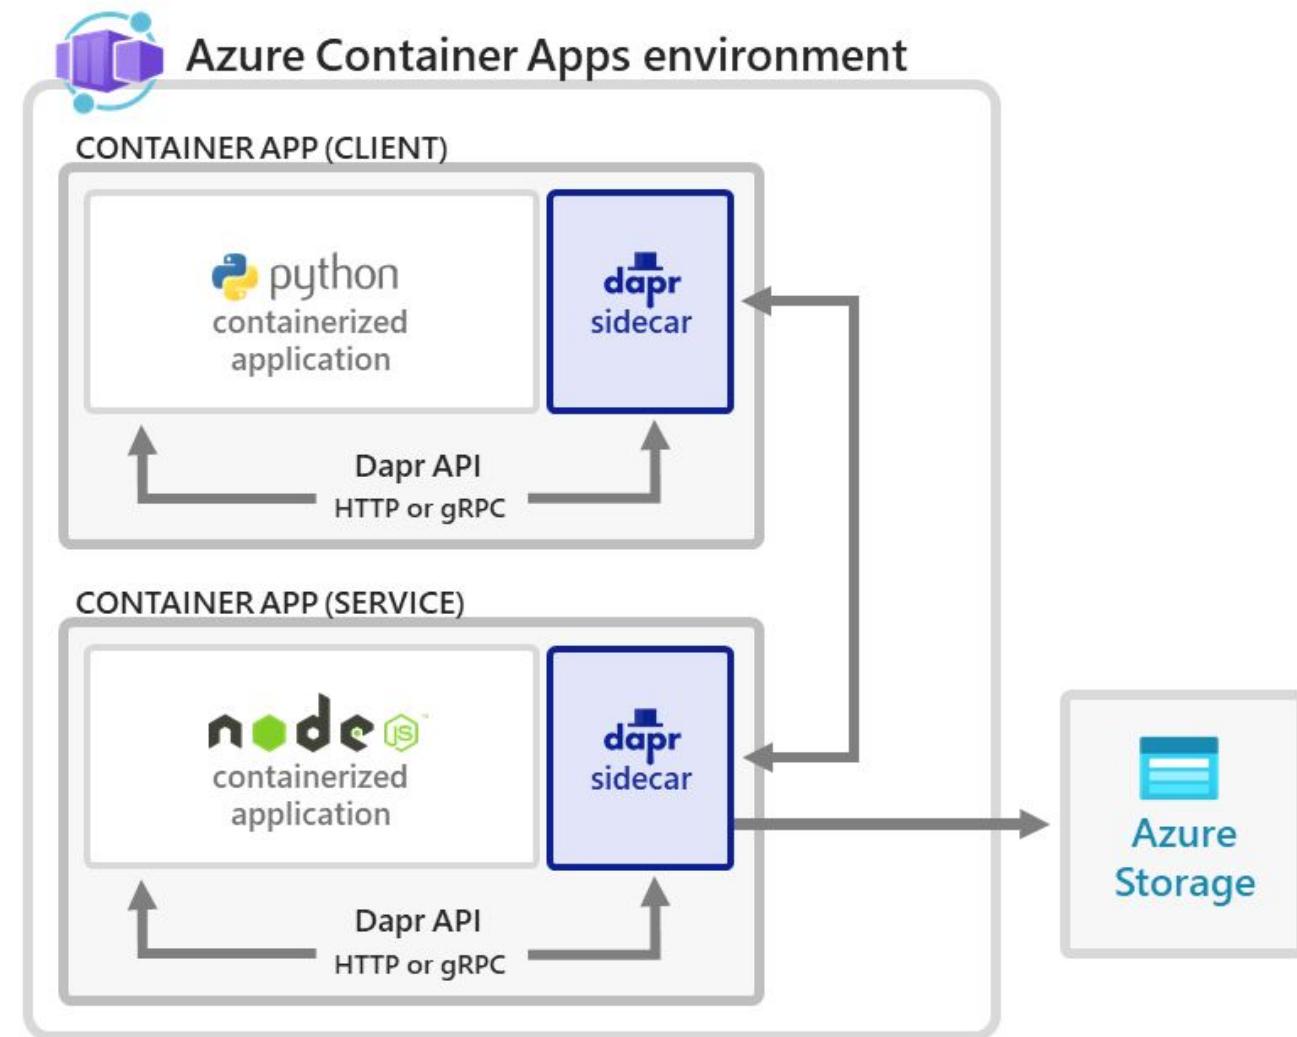
# Container Instances (ACI)



- Started in 2017
- Similar scope as AWS Fargate
- No scaling or load balancing
- Unopinionated approach to start containers
- Simpler to use than ECS / Fargate on AWS
- Can be added as virtual nodes to AKS

# Container Apps

- Started late 2021, since May 2022 GA
- Multiple revisions with traffic split
- Autoscaling options: HTTP, CPU/Memory, KEDA
- Native Dapr integration



A whale is captured mid-leap, its dark, textured body and white underbelly creating a dramatic contrast against the bright spray of water it has just disturbed. The whale's tail is visible above the surface, and its long, curved body extends downwards into the deep blue ocean. The background shows a vast expanse of blue water under a clear sky.

# OpenFaaS

# OpenFaaS

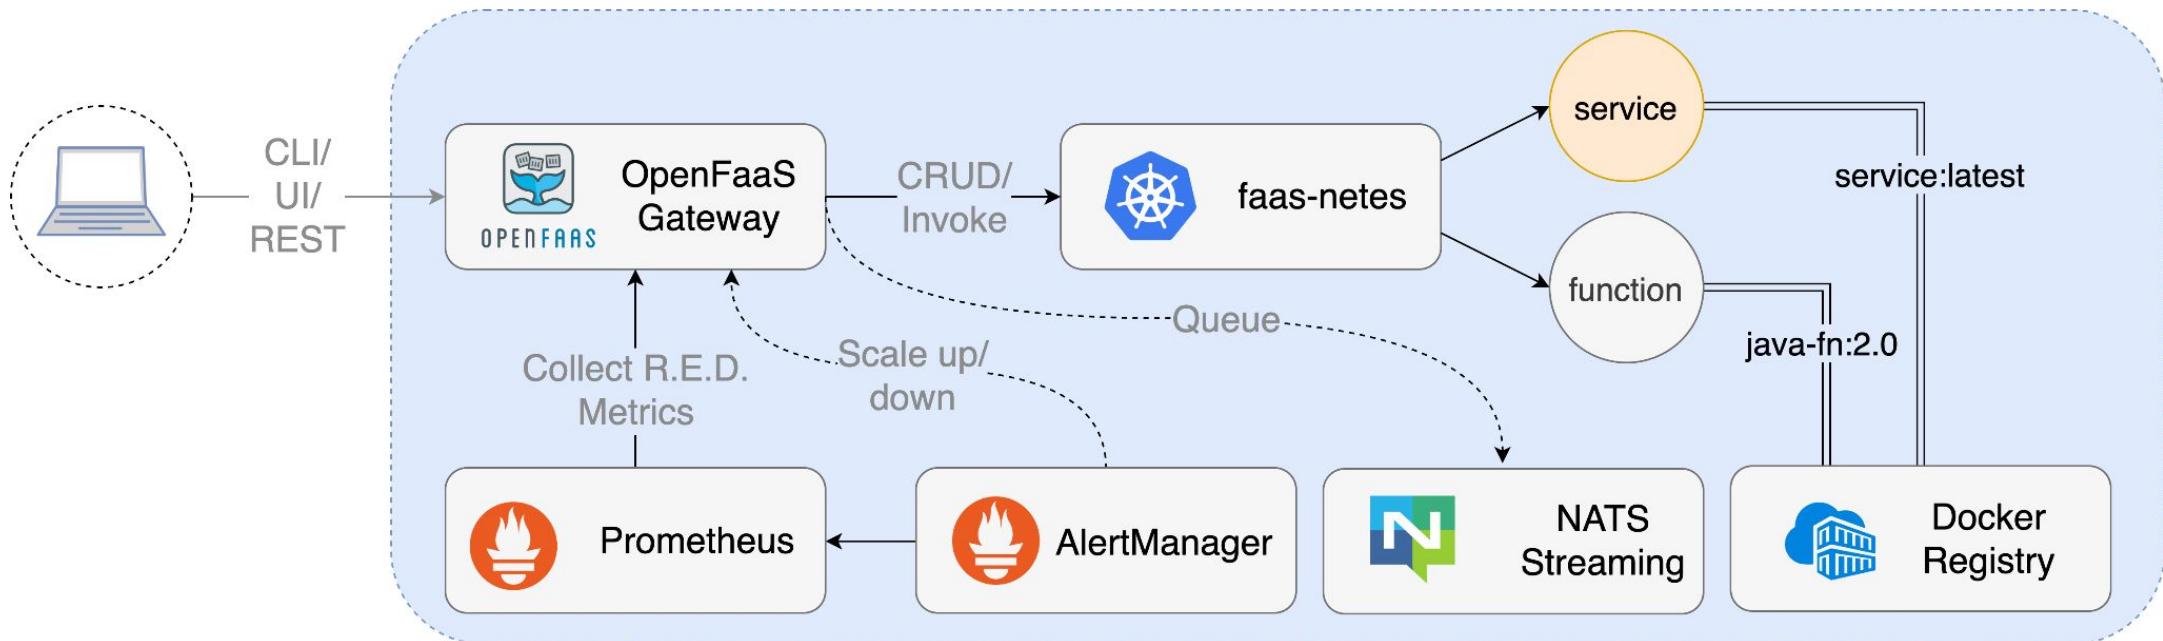
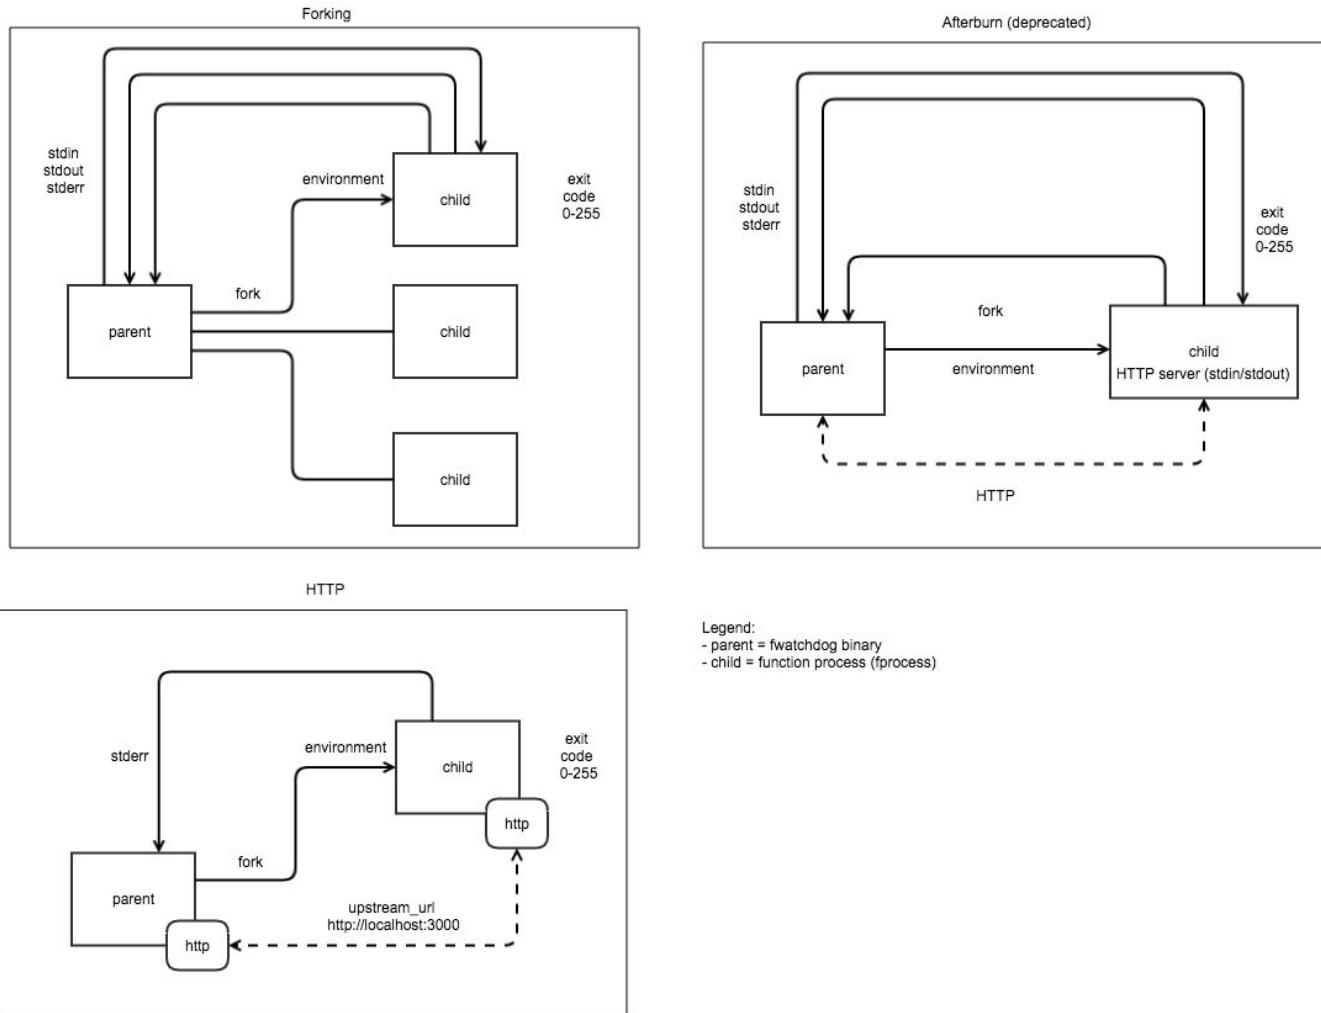


diagram from <https://docs.openfaas.com/architecture/stack/>

# WatchDogs



# Open Core

- Exclusive Features of OpenFaaS Pro:
  - New Autoscaling engine including scale-to-zero
  - Retry of failed functions
  - New User Interface
  - Support for EDA application via triggers (Kafka, AWS SQS)
- No free trial (?)

A photograph of a traditional wooden longtail boat sailing on the ocean. The boat's hull is dark wood, and its deck is made of light-colored planks. A large wooden pole (steering oar) is visible at the stern. The boat is moving towards a group of green, craggy limestone islands under a blue sky with white clouds.

# Knative

**Kubernetes-based platform to  
deploy and manage modern  
serverless workloads.**

<https://knative.dev>

---

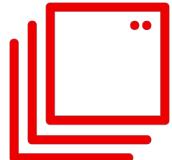
# Modern Event-driven Workloads with Knative

W-JAX, November 10th, 12:00 - 13:00  
Room Atlanta

# Bringing Serverless applications to Kubernetes

## Serving

A request-driven model that serves the container with your application and can "scale to zero".



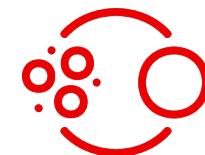
## Eventing

Common infrastructure for consuming and producing events that will stimulate applications.



## Functions

A programming model that lets you focus on just your code for faster iterations.



# Background Information

- Started as an **Open Source** Project mid-2018 by Google
- CNCF Incubating Project since March 2022
- Community driven with a lot of vendor backing
  - <https://github.com/knative>
  - <https://knative.dev>
  - Support by Google\*, Red Hat, IBM, VMware, Triggermesh and more
  - Organized in multiple Working Groups with weekly meetings
- Releases
  - Current: **v1.8**
  - quarterly release cadence (starting 2023)

# Serving Concepts

- **Demand-based autoscaling**, including scale-to-zero
- **Separation** of code and configuration
- Simplified deployment model catered for **stateless applications**
- Rich **traffic split capabilities** to enable custom rollout strategies of new versions

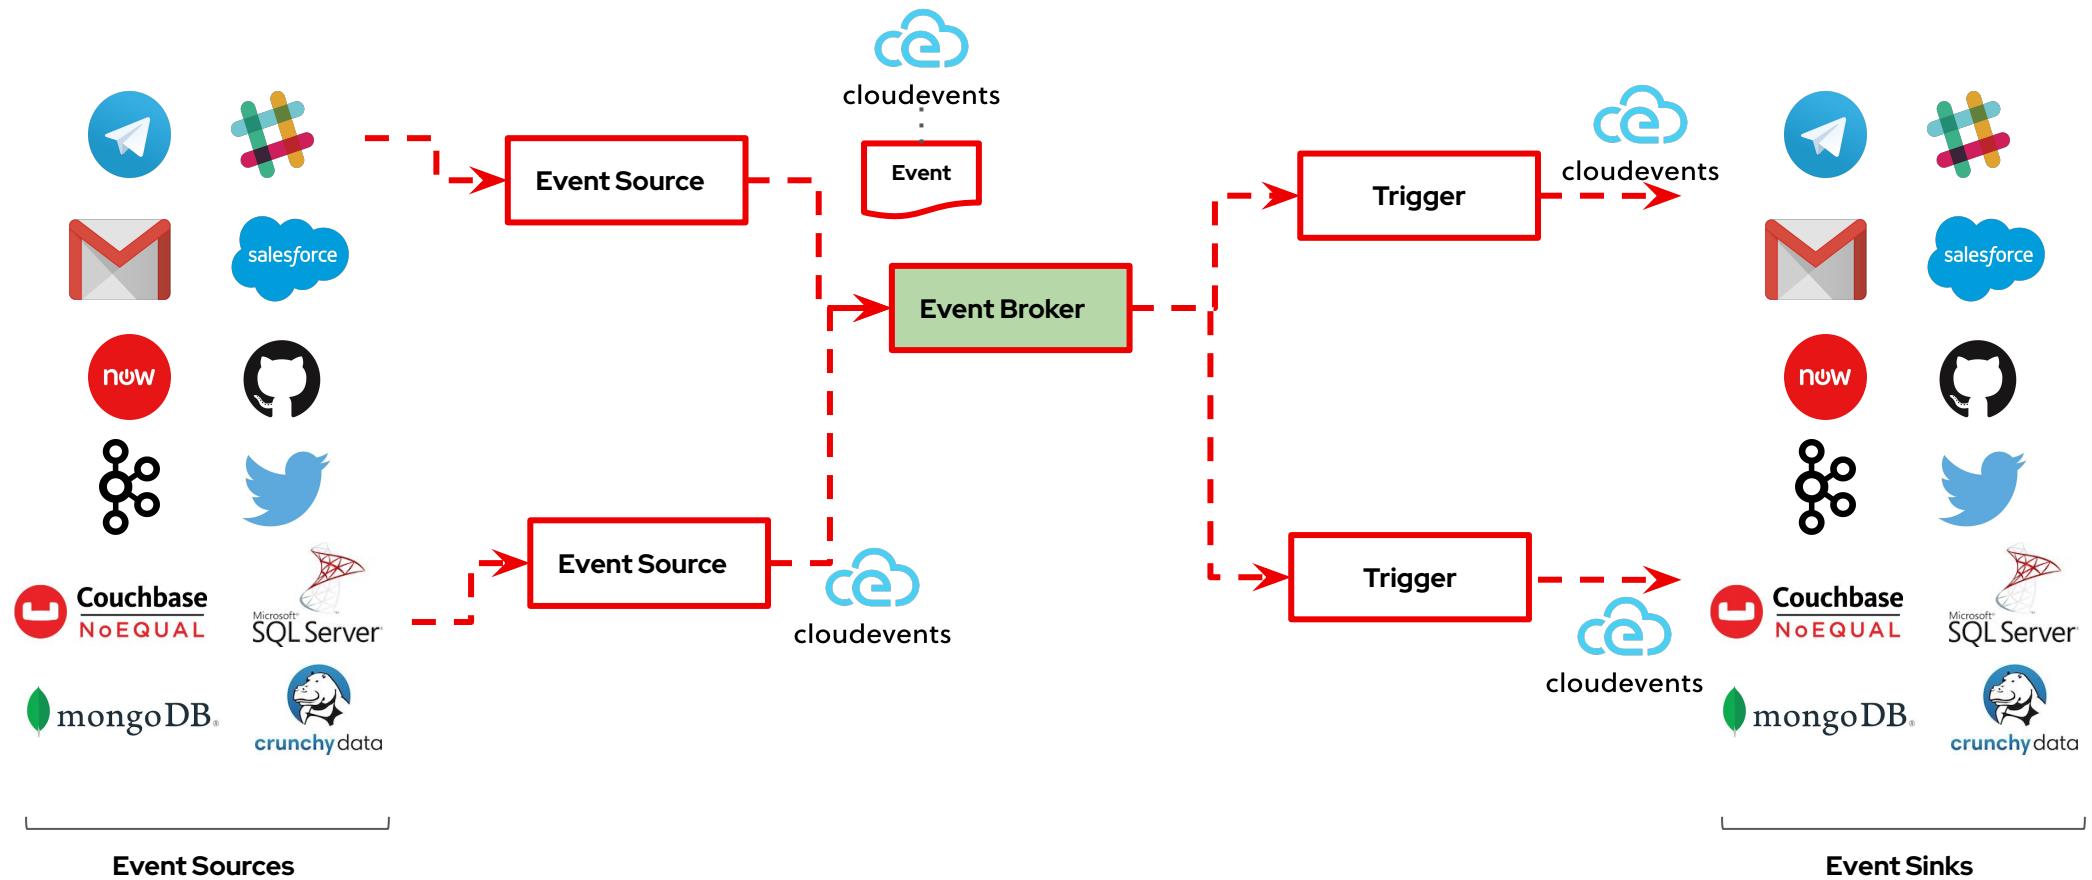
# From Deployment to KService

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: random
spec:
  replicas: 1
  selector:
    matchLabels:
      app: random
  template:
    metadata:
      labels:
        app: random
    spec:
      containers:
        - image: rhuss/random
          name: random
      ports:
        - containerPort: 8080
```

```
apiVersion: serving.knative.dev/v1
kind: Service
metadata:
  name: random
spec:
  replicas: 1
  selector:
    matchLabels:
      app: random
  template:
    metadata:
      labels:
        app: random
    spec:
      containers:
        - image: rhuss/random
          name: random
      ports:
        - containerPort: 8080
```

No more K8s  
**Service** or  
**Ingress**  
required!

# Event Driven Applications



# Google Cloud Run



- Two variations:
  - Google Cloud Run → Knative API implemented on App Engine
  - Google Cloud Run For Anthos → Knative on Kubernetes
- Events for Cloud Run on Anthos
  - Knative Eventing on GKE
- Eventarc (inspired by Eventing)
- Not all Knative features implemented
  - see <https://ahmet.im/blog/cloud-run-is-a-knative> (?)

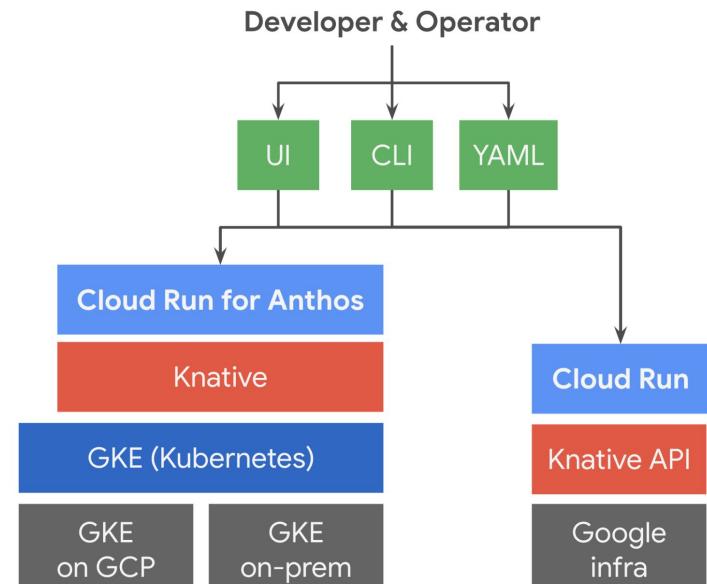
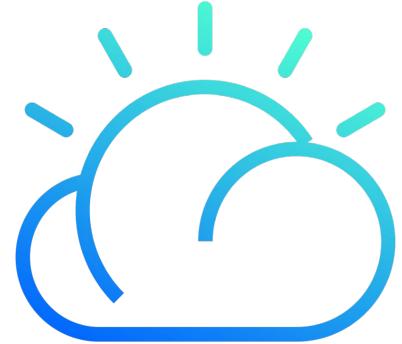


Diagram: @ahmetb

# IBM Cloud Code Engine



- Knative on the IBM Cloud
- Runs on IKS
- KafkaSource support
- Support for Batch Jobs
- In-Cluster container builds via Cloud Native BuildPacks

# Cloud Native Runtimes for VMware Tanzu

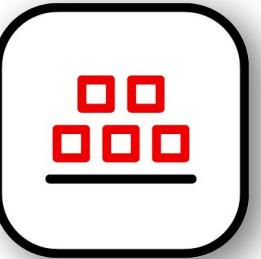


- Knative distribution on-top of Kubernetes
- Integral part of the Tanzu Application Platform (TAP)
- Eventing Broker is backed by RabbitMQ
- Includes TriggerMesh Sources for Amazon Web Services (SAWS)



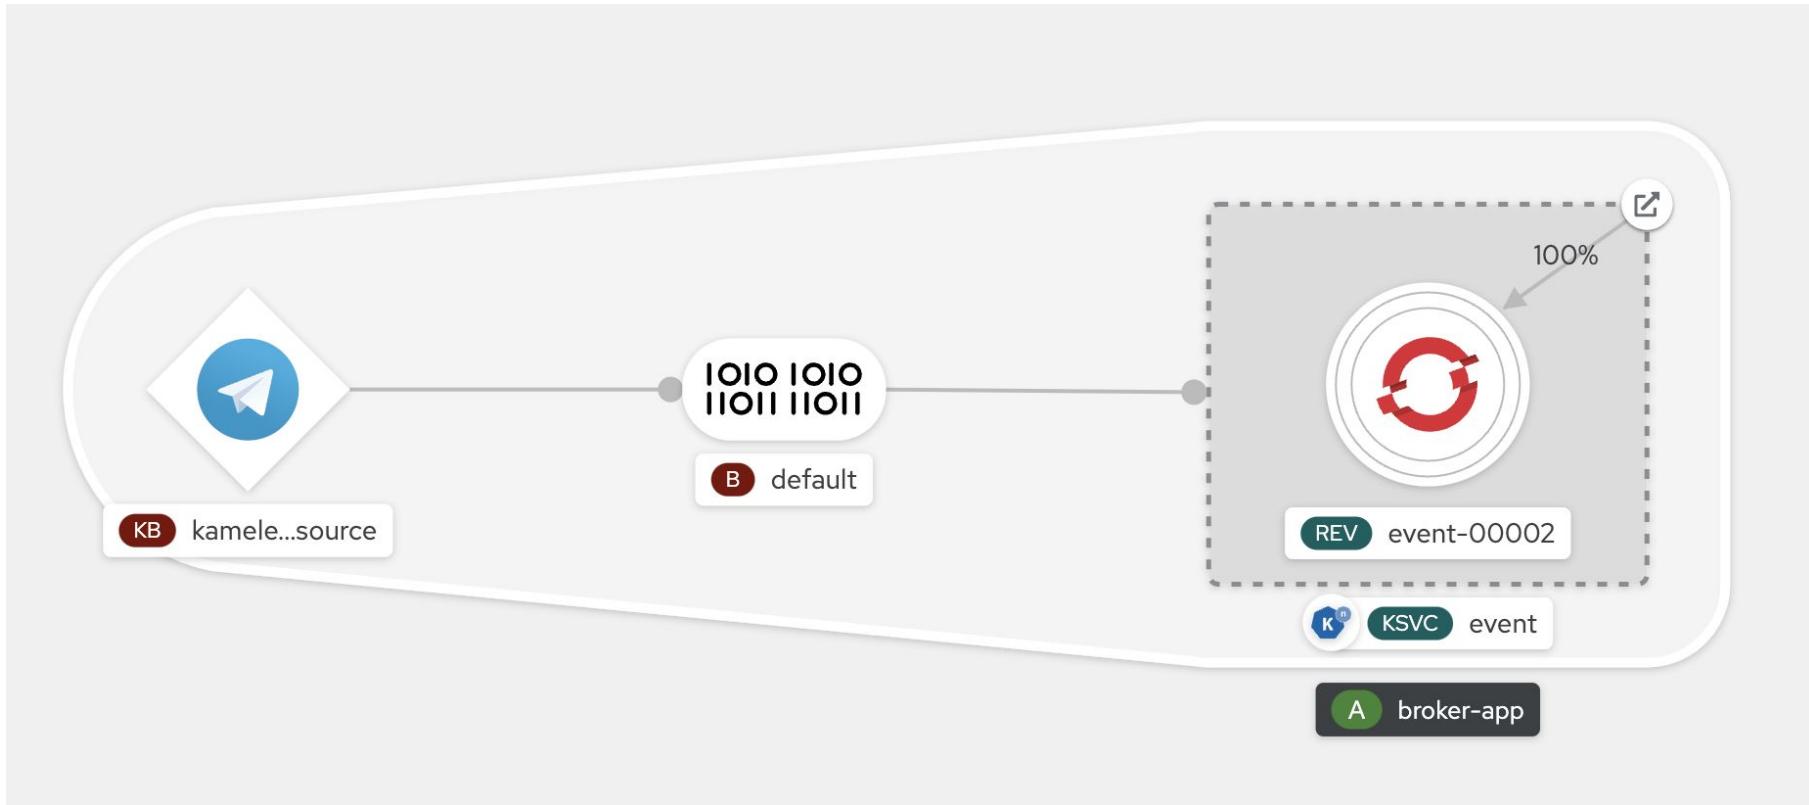
# Red Hat OpenShift Serverless

- Knative bundled with Red Hat OpenShift
- Runs everywhere where OpenShift runs
  - on-prem
  - AWS (ROSA)
  - Azure (ARO)
  - IBM (ROKS)
- Adds operator-based installation experience
- Kafka-backed Knative broker
- Kafka EventSource and Kafka EventSink
- 50+ sources via Apache Camel connectors (Kamelets)



# Red Hat OpenShift Serverless

- Great visualization in the OpenShift Console



# Challenges



# Coldstart

- Application startup times in Kubernetes:

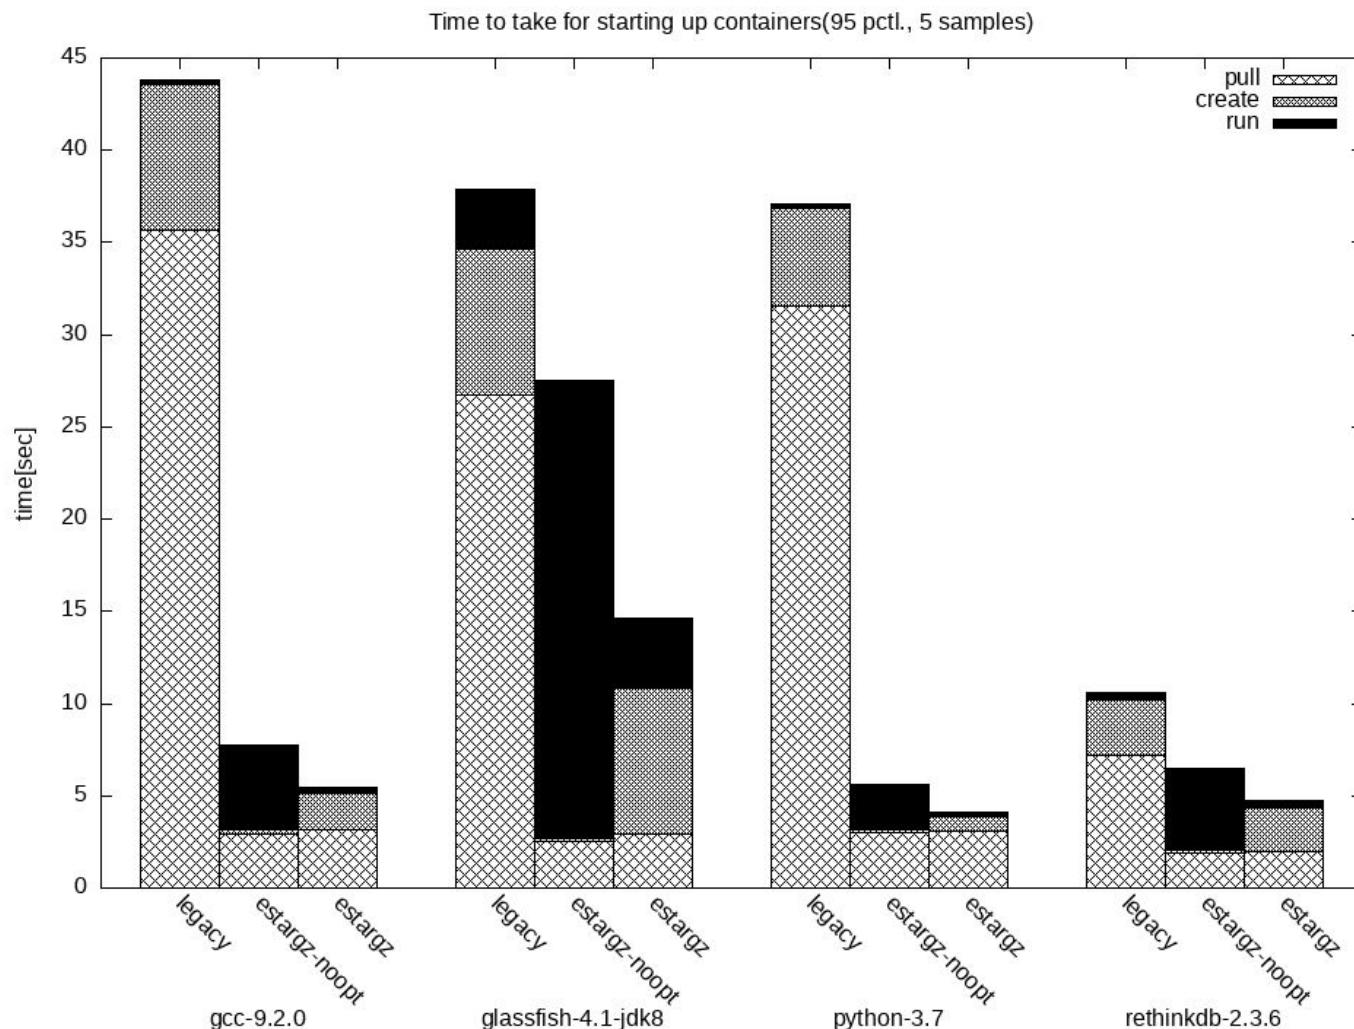
Pod Scheduling 0.1 - 10s	Image Pull 0- 10s	Container Startup <1s	Application Startup 0.01 - 5s	Pod Ready 0.1-1s
-----------------------------	----------------------	-----------------------------	----------------------------------	---------------------

- Scheduling time depends on cluster size and can take [10s for a 1000 nodes cluster](#)
- Application startup depends on the application stack
  - Fast for optimized frameworks like Quarkus or Node
  - JavaEE servers are no good fit
- Pod serves requests when readiness probes succeed
  - Kubelet has a 1s granularity for checking probes

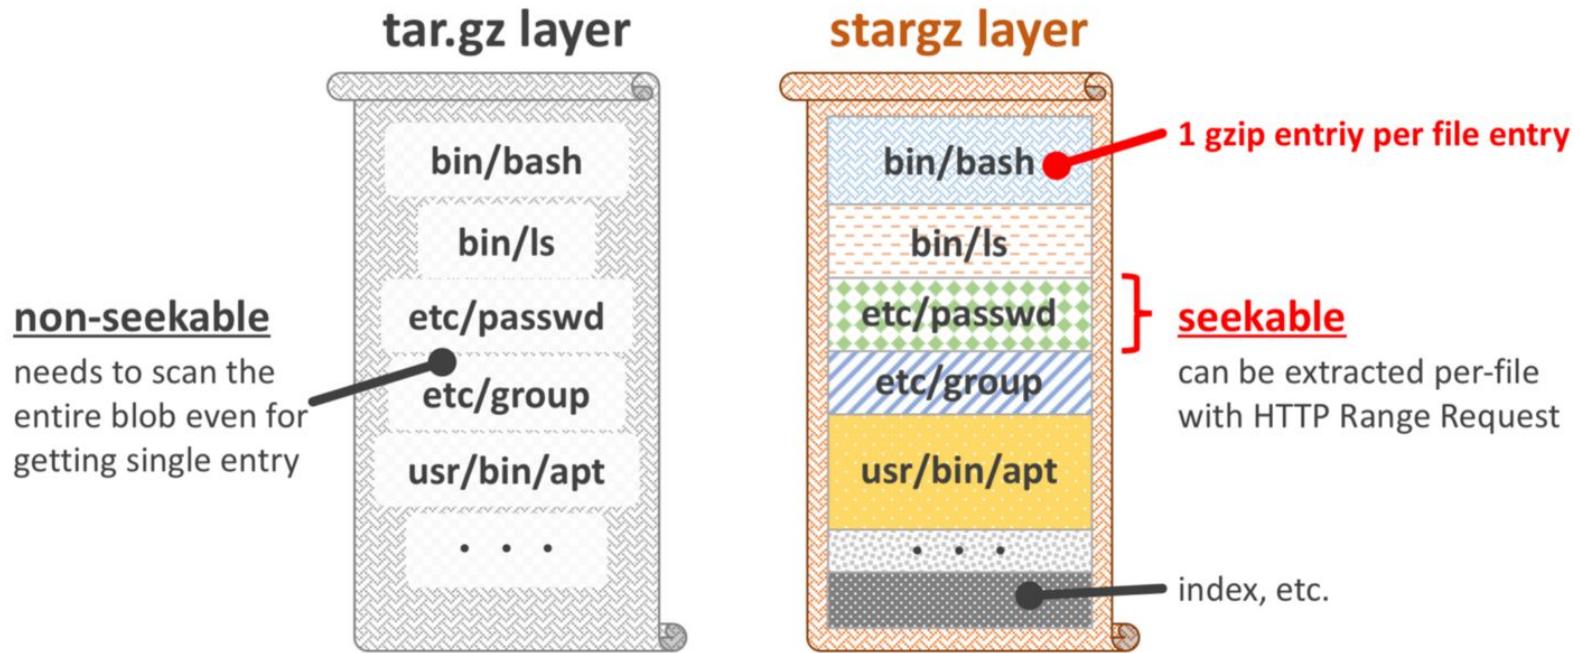
# Image Pull

- Application image needs to be pulled from a container image registry
  - Cached at the Container Runtime for future usage
  - Workarounds: Pre-warm with a Daemonset
- Lazy pulling with stargz snapshotter
  - Files from layers only loaded when referenced
  - eStargz: Measure order of files access and rearrange the image

# containerd/stargz-snapshotter

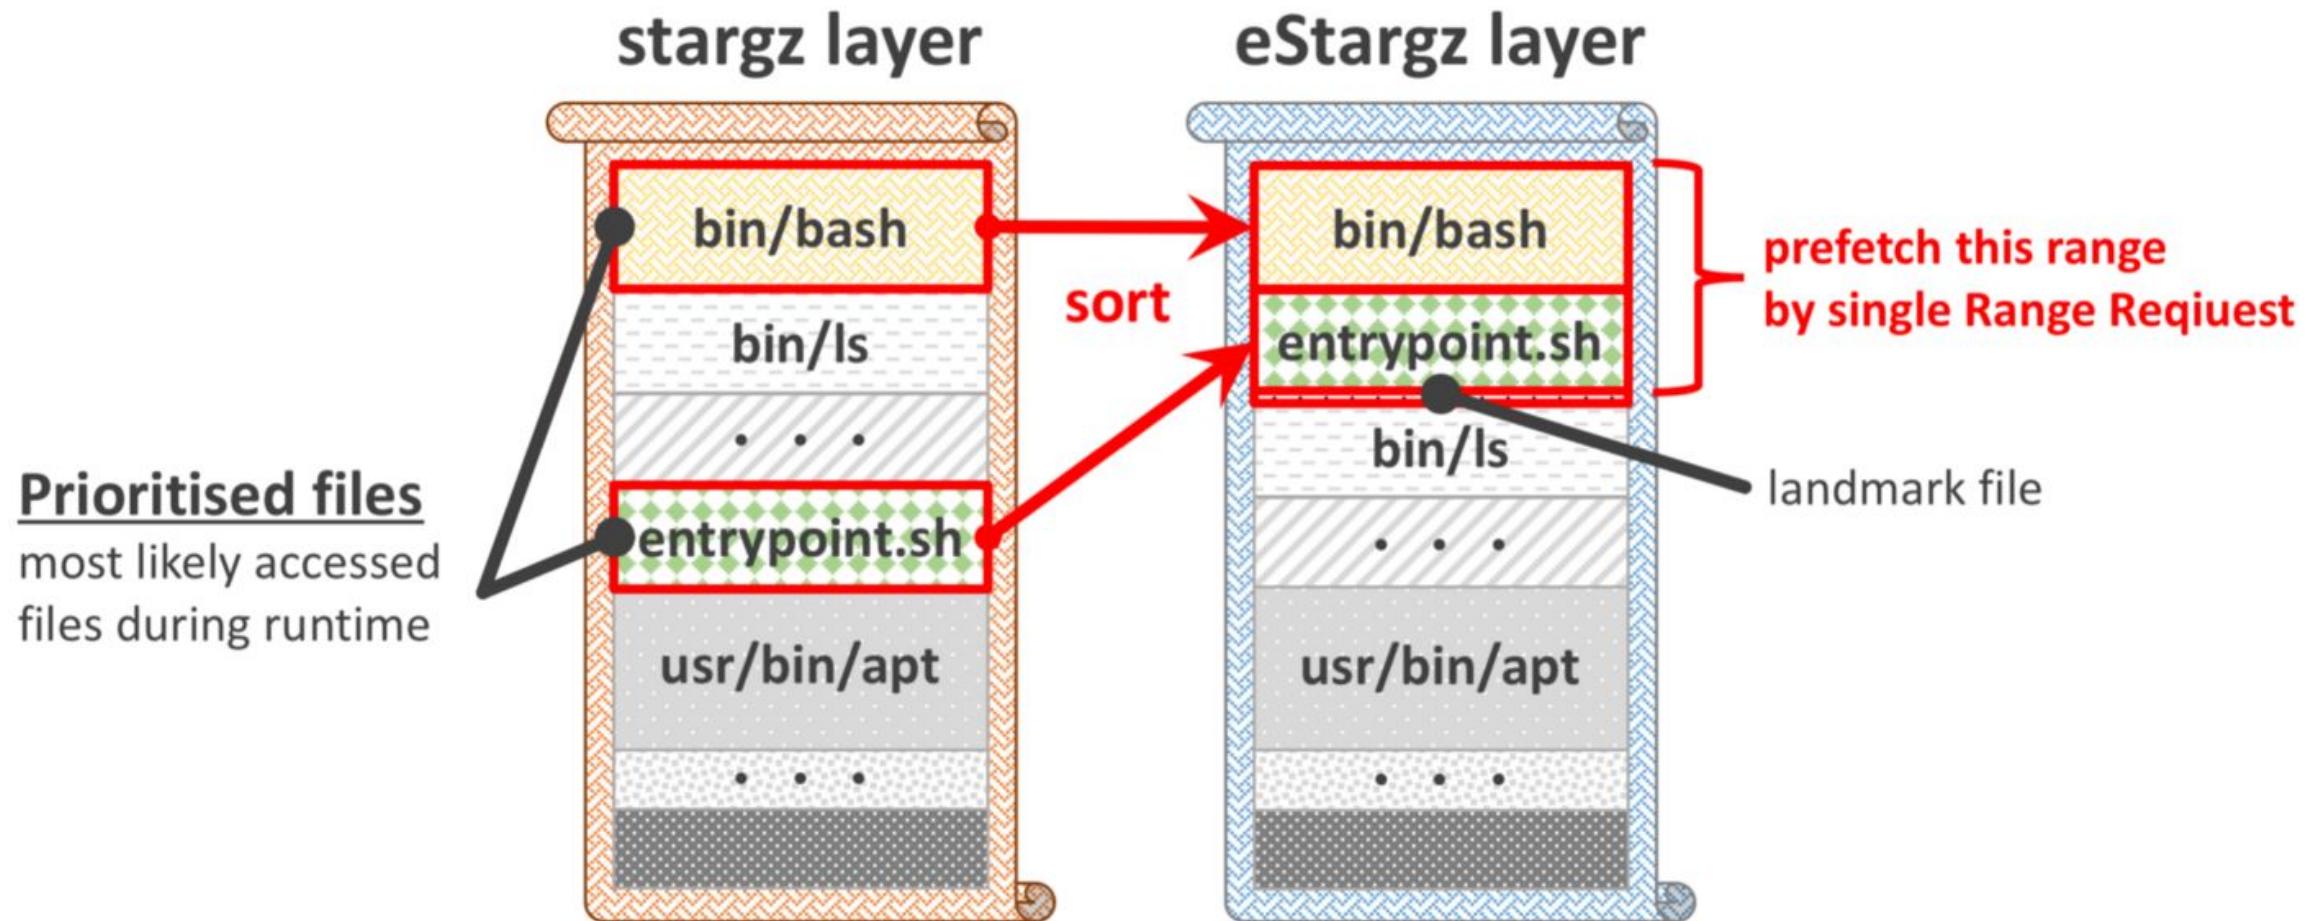


# Stargz (“seekable tar.gz”)



- tar: Non-indexed, gz: Non-Seekable
- Solution: Gzip tar-entries individually, concatenate and add gzipped-index file and tar footer

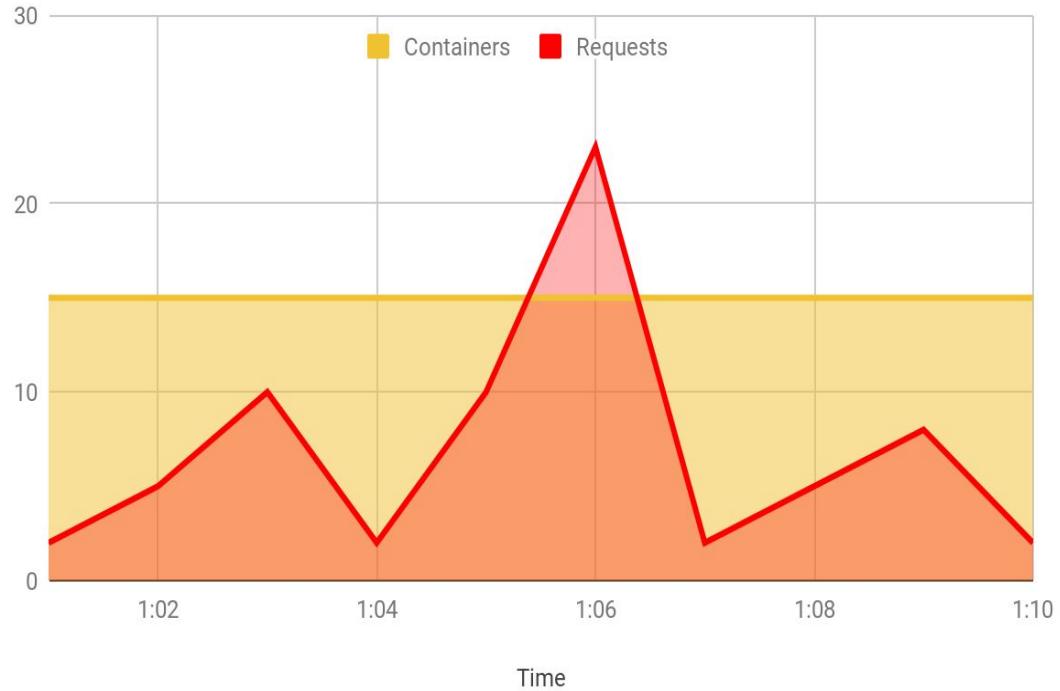
# eStargz Optimization



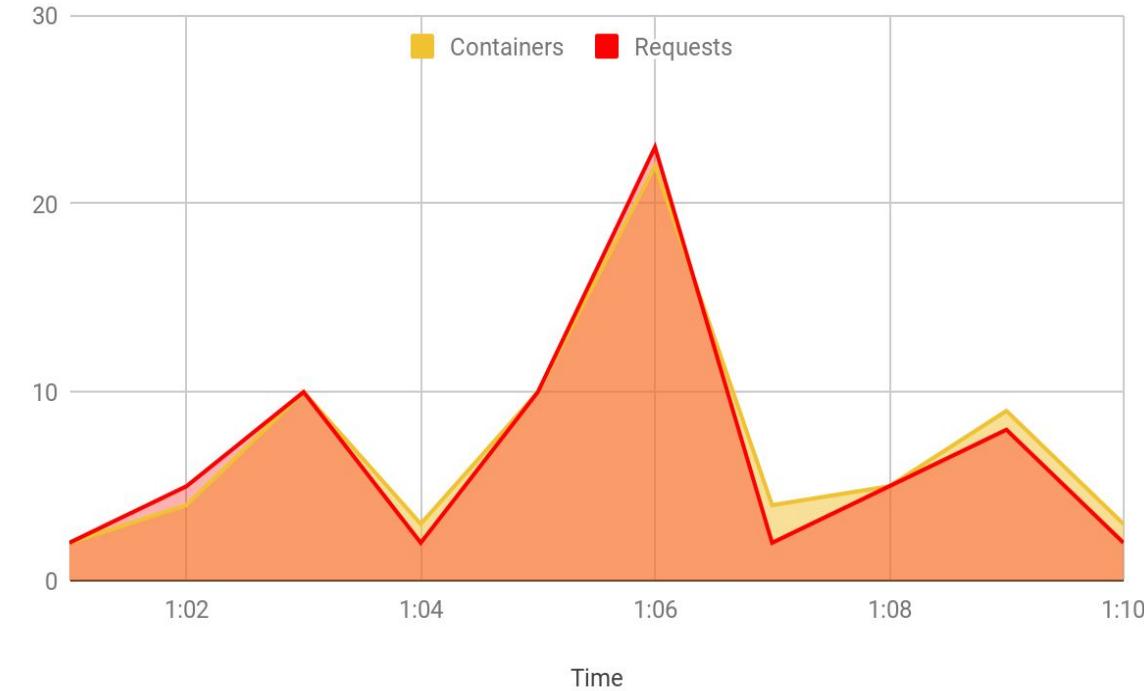
# Possible Startup Optimizations

- Pooling of network interfaces
- Move the Pod Lifecycle Event Generator (PLEG) to an event-driven model instead of 1s polling (requires change to the CRI)
- Precreated pool of “Pod hulls”
- More pooling in general
- See [Cold start latency improvements in Kubernetes and Knative](#) for more details

# Autoscaling



NOT Serverless

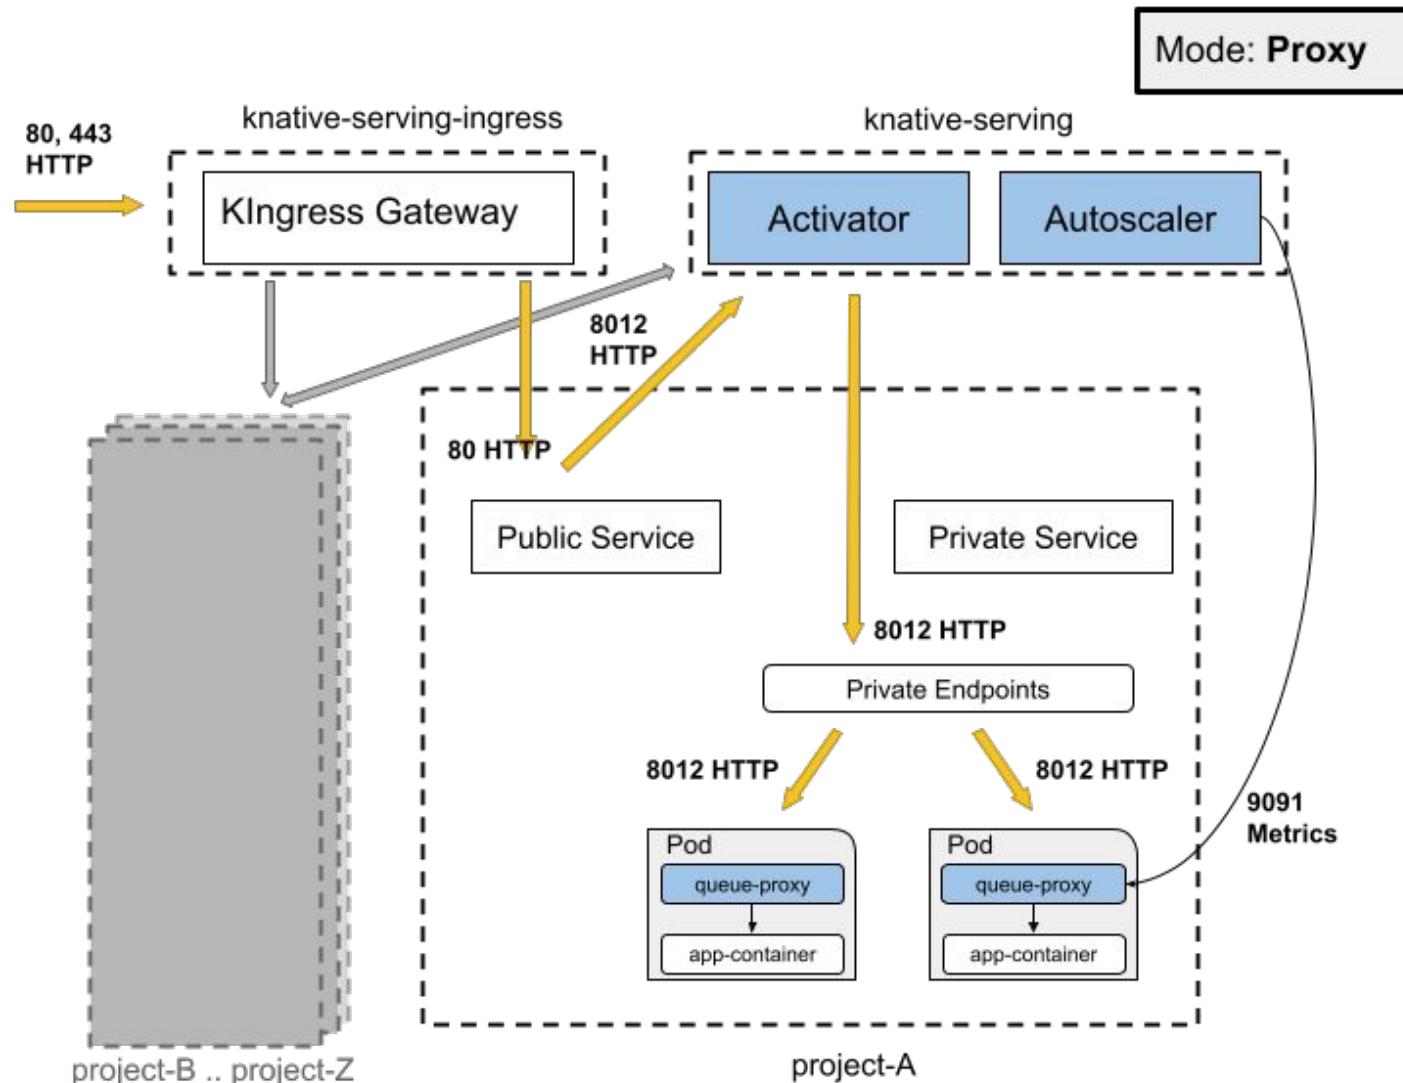


with Serverless

# Autoscaling on Kubernetes

- Kubernetes Native:
  - Horizontal Pod Autoscaler (HPA)
  - Vertical Pod Autoscaler (VPA)
  - Cluster Autoscaler (CA)
- Knative Autoscaler
- KEDA

# Knative Autoscaling



# Pros & Cons



# Container Drawbacks

- Container Image needs to be **built**
  - Explicit via Docker/Buildah tooling
  - Automatically by language specific tooling
    - [Cloud Native Buildpacks](#), [Shipwright](#), [Eclipse JKube](#) ...
  - Knative Functions
- **Coldstart** times
- Extra overhead because of **additional abstraction**

# Container Benefits

- **Standard** format (OCI)
- **No vendor locking** per se
  - Multitude cloud offering for container deployments
- No language restrictions
- Rich tooling landscape
- Hybrid scenarios (on-prem combined with cloud) easily possible
- Versioning
- Powerful orchestration platforms underneath (Kubernetes)

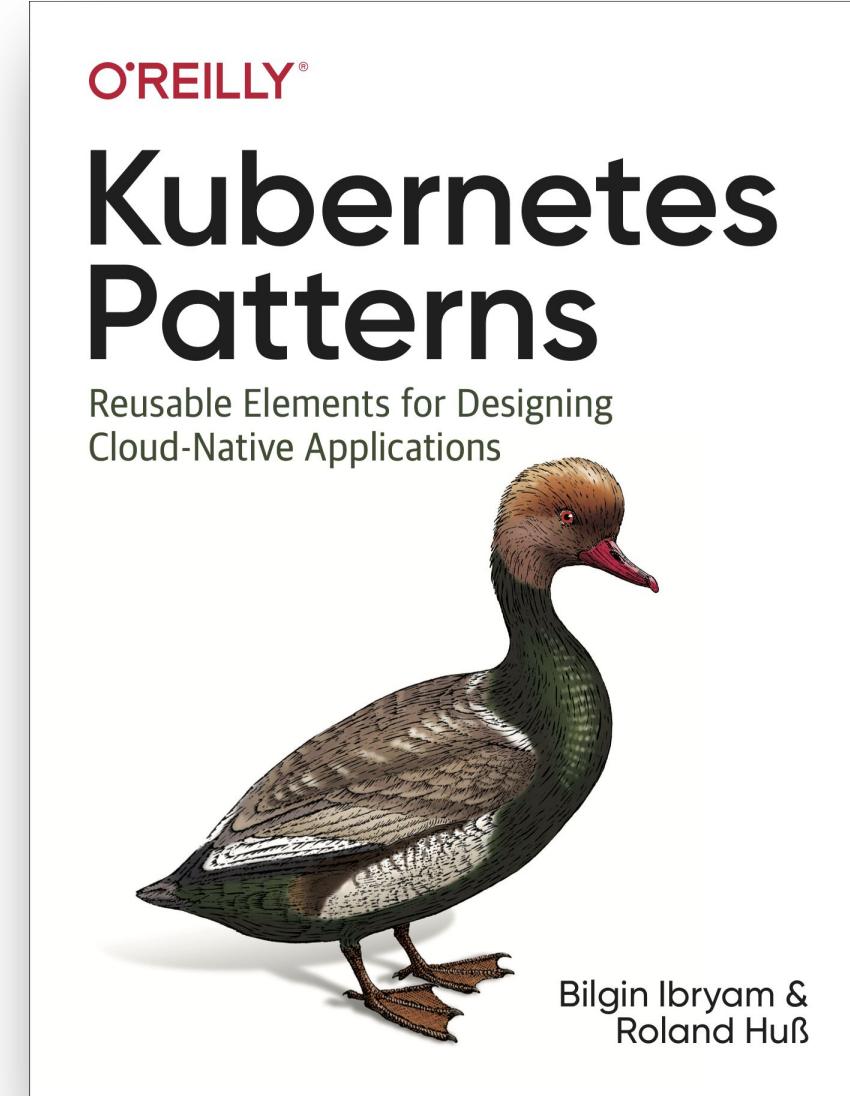
# The Future



# What's (maybe) next ?

- **Web Assembly** (WASM) as a popular deployment format for serverless functions
- **Serverless Kubernetes Control-Plane**
  - so that optional features don't use resources
- **Sub-second cold start** times become a reality
- **Scale-to-Zero** as a **core** Kubernetes feature
- **Hybrid Serverless** with multi tenancy

# Thank you



<https://k8spatterns.io>

# Picture Credits

<https://www.pexels.com/photo/boat-island-ocean-sea-218999/>

<https://unsplash.com/photos/UGMf30W28qc>

<https://pixabay.com/photos/hamburg-speicherstadt-channel-297671/>

<https://pixabay.com/photos/beer-machine-alcohol-brewery-1513436/>

<https://www.pexels.com/photo/people-art-abstract-blur-6077519/>

<https://www.pexels.com/photo/person-holding-a-glass-ball-1252893/>

<https://me.me/i/aws-lambda-is-just-glorified-cgi-bin-imgflip-com-change-my-mind-d0b715592ba34b08b79452ad02783ca2>

<https://unsplash.com/photos/vkQgb1lZZPQ>

<https://unsplash.com/photos/9JrBiphz0e0>

<https://www.pexels.com/photo/person-rock-climbing-3077882/>