



Elastic Kubernetes

Zbyněk Roubalík **@zroubalik**

Principal Software Engineer, Red Hat

OpenShift Serverless & KEDA

Roland Huß **@ro14nd@hachyderm.io**

Senior Principal Software Engineer, Red Hat

OpenShift Serverless Architect



Why “Elastic Kubernetes”?



- Workloads have different size and resources requirements
 - **Under provisioning** - not enough resources might cause outages, latency and overall bad UX
 - **Over provisioning** - too many resources, expensive, wasteful
- It's hard to specify resource and pod allocation, as resources requirements change across the time -> we need elasticity
 - Applications autoscaling
 - Cluster autoscaling

Why "Elastic Kubernetes"?

Resources defined statically

vs.

Enabled autoscaling



Resource requests & limits

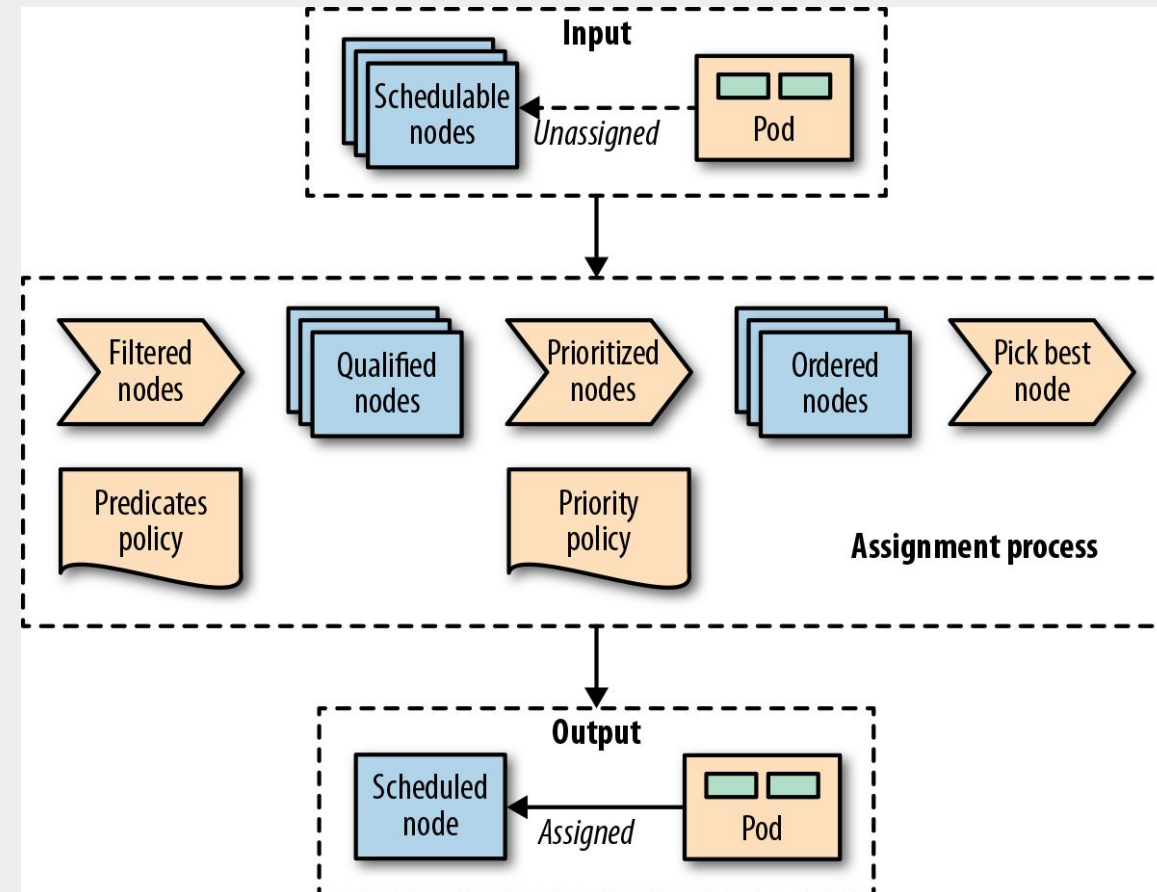
- It's recommended to **always set memory resource requests and limits**
 - and requests == limits
- It's recommended to **always set cpu resource requests**
 - and no limits
- Every workload has a different set of requirements
- Capacity planning is important aspect
- If not sure about correct values -> collect metrics over time and modify gradually

Pods:

- Resource **Requests** - the minimal amount of resources needed for the Pod
- Resource **Limits** - the maximal amount of resources that can be used by Pod

Kubernetes scheduler

- Responsible for assigning Pods to nodes
- **Predicates:** Filter of suitable nodes
- **Priorities:** Ordering of nodes according to preference
- Predicates + Priorities = Scheduler Policy






A photograph of a wooden boat's interior, looking towards the bow. The boat is on a dark blue sea. In the background, there are several large, rocky islands under a blue sky with white clouds. The text "Horizontal Pod Scaling" is overlaid in white.

Horizontal Pod Scaling

Horizontal Pod Scaling

- **Scale out/in** operation
- Increasing/decreasing the number of replicas (Pods)
- Application deployment or resource requests & limits don't change



Manual Horizontal Pod Scaling

- Specifying the number of replicas (Pods) manually
- Can be declared:
 - Imperatively (kubectl)
 - Declaratively (yaml)

```
$ kubectl scale my-app --replicas=5
```


Horizontal Pod Autoscaling

- **Horizontal Pod Autoscaler (HPA)**
 - Built-in Kubernetes component
- <https://kubernetes.io/docs/tasks/run-application/horizontal-pod-autoscale/>

Example:

Application performing a resource expensive task


- Application is deployed as standard Kubernetes Deployment



Example:

Application redesigned to utilize Horizontal Pod Autoscaler

- Application remains the same and is being deployed the same way
- Autoscaling via HPA: based on CPU & Memory consumption



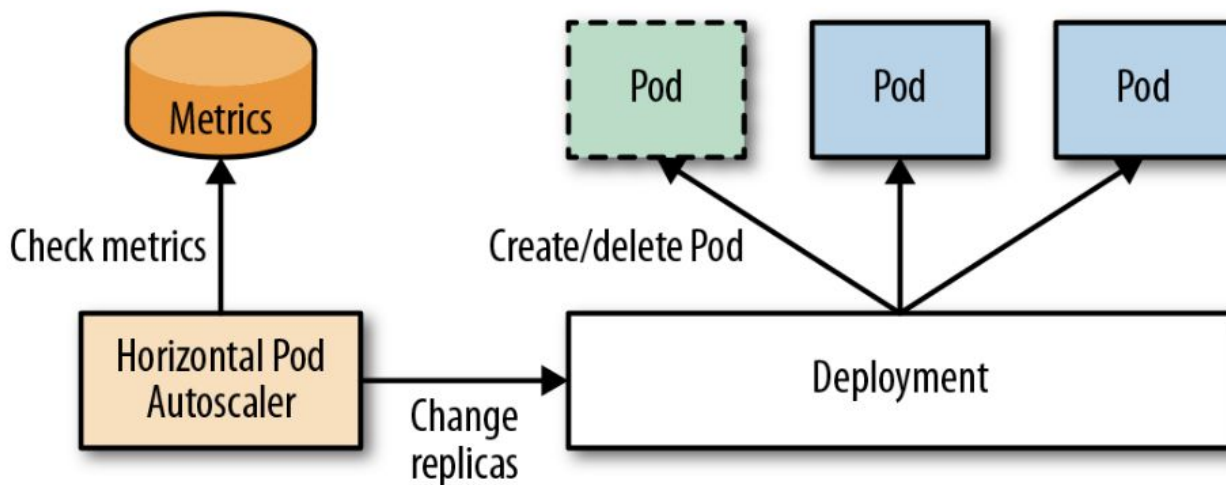
Horizontal Pod Autoscaler

- Needs Kubernetes Metrics Server enabled
 - it is a cluster wide aggregator of resource usage data
- Scales Deployments, StatefulSets and Custom Resources that enable the **/scale** subresource
- Target workload needs to specify **resource limits**
- Metric Types:
 - **Resource metrics** – cpu/memory utilization (built-in)
 - **Custom / External Metrics** – metrics about custom resources
- Can **not scale to 0** (for custom metrics currently in development)

Horizontal Pod Autoscaler

- HPA operates on the ratio between desired metric value and current metric value:

```
desiredReplicas = ceil(  
    currentReplicas * (currentMetricValue / desiredMetricValue)  
)
```



HPA

- Can be declared:
 - Imperatively (kubectl)
 - Declaratively (yaml)
- Multiple metrics can be defined

```
apiVersion: autoscaling/v2
kind: HorizontalPodAutoscaler
metadata:
  name: example-hpa
spec:
  scaleTargetRef:
    apiVersion: apps/v1
    kind: Deployment
    name: my-app
  minReplicas: 1
  maxReplicas: 10
  metrics:
  - type: Resource
    resource:
      name: cpu
      target:
        type: Utilization
        averageUtilization: 50
```


```
$ kubectl autoscale deployment my-app --cpu-percent=50 --min=1 --max=10
```




Vertical Pod Scaling

Vertical Pod Scaling

- **Scale up/down** operation
- Increasing/decreasing the amount of resources assigned to the Pod
- Number of replicas doesn't change



Vertical Pod Autoscaling

- **Vertical Pod Autoscaler (VPA)**
 - Add-on that needs to be installed on cluster
- <https://github.com/kubernetes/autoscaler/tree/master/vertical-pod-autoscaler>

Example:

Application performing a resource expensive task


- Application is deployed as standard Kubernetes Deployment



Example:

Application redesigned to utilize Vertical Pod Autoscaler (VPA)

- Resources assigned to the application has been controlled by VPA based on resource utilization
- Autoscaling via VPA: based on CPU & Memory consumption

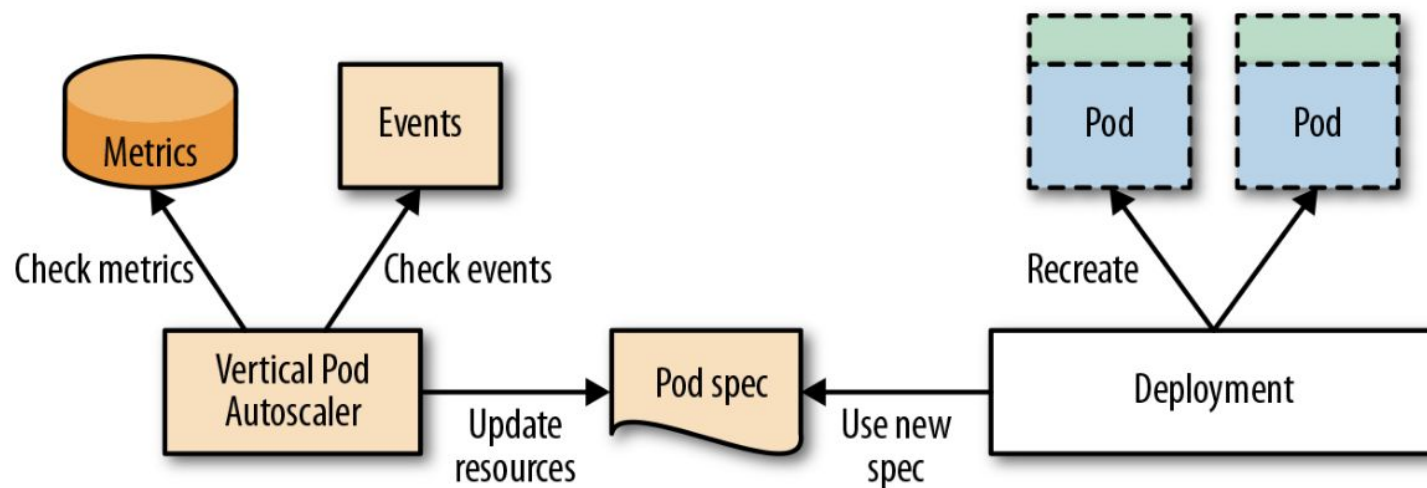


Vertical Pod Autoscaler

- **Determines** resource limits/requests based on historic and current metrics
- Containers might be **rescheduled** on a different node based on VPA limit recommendations
- Can **not be used** together with **HPA**
- Four modes:
 - **Auto/Recreate** - automatically apply the VPA resource recommendations
 - **Initial** - apply the VPA recommendations only at pod creation
 - **Off** - only provides the VPA recommendations in the status section

Vertical Pod Autoscaler

- **Recommender** – monitors resource utilization and computes target values
- **Updater** – evicts those pods that need the new resource limits
- **Admission Plugin** – sets the correct resource requests on new pods



VPA

```
apiVersion: autoscaling.k8s.io/v1
kind: VerticalPodAutoscaler
metadata:
  name: example-vpa
spec:
  targetRef:
    apiVersion: apps/v1
    kind: Deployment
    name: my-app
  updatePolicy:
    updateMode: "Off"
```

```
$ kubectl describe vpa example-vpa
```

```
...
status:
  recommendation:
    containerRecommendations:
      - containerName: my-container
        lowerBound:
          cpu: 25m
          memory: 262144k
        target:
          cpu: 25m
          memory: 262144k
        uncappedTarget:
          cpu: 25m
          memory: 262144k
        upperBound:
          cpu: 262m
          memory: "274357142"
```




Demand based Autoscaling

Demand Based Autoscaling




- **Knative Autoscaler**
 - Core component of Knative Serving
 - CNCF project
- <https://knative.dev/docs/serving/autoscaling/autoscaler-types/#knative-pod-autoscaler-kpa>

Example:

Application serving HTTP requests

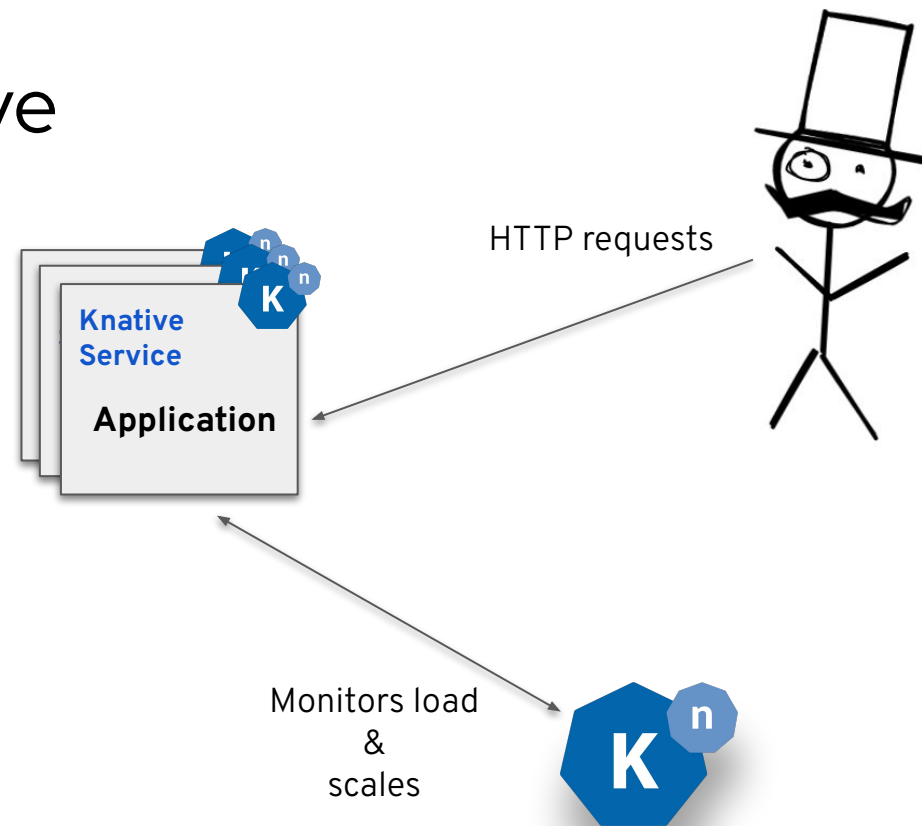
- Application is deployed as standard Kubernetes Deployment
- Can be autoscaled only via standard k8s HPA: CPU & Memory
- No demand based autoscaling



Example:

Application redesigned to utilize Knative

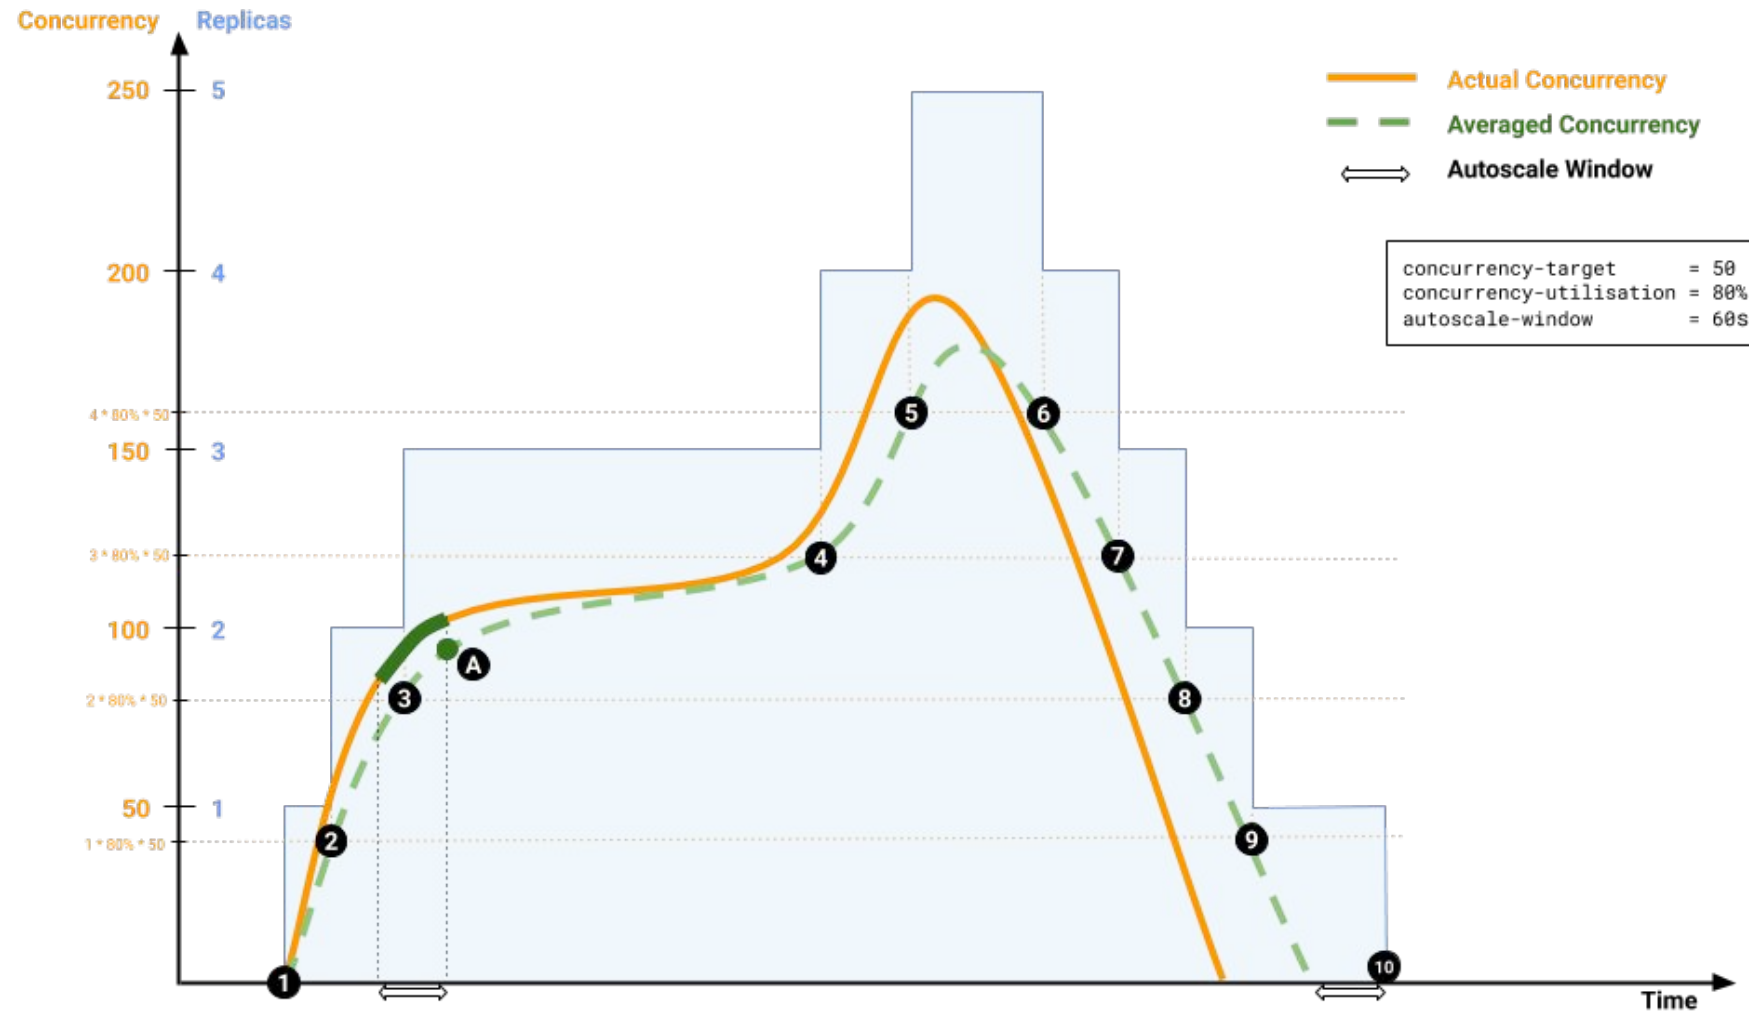
- Application is deployed as Knative Service
- Knative Autoscaler monitors load and scales the application based on demand



Knative Autoscaler Concepts

- Knative Autoscaler scales **Knative Service**, a CR representing the workload, it manages needed Kubernetes resources (Deployment, Service, Ingress,...)
- **Activator** component enables scale to 0
 - Incoming requests are being hold until the app is scaled to 1 replica
- Autoscaler itself has 3 components:
 - **PodAutoscaler Reconciler** - ensures that all components are up to date
 - **Collector** - collect metrics from various sources
 - **Decider** - based on metrics decides how the app should be scaled
 - `want = concurrencyInSystem/targetConcurrencyPerInstance`

Knative Autoscaler



1 Scale up from 0 to 1 replica on first request.

2 Scale from 1 to 2 replicas if the utilisation 80% of the concurrency target 50 is reached for the **averaged concurrency**.

3 ... 9 Up- and downscaling events when **averaged concurrency** crosses the utilisation threshold counted across the current number of replicas. ($2 * 80\% * 50 = 80$, $3 * 80\% * 50 = 120$, ...)

10 Scale down to 0 when **averaged concurrency** is going down to 0 for the length of the autoscale window.

A The **averaged concurrency** is calculated every 2 seconds by averaging **concurrent requests** for the past auto-scale window length (default: 60s)

Knative Service

- Application deployed as Knative Service
- Can be declared:
 - Imperatively (kn cli)
 - Declaratively (yaml)

```
apiVersion: serving.knative.dev/v1
kind: Service
metadata:
  name: example
spec:
  template:
    spec:
      containers:
        - image: johndoe/my-image
          ports:
            - containerPort: 8080
```

```
$ kn service create example --image johndoe/my-image --port 8080
```


A large crowd of people is seen from behind, with their arms raised in the air, facing a stage. The stage is illuminated with bright yellow and blue spotlights that create a hazy atmosphere. A semi-transparent dark rectangle is overlaid in the center of the image, containing the text "Event Driven Autoscaling" in white. On the left and right sides of the stage, there are circular signs with text: "ES RDEN ES" on the left and "LE ARD T E" on the right.

Event Driven Autoscaling

Event Driven Autoscaling

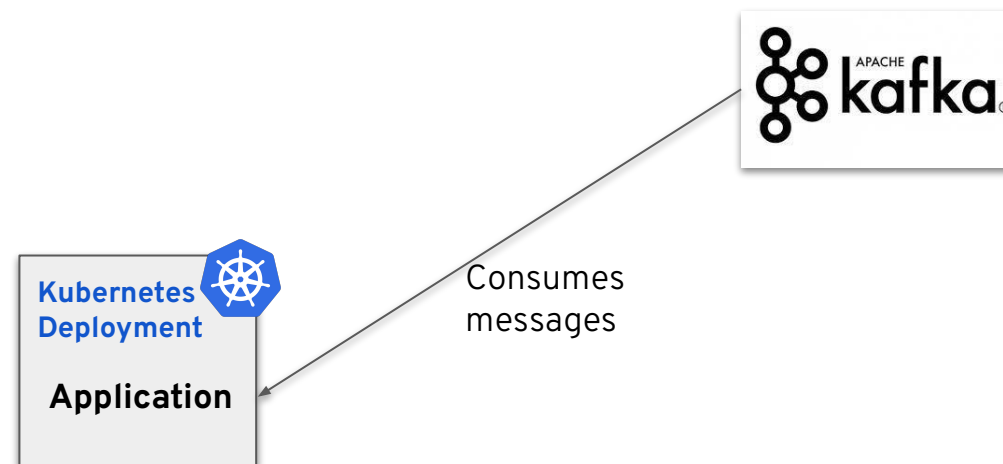


- **K**ubernetes **E**vent **D**riven **A**utoscaling dead simple
 - CNCF project
- <https://keda.sh>

Example:

Application consuming messages from Kafka topic

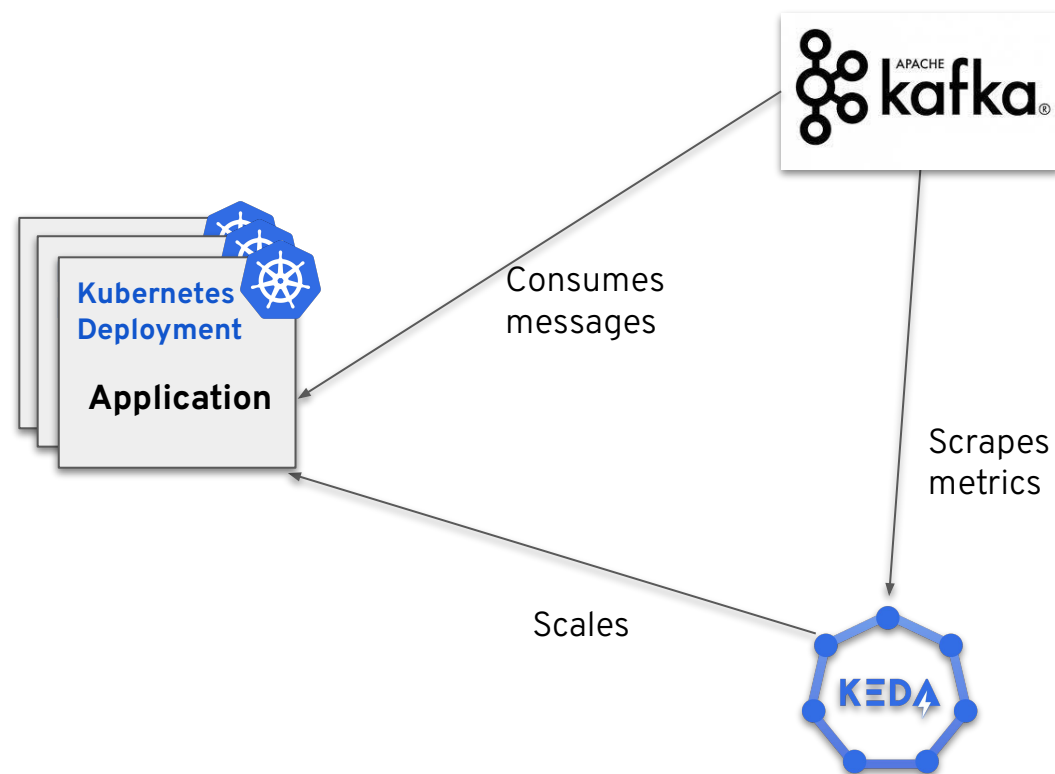
- Application is deployed as standard Kubernetes Deployment
- Can be autoscaled only via standard k8s HPA: CPU & Memory
- No event-driven autoscaling



Example:

Application redesigned to utilize **KEDA**

- Application remains the same and is being deployed the same way
- Event driven autoscaling enabled through KEDA

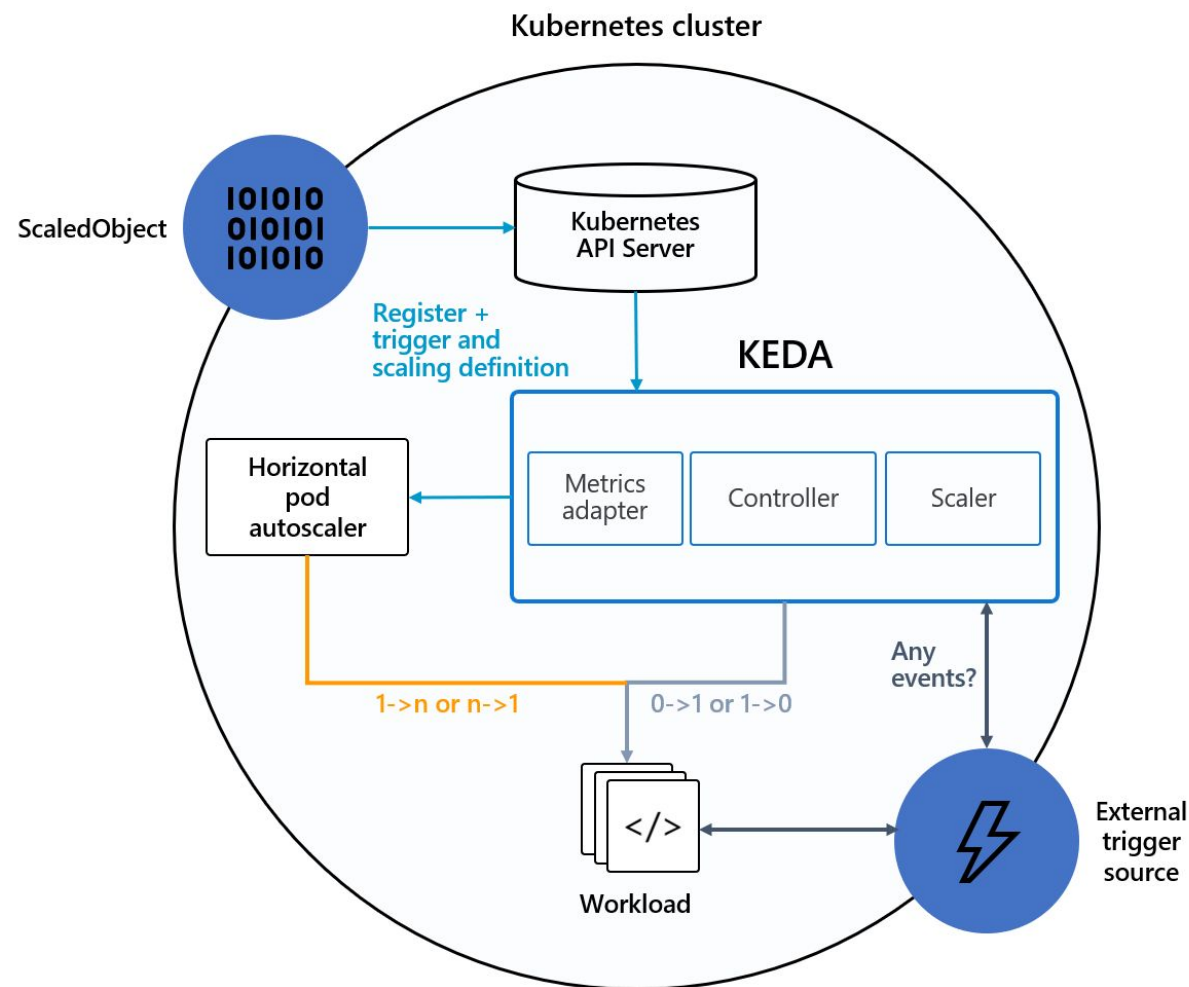


KEDA Concepts

- Automatically scale Kubernetes Deployments, Jobs & Custom Resources
- Provides **50+** built-in scalers, but users can build own external scalers
 - Kafka, Prometheus, RabbitMQ, AWS services, Azure Services,...
- Scale resources **based on events** in the target scalers, eg. messages in Kafka topic
- KEDA does not scale on HTTP requests
- Save resources by **scale to 0**
- KEDA **does not** manipulate the data, just scales the workload

KEDA Architecture

- KEDA is built on top of Kubernetes
- Use **ScaledObject/ScaledJob** to define scaling metadata
- Manages workloads to scale to 0
- Registers itself as Kubernetes Metric Adapter
- Provides external metrics for Horizontal Pod Autoscaler (HPA)



ScaledObject

- Can target Deployment, StatefulSet or Custom Resource with /scale
- Multiple scalers can be defined as triggers for the target workload
- User can specify HPA related settings to tweak the scaling behavior

```
apiVersion: keda.sh/v1alpha1
kind: ScaledObject
metadata:
  name: example-so
spec:
  scaleTargetRef:
    name: example-deployment
  minReplicaCount: 0
  maxReplicaCount: 100
  triggers:
  - type: kafka
    metadata:
      bootstrapServers: kafka.svc:9092
      consumerGroup: my-group
      topic: test-topic
      lagThreshold: '5'
```

Demo

Cluster Autoscaling


A scenic view of a city canal, likely in Copenhagen, featuring historic brick buildings with arched windows and a bridge over the water. The sky is a warm orange and yellow, suggesting sunset or sunrise. The text 'Cluster Autoscaling' is overlaid in white on a semi-transparent dark band across the middle of the image.

Cluster Autoscaling

- When autoscaling of applications is not enough
- Autoscaling of applications (more Pods and Containers) makes pressure on the infrastructure as well
- We should scale Kubernetes Nodes together with applications to really achieve “Elastic Kubernetes”

Cluster API & Cluster Autoscaler

- Standardized around **Cluster API**:
 - To manage the lifecycle of Kubernetes conformant clusters using a declarative API.
 - To work in different environments, both on-premise and cloud.
- There are cloud vendor specific implementations of cloud autoscalers
- <https://cluster-api.sigs.k8s.io/>




Cluster Autoscaler

- Adjusts the number of Nodes in the cluster when Pods fail to schedule or when nodes are underutilized
- Makes decision based on requests and limits for CPU and memory resources, **not the actual load!**
- Best practices:
 - Set correct resources requests & limits
 - Use PodDisruptionBudgets to prevent pods from being deleted too abruptly
 - Check if your cloud provider's quota is big enough

Cluster Autoscaler

- Unschedulable pods makes pressure on Cluster Autoscaler -> new Node is provisioned
- In case Node is not needed -> Cluster Autoscaler terminates the underlying instance in a cloud-provider-dependent manner



Cluster Autoscaler

- Technique: Keep an empty spare node in the cluster to reduce provisioning time
- Shrinking the cluster is hard as it requires rebalancing the cluster.
- Works on a different time-scale than application auto-scaling
 - it takes much more time to spin up a new node than a new pod



Summary

Summary

- **Why Elastic Kubernetes**
- **Application Autoscaling**
 - HPA
 - VPA
 - Knative Autoscaler
 - KEDA
- **Cluster Autoscaling**

Thank you

 @ro14nd@hachyderm.io

 @zroubalik

Picture Credits

<https://kubernetes.io/docs/concepts/scheduling-eviction/scheduling-framework/>

<https://www.pexels.com/photo/boat-island-ocean-sea-218999/>

<https://pixabay.com/photos/hamburg-speicherstadt-channel-2976711/>

<https://pixabay.com/photos/beer-machine-alcohol-brewery-1513436/>

<https://pixabay.com/photos/audience-band-concert-crowd-1853662/>

<https://pixabay.com/photos/metal-tube-chair-stack-stacked-1756245/>

<https://pixabay.com/photos/above-summary-the-activities-of-the-1956156/>